

# Bibliography

The following abbreviations are used for frequently cited conferences and journals:

<b>AAAI</b>	Proceedings of the AAAI Conference on Artificial Intelligence
<b>AAMAS</b>	Proceedings of the International Conference on Autonomous Agents and Multi-agent Systems
<b>ACL</b>	Proceedings of the Annual Meeting of the Association for Computational Linguistics
<b>AIJ</b>	Artificial Intelligence (Journal)
<b>AIMag</b>	AI Magazine
<b>AIPS</b>	Proceedings of the International Conference on AI Planning Systems
<b>AISTATS</b>	Proceedings of the International Conference on Artificial Intelligence and Statistics
<b>BBS</b>	Behavioral and Brain Sciences
<b>CACM</b>	Communications of the Association for Computing Machinery
<b>COGSCI</b>	Proceedings of the Annual Conference of the Cognitive Science Society
<b>COLING</b>	Proceedings of the International Conference on Computational Linguistics
<b>COLT</b>	Proceedings of the Annual ACM Workshop on Computational Learning Theory
<b>CP</b>	Proceedings of the International Conference on Principles and Practice of Constraint Programming
<b>CVPR</b>	Proceedings of the IEEE Conference on Computer Vision and Pattern Recognition
<b>EC</b>	Proceedings of the ACM Conference on Electronic Commerce
<b>ECAI</b>	Proceedings of the European Conference on Artificial Intelligence
<b>ECCV</b>	Proceedings of the European Conference on Computer Vision
<b>ECML</b>	Proceedings of the The European Conference on Machine Learning
<b>ECP</b>	Proceedings of the European Conference on Planning
<b>EMNLP</b>	Proceedings of the Conference on Empirical Methods in Natural Language Processing
<b>FGCS</b>	Proceedings of the International Conference on Fifth Generation Computer Systems
<b>FOCS</b>	Proceedings of the Annual Symposium on Foundations of Computer Science
<b>GECCO</b>	Proceedings of the Genetics and Evolutionary Computing Conference
<b>HRI</b>	Proceedings of the International Conference on Human-Robot Interaction
<b>ICAPS</b>	Proceedings of the International Conference on Automated Planning and Scheduling
<b>ICASSP</b>	Proceedings of the International Conference on Acoustics, Speech, and Signal Processing
<b>ICCV</b>	Proceedings of the International Conference on Computer Vision
<b>ICLP</b>	Proceedings of the International Conference on Logic Programming
<b>ICLR</b>	Proceedings of the International Conference on Learning Representations
<b>ICML</b>	Proceedings of the International Conference on Machine Learning
<b>ICPR</b>	Proceedings of the International Conference on Pattern Recognition
<b>ICRA</b>	Proceedings of the IEEE International Conference on Robotics and Automation
<b>ICSLP</b>	Proceedings of the International Conference on Speech and Language Processing
<b>IJAR</b>	International Journal of Approximate Reasoning
<b>IJCAI</b>	Proceedings of the International Joint Conference on Artificial Intelligence
<b>IJCNN</b>	Proceedings of the International Joint Conference on Neural Networks
<b>IJCV</b>	International Journal of Computer Vision
<b>ILP</b>	Proceedings of the International Workshop on Inductive Logic Programming
<b>IROS</b>	Proceedings of the International Conference on Intelligent Robots and Systems
<b>ISMIS</b>	Proceedings of the International Symposium on Methodologies for Intelligent Systems
<b>ISRR</b>	Proceedings of the International Symposium on Robotics Research
<b>JACM</b>	Journal of the Association for Computing Machinery
<b>JAIR</b>	Journal of Artificial Intelligence Research
<b>JAR</b>	Journal of Automated Reasoning
<b>JASA</b>	Journal of the American Statistical Association
<b>JMLR</b>	Journal of Machine Learning Research
<b>JSL</b>	Journal of Symbolic Logic
<b>KDD</b>	Proceedings of the International Conference on Knowledge Discovery and Data Mining
<b>KR</b>	Proceedings of the International Conference on Principles of Knowledge Representation and Reasoning
<b>LICS</b>	Proceedings of the IEEE Symposium on Logic in Computer Science
<b>NeurIPS</b>	Advances in Neural Information Processing Systems
<b>PAMI</b>	IEEE Transactions on Pattern Analysis and Machine Intelligence
<b>PNAS</b>	Proceedings of the National Academy of Sciences of the United States of America
<b>PODS</b>	Proceedings of the ACM International Symposium on Principles of Database Systems
<b>RSS</b>	Proceedings of the Conference on Robotics: Science and Systems
<b>SIGIR</b>	Proceedings of the Special Interest Group on Information Retrieval
<b>SIGMOD</b>	Proceedings of the ACM SIGMOD International Conference on Management of Data
<b>SODA</b>	Proceedings of the Annual ACM-SIAM Symposium on Discrete Algorithms
<b>STOC</b>	Proceedings of the Annual ACM Symposium on Theory of Computing
<b>TARK</b>	Proceedings of the Conference on Theoretical Aspects of Reasoning about Knowledge
<b>UAI</b>	Proceedings of the Conference on Uncertainty in Artificial Intelligence

- Aaronson, S.** (2014). My conversation with "Eugene Goostman," the chatbot that's all over the news for allegedly passing the Turing test. Shtetl-Optimized, [www.scottaaronson.com/blog/?p=1858](http://www.scottaaronson.com/blog/?p=1858).
- Aarts, E.** and Lenstra, J. K. (2003). *Local Search in Combinatorial Optimization*. Princeton University Press.
- Aarup, M.**, Arentoft, M. M., Parrod, Y., Stader, J., and Stokes, I. (1994). OPTIMUM-AIV: A knowledge-based planning and scheduling system for spacecraft AIV. In Fox, M. and Zweber, M. (Eds.), *Knowledge Based Scheduling*. Morgan Kaufmann.
- Abbas, A.** (2018). *Foundations of Multiattribute Utility*. Cambridge University Press.
- Abbeel, P.** and Ng, A. Y. (2004). Apprenticeship learning via inverse reinforcement learning. In *ICML-04*.
- Abney, S.**, McAllester, D. A., and Pereira, F. (1999). Relating probabilistic grammars and automata. In *ACL-99*.
- Abramson, B.** (1987). *The expected-outcome model of two-player games*. Ph.D. thesis, Columbia University.
- Abramson, B.** (1990). Expected-outcome: A general model of static evaluation. *PAMI*, 12, 182–193.
- Abreu, D.** and Rubinstein, A. (1988). The structure of Nash equilibrium in repeated games with finite automata. *Econometrica*, 56, 1259–1281.
- Achlioptas, D.** (2009). Random satisfiability. In Biere, A., Heule, M., van Maaren, H., and Walsh, T. (Eds.), *Handbook of Satisfiability*. IOS Press.
- Ackerman, E.** and Guizzo, E. (2016). The next generation of Boston Dynamics' Atlas robot is quiet, robust, and tether free. *IEEE Spectrum*, 24, 2016.
- Ackerman, N.**, Freer, C., and Roy, D. (2013). On the computability of conditional probability. arXiv 1005.3014.
- Ackley, D. H.** and Littman, M. L. (1991). Interactions between learning and evolution. In Langton, C., Taylor, C., Farmer, J. D., and Rasmussen, S. (Eds.), *Artificial Life II*. Addison-Wesley.
- Adida, B.** and Birbeck, M. (2008). RDFa primer. Tech. rep., W3C.
- Adolph, K. E.**, Kretch, K. S., and LoBue, V. (2014). Fear of heights in infants? *Current Directions in Psychological Science*, 23, 60–66.
- Agerbeck, C.** and Hansen, M. O. (2008). A multi-agent approach to solving *NP*-complete problems. Master's thesis, Technical Univ. of Denmark.
- Aggarwal, G.**, Goel, A., and Motwani, R. (2006). Truthful auctions for pricing search keywords. In *EC-06*.
- Agha, G.** (1986). *ACTORS: A Model of Concurrent Computation in Distributed Systems*. MIT Press.
- Agichtein, E.** and Gravano, L. (2003). Querying text databases for efficient information extraction. In *Proc. IEEE Conference on Data Engineering*.
- Agmon, S.** (1954). The relaxation method for linear inequalities. *Canadian Journal of Mathematics*, 6, 382–392.
- Agostinelli, F.**, McAleer, S., Shmakov, A., and Baldi, P. (2019). Solving the Rubik's Cube with deep reinforcement learning and search. *Nature Machine Intelligence*, 1, 356–363.
- Agrawal, P.**, Nair, A. V., Abbeel, P., Malik, J., and Levine, S. (2017). Learning to poke by poking: Experiential learning of intuitive physics. In *NeurIPS-29*.
- Agre, P. E.** and Chapman, D. (1987). Pengi: an implementation of a theory of activity. In *IJCAI-87*.
- Aizerman, M.**, Braverman, E., and Rozonoer, L. (1964). Theoretical foundations of the potential function method in pattern recognition learning. *Automation and Remote Control*, 25, 821–837.
- Akametalu, A. K.**, Fisac, J. F., Gillula, J. H., Kaynama, S., Zeilinger, M. N., and Tomlin, C. J. (2014). Reachability-based safe learning with Gaussian processes. In *53rd IEEE Conference on Decision and Control*.
- Akgun, B.**, Cakmak, M., Jiang, K., and Thomaz, A. (2012). Keyframe-based learning from demonstration. *International Journal of Social Robotics*, 4, 343–355.
- Aldous, D.** and Vazirani, U. (1994). "Go with the winners" algorithms. In *FOCS-94*.
- Alemi, A. A.**, Chollet, F., Een, N., Irving, G., Szegedy, C., and Urban, J. (2017). DeepMath - Deep sequence models for premise selection. In *NeurIPS-29*.
- Allais, M.** (1953). Le comportement de l'homme rationnel devant la risque: critique des postulats et axiomes de l'école Américaine. *Econometrica*, 21, 503–546.
- Allan, J.**, Harman, D., Kanoulas, E., Li, D., Van Gysel, C., and Vorhees, E. (2017). Trec 2017 common core track overview. In *Proc. TREC*.
- Allen, J. F.** (1983). Maintaining knowledge about temporal intervals. *CACM*, 26, 832–843.
- Allen, J. F.** (1984). Towards a general theory of action and time. *AIJ*, 23, 123–154.
- Allen, J. F.** (1991). Time and time again: The many ways to represent time. *Int. J. Intelligent Systems*, 6, 341–355.
- Allen, J. F.**, Hendler, J., and Tate, A. (Eds.). (1990). *Readings in Planning*. Morgan Kaufmann.
- Allen, P.** and Greaves, M. (2011). The singularity isn't near. *Technology review*, 12, 7–8.
- Allen-Zhu, Z.**, Li, Y., and Song, Z. (2018). A convergence theory for deep learning via over-parameterization. arXiv:1811.03962.
- Alterman, R.** (1988). Adaptive planning. *Cognitive Science*, 12, 393–422.
- Amarel, S.** (1967). An approach to heuristic problem-solving and theorem proving in the propositional calculus. In Hart, J. and Takasu, S. (Eds.), *Systems and Computer Science*. University of Toronto Press.
- Amarel, S.** (1968). On representations of problems of reasoning about actions. In Michie, D. (Ed.), *Machine Intelligence 3*, Vol. 3. Elsevier.
- Amir, E.** and Russell, S. J. (2003). Logical filtering. In *IJCAI-03*.
- Amit, Y.** and Geman, D. (1997). Shape quantization and recognition with randomized trees. *Neural Computation*, 9, 1545–1588.
- Amodei, D.** and Hernandez, D. (2018). AI and compute. OpenAI blog, [blog.openai.com/ai-and-compute/](https://blog.openai.com/ai-and-compute/).
- Amodei, D.**, Olah, C., Steinhardt, J., Christiano, P., Schulman, J., and Mané, D. (2016). Concrete problems in AI safety. arXiv:1606.06565.
- Andersen, S. K.**, Olesen, K. G., Jensen, F. V., and Jensen, F. (1989). HUGIN—A shell for building Bayesian belief universes for expert systems. In *IJCAI-89*.
- Anderson, J. R.** (1980). *Cognitive Psychology and Its Implications*. W. H. Freeman.
- Anderson, J. R.** (1983). *The Architecture of Cognition*. Harvard University Press.
- Anderson, K.**, Sturtevant, N. R., Holte, R. C., and Schaeffer, J. (2008). Coarse-to-fine search techniques. Tech. rep., University of Alberta.
- Andoni, A.** and Indyk, P. (2006). Near-optimal hashing algorithms for approximate nearest neighbor in high dimensions. In *FOCS-06*.
- Andor, D.**, Alberti, C., Weiss, D., Severyn, A., Presta, A., Ganchev, K., Petrov, S., and Collins, M. (2016). Globally normalized transition-based neural networks. arXiv:1603.06042.
- Andre, D.**, Friedman, N., and Parr, R. (1998). Generalized prioritized sweeping. In *NeurIPS-10*.
- Andre, D.** and Russell, S. J. (2002). State abstraction for programmable reinforcement learning agents. In *AAAI-02*.
- Andreae, P.** (1985). *Justified Generalisation: Learning Procedures from Examples*. Ph.D. thesis, MIT.
- Andrieu, C.**, Doucet, A., and Holenstein, R. (2010). Particle Markov chain Monte Carlo methods. *J. Royal Statistical Society, Series B*, 72, 269–342.
- Andrychowicz, M.**, Baker, B., Chociej, M., Jozefowicz, R., McGrew, B., Pachocki, J., Petron, A., Plappert, M., Powell, G., Ray, A., et al. (2018a). Learning dexterous in-hand manipulation. arXiv:1808.00177.
- Andrychowicz, M.**, Wolski, F., Ray, A., Schneider, J., Fong, R., Welinder, P., McGrew, B., Tobin, J., Abbeel, P., and Zaremba, W. (2018b). Hindsight experience replay. In *NeurIPS-30*.
- Aneja, J.**, Deshpande, A., and Schwing, A. (2018). Convolutional image captioning. In *CVPR-18*.
- Aoki, M.** (1965). Optimal control of partially observable Markov systems. *J. Franklin Institute*, 280, 367–386.
- Appel, K.** and Haken, W. (1977). Every planar map is four colorable: Part I: Discharging. *Illinois J. Math.*, 21, 429–490.
- Appelt, D.** (1999). Introduction to information extraction. *AJ Communications*, 12, 161–172.
- Apt, K. R.** (1999). The essence of constraint propagation. *Theoretical Computer Science*, 221, 179–210.
- Apt, K. R.** (2003). *Principles of Constraint Programming*. Cambridge University Press.
- Apté, C.**, Damerau, F., and Weiss, S. (1994). Automated learning of decision rules for text categorization. *ACM Transactions on Information Systems*, 12, 233–251.
- Arbuthnot, J.** (1692). *Of the Laws of Chance*. Motte, London. Translation into English, with additions, of Huygens (1657).
- Archibald, C.**, Altman, A., and Shoham, Y. (2009). Analysis of a winning computational billiards player. In *IJCAI-09*.
- Arfaee, S. J.**, Zilles, S., and Holte, R. C. (2010). Bootstrap learning of heuristic functions. In *Third Annual Symposium on Combinatorial Search*.
- Argall, B. D.**, Chernova, S., Veloso, M., and Browning, B. (2009). A survey of robot learning from demonstration. *Robotics and autonomous systems*, 57, 469–483.
- Ariely, D.** (2009). *Predictably Irrational* (Revised edition). Harper.
- Arkin, R.** (1998). *Behavior-Based Robotics*. MIT Press.
- Arkin, R.** (2015). The case for banning killer robots: Counterpoint. *CACM*, 58.
- Armando, A.**, Carbone, R., Compagna, L., Cuellar, J., and Tobarra, L. (2008). Formal analysis of SAML 2.0 web browser single sign-on: Breaking the SAML-based single sign-on for Google apps. In *Proc. 6th ACM Workshop on Formal Methods in Security Engineering*.
- Armstrong, S.** and Levinstein, B. (2017). Low impact artificial intelligences. arXiv:1705.10720.
- Arnauld, A.** (1662). *La logique, ou l'art de penser*. Chez Charles Savreux, Paris.
- Arora, N. S.**, Russell, S. J., and Sudderth, E. (2013). NET-VISA: Network processing vertically integrated seismic analysis. *Bull. Seism. Soc. Amer.*, 103, 709–729.
- Arora, S.** (1998). Polynomial time approximation schemes for Euclidean traveling salesman and other geometric problems. *JACM*, 45, 753–782.

- Arpit**, D., Jastrzebski, S., Ballas, N., Krueger, D., Bengio, E., Kanwal, M. S., Maharaj, T., Fischer, A., Courville, A., Bengio, Y., and Lacoste-Julien, S. (2017). A closer look at memorization in deep networks. arXiv:1706.05394.
- Arrow**, K. J. (1951). *Social Choice and Individual Values*. Wiley.
- Arulampalam**, M. S., Maskell, S., Gordon, N., and Clapp, T. (2002). A tutorial on particle filters for online nonlinear/non-Gaussian Bayesian tracking. *IEEE Transactions on Signal Processing*, 50, 174–188.
- Arulkumaran**, K., Deisenroth, M. P., Brundage, M., and Bharath, A. A. (2017). Deep reinforcement learning: A brief survey. *IEEE Signal Processing Magazine*, 34, 26–38.
- Arunachalam**, R. and Sadeh, N. M. (2005). The supply chain trading agent competition. *Electronic Commerce Research and Applications*, Spring, 66–84.
- Ashby**, W. R. (1940). Adaptiveness and equilibrium. *J. Mental Science*, 86, 478–483.
- Ashby**, W. R. (1948). Design for a brain. *Electronic Engineering, December*, 379–383.
- Ashby**, W. R. (1952). *Design for a Brain*. Wiley.
- Asimov**, I. (1942). Runaround. *Astounding Science Fiction, March*.
- Asimov**, I. (1950). *I, Robot*. Doubleday.
- Asimov**, I. (1958). The feeling of power. *If: Worlds of Science Fiction, February*.
- Astrom**, K. J. (1965). Optimal control of Markov decision processes with incomplete state estimation. *J. Math. Anal. Appl.*, 10, 174–205.
- Atkeson**, C. G., Moore, A. W., and Schaal, S. (1997). Locally weighted learning for control. In *Lazy learning*. Springer.
- Audi**, R. (Ed.). (1999). *The Cambridge Dictionary of Philosophy*. Cambridge University Press.
- Auer**, P., Cesa-Bianchi, N., and Fischer, P. (2002). Finite-time analysis of the multiarmed bandit problem. *Machine Learning*, 47, 235–256.
- Aumann**, R. and Brandenburger, A. (1995). Epistemic conditions for nash equilibrium. *Econometrica*, 67, 1161–1180.
- Axelrod**, R. (1985). *The Evolution of Cooperation*. Basic Books.
- Ba**, J. L., Kiros, J. R., and Hinton, G. E. (2016). Layer normalization. arXiv:1607.06450.
- Baader**, F., Calvanese, D., McGuinness, D., Nardi, D., and Patel-Schneider, P. (2007). *The Description Logic Handbook* (2nd edition). Cambridge University Press.
- Baader**, F. and Snyder, W. (2001). Unification theory. In Robinson, J. and Voronkov, A. (Eds.), *Handbook of Automated Reasoning*. Elsevier.
- Bacchus**, F. (1990). *Representing and Reasoning with Probabilistic Knowledge*. MIT Press.
- Bacchus**, F. and Grove, A. (1995). Graphical models for preference and utility. In *UAI-95*.
- Bacchus**, F. and Grove, A. (1996). Utility independence in a qualitative decision theory. In *KR-96*.
- Bacchus**, F., Grove, A., Halpern, J. Y., and Koller, D. (1992). From statistics to beliefs. In *AAAI-92*.
- Bacchus**, F. and van Beek, P. (1998). On the conversion between non-binary and binary constraint satisfaction problems. In *AAAI-98*.
- Bacchus**, F. and van Run, P. (1995). Dynamic variable ordering in CSPs. In *CP-95*.
- Bacchus**, F., Dalmao, S., and Pitassi, T. (2003). Value elimination: Bayesian inference via backtracking search. In *UAI-03*.
- Bachmann**, P. G. H. (1894). *Die analytische Zahlentheorie*. B. G. Teubner, Leipzig.
- Backus**, J. W. (1959). The syntax and semantics of the proposed international algebraic language of the Zurich ACM-GAMM conference. *Proc. Int'l Conf. on Information Processing*.
- Bacon**, F. (1609). *Wisdom of the Ancients*. Cassell and Company.
- Baeza-Yates**, R. and Ribeiro-Neto, B. (2011). *Modern Information Retrieval* (2nd edition). Addison-Wesley.
- Bagdasaryan**, E., Veit, A., Hua, Y., Estrin, D., and Shmatikov, V. (2018). How to backdoor federated learning. arXiv:1807.00459.
- Bagnell**, J. A. and Schneider, J. (2001). Autonomous helicopter control using reinforcement learning policy search methods. In *ICRA-01*.
- Bahdanau**, D., Cho, K., and Bengio, Y. (2015). Neural machine translation by jointly learning to align and translate. In *ICLR-15*.
- Bahubalendruni**, M. R. and Biswal, B. B. (2016). A review on assembly sequence generation and its automation. *Proc. Institution of Mechanical Engineers, Part C: Journal of Mechanical Engineering Science*, 230, 824–838.
- Bai**, A. and Russell, S. J. (2017). Efficient reinforcement learning with hierarchies of machines by leveraging internal transitions. In *IJCAI-17*.
- Bai**, H., Cai, S., Ye, N., Hsu, D., and Lee, W. S. (2015). Intention-aware online POMDP planning for autonomous driving in a crowd. In *ICRA-15*.
- Bajcsy**, A., Losey, D. P., O’Malley, M. K., and Dragan, A. D. (2017). Learning robot objectives from physical human interaction. *Proceedings of Machine Learning Research*, 78, 217–226.
- Baker**, C. L., Saxe, R., and Tenenbaum, J. B. (2009). Action understanding as inverse planning. *Cognition*, 113, 329–349.
- Baker**, J. (1975). The Dragon system—An overview. *IEEE Transactions on Acoustics, Speech, and Signal Processing*, 23, 24–29.
- Baker**, J. (1979). Trainable grammars for speech recognition. In *Speech Communication Papers for the 97th Meeting of the Acoustical Society of America*.
- Baldi**, P., Chauvin, Y., Hunkapiller, T., and McClure, M. (1994). Hidden Markov models of biological primary sequence information. *PNAS*, 91, 1059–1063.
- Baldwin**, J. M. (1896). A new factor in evolution. *American Naturalist*, 30, 441–451. Continued on pages 536–553.
- Ballard**, B. W. (1983). The \*-minimax search procedure for trees containing chance nodes. *AIJ*, 21, 327–350.
- Baluja**, S. (1997). Genetic algorithms and explicit search statistics. In *NeurIPS 9*.
- Bancilhon**, F., Maier, D., Sagiv, Y., and Ullman, J. D. (1986). Magic sets and other strange ways to implement logic programs. In *PODS-86*.
- Banko**, M. and Brill, E. (2001). Scaling to very very large corpora for natural language disambiguation. In *ACL-01*.
- Banko**, M., Brill, E., Dumais, S. T., and Lin, J. (2002). AskMSR: Question answering using the worldwide web. In *Proc. AAAI Spring Symposium on Mining Answers from Texts and Knowledge Bases*.
- Banko**, M., Cafarella, M. J., Soderland, S., Broadhead, M., and Etzioni, O. (2007). Open information extraction from the web. In *IJCAI-07*.
- Banko**, M. and Etzioni, O. (2008). The tradeoffs between open and traditional relation extraction. In *ACL-08*.
- Bansal**, K., Loos, S., Rabe, M. N., Szegedy, C., and Wilcox, S. (2019). HOList: An environment for machine learning of higher-order theorem proving (extended version). arXiv:1904.03241.
- Bar-Hillel**, Y. (1954). Indexical expressions. *Mind*, 63, 359–379.
- Bar-Shalom**, Y. (Ed.). (1992). *Multitarget-Multisensor Tracking: Advanced Applications*. Artech House.
- Bar-Shalom**, Y. and Fortmann, T. E. (1988). *Tracking and Data Association*. Academic Press.
- Bar-Shalom**, Y., Li, X.-R., and Kirubarajan, T. (2001). *Estimation, Tracking and Navigation: Theory, Algorithms and Software*. Wiley.
- Barber**, D. (2012). *Bayesian Reasoning and Machine Learning*. Cambridge University Press.
- Barr**, A. and Feigenbaum, E. A. (Eds.). (1981). *The Handbook of Artificial Intelligence*, Vol. 1. HeurisTech Press and William Kaufmann.
- Barreiro**, J., Boyce, M., Do, M., Frank, J., Iatauro, M., Kichkaylo, T., Morris, P., Ong, J., Remolina, E., Smith, T., et al. (2012). EUROPA: A platform for AI planning, scheduling, constraint programming, and optimization. *4th International Competition on Knowledge Engineering for Planning and Scheduling (ICKEPS)*.
- Barreno**, M., Nelson, B., Joseph, A. D., and Tygar, J. D. (2010). The security of machine learning. *Machine Learning*, 81, 121–148.
- Barrett**, S. and Stone, P. (2015). Cooperating with unknown teammates in complex domains: A robot soccer case study of ad hoc teamwork. In *AAAI-15*.
- Barták**, R., Salido, M. A., and Rossi, F. (2010). New trends in constraint satisfaction, planning, and scheduling: A survey. *The Knowledge Engineering Review*, 25, 249–279.
- Bartholdi**, J. J., Tovey, C. A., and Trick, M. A. (1989). The computational difficulty of manipulating an election. *Social Choice and Welfare*, 6, 227–241.
- Barto**, A. G., Bradtko, S. J., and Singh, S. (1995). Learning to act using real-time dynamic programming. *AIJ*, 73, 81–138.
- Barto**, A. G., Sutton, R. S., and Brouwer, P. S. (1981). Associative search network: A reinforcement learning associative memory. *Biological Cybernetics*, 40, 201–211.
- Barwise**, J. and Etchemendy, J. (2002). *Language, Proof and Logic*. CSLI Press.
- Baum**, E., Boneh, D., and Garrett, C. (1995). On genetic algorithms. In *COLT-95*.
- Baum**, E. and Smith, W. D. (1997). A Bayesian approach to relevance in game playing. *AIJ*, 97, 195–242.
- Baum**, L. E. and Petrie, T. (1966). Statistical inference for probabilistic functions of finite state Markov chains. *Annals of Mathematical Statistics*, 41, 1554–1563.
- Baxter**, J. and Bartlett, P. (2000). Reinforcement learning in POMDPs via direct gradient ascent. In *ICML-00*.
- Bayardo**, R. J. and Agrawal, R. (2005). Data privacy through optimal k-anonymization. In *Proc. 21st Int'l Conf. on Data Engineering*.
- Bayardo**, R. J. and Miranker, D. P. (1994). An optimal backtrack algorithm for tree-structured constraint satisfaction problems. *AIJ*, 71, 159–181.
- Bayardo**, R. J. and Schrag, R. C. (1997). Using CSP look-back techniques to solve real-world SAT instances. In *AAAI-97*.
- Bayes**, T. (1763). An essay towards solving a problem in the doctrine of chances. *Phil. Trans. Roy. Soc.*, 53, 370–418.

- Beal, J.** and Winston, P. H. (2009). The new frontier of human-level artificial intelligence. *IEEE Intelligent Systems*, 24, 21–23.
- Beardon, A. F.**, Candeal, J. C., Herden, G., Induráin, E., and Mehta, G. B. (2002). The non-existence of a utility function and the structure of non-representable preference relations. *Journal of Mathematical Economics*, 37, 17 – 38.
- Beattie, C.**, Leibo, J. Z., Teplyashin, D., Ward, T., Wainwright, M., Küttler, H., Lefrancq, A., Green, S., Valdés, V., Sadik, A., Schriftwieser, J., Anderson, K., York, S., Cant, M., Cain, A., Bolton, A., Gaffney, S., King, H., Hassabis, D., Legg, S., and Petersen, S. (2016). DeepMind lab. arXiv:1612.03801.
- Bechhofer, R.** (1954). A single-sample multiple decision procedure for ranking means of normal populations with known variances. *Annals of Mathematical Statistics*, 25, 16–39.
- Beck, J. C.**, Feng, T. K., and Watson, J.-P. (2011). Combining constraint programming and local search for job-shop scheduling. *INFORMS Journal on Computing*, 23, 1–14.
- Beckert, B.** and Posegga, J. (1995). Leantap: Lean, tableau-based deduction. *JAR*, 15, 339–358.
- Beeri, C.**, Fagin, R., Maier, D., and Yannakakis, M. (1983). On the desirability of acyclic database schemes. *JACM*, 30, 479–513.
- Bekey, G.** (2008). *Robotics: State Of The Art And Future Challenges*. Imperial College Press.
- Belkin, M.**, Hsu, D., Ma, S., and Mandal, S. (2019). Reconciling modern machine-learning practice and the classical bias-variance trade-off. *PNAS*, 116, 15849–15854.
- Bell, C.** and Tate, A. (1985). Using temporal constraints to restrict search in a planner. In *Proc. Third Alvey IKBS SIG Workshop*.
- Bell, J. L.** and Machover, M. (1977). *A Course in Mathematical Logic*. Elsevier.
- Bellamy, E.** (2003). *Looking Backward: 2000–1887*. Broadview Press.
- Bellamy, R. K. E.**, Dey, K., Hind, M., Hoffman, S. C., Houde, S., Kannan, K., Lohia, P., Martino, J., Mehta, S., Mojsilovic, A., Nagar, S., Ramamurthy, K. N., Richards, J. T., Saha, D., Sattigeri, P., Singh, M., Varshney, K. R., and Zhang, Y. (2018). AI fairness 360: An extensible toolkit for detecting, understanding, and mitigating unwanted algorithmic bias. arXiv:1810.01943.
- Bellemare, M. G.**, Naddaf, Y., Veness, J., and Bowling, M. (2013). The arcade learning environment: An evaluation platform for general agents. *JAIR*, 47, 253–279.
- Bellman, R. E.** (1952). On the theory of dynamic programming. *PNAS*, 38, 716–719.
- Bellman, R. E.** (1958). On a routing problem. *Quarterly of Applied Mathematics*, 16.
- Bellman, R. E.** (1961). *Adaptive Control Processes: A Guided Tour*. Princeton University Press.
- Bellman, R. E.** (1965). On the application of dynamic programming to the determination of optimal play in chess and checkers. *PNAS*, 53, 244–246.
- Bellman, R. E.** (1984). *Eye of the Hurricane*. World Scientific.
- Bellman, R. E.** and Dreyfus, S. E. (1962). *Applied Dynamic Programming*. Princeton University Press.
- Bellman, R. E.** (1957). *Dynamic Programming*. Princeton University Press.
- Ben-Tal, A.** and Nemirovski, A. (2001). *Lectures on Modern Convex Optimization: Analysis, Algorithms, and Engineering Applications*. SIAM (Society for Industrial and Applied Mathematics).
- Bengio, Y.**, Simard, P., and Frasconi, P. (1994). Learning long-term dependencies with gradient descent is difficult. *IEEE Transactions on Neural Networks*, 5, 157–166.
- Bengio, Y.** and Bengio, S. (2001). Modeling high-dimensional discrete data with multi-layer neural networks. In *NeurIPS 13*.
- Bengio, Y.**, Ducharme, R., Vincent, P., and Jauvin, C. (2003). A neural probabilistic language model. *JMLR*, 3, 1137–1155.
- Bengio, Y.** and LeCun, Y. (2007). Scaling learning algorithms towards AI. In Bottou, L., Chapelle, O., DeCoste, D., and Weston, J. (Eds.), *Large-Scale Kernel Machines*. MIT Press.
- Benjamin, M.** (2013). *Drone Warfare: Killing by Remote Control*. Verso Books.
- Bentham, J.** (1823). *Principles of Morals and Legislation*. Oxford University Press, Oxford. Original work published in 1789.
- Benzmüller, C.** and Paleo, B. W. (2013). Formalization, mechanization and automation of Gödel's proof of God's existence. arXiv:1308.4526.
- Beresniak, A.**, Medina-Lara, A., Auray, J. P., De Wever, A., Praet, J.-C., Tarricone, R., Torbica, A., Dupont, D., Lamure, M., and Duru, G. (2015). Validation of the underlying assumptions of the quality-adjusted life-years outcome: Results from the ECHOOUTCOME European project. *PharmacoEconomics*, 33, 61–69.
- Berger, J. O.** (1985). *Statistical Decision Theory and Bayesian Analysis*. Springer Verlag.
- Bergstra, J.** and Bengio, Y. (2012). Random search for hyper-parameter optimization. *JMLR*, 13, 281–305.
- Berk, R.**, Heidari, H., Jabbari, S., Kearns, M., and Roth, A. (2017). Fairness in criminal justice risk assessments: The state of the art. arXiv:1703.09207.
- Berkson, J.** (1944). Application of the logistic function to bio-assay. *JASA*, 39, 357–365.
- Berleur, J.** and Brunnstein, K. (2001). *Ethics of Computing: Codes, Spaces for Discussion and Law*. Chapman and Hall.
- Berlin, K.**, Koren, S., Chin, C.-S., Drake, J. P., Landolin, J. M., and Phillippe, A. M. (2015). Assembling large genomes with single-molecule sequencing and locality-sensitive hashing. *Nature Biotechnology*, 33, 623.
- Berliner, H. J.** (1979). The B\* tree search algorithm: A best-first proof procedure. *AIJ*, 12, 23–40.
- Berliner, H. J.** (1980a). Backgammon computer program beats world champion. *AIJ*, 14, 205–220.
- Berliner, H. J.** (1980b). Computer backgammon. *Scientific American*, 249, 64–72.
- Bermúdez-Chacón, R.**, Gonnet, G. H., and Smith, K. (2015). Automatic problem-specific hyperparameter optimization and model selection for supervised machine learning. Tech. rep., ETH Zurich.
- Bernardo, J. M.** and Smith, A. (1994). *Bayesian Theory*. Wiley.
- Berners-Lee, T.**, Hendler, J., and Lassila, O. (2001). The semantic web. *Scientific American*, 284, 34–43.
- Bernoulli, D.** (1738). Specimen theoriae novae de mensura sortis. *Proc. St. Petersburg Imperial Academy of Sciences*, 5, 175–192.
- Bernstein, P. L.** (1996). *Against the Gods: The Remarkable Story of Risk*. Wiley.
- Berrada, L.**, Zisserman, A., and Kumar, M. P. (2019). Training neural networks for and by interpolation. arXiv:1906.05661.
- Berrou, C.**, Glavieux, A., and Thitimajshima, P. (1993). Near Shannon limit error control-correcting coding and decoding: Turbo-codes. I. In *Proc. IEEE International Conference on Communications*.
- Berry, D. A.** and Fristedt, B. (1985). *Bandit Problems: Sequential Allocation of Experiments*. Chapman and Hall.
- Bertele, U.** and Brioschi, F. (1972). *Nonserial Dynamic Programming*. Academic Press.
- Bertoli, P.**, Cimatti, A., and Roveri, M. (2001a). Heuristic search + symbolic model checking = efficient conformant planning. In *IJCAI-01*.
- Bertoli, P.**, Cimatti, A., Roveri, M., and Traverso, P. (2001b). Planning in nondeterministic domains under partial observability via symbolic model checking. In *IJCAI-01*.
- Bertot, Y.**, Casteran, P., Huet, G., and Paulin-Mohring, C. (2004). *Interactive Theorem Proving and Program Development*. Springer.
- Bertsekas, D.** (1987). *Dynamic Programming: Deterministic and Stochastic Models*. Prentice-Hall.
- Bertsekas, D.** and Tsitsiklis, J. N. (1996). *Neuro-Dynamic Programming*. Athena Scientific.
- Bertsekas, D.** and Tsitsiklis, J. N. (2008). *Introduction to Probability* (2nd edition). Athena Scientific.
- Bertsekas, D.** and Shreve, S. E. (2007). *Stochastic Optimal Control: The Discrete-Time Case*. Athena Scientific.
- Bertsimas, D.**, Delarue, A., and Martin, S. (2019). Optimizing schools' start time and bus routes. *PNAS*, 116, 13, 5943–5948.
- Bertsimas, D.** and Dunn, J. (2017). Optimal classification trees. *Machine Learning*, 106, 1039–1082.
- Bessen, J.** (2015). *Learning by Doing: The Real Connection between Innovation, Wages, and Wealth*. Yale University Press.
- Bessière, C.** (2006). Constraint propagation. In Rossi, F., van Beek, P., and Walsh, T. (Eds.), *Handbook of Constraint Programming*. Elsevier.
- Beutel, A.**, Chen, J., Doshi, T., Qian, H., Woodruff, A., Luu, C., Kreitmann, P., Bischof, J., and Chi, E. H. (2019). Putting fairness principles into practice: Challenges, metrics, and improvements. arXiv:1901.04562.
- Bhar, R.** and Hamori, S. (2004). *Hidden Markov Models: Applications to Financial Economics*. Springer.
- Bibel, W.** (1993). *Deduction: Automated Logic*. Academic Press.
- Bien, J.**, Tibshirani, R., et al. (2011). Prototype selection for interpretable classification. *Annals of Applied Statistics*, 5, 2403–2424.
- Biere, A.**, Heule, M., van Maaren, H., and Walsh, T. (Eds.). (2009). *Handbook of Satisfiability*. IOS Press.
- Bies, A.**, Mott, J., and Warner, C. (2015). English news text treebank: Penn treebank revised. Linguistic Data Consortium.
- Billings, D.**, Burch, N., Davidson, A., Holte, R. C., Schaeffer, J., Schauenberg, T., and Szafron, D. (2003). Approximating game-theoretic optimal strategies for full-scale poker. In *IJCAI-03*.
- Billingsley, P.** (2012). *Probability and Measure* (4th edition). Wiley.
- Binder, J.**, Koller, D., Russell, S. J., and Kanazawa, K. (1997a). Adaptive probabilistic networks with hidden variables. *Machine Learning*, 29, 213–244.
- Binder, J.**, Murphy, K., and Russell, S. J. (1997b). Space-efficient inference in dynamic probabilistic networks. In *IJCAI-97*.
- Bingham, E.**, Chen, J., Jankowiak, M., Obermeyer, F., Pradhan, N., Karaletsos, T., Singh, R., Szerlip, P., Horsfall, P., and Goodman, N. D. (2019). Pyro: Deep universal probabilistic programming. *JMLR*, 20, 1–26.
- Binmore, K.** (1982). *Essays on Foundations of Game Theory*. Pitman.

- Biran, O.** and Cotton, C. (2017). Explanation and justification in machine learning: A survey. In *Proc. IJCAI-17 Workshop on Explainable AI*.
- Bishop, C. M.** (1995). *Neural Networks for Pattern Recognition*. Oxford University Press.
- Bishop, C. M.** (2007). *Pattern Recognition and Machine Learning*. Springer-Verlag.
- Bisson, T.** (1990). They're made out of meat. *Omni Magazine*.
- Bistarelli, S.**, Montanari, U., and Rossi, F. (1997). Semiring-based constraint satisfaction and optimization. *JACM*, 44, 201–236.
- Bitner, J. R.** and Reingold, E. M. (1975). Backtrack programming techniques. *CACM*, 18, 651–656.
- Bizer, C.**, Auer, S., Kobilarov, G., Lehmann, J., and Cyganiak, R. (2007). DBpedia – querying Wikipedia like a database. In *16th International Conference on World Wide Web*.
- Blazewicz, J.**, Ecker, K., Pesch, E., Schmidt, G., and Weglarz, J. (2007). *Handbook on Scheduling: Models and Methods for Advanced Planning*. Springer-Verlag.
- Blei, D. M.**, Ng, A. Y., and Jordan, M. I. (2002). Latent Dirichlet allocation. In *NeurIPS 14*.
- Bliss, C. I.** (1934). The method of probits. *Science*, 79, 38–39.
- Block, H. D.**, Knight, B., and Rosenblatt, F. (1962). Analysis of a four-layer series-coupled perceptron. *Rev. Modern Physics*, 34, 275–282.
- Block, N.** (2009). Comparing the major theories of consciousness. In Gazzaniga, M. S. (Ed.), *The Cognitive Neurosciences*. MIT Press.
- Blum, A. L.** and Furst, M. (1997). Fast planning through planning graph analysis. *AIJ*, 90, 281–300.
- Blum, A. L.** (1996). On-line algorithms in machine learning. In *Proc. Workshop on On-Line Algorithms, Dagstuhl*.
- Blum, A. L.**, Hopcroft, J., and Kannan, R. (2020). *Foundations of Data Science*. Cambridge University Press.
- Blum, A. L.** and Mitchell, T. M. (1998). Combining labeled and unlabeled data with co-training. In *COLT-98*.
- Blumer, A.**, Ehrenfeucht, A., Haussler, D., and Warmuth, M. (1989). Learnability and the Vapnik-Chervonenkis dimension. *JACM*, 36, 929–965.
- Bobrow, D. G.** (1967). Natural language input for a computer problem solving system. In Minsky, M. L. (Ed.), *Semantic Information Processing*. MIT Press.
- Bod, R.** (2008). The data-oriented parsing approach: Theory and application. In *Computational Intelligence: A Compendium*. Springer-Verlag.
- Bod, R.**, Scha, R., and Sima'an, K. (2003). *Data-Oriented Parsing*. CSLI Press.
- Boddington, P.** (2017). *Towards a Code of Ethics for Artificial Intelligence*. Springer-Verlag.
- Boden, M. A.** (Ed.). (1990). *The Philosophy of Artificial Intelligence*. Oxford University Press.
- Bolognesi, A.** and Ciancarini, P. (2003). Computer programming of kriegspiel endings: The case of KR vs. K. In *Advances in Computer Games 10*.
- Bolton, R. J.** and Hand, D. J. (2002). Statistical fraud detection: A review. *Statistical science*, 17, 235–249.
- Bonawitz, K.**, Ivanov, V., Kreuter, B., Marcedone, A., McMahan, H. B., Patel, S., Ramage, D., Segal, A., and Seth, K. (2017). Practical secure aggregation for privacy-preserving machine learning. In *Proc. ACM SIGSAC Conference on Computer and Communications Security*.
- Bond, A. H.** and Gasser, L. (Eds.). (1988). *Readings in Distributed Artificial Intelligence*. Morgan Kaufmann.
- Bonet, B.** (2002). An epsilon-optimal grid-based algorithm for partially observable Markov decision processes. In *ICML-02*.
- Bonet, B.** and Geffner, H. (1999). Planning as heuristic search: New results. In *ECP-99*.
- Bonet, B.** and Geffner, H. (2000). Planning with incomplete information as heuristic search in belief space. In *ICAPS-00*.
- Bonet, B.** and Geffner, H. (2005). An algorithm better than AO\*? In *AAAI-05*.
- Boole, G.** (1847). *The Mathematical Analysis of Logic: Being an Essay towards a Calculus of Deductive Reasoning*. Macmillan, Barclay, and Macmillan.
- Booth, T. L.** (1969). Probabilistic representation of formal languages. In *IEEE Conference Record of the 1969 Tenth Annual Symposium on Switching and Automata Theory*.
- Borel, E.** (1921). La théorie du jeu et les équations intégrales à noyau symétrique. *Comptes Rendus Hebdomadaires des Séances de l'Académie des Sciences*, 173, 1304–1308.
- Borenstein, J.**, Everett, B., and Feng, L. (1996). *Navigating Mobile Robots: Systems and Techniques*. A. K. Peters, Ltd.
- Borenstein, J.** and Koren, Y. (1991). The vector field histogram—Fast obstacle avoidance for mobile robots. *IEEE Transactions on Robotics and Automation*, 7, 278–288.
- Borgida, A.**, Brachman, R. J., McGuinness, D., and Alperin Resnick, L. (1989). CLASSIC: A structural data model for objects. *SIGMOD Record*, 18, 58–67.
- Boroditsky, L.** (2003). Linguistic relativity. In Nadel, L. (Ed.), *Encyclopedia of Cognitive Science*. Macmillan.
- Boser, B.**, Guyon, I., and Vapnik, V. N. (1992). A training algorithm for optimal margin classifiers. In *COLT-92*.
- Bosse, M.**, Newman, P., Leonard, J., Soika, M., Feiten, W., and Teller, S. (2004). Simultaneous localization and map building in large-scale cyclic environments using the Atlas framework. *Int. J. Robotics Research*, 23, 1113–1139.
- Bostrom, N.** (2005). A history of transhumanist thought. *Journal of Evolution and Technology*, 14, 1–25.
- Bostrom, N.** (2014). *Superintelligence: Paths, Dangers, Strategies*. Oxford University Press.
- Bottou, L.** and Bousquet, O. (2008). The tradeoffs of large scale learning. In *NeurIPS 20*.
- Bottou, L.**, Curtis, F. E., and Nocedal, J. (2018). Optimization methods for large-scale machine learning. *SIAM Review*, 60, 223–311.
- Boué, L.** (2019). Real numbers, data science and chaos: How to fit any dataset with a single parameter. arXiv:1904.12320.
- Bousmalis, K.**, Irpan, A., Wohlhart, P., Bai, Y., Kelcey, M., Kalakrishnan, M., Downs, L., Ibarz, J., Pastor, P., Konolige, K., Levine, S., and Vanhoucke, V. (2017). Using simulation and domain adaptation to improve efficiency of deep robotic grasping. arXiv:1709.07857.
- Boutilier, C.** (2002). A POMDP formulation of preference elicitation problems. In *AAAI-02*.
- Boutilier, C.** and Brafman, R. I. (2001). Partial-order planning with concurrent interacting actions. *JAIR*, 14, 105–136.
- Boutilier, C.**, Dearden, R., and Goldszmidt, M. (2000). Stochastic dynamic programming with factored representations. *AIJ*, 121, 49–107.
- Boutilier, C.**, Reiter, R., and Price, B. (2001). Symbolic dynamic programming for first-order MDPs. In *IJCAI-01*.
- Boutilier, C.**, Brafman, R. I., Domshlak, C., Hoos, H. H., and Poole, D. (2004). CP-nets: A tool for representing and reasoning with conditional *ceteris paribus* preference statements. *JAIR*, 21, 135–191.
- Boutilier, C.**, Friedman, N., Goldszmidt, M., and Koller, D. (1996). Context-specific independence in Bayesian networks. In *UAI-96*.
- Bouzy, B.** and Cazenave, T. (2001). Computer Go: An AI oriented survey. *AIJ*, 132, 39–103.
- Bowling, M.**, Burch, N., Johanson, M., and Tammelin, O. (2015). Heads-up limit hold'em poker is solved. *Science*, 347, 145–149.
- Bowling, M.**, Johanson, M., Burch, N., and Szafrański, D. (2008). Strategy evaluation in extensive games with importance sampling. In *ICML-08*.
- Bowman, S.**, Angeli, G., Potts, C., and Manning, C. (2015). A large annotated corpus for learning natural language inference. In *EMNLP-15*.
- Box, G. E. P.** (1957). Evolutionary operation: A method of increasing industrial productivity. *Applied Statistics*, 6, 81–101.
- Box, G. E. P.**, Jenkins, G., Reinsel, G., and Ljung, G. M. (2016). *Time Series Analysis: Forecasting and Control* (5th edition). Wiley.
- Box, G. E. P.** and Tiao, G. C. (1973). *Bayesian Inference in Statistical Analysis*. Addison-Wesley.
- Boyan, J. A.** and Moore, A. W. (1998). Learning evaluation functions for global optimization and Boolean satisfiability. In *AAAI-98*.
- Boyd, S.** and Vandenberghe, L. (2004). *Convex Optimization*. Cambridge University Press.
- Boyer, X.**, Friedman, N., and Koller, D. (1999). Discovering the hidden structure of complex dynamic systems. In *UAI-99*.
- Boyer, R. S.** and Moore, J. S. (1979). *A Computational Logic*. Academic Press.
- Boyer, R. S.** and Moore, J. S. (1984). Proof checking the RSA public key encryption algorithm. *American Mathematical Monthly*, 91, 181–189.
- Brachman, R. J.** (1979). On the epistemological status of semantic networks. In Findler, N. V. (Ed.), *Associative Networks: Representation and Use of Knowledge by Computers*. Academic Press.
- Brachman, R. J.** and Levesque, H. J. (Eds.). (1985). *Readings in Knowledge Representation*. Morgan Kaufmann.
- Bradt, R. N.**, Johnson, S. M., and Karlin, S. (1956). On sequential designs for maximizing the sum of n observations. *Ann. Math. Statist.*, 27, 1060–1074.
- Brafman, O.** and Brafman, R. (2009). *Sway: The Irresistible Pull of Irrational Behavior*. Broadway Business.
- Brafman, R. I.** and Domshlak, C. (2008). From one to many: Planning for loosely coupled multi-agent systems. In *ICAPS-08*.
- Brafman, R. I.** and Tennenholtz, M. (2000). A near optimal polynomial time algorithm for learning in certain classes of stochastic games. *AIJ*, 121, 31–47.
- Braitenberg, V.** (1984). *Vehicles: Experiments in Synthetic Psychology*. MIT Press.
- Brandt, F.**, Conitzer, V., Endriss, U., Lang, J., and Procaccia, A. D. (Eds.). (2016). *Handbook of Computational Social Choice*. Cambridge University Press.
- Brants, T.** (2000). TnT: A statistical part-of-speech tagger. In *Proc. Sixth Conference on Applied Natural Language Processing*.
- Brants, T.**, Popat, A. C., Xu, P., Och, F. J., and Dean, J. (2007). Large language models in machine translation. In *EMNLP-CoNLL-07*.
- Bratko, I.** (2009). *Prolog Programming for Artificial Intelligence* (4th edition). Addison-Wesley.

- Bratman, M. E.** (1987). *Intention, Plans, and Practical Reason*. Harvard University Press.
- Breck, E., Cai, S., Nielsen, E., Salib, M., and Sculley, D.** (2016). What's your ML test score? A rubric for ML production systems. In *Proc. NIPS 2016 Workshop on Reliable Machine Learning in the Wild*.
- Breese, J. S.** (1992). Construction of belief and decision networks. *Computational Intelligence*, 8, 624–647.
- Breese, J. S. and Heckerman, D.** (1996). Decision-theoretic troubleshooting: A framework for repair and experiment. In *UAI-96*.
- Breiman, L., Friedman, J., Olshen, R. A., and Stone, C. J.** (1984). *Classification and Regression Trees*. Wadsworth International Group.
- Breiman, L.** (2001). Random forests. *Machine Learning*, 45(1), 5–32.
- Brelaz, D.** (1979). New methods to color the vertices of a graph. *CACM*, 22, 251–256.
- Brent, R. P.** (1973). *Algorithms for Minimization without Derivatives*. Prentice-Hall.
- Bresnan, J.** (1982). *The Mental Representation of Grammatical Relations*. MIT Press.
- Brewka, G., Dix, J., and Konolige, K.** (1997). *Nonmonotonic Reasoning: An Overview*. Center for the Study of Language and Information (CSLI).
- Brickley, D. and Guha, R. V.** (2004). RDF vocabulary description language 1.0: RDF schema. Tech. rep., W3C.
- Briggs, R.** (1985). Knowledge representation in Sanskrit and artificial intelligence. *AIMag*, 6, 32–39.
- Brill, E.** (1992). A simple rule-based part of speech tagger. In *Proc. Third Conference on Applied Natural Language Processing*.
- Brin, D.** (1998). *The Transparent Society*. Perseus.
- Brin, S. and Page, L.** (1998). The anatomy of a large-scale hypertextual web search engine. In *Proc. Seventh World Wide Web Conference*.
- Bringsjord, S.** (2008). If I were judge. In Epstein, R., Roberts, G., and Beber, G. (Eds.), *Parsing the Turing Test*. Springer.
- Broadbent, D. E.** (1958). *Perception and Communication*. Pergamon.
- Brockman, G., Cheung, V., Pettersson, L., Schneider, J., Schulman, J., Tang, J., and Zaremba, W.** (2016). OpenAI gym. arXiv:1606.01540.
- Brooks, R. A.** (1986). A robust layered control system for a mobile robot. *IEEE J. of Robotics and Automation*, 2, 14–23.
- Brooks, R. A.** (1989). Engineering approach to building complete, intelligent beings. *Proc. SPIE—the International Society for Optical Engineering*, 1002, 618–625.
- Brooks, R. A.** (1991). Intelligence without representation. *AIJ*, 47, 139–159.
- Brooks, R. A. and Lozano-Perez, T.** (1985). A subdivision algorithm in configuration space for findpath with rotation. *IEEE Transactions on Systems, Man and Cybernetics*, 15, 224–233.
- Brooks, R. A.** (2017). The seven deadly sins of AI predictions. *MIT Technology Review*, Oct 6.
- Brooks, S., Gelman, A., Jones, G., and Meng, X.-L.** (2011). *Handbook of Markov Chain Monte Carlo*. Chapman & Hall/CRC.
- Brown, C., Finkelstein, L., and Purdom, P.** (1988). Backtrack searching in the presence of symmetry. In Mora, T. (Ed.), *Applied Algebra, Algebraic Algorithms and Error-Correcting Codes*. Springer-Verlag.
- Brown, K. C.** (1974). A note on the apparent bias of net revenue estimates. *J. Finance*, 29, 1215–1216.
- Brown, N. and Sandholm, T.** (2017). Libratus: The superhuman AI for no-limit poker. In *IJCAI-17*.
- Brown, N. and Sandholm, T.** (2019). Superhuman AI for multiplayer poker. *Science*, 365, 885–890.
- Brown, P. F., Cocke, J., Della Pietra, S. A., Della Pietra, V. J., Jelinek, F., Mercer, R. L., and Roossin, P.** (1988). A statistical approach to language translation. In *COLING-88*.
- Brown, P. F., Desouza, P. V., Mercer, R. L., Pietra, V. J. D., and Lai, J. C.** (1992). Class-based n-gram models of natural language. *Computational linguistics*, 18(4).
- Browne, C., Powley, E. J., Whitehouse, D., Lucas, S. M., Cowling, P. I., Rohlfshagen, P., Tavener, S., Liebana, D. P., Samothrakis, S., and Colton, S.** (2012). A survey of Monte Carlo tree search methods. *IEEE Transactions on Computational Intelligence and AI in Games*, 4, 1–43.
- Brownston, L., Farrell, R., Kant, E., and Martin, N.** (1985). *Programming Expert Systems in OPS5: An Introduction to Rule-Based Programming*. Addison-Wesley.
- Bruce, V., Green, P., and Georgeson, M.** (2003). *Visual Perception: Physiology, Psychology and Ecology*. Routledge and Kegan Paul.
- Brügmann, B.** (1993). Monte Carlo Go. Tech. rep., Department of Physics, Syracuse University.
- Bryce, D. and Kambhampati, S.** (2007). A tutorial on planning graph-based reachability heuristics. *AIMag, Spring*, 47–83.
- Bryce, D., Kambhampati, S., and Smith, D. E.** (2006). Planning graph heuristics for belief space search. *JAIR*, 26, 35–99.
- Brynjolfsson, E. and McAfee, A.** (2011). *Race Against the Machine*. Digital Frontier Press.
- Brynjolfsson, E. and McAfee, A.** (2014). *The Second Machine Age*. W. W. Norton.
- Brynjolfsson, E., Rock, D., and Syverson, C.** (2018). Artificial intelligence and the modern productivity paradox: A clash of expectations and statistics. In Agrawal, A., Gans, J., and Goldfarb, A. (Eds.), *The Economics of Artificial Intelligence: An Agenda*. University of Chicago Press.
- Bryson, A. E. and Ho, Y.-C.** (1969). *Applied Optimal Control*. Blaisdell.
- Bryson, A. E.** (1962). A gradient method for optimizing multi-stage allocation processes. In *Proc. of a Harvard Symposium on Digital Computers and Their Applications*.
- Bryson, J. J.** (2012). A role for consciousness in action selection. *International Journal of Machine Consciousness*, 4, 471–482.
- Bryson, J. J. and Winfield, A.** (2017). Standardizing ethical design for artificial intelligence and autonomous systems. *Computer*, 50, 116–119.
- Buchanan, B. G., Mitchell, T. M., Smith, R. G., and Johnson, C. R.** (1978). Models of learning systems. In *Encyclopedia of Computer Science and Technology*, Vol. 11. Dekker.
- Buchanan, B. G. and Shortliffe, E. H.** (Eds.). (1984). *Rule-Based Expert Systems: The MYCIN Experiments of the Stanford Heuristic Programming Project*. Addison-Wesley.
- Buchanan, B. G., Sutherland, G. L., and Feigenbaum, E. A.** (1969). Heuristic DENDRAL: A program for generating explanatory hypotheses in organic chemistry. In Meltzer, B., Michie, D., and Swann, M. (Eds.), *Machine Intelligence 4*. Edinburgh University Press.
- Buck, C., Heafield, K., and Van Ooyen, B.** (2014). N-gram counts and language models from the common crawl. In *Proc. International Conference on Language Resources and Evaluation*.
- Buehler, M., Iagnemma, K., and Singh, S.** (Eds.). (2006). *The 2005 DARPA Grand Challenge: The Great Robot Race*. Springer-Verlag.
- Buffon, G.** (1777). *Essai d'arithmetique morale*. Suplement to *Histoire naturelle*, vol. IV.
- Bunt, H. C.** (1985). The formal representation of (quasi-) continuous concepts. In Hobbs, J. R. and Moore, R. C. (Eds.), *Formal Theories of the Commonsense World*. Ablex.
- Buolamwini, J. and Gebru, T.** (2018). Gender shades: Intersectional accuracy disparities in commercial gender classification. In *Conference on Fairness, Accountability and Transparency*.
- Burgard, W., Cremers, A. B., Fox, D., Hahnel, D., Lakemeyer, G., Schulz, D., Steiner, W., and Thrun, S.** (1999). Experiences with an interactive museum tour-guide robot. *AIJ*, 114, 3–55.
- Burkov, A.** (2019). *The Hundred-Page Machine Learning Book*. Burkov.
- Burns, E., Hatem, M., Leighton, M. J., and Rumel, W.** (2012). Implementing fast heuristic search code. In *Symposium on Combinatorial Search*.
- Buro, M.** (1995). ProbCut: An effective selective extension of the alpha-beta algorithm. *J. International Computer Chess Association*, 18, 71–76.
- Buro, M.** (2002). Improving heuristic mini-max search by supervised learning. *AIJ*, 134, 85–99.
- Burstein, J., Leacock, C., and Swartz, R.** (2001). Automated evaluation of essays and short answers. In *Fifth International Computer Assisted Assessment Conference*.
- Burton, R.** (2009). *On Being Certain: Believing You Are Right Even When You're Not*. St. Martin's Griffin.
- Buss, D. M.** (2005). *Handbook of Evolutionary Psychology*. Wiley.
- Butler, S.** (1863). Darwin among the machines. *The Press (Christchurch, New Zealand)*, June 13.
- Bylander, T.** (1994). The computational complexity of propositional STRIPS planning. *AIJ*, 69, 165–204.
- Byrd, R. H., Lu, P., Nocedal, J., and Zhu, C.** (1995). A limited memory algorithm for bound constrained optimization. *SIAM Journal on Scientific and Statistical Computing*, 16, 1190–1208.
- Cabeza, R. and Nyberg, L.** (2001). Imaging cognition II: An empirical review of 275 PET and fMRI studies. *J. Cognitive Neuroscience*, 12, 1–47.
- Cafarella, M. J., Halevy, A., Zhang, Y., Wang, D. Z., and Wu, E.** (2008). Webtables: Exploring the power of tables on the web. In *VLDB-08*.
- Calvanese, D., Lenzerini, M., and Nardi, D.** (1999). Unifying class-based representation formalisms. *JAIR*, 11, 199–240.
- Camacho, R. and Michie, D.** (1995). Behavioral cloning: A correction. *AIMag*, 16, 92.
- Campbell, D. E. and Kelly, J.** (2002). Impossibility theorems in the Arrowian framework. In Arrow, K. J., Sen, A. K., and Suzumura, K. (Eds.), *Handbook of Social Choice and Welfare Volume 1*. Elsevier Science.
- Campbell, M. S., Hoane, A. J., and Hsu, F.-H.** (2002). Deep Blue. *AIJ*, 134, 57–83.
- Cannings, C., Thompson, E., and Skolnick, M. H.** (1978). Probability functions on complex pedigrees. *Advances in Applied Probability*, 10, 26–61.
- Canny, J. and Reif, J.** (1987). New lower bound techniques for robot motion planning problems. In *FOCS-87*.
- Canny, J.** (1986). A computational approach to edge detection. *PAMI*, 8, 679–698.
- Canny, J.** (1988). *The Complexity of Robot Motion Planning*. MIT Press.

- Capen**, E., Clapp, R., and Campbell, W. (1971). Competitive bidding in high-risk situations. *J. Petroleum Technology*, 23, 641–653.
- Carbonell**, J. G. (1983). Derivational analogy and its role in problem solving. In *AAAI-83*.
- Carbonell**, J. G., Knoblock, C. A., and Minton, S. (1989). PRODIGY: An integrated architecture for planning and learning. Technical report, Computer Science Department, Carnegie-Mellon University.
- Carbonnel**, C. and Cooper, M. C. (2016). Tractability in constraint satisfaction problems: A survey. *Constraints*, 21(2), 115–144.
- Cardano**, G. (1663). *Liber de ludo aleae*. Lyons.
- Carlini**, N., Athalye, A., Papernot, N., Brendel, W., Rauber, J., Tsipras, D., Goodfellow, I., Madry, A., and Kurakin, A. (2019). On evaluating adversarial robustness. arXiv:1902.06705.
- Carnap**, R. (1928). *Der logische Aufbau der Welt*. Weltkreis-verlag. Translated into English as The Logical Structure of the World (Carnap, 1967).
- Carnap**, R. (1948). On the application of inductive logic. *Philosophy and Phenomenological Research*, 8, 133–148.
- Carnap**, R. (1950). *Logical Foundations of Probability*. University of Chicago Press.
- Carpenter**, B., Gelman, A., Hoffman, M., Lee, D., Goodrich, B., Betancourt, M., Brubaker, M., Guo, J., Li, P., and Riddell, A. (2017). Stan: A probabilistic programming language. *Journal of Statistical Software*, 76, 1–32.
- Carroll**, S. (2007). *The Making of the Fittest: DNA and the Ultimate Forensic Record of Evolution*. Norton.
- Casati**, R. and Varzi, A. (1999). *Parts and Places: The Structures of Spatial Representation*. MIT Press.
- Cassandra**, A. R., Kaelbling, L. P., and Littman, M. L. (1994). Acting optimally in partially observable stochastic domains. In *AAAI-94*.
- Cassandras**, C. G. and Lygeros, J. (2006). *Stochastic Hybrid Systems*. CRC Press.
- Castro**, R., Coates, M., Liang, G., Nowak, R., and Yu, B. (2004). Network tomography: Recent developments. *Statistical Science*, 19, 499–517.
- Cauchy**, A. (1847). Méthode générale pour la résolution des systèmes d'équations simultanées. *Comp. Rend. Sci. Paris*, 25, 536–538.
- Cesa-Bianchi**, N. and Lugosi, G. (2006). *Prediction, Learning, and Games*. Cambridge University Press.
- Chajewska**, U., Koller, D., and Parr, R. (2000). Making rational decisions using adaptive utility elicitation. In *AAAI-00*.
- Chakrabarti**, P. P., Ghose, S., Acharya, A., and de Sarkar, S. C. (1989). Heuristic search in restricted memory. *AII*, 41, 197–222.
- Chalkiadakis**, G., Elkind, E., and Wooldridge, M. (2011). *Computational Aspects of Cooperative Game Theory*. Morgan Kaufmann.
- Chalmers**, D. J. (1992). Subsymbolic computation and the Chinese room. In Dinsmore, J. (Ed.), *The symbolic and connectionist paradigms: Closing the gap*. Lawrence Erlbaum.
- Chandola**, V., Banerjee, A., and Kumar, V. (2009). Anomaly detection: A survey. *ACM Computing Surveys*, 41.
- Chandra**, A. K. and Harel, D. (1980). Computable queries for relational data bases. *J. Computer and System Sciences*, 21, 156–178.
- Chang**, C.-L. and Lee, R. C.-T. (1973). *Symbolic Logic and Mechanical Theorem Proving*. Academic Press.
- Chang**, H. S., Fu, M. C., Hu, J., and Marcus, S. I. (2005). An adaptive sampling algorithm for solving Markov decision processes. *Operations Research*, 53, 126–139.
- Chao**, W.-L., Hu, H., and Sha, F. (2018). Being negative but constructively: Lessons learnt from creating better visual question answering datasets. In *ACL-18*.
- Chapman**, D. (1987). Planning for conjunctive goals. *AII*, 32, 333–377.
- Charniak**, E. (1993). *Statistical Language Learning*. MIT Press.
- Charniak**, E. (1996). Tree-bank grammars. In *AAAI-96*.
- Charniak**, E. (1997). Statistical parsing with a context-free grammar and word statistics. In *AAAI-97*.
- Charniak**, E. and Goldman, R. (1992). A Bayesian model of plan recognition. *AII*, 64, 53–79.
- Charniak**, E., Riesbeck, C., McDermott, D., and Meehan, J. (1987). *Artificial Intelligence Programming* (2nd edition). Lawrence Erlbaum.
- Charniak**, E. (1991). Bayesian networks without tears. *AI Mag*, 12, 50–63.
- Charniak**, E. (2018). *Introduction to Deep Learning*. MIT Press.
- Chaslot**, G., Bakkes, S., Szita, I., and Spronck, P. (2008). Monte-Carlo tree search: A new framework for game AI. In *Proc. Fourth Artificial Intelligence and Interactive Digital Entertainment Conference*.
- Chater**, N. and Oaksford, M. (Eds.). (2008). *The Probabilistic Mind: Prospects for Bayesian Cognitive Science*. Oxford University Press.
- Chatfield**, C. (1989). *The Analysis of Time Series: An Introduction* (4th edition). Chapman and Hall.
- Chavira**, M. and Darwiche, A. (2008). On probabilistic inference by weighted model counting. *AII*, 172, 772–799.
- Chawla**, N. V., Bowyer, K. W., Hall, L. O., and Kegelmeyer, W. P. (2002). SMOTE: Synthetic minority over-sampling technique. *AII*, 16, 321–357.
- Cheeseman**, P. (1985). In defense of probability. In *IJCAI-85*.
- Cheeseman**, P. (1988). An inquiry into computer understanding. *Computational Intelligence*, 4, 58–66.
- Cheeseman**, P., Kanefsky, B., and Taylor, W. (1991). Where the really hard problems are. In *IJCAI-91*.
- Cheeseman**, P., Self, M., Kelly, J., and Stutz, J. (1988). Bayesian classification. In *AAAI-88*.
- Cheeseman**, P. and Stutz, J. (1996). Bayesian classification (AutoClass): Theory and results. In Fayyad, U., Piatesky-Shapiro, G., Smyth, P., and Uthurusamy, R. (Eds.), *Advances in Knowledge Discovery and Data Mining*. AAAI Press/MIT Press.
- Chen**, D. and Manning, C. (2014). A fast and accurate dependency parser using neural networks. In *EMNLP-14*.
- Chen**, J., Holte, R. C., Zilles, S., and Sturtevant, N. R. (2017). Front-to-end bidirectional heuristic search with near-optimal node expansions. *IJCAI-17*.
- Chen**, M. X., Firat, O., Bapna, A., Johnson, M., Macherey, W., Foster, G., Jones, L., Parmar, N., Schuster, M., Chen, Z., Wu, Y., and Hughes, M. (2018). The best of both worlds: Combining recent advances in neural machine translation. In *ACL-18*.
- Chen**, S. F. and Goodman, J. (1996). An empirical study of smoothing techniques for language modeling. In *ACL-96*.
- Chen**, T. and Guestrin, C. (2016). XGBoost: A scalable tree boosting system. In *KDD-16*.
- Cheng**, J. and Druzdzel, M. J. (2000). AIS-BN: An adaptive importance sampling algorithm for evidential reasoning in large Bayesian networks. *AII*, 13, 155–188.
- Cheng**, J., Greiner, R., Kelly, J., Bell, D. A., and Liu, W. (2002). Learning Bayesian networks from data: An information-theory based approach. *AII*, 137, 43–90.
- Chiu**, C., Sainath, T., Wu, Y., Prabhavalkar, R., Nguyen, P., Chen, Z., Kannan, A., Weiss, R., Rao, K., Gonina, K., Jaityl, N., Li, B., Chorowski, J., and Baccianini, M. (2017). State-of-the-art speech recognition with sequence-to-sequence models. arXiv:1712.01769.
- Chklovski**, T. and Gil, Y. (2005). Improving the design of intelligent acquisition interfaces for collecting world knowledge from web contributors. In *Proc. Third International Conference on Knowledge Capture*.
- Chollet**, F. (2019). On the measure of intelligence. arXiv:1911.01547.
- Chollet**, F. (2017). *Deep Learning with Python*. Manning.
- Chomsky**, N. (1956). Three models for the description of language. *IRE Transactions on Information Theory*, 2, 113–124.
- Chomsky**, N. (1957). *Syntactic Structures*. Mouton.
- Choromanska**, A., Henaff, M., Mathieu, M., Arous, G. B., and LeCun, Y. (2014). The loss surface of multilayer networks. arXiv:1412.0233.
- Choset**, H. (1996). *Sensor Based Motion Planning: The Hierarchical Generalized Voronoi Graph*. Ph.D. thesis, California Institute of Technology.
- Choset**, H., Hutchinson, S., Lynch, K., Kantor, G., Burgard, W., Kavraki, L., and Thrun, S. (2005). *Principles of Robot Motion: Theory, Algorithms, and Implementation*. MIT Press.
- Chouldechova**, A. (2017). Fair prediction with disparate impact: A study of bias in recidivism prediction instruments. *Big Data*, 5, 153–163.
- Chouldechova**, A. and Roth, A. (2018). The frontiers of fairness in machine learning. arXiv:1810.08810.
- Christian**, B. (2011). *The Most Human Human*. Doubleday.
- Christin**, A., Rosenblat, A., and Boyd, D. (2015). Courts and predictive algorithms. *Data & Civil Rights*.
- Chung**, K. L. (1979). *Elementary Probability Theory with Stochastic Processes* (3rd edition). Springer-Verlag.
- Church**, A. (1936). A note on the Entscheidungsproblem. *JSL*, 1, 40–41 and 101–102.
- Church**, A. (1956). *Introduction to Mathematical Logic*. Princeton University Press.
- Church**, K. (1988). A stochastic parts program and noun phrase parser for unrestricted texts. In *Proc. Second Conference on Applied Natural Language Processing*.
- Church**, K. and Patil, R. (1982). Coping with syntactic ambiguity or how to put the block in the box on the table. *Computational Linguistics*, 8, 139–149.
- Church**, K. (2004). Speech and language processing: Can we use the past to predict the future. In *Proc. Conference on Text, Speech, and Dialogue*.
- Church**, K. and Gale, W. A. (1991). A comparison of the enhanced Good-Turing and deleted estimation methods for estimating probabilities of English bigrams. *Computer Speech and Language*, 5, 19–54.
- Church**, K. and Hestness, J. (2019). A survey of 25 years of evaluation. *Natural Language Engineering*, 25, 753–767.
- Churchland**, P. M. (2013). *Matter and Consciousness* (3rd edition). MIT Press.

- Ciancarini, P.** and Favini, G. P. (2010). Monte Carlo tree search in Kriegspiel. *AIJ*, 174, 670–684.
- Ciancarini, P.** and Wooldridge, M. (2001). *Agent-Oriented Software Engineering*. Springer-Verlag.
- Cimatti, A.**, Roveri, M., and Traverso, P. (1998). Automatic OBDD-based generation of universal plans in non-deterministic domains. In *AAAI-98*.
- Claret, G.**, Rajamani, S. K., Nori, A. V., Gordon, A. D., and Borgström, J. (2013). Bayesian inference using data flow analysis. In *Proc. 9th Joint Meeting on Foundations of Software Engineering*.
- Clark, A.** (1998). *Being There: Putting Brain, Body, and World Together Again*. MIT Press.
- Clark, A.** (2015). *Surfing Uncertainty: Prediction, Action, and the Embodied Mind*. Oxford University Press.
- Clark, K. L.** (1978). Negation as failure. In Gallois, H. and Minker, J. (Eds.), *Logic and Data Bases*. Plenum.
- Clark, P.**, Cowhey, I., Etzioni, O., Khot, T., Sabharwal, A., Schoenick, C., and Tafjord, O. (2018). Think you have solved question answering? Try ARC, the AI2 reasoning challenge. arXiv:1803.05457.
- Clark, P.**, Etzioni, O., Khot, T., Mishra, B. D., Richardson, K., et al. (2019). From 'F' to 'A' on the NY Regents science exams: An overview of the Aristo project. arXiv:1909.01958.
- Clark, S.** and Curran, J. R. (2004). Parsing the WSJ using CCG and log-linear models. In *ACL-04*.
- Clarke, A. C.** (1968). *2001: A Space Odyssey*. Signet.
- Clarke, E.** and Grumberg, O. (1987). Research on automatic verification of finite-state concurrent systems. *Annual Review of Computer Science*, 2, 269–290.
- Clearwater, S. H.** (Ed.). (1996). *Market-Based Control*. World Scientific.
- Clocksin, W. F.** and Mellish, C. S. (2003). *Programming in Prolog* (5th edition). Springer-Verlag.
- Clocksin, W. F.** (2003). *Clause and Effect: Prolog Programming for the Working Programmer*. Springer.
- Coase, R. H.** (1960). The problem of social cost. *Journal of Law and Economics*, pp. 1–44.
- Coates, A.**, Abbeel, P., and Ng, A. Y. (2009). Apprenticeship learning for helicopter control. *Association for Computing Machinery*, 52(7).
- Cobham, A.** (1964). The intrinsic computational difficulty of functions. In *Proc. International Congress for Logic, Methodology, and Philosophy of Science*.
- Cohen, P. R.** (1995). *Empirical Methods for Artificial Intelligence*. MIT Press.
- Cohen, P. R.** and Levesque, H. J. (1990). Intention is choice with commitment. *AIJ*, 42, 213–261.
- Cohen, P. R.**, Morgan, J., and Pollack, M. E. (1990). *Intentions in Communication*. MIT Press.
- Cohen, P. R.** and Perrault, C. R. (1979). Elements of a plan-based theory of speech acts. *Cognitive Science*, 3, 177–212.
- Cohn, A. G.**, Bennett, B., Gooday, J. M., and Gotts, N. (1997). RCC: A calculus for region based qualitative spatial reasoning. *GeoInformatica*, 1, 275–316.
- Collin, Z.**, Dechter, R., and Katz, S. (1999). Self-stabilizing distributed constraint satisfaction. *Chicago J. of Theoretical Computer Science*, 1999.
- Collins, M.** (1999). *Head-driven Statistical Models for Natural Language Processing*. Ph.D. thesis, University of Pennsylvania.
- Collins, M.** and Duffy, K. (2002). New ranking algorithms for parsing and tagging: Kernels over discrete structures, and the voted perceptron. In *ACL-02*.
- Colmerauer, A.** and Roussel, P. (1993). The birth of Prolog. *SIGPLAN Notices*, 28, 37–52.
- Colmerauer, A.**, Kanoui, H., Pasero, R., and Roussel, P. (1973). Un système de communication homme-machine en Français. Rapport, Groupe d'Intelligence Artificielle, Université d'Aix-Marseille II.
- Condon, J. H.** and Thompson, K. (1982). Belle chess hardware. In Clarke, M. R. B. (Ed.), *Advances in Computer Chess 3*. Pergamon.
- Congdon, C. B.**, Huber, M., Kortenkamp, D., Bidlack, C., Cohen, C., Huffman, S., Koss, F., Raschke, U., and Weymouth, T. (1992). CARMEL versus Flakey: A comparison of two robots. Tech. rep., American Association for Artificial Intelligence.
- Conlisk, J.** (1989). Three variants on the Allais example. *American Economic Review*, 79, 392–407.
- Connell, J.** (1989). *A Colony Architecture for an Artificial Creature*. Ph.D. thesis, Artificial Intelligence Laboratory, MIT.
- Conway, D.** and White, J. (2012). *Machine Learning for Hackers*. O'Reilly.
- Cook, S. A.** (1971). The complexity of theorem-proving procedures. In *STOC-71*.
- Cook, S. A.** and Mitchell, D. (1997). Finding hard instances of the satisfiability problem: A survey. In Du, D., Gu, J., and Pardalos, P. (Eds.), *Satisfiability problems: Theory and applications*. American Mathematical Society.
- Cooper, G.** (1990). The computational complexity of probabilistic inference using Bayesian belief networks. *AIJ*, 42, 393–405.
- Cooper, G.** and Herskovits, E. (1992). A Bayesian method for the induction of probabilistic networks from data. *Machine Learning*, 9, 309–347.
- Copeland, J.** (1993). *Artificial Intelligence: A Philosophical Introduction*. Blackwell.
- Corbett-Davies, S.** and Goel, S. (2018). The measure and mismeasure of fairness: A critical review of fair machine learning. arXiv:1808.00023.
- Corbett-Davies, S.**, Pierson, E., Feller, A., Goel, S., and Huq, A. (2017). Algorithmic decision making and the cost of fairness. arXiv:1701.08230.
- Cormen, T. H.**, Leiserson, C. E., Rivest, R., and Stein, C. (2009). *Introduction to Algorithms* (3rd edition). MIT Press.
- Cortes, C.** and Vapnik, V. N. (1995). Support vector networks. *Machine Learning*, 20, 273–297.
- Cournot, A.** (Ed.). (1838). *Recherches sur les principes mathématiques de la théorie des richesses*. L. Hachette, Paris.
- Cover, T.** and Thomas, J. (2006). *Elements of Information Theory* (2nd edition). Wiley.
- Cowan, J. D.** and Sharp, D. H. (1988a). Neural nets. *Quarterly Reviews of Biophysics*, 21, 365–427.
- Cowan, J. D.** and Sharp, D. H. (1988b). Neural nets and artificial intelligence. *Daedalus*, 117, 85–121.
- Cowell, R.**, Dawid, A. P., Lauritzen, S., and Spiegelhalter, D. J. (2002). *Probabilistic Networks and Expert Systems*. Springer.
- Cox, I.** (1993). A review of statistical data association techniques for motion correspondence. *IJCV*, 10, 53–66.
- Cox, I.** and Hingorani, S. L. (1994). An efficient implementation and evaluation of Reid's multiple hypothesis tracking algorithm for visual tracking. In *ICPR-94*.
- Cox, I.** and Wilfong, G. T. (Eds.). (1990). *Autonomous Robot Vehicles*. Springer Verlag.
- Cox, R. T.** (1946). Probability, frequency, and reasonable expectation. *American Journal of Physics*, 14, 1–13.
- Craig, J.** (1989). *Introduction to Robotics: Mechanics and Control* (2nd edition). Addison-Wesley.
- Craik, K.** (1943). *The Nature of Explanation*. Cambridge University Press.
- Cramton, P.**, Shoham, Y., and Steinberg, R. (Eds.). (2006). *Combinatorial Auctions*. MIT Press.
- Craven, M.**, DiPasquo, D., Freitag, D., McCallum, A., Mitchell, T. M., Nigam, K., and Slattery, S. (2000). Learning to construct knowledge bases from the World Wide Web. *AIJ*, 118, 69–113.
- Crawford, J. M.** and Auton, L. D. (1993). Experimental results on the crossover point in satisfiability problems. In *AAAI-93*.
- Crick, F.** (1999). The impact of molecular biology on neuroscience. *Phil. Trans. Roy. Soc., B*, 354, 2021–2025.
- Crick, F.** and Koch, C. (2003). A framework for consciousness. *Nature Neuroscience*, 6, 119.
- Crisan, D.** and Doucet, A. (2002). A survey of convergence results on particle filtering methods for practitioners. *IEEE Trans. Signal Processing*, 50, 736–746.
- Cristianini, N.** and Hahn, M. (2007). *Introduction to Computational Genomics: A Case Studies Approach*. Cambridge University Press.
- Cristianini, N.** and Schölkopf, B. (2002). Support vector machines and kernel methods: The new generation of learning machines. *AI Mag*, 23, 31–41.
- Cristianini, N.** and Shawe-Taylor, J. (2000). *An Introduction to Support Vector Machines and Other Kernel-Based Learning Methods*. Cambridge University Press.
- Crockett, L.** (1994). *The Turing Test and the Frame Problem: AI's Mistaken Understanding of Intelligence*. Ablex.
- Croft, W. B.**, Metzler, D., and Strohman, T. (2010). *Search Engines: Information Retrieval in Practice*. Addison-Wesley.
- Cross, S. E.** and Walker, E. (1994). DART: Applying knowledge based planning and scheduling to crisis action planning. In Zweber, M. and Fox, M. S. (Eds.), *Intelligent Scheduling*. Morgan Kaufmann.
- Cruse, A.** (2011). *Meaning in Language: An Introduction to Semantics and Pragmatics*. Oxford University Press.
- Culberson, J.** and Schaeffer, J. (1996). Searching with pattern databases. In *Advances in Artificial Intelligence (Lecture Notes in Artificial Intelligence 1081)*. Springer-Verlag.
- Culberson, J.** and Schaeffer, J. (1998). Pattern databases. *Computational Intelligence*, 14, 318–334.
- Cummins, D.** and Allen, C. (1998). *The Evolution of Mind*. Oxford University Press.
- Cushing, W.**, Kambhampati, S., Mausam, and Weld, D. S. (2007). When is temporal planning really temporal? In *IJCAI-07*.
- Cusumano-Towner, M. F.**, Saad, F., Lew, A. K., and Mansingka, V. K. (2019). Gen: A general-purpose probabilistic programming system with programmable inference. In *PLDI-19*.
- Cybenko, G.** (1988). Continuous valued neural networks with two hidden layers are sufficient. Technical report, Department of Computer Science, Tufts University.
- Cybenko, G.** (1989). Approximation by superpositions of a sigmoidal function. *Mathematics of Controls, Signals, and Systems*, 2, 303–314.
- Cyert, R.** and de Groot, M. (1979). Adaptive utility. In Allais, M. and Hagen, O. (Eds.), *Expected Utility Hypothesis and the Allais Paradox*. D. Reidel.
- Dagan, I.**, Glickman, O., and Magnini, B. (2005). The PASCAL recognising textual entailment challenge. In *Machine Learning Challenges Workshop*.
- Daganzo, C.** (1979). *Multinomial Probit: The Theory and Its Application to Demand Forecasting*. Academic Press.

- Dagum**, P. and Luby, M. (1993). Approximating probabilistic inference in Bayesian belief networks is NP-hard. *AIJ*, 60, 141–153.
- Dagum**, P. and Luby, M. (1997). An optimal approximation algorithm for Bayesian inference. *AIJ*, 93, 1–27.
- Dai**, A. M. and Le, Q. V. (2016). Semi-supervised sequence learning. In *NeurIPS* 28.
- Dalal**, N. and Triggs, B. (2005). Histograms of oriented gradients for human detection. In *CVPR-05*.
- Dalvi**, N. N., Ré, C., and Suciu, D. (2009). Probabilistic databases. *CACM*, 52, 86–94.
- Daly**, R., Shen, Q., and Aitken, S. (2011). Learning Bayesian networks: Approaches and issues. *Knowledge Engineering Review*, 26, 99–157.
- Damasio**, A. R. (1999). *The Feeling of What Happens: Body and Emotion in the Making of Consciousness*. Houghton Mifflin.
- Danaher**, J. and McArthur, N. (2017). *Robot Sex: Social and Ethical Implications*. MIT Press.
- Dantzig**, G. B. (1949). Programming of interdependent activities: II. Mathematical model. *Econometrica*, 17, 200–211.
- Darwiche**, A. (2001). Recursive conditioning. *AIJ*, 126, 5–41.
- Darwiche**, A. and Ginsberg, M. L. (1992). A symbolic generalization of probability theory. In *AAAI-92*.
- Darwiche**, A. (2009). *Modeling and reasoning with Bayesian networks*. Cambridge University Press.
- Darwin**, C. (1859). *On The Origin of Species by Means of Natural Selection*. J. Murray.
- Dasgupta**, P., Chakrabarti, P. P., and de Sarkar, S. C. (1994). Agent searching in a tree and the optimality of iterative deepening. *AIJ*, 71, 195–208.
- Dasgupta**, P. and Maskin, E. (2008). On the robustness of majority rule. *Journal of the European Economic Association*, 6, 949–973.
- Dauphin**, Y., Pascanu, R., Gulcehre, C., Cho, K., Ganguli, S., and Bengio, Y. (2015). Identifying and attacking the saddle point problem in high-dimensional non-convex optimization. In *NeurIPS* 27.
- Davidson**, D. (1980). *Essays on Actions and Events*. Oxford University Press.
- Davidson**, D. (1986). A nice derangement of epitaphs. *Philosophical Grounds of Rationality*, 4, 157–174.
- Davis**, E. (1986). *Representing and Acquiring Geographic Knowledge*. Pitman and Morgan Kaufmann.
- Davis**, E. (1990). *Representations of Commonsense Knowledge*. Morgan Kaufmann.
- Davis**, E. (2005). Knowledge and communication: A first-order theory. *AIJ*, 166, 81–140.
- Davis**, E. (2006). The expressivity of quantifying over regions. *J. Logic and Computation*, 16, 891–916.
- Davis**, E. (2007). Physical reasoning. In van Harmelen, F., Lifschitz, V., and Porter, B. (Eds.), *The Handbook of Knowledge Representation*. Elsevier.
- Davis**, E. (2008). Pouring liquids: A study in commonsense physical reasoning. *AIJ*, 172.
- Davis**, E. (2017). Logical formalizations of commonsense reasoning: A survey. *JAIR*, 59, 651–723.
- Davis**, E. and Morgenstern, L. (2004). Introduction: Progress in formal commonsense reasoning. *AIJ*, 153, 1–12.
- Davis**, E. and Morgenstern, L. (2005). A first-order theory of communication and multi-agent plans. *J. Logic and Computation*, 15, 701–749.
- Davis**, M. (1957). A computer program for Presburger's algorithm. In *Proving Theorems (as Done by Man, Logician, or Machine)*. Proc. Summer Institute for Symbolic Logic. Second edition; publication date is 1960.
- Davis**, M., Logemann, G., and Loveland, D. (1962). A machine program for theorem-proving. *CACM*, 5, 394–397.
- Davis**, M. and Putnam, H. (1960). A computing procedure for quantification theory. *JACM*, 7, 201–215.
- Dayan**, P. (1992). The convergence of TD( $\lambda$ ) for general  $\lambda$ . *Machine Learning*, 8, 341–362.
- Dayan**, P. and Abbott, L. F. (2001). *Theoretical Neuroscience: Computational and Mathematical Modeling of Neural Systems*. MIT Press.
- Dayan**, P. and Hinton, G. E. (1993). Feudal reinforcement learning. In *NeurIPS* 5.
- Dayan**, P. and Niv, Y. (2008). Reinforcement learning and the brain: The good, the bad and the ugly. *Current Opinion in Neurobiology*, 18, 185–196.
- de Condorcet**, M. (1785). *Essay on the Application of Analysis to the Probability of Majority Decisions*. Imprimerie Royale.
- de Dombal**, F. T., Leaper, D. J., Horrocks, J. C., and Staniland, J. R. (1974). Human and computer-aided diagnosis of abdominal pain: Further report with emphasis on performance of clinicians. *British Medical Journal*, 1, 376–380.
- de Dombal**, F. T., Staniland, J. R., and Clamp, S. E. (1981). Geographical variation in disease presentation. *Medical Decision Making*, 1, 59–69.
- de Farias**, D. P. and Roy, B. V. (2003). The linear programming approach to approximate dynamic programming. *Operations Research*, 51, 839–1016.
- de Finetti**, B. (1937). Le prévision: ses lois logiques, ses sources subjectives. *Ann. Inst. Poincaré*, 7, 1–68.
- de Finetti**, B. (1993). On the subjective meaning of probability. In Monari, P. and Cocchi, D. (Eds.), *Probabilità e Induzione*. Clueb.
- de Freitas**, J. F. G., Nirajan, M., and Gee, A. H. (2000). Sequential Monte Carlo methods to train neural network models. *Neural Computation*, 12, 933–953.
- de Ghellinck**, G. (1960). Les problèmes de décisions séquentielles. *Cahiers du Centre d'Études de Recherche Opérationnelle*, 2, 161–179.
- de Kleer**, J. (1975). Qualitative and quantitative knowledge in classical mechanics. Tech. rep., MIT Artificial Intelligence Laboratory.
- de Kleer**, J. (1989). A comparison of ATMS and CSP techniques. In *IJCAI-89*.
- de Kleer**, J. and Brown, J. S. (1985). A qualitative physics based on confluences. In Hobbs, J. R. and Moore, R. C. (Eds.), *Formal Theories of the Commonsense World*. Ablex.
- de Marcken**, C. (1996). *Unsupervised Language Acquisition*. Ph.D. thesis, MIT.
- De Morgan**, A. (1864). On the syllogism, No. IV, and on the logic of relations. *Transaction of the Cambridge Philosophical Society*, X, 331–358.
- de Salvo Braz**, R., Amir, E., and Roth, D. (2007). Lifted first-order probabilistic inference. In Getoor, L. and Taskar, B. (Eds.), *Introduction to Statistical Relational Learning*. MIT Press.
- Deacon**, T. W. (1997). *The Symbolic Species: The Co-evolution of Language and the Brain*. W. W. Norton.
- Deale**, M., Yvanovich, M., Schmitzus, D., Kautz, D., Carpenter, M., Zweben, M., Davis, G., and Daun, B. (1994). The space shuttle ground processing scheduling system. In Zweben, M. and Fox, M. (Eds.), *Intelligent Scheduling*. Morgan Kaufmann.
- Dechter**, R., Patterson, D. A., and Young, C. (2018). A new golden age in computer architecture: Empowering the machine-learning revolution. *IEEE Micro*, 38, 21–29.
- Dechter**, R., Basye, K., Chekaluk, R., and Hyun, S. (1990). Coping with uncertainty in a control system for navigation and exploration. In *AAAI-90*.
- Dechter**, R. and Boddy, M. (1988). An analysis of time-dependent planning. In *AAAI-88*.
- Dechter**, R., Firby, R. J., and Miller, D. (1990). Hierarchical planning involving deadlines, travel time, and resources. *Computational Intelligence*, 6, 381–398.
- Dechter**, R., Kaelbling, L. P., Kirman, J., and Nicholson, A. (1993). Planning with deadlines in stochastic domains. In *AAAI-93*.
- Dean**, T. and Kanazawa, K. (1989a). A model for projection and action. In *IJCAI-89*.
- Dean**, T. and Kanazawa, K. (1989b). A model for reasoning about persistence and causation. *Computational Intelligence*, 5, 142–150.
- Dean**, T. and Wellman, M. P. (1991). *Planning and Control*. Morgan Kaufmann.
- Dearden**, R., Friedman, N., and Andre, D. (1999). Model-based Bayesian exploration. In *UAI-99*.
- Dearden**, R., Friedman, N., and Russell, S. J. (1998). Bayesian Q-learning. In *AAAI-98*.
- Debevec**, P., Taylor, C., and Malik, J. (1996). Modeling and rendering architecture from photographs: A hybrid geometry- and image-based approach. In *Proc. 23rd Annual Conference on Computer Graphics (SIGGRAPH)*.
- Debreu**, G. (1960). Topological methods in cardinal utility theory. In Arrow, K. J., Karlin, S., and Suppes, P. (Eds.), *Mathematical Methods in the Social Sciences*, 1959. Stanford University Press.
- Dechter**, A. and Dechter, R. (1987). Removing redundancies in constraint networks. In *AAAI-87*.
- Dechter**, R. (1990a). Enhancement schemes for constraint processing: Backjumping, learning and cutset decomposition. *AIJ*, 41, 273–312.
- Dechter**, R. (1990b). On the expressiveness of networks with hidden variables. In *AAAI-90*.
- Dechter**, R. (1999). Bucket elimination: A unifying framework for reasoning. *AIJ*, 113, 41–85.
- Dechter**, R. and Pearl, J. (1985). Generalized best-first search strategies and the optimality of A\*. *JACM*, 32, 505–536.
- Dechter**, R. and Pearl, J. (1987). Network-based heuristics for constraint-satisfaction problems. *AIJ*, 34, 1–38.
- Dechter**, R. and Pearl, J. (1989). Tree clustering for constraint networks. *AIJ*, 38, 353–366.
- Dechter**, R. and Rish, I. (2003). Mini-buckets: A general scheme for bounded inference. *JACM*, 50, 107–153.
- Dechter**, R. (2003). *Constraint Processing*. Morgan Kaufmann.
- Dechter**, R. (2019). *Reasoning with Probabilistic and Deterministic Graphical Models: Exact Algorithms* (2nd edition). Morgan & Claypool.
- Dechter**, R. and Frost, D. (2002). Backjump-based backtracking for constraint satisfaction problems. *AIJ*, 136, 147–188.
- Dechter**, R. and Mateescu, R. (2007). AND/OR search spaces for graphical models. *AIJ*, 171, 73–106.
- DeCoste**, D. and Schölkopf, B. (2002). Training invariant support vector machines. *Machine Learning*, 46, 161–190.
- Dedekind**, R. (1888). *Was sind und was sollen die Zahlen*. Braunschweig, Germany.

- Deerwester**, S. C., Dumais, S. T., Landauer, T. K., Furnas, G. W., and Harshman, R. A. (1990). Indexing by latent semantic analysis. *J. American Society for Information Science*, 41, 391–407.
- DeGroot**, M. H. (1970). *Optimal Statistical Decisions*. McGraw-Hill.
- DeGroot**, M. H. and Schervish, M. J. (2001). *Probability and Statistics* (3rd edition). Addison Wesley.
- Dehaene**, S. (2014). *Consciousness and the Brain: Deciphering How the Brain Codes Our Thoughts*. Penguin Books.
- Del Moral**, P., Doucet, A., and Jasra, A. (2006). Sequential Monte Carlo samplers. *J. Royal Statistical Society*, 68, 411–436.
- Del Moral**, P. (2004). *Feynman–Kac Formulae, Geometrical and Interacting Particle Systems with Applications*. Springer-Verlag.
- Delgrande**, J. and Schaub, T. (2003). On the relation between Reiter's default logic and its (major) variants. In *Seventh European Conference on Symbolic and Quantitative Approaches to Reasoning with Uncertainty*.
- Delling**, D., Sanders, P., Schultes, D., and Wagner, D. (2009). Engineering route planning algorithms. In Lerner, J., Wagner, D., and Zweig, K. (Eds.), *Algorithmics, LNCS*. Springer-Verlag.
- Dempster**, A. P. (1968). A generalization of Bayesian inference. *J. Royal Statistical Society*, 30 (Series B), 205–247.
- Dempster**, A. P., Laird, N., and Rubin, D. (1977). Maximum likelihood from incomplete data via the EM algorithm. *J. Royal Statistical Society*, 39 (Series B), 1–38.
- Denardo**, E. V. (1967). Contraction mappings in the theory underlying dynamic programming. *SIAM Review*, 9, 165–177.
- Deng**, J., Dong, W., Socher, R., Li, L.-J., Li, K., and Fei-Fei, L. (2009). Imagenet: A large-scale hierarchical image database. In *CVPR-09*.
- Deng**, L. (2016). Deep learning: From speech recognition to language and multimodal processing. *APSIPA Transactions on Signal and Information Processing*, 5.
- Deng**, L., Yu, D., et al. (2014). Deep learning: Methods and applications. *Foundations and Trends in Signal Processing*, 7, 197–387.
- Deng**, X. and Papadimitriou, C. H. (1990). Exploring an unknown graph. In *FOCS-90*.
- Deng**, X. and Papadimitriou, C. H. (1994). On the complexity of cooperative solution concepts. *Mathematics of Operations Research*, 19, 257–266.
- Denney**, E., Fischer, B., and Schumann, J. (2006). An empirical evaluation of automated theorem provers in software certification. *Int. J. AI Tools*, 15, 81–107.
- D'Épenoux**, F. (1963). A probabilistic production and inventory problem. *A probabilistic production and inventory problem*, 10, 98–108.
- Dervovic**, D., Herbster, M., Mountney, P., Severini, S., Usher, N., and Wossnig, L. (2018). Quantum linear systems algorithms: A primer. arXiv:1802.08227.
- Descartes**, R. (1637). Discourse on method. In Cottingham, J., Stoothoff, R., and Murdoch, D. (Eds.), *The Philosophical Writings of Descartes*, Vol. I. Cambridge University Press, Cambridge.
- Descofte**, Y. and Latombe, J.-C. (1985). Making compromises among antagonist constraints in a planner. *AIJ*, 27, 183–217.
- Deshpande**, I., Hu, Y.-T., Sun, R., Pyrros, A., Sidiqui, N., Koyejo, S., Zhao, Z., Forsyth, D., and Schwing, A. (2019). Max-sliced Wasserstein distance and its use for GANs. In *CVPR-19*.
- Deutscher**, G. (2010). *Through the Language Glass: Why the World Looks Different in Other Languages*. Metropolitan Books.
- Devlin**, J., Chang, M.-W., Lee, K., and Toutanova, K. (2018). Bert: Pre-training of deep bidirectional transformers for language understanding. arXiv:1810.04805.
- Devlin**, J. (2018). *Turned On: Science, Sex and Robots*. Bloomsbury.
- Devroye**, L. (1987). *A course in density estimation*. Birkhäuser.
- Dias**, M. B., Zlot, R., Kalra, N., and Stentz, A. (2006). Market-based multirobot coordination: A survey and analysis. *Proc. IEEE*, 94, 1257–1270.
- Dickmanns**, E. D. and Zapp, A. (1987). Autonomous high speed road vehicle guidance by computer vision. In *Automatic Control—World Congress, 1987: Selected Papers from the 10th Triennial World Congress of the International Federation of Automatic Control*.
- Dieterich**, T. (2000). Hierarchical reinforcement learning with the MAXQ value function decomposition. *JAIR*, 13, 227–303.
- Dijkstra**, E. W. (1959). A note on two problems in connexion with graphs. *Numerische Mathematik*, 1, 269–271.
- Dijkstra**, E. W. (1984). The threats to computing science. In *ACM South Central Regional Conference*.
- Ding**, Y., Sohn, J. H., Kawczynski, M. G., Trivedi, H., Harnish, R., Jenkins, N. W., Litviev, D., Copeland, T. P., Aboian, M. S., Mari Aparici, C., et al. (2018). A deep learning model to predict a diagnosis of alzheimer disease by using 18F-FDG PET of the brain. *Radiology*, p. 180958.
- Dinh**, H., Russell, A., and Su, Y. (2007). On the value of good advice: The complexity of A\* with accurate heuristics. In *AAAI-07*.
- Dissanayake**, G., Newman, P., Clark, S., Durrant-Whyte, H., and Csorba, M. (2001). A solution to the simultaneous localisation and map building (SLAM) problem. *IEEE Transactions on Robotics and Automation*, 17, 229–241.
- Dittmer**, S. and Jensen, F. (1997). Myopic value of information in influence diagrams. In *UAI-97*.
- Do**, M. and Kambhampati, S. (2003). Planning as constraint satisfaction: solving the planning graph by compiling it into CSP. *AIJ*, 132, 151–182.
- Do**, M. B. and Kambhampati, S. (2001). Sapa: A domain-independent heuristic metric temporal planner. In *ECP-01*.
- Doctorow**, C. (2001). Metacrap: Putting the torch to seven straw-men of the meta-utopia. www.well.com/doctorow/metacrap.htm.
- Doctorow**, C. and Stross, C. (2012). *The Rapture of the Nerds: A Tale of the Singularity, Posthumanity, and Awkward Social Situations*. Tor Books.
- Dodd**, L. (1988). The inside/outside algorithm: Grammatical inference applied to stochastic context-free grammars. Tech. rep., Royal Signals and Radar Establishment, Malvern.
- Domingos**, P. and Pazzani, M. (1997). On the optimality of the simple Bayesian classifier under zero-one loss. *Machine Learning*, 29, 103–30.
- Domingos**, P. (2012). A few useful things to know about machine learning. *Commun. ACM*, 55(10), 78–87.
- Domingos**, P. (2015). *The Master Algorithm: How the Quest for the Ultimate Learning Machine Will Remake Our World*. Basic Books.
- Dong**, X., Gabrilovich, E., Heitz, G., Horn, W., Lao, N., Murphy, K., Strohmann, T., Sun, S., and Zhang, W. (2014). Knowledge vault: A web-scale approach to probabilistic knowledge fusion. In *KDD-14*.
- Doorenbos**, R. (1994). Combining left and right unlinking for matching a large number of learned rules. In *AAAI-94*.
- Doran**, J. and Michie, D. (1966). Experiments with the graph traverser program. *Proc. Roy. Soc.*, 294, Series A, 235–259.
- Dorf**, R. C. and Bishop, R. H. (2004). *Modern Control Systems* (10th edition). Prentice-Hall.
- Dorigo**, M., Birattari, M., Blum, C., Clerc, M., Stützle, T., and Winfield, A. (2008). *Ant Colony Optimization and Swarm Intelligence: 6th International Conference, ANTS 2008, Brussels, Belgium, September 22–24, 2008, Proceedings*, Vol. 5217. Springer-Verlag.
- Doshi-Velez**, F. and Kim, B. (2017). Towards a rigorous science of interpretable machine learning. arXiv:1702.08608.
- Doucet**, A. (1997). *Monte Carlo methods for Bayesian estimation of hidden Markov models: Application to radiation signals*. Ph.D. thesis, Université de Paris-Sud.
- Doucet**, A., de Freitas, J. F. G., and Gordon, N. (2001). *Sequential Monte Carlo Methods in Practice*. Springer-Verlag.
- Doucet**, A., de Freitas, J. F. G., Murphy, K., and Russell, S. J. (2000). Rao-Blackwellised particle filtering for dynamic Bayesian networks. In *UAI-00*.
- Doucet**, A. and Johansen, A. M. (2011). A tutorial on particle filtering and smoothing: Fifteen years later. In Crisan, D. and Rozovskii, B. (Eds.), *Oxford Handbook of Nonlinear Filtering*. Oxford.
- Dowty**, D., Wall, R., and Peters, S. (1991). *Introduction to Montague Semantics*. D. Reidel.
- Doyle**, J. (1979). A truth maintenance system. *AIJ*, 12, 231–272.
- Doyle**, J. (1983). What is rational psychology? Toward a modern mental philosophy. *AIMag*, 4, 50–53.
- Drabble**, B. (1990). Mission scheduling for spacecraft: Diaries of T-SCHED. In *Expert Planning Systems*. Institute of Electrical Engineers.
- Dragan**, A. D., Lee, K. C., and Srinivasa, S. (2013). Legibility and predictability of robot motion. In *HRI-13*.
- Dredze**, M., Crammer, K., and Pereira, F. (2008). Confidence-weighted linear classification. In *ICML-08*.
- Dressel**, J. and Farid, H. (2018). The accuracy, fairness, and limits of predicting recidivism. *Science Advances*, 4, eaao5580.
- Dreyfus**, H. L. (1972). *What Computers Can't Do: A Critique of Artificial Reason*. Harper and Row.
- Dreyfus**, H. L. (1992). *What Computers Still Can't Do: A Critique of Artificial Reason*. MIT Press.
- Dreyfus**, H. L. and Dreyfus, S. E. (1986). *Mind over Machine: The Power of Human Intuition and Expertise in the Era of the Computer*. Blackwell.
- Dreyfus**, S. E. (1962). The numerical solution of variational problems. *J. Math. Anal. and Appl.*, 5, 30–45.
- Dreyfus**, S. E. (1969). An appraisal of some shortest-paths algorithms. *Operations Research*, 17, 395–412.
- Dreyfus**, S. E. (1990). Artificial neural networks, back propagation, and the Kelley–Bryson gradient procedure. *J. Guidance, Control, and Dynamics*, 13, 926–928.
- Du**, S. S., Lee, J. D., Li, H., Wang, L., and Zhai, X. (2018). Gradient descent finds global minima of deep neural networks. arXiv:1811.03804.
- Dubois**, D. and Prade, H. (1994). A survey of belief revision and updating rules in various uncertainty models. *Int. J. Intelligent Systems*, 9, 61–100.
- Duda**, R. O. and Hart, P. E. (1973). *Pattern classification and scene analysis*. Wiley.

- Duda**, R. O., Hart, P. E., and Stork, D. G. (2001). *Pattern Classification* (2nd edition). Wiley.
- Dudek**, G. and Jenkin, M. (2000). *Computational Principles of Mobile Robotics*. Cambridge University Press.
- Duffy**, D. (1991). *Principles of Automated Theorem Proving*. John Wiley & Sons.
- Dunn**, H. L. (1946). Record linkage". *Am. J. Public Health*, 36, 1412–1416.
- Durfee**, E. H. and Lesser, V. R. (1989). Negotiating task decomposition and allocation using partial global planning. In Huhns, M. and Gasser, L. (Eds.), *Distributed AI*, Vol. 2. Morgan Kaufmann.
- Durme**, B. V. and Pasca, M. (2008). Finding cars, goddesses and enzymes: Parametrizable acquisition of labeled instances for open-domain information extraction. In *AAAI-08*.
- Dwork**, C. (2008). Differential privacy: A survey of results. In *International Conference on Theory and Applications of Models of Computation*.
- Dwork**, C., Hardt, M., Pitassi, T., Reingold, O., and Zemel, R. (2012). Fairness through awareness. In *Proc. 3rd innovations in theoretical computer science conference*.
- Dwork**, C., Roth, A., et al. (2014). The algorithmic foundations of differential privacy. *Foundations and Trends in Theoretical Computer Science*, 9, 211–407.
- Dyson**, F. (2004). A meeting with Enrico Fermi. *Nature*, 427, 297.
- Dyson**, G. (1998). *Darwin among the machines : the evolution of global intelligence*. Perseus Books.
- Earley**, J. (1970). An efficient context-free parsing algorithm. *CACM*, 13, 94–102.
- Ebendt**, R. and Drechsler, R. (2009). Weighted A\* search—unifying view and application. *AIJ*, 173, 1310–1342.
- Eckerle**, J., Chen, J., Sturtevant, N. R., Zilles, S., and Holte, R. C. (2017). Sufficient conditions for node expansion in bidirectional heuristic search. In *ICAPS-17*.
- Eckhouse**, L., Lum, K., Conti-Cook, C., and Ciccolini, J. (2019). Layers of bias: A unified approach for understanding problems with risk assessment. *Criminal Justice and Behavior*, 46, 185–209.
- Edelkamp**, S. (2009). Scaling search with symbolic pattern databases. In *Model Checking and Artificial Intelligence (MOCHART)*.
- Edelkamp**, S. and Schrödl, S. (2012). *Heuristic Search*. Morgan Kaufmann.
- Edmonds**, J. (1965). Paths, trees, and flowers. *Canadian J. of Mathematics*, 17, 449–467.
- Edwards**, P. (Ed.). (1967). *The Encyclopedia of Philosophy*. Macmillan.
- Eiter**, T., Leone, N., Mateis, C., Pfeifer, G., and Scarcello, F. (1998). The KR system dlv: Progress report, comparisons and benchmarks. In *KR-98*.
- Elio**, R. (Ed.). (2002). *Common Sense, Reasoning, and Rationality*. Oxford University Press.
- Elkan**, C. (1997). Boosting and naive Bayesian learning. Tech. rep., Department of Computer Science and Engineering, University of California, San Diego.
- Ellsberg**, D. (1962). *Risk, Ambiguity, and Decision*. Ph.D. thesis, Harvard University.
- Elman**, J. L. (1990). Finding structure in time. *Cognitive Science*, 14, 179–211.
- Elman**, J. L., Bates, E., Johnson, M., Karmiloff-Smith, A., Parisi, D., and Plunkett, K. (1997). *Rethinking Innateness*. MIT Press.
- Elo**, A. E. (1978). *The rating of chess players: Past and present*. Arco Publishing.
- Elsken**, T., Metzen, J. H., and Hutter, F. (2018). Neural architecture search: A survey. arXiv:1808.05377.
- Empson**, W. (1953). *Seven Types of Ambiguity*. New Directions.
- Enderton**, H. B. (1972). *A Mathematical Introduction to Logic*. Academic Press.
- Engel**, J., Resnick, C., Roberts, A., Dieleman, S., Norouzi, M., Eck, D., and Simonyan, K. (2017). Neural audio synthesis of musical notes with wavenet autoencoders. In *Proc. 34th International Conference on Machine Learning-Volume 70*.
- Epstein**, R., Roberts, G., and Beber, G. (Eds.). (2008). *Parsing the Turing test*. Springer.
- Erdmann**, M. A. and Mason, M. (1988). An exploration of sensorless manipulation. *IEEE Journal of Robotics and Automation*, 4, 369–379.
- Ernst**, H. A. (1961). *MH-1, a Computer-Operated Mechanical Hand*. Ph.D. thesis, MIT.
- Ernst**, M., Millstein, T., and Weld, D. S. (1997). Automatic SAT-compilation of planning problems. In *IJCAI-97*.
- Erol**, K., Hendler, J., and Nau, D. S. (1994). HTN planning: Complexity and expressivity. In *AAAI-94*.
- Erol**, K., Hendler, J., and Nau, D. S. (1996). Complexity results for HTN planning. *AIJ*, 18, 69–93.
- Erol**, Y., Li, L., Ramsundar, B., and Russell, S. J. (2013). The extended parameter filter. In *ICML-13*.
- Erol**, Y., Wu, Y., Li, L., and Russell, S. J. (2017). A nearly-black-box online algorithm for joint parameter and state estimation in temporal models. In *AAAI-17*.
- Esteva**, A., Kuprel, B., Novoa, R. A., Ko, J., Swetter, S. M., Blau, H. M., and Thrun, S. (2017). Dermatologist-level classification of skin cancer with deep neural networks. *Nature*, 542, 115.
- Etzioni**, A. (2004). *From Empire to Community: A New Approach to International Relation*. Palgrave Macmillan.
- Etzioni**, A. and Etzioni, O. (2017a). Incorporating ethics into artificial intelligence. *The Journal of Ethics*, 21, 403–418.
- Etzioni**, A. and Etzioni, O. (2017b). Should artificial intelligence be regulated? *Issues in Science and Technology*, Summer.
- Etzioni**, O. (1989). Tractable decision-analytic control. In *Proc. First International Conference on Knowledge Representation and Reasoning*.
- Etzioni**, O., Banko, M., Soderland, S., and Weld, D. S. (2008). Open information extraction from the web. *CACM*, 51.
- Etzioni**, O., Hanks, S., Weld, D. S., Draper, D., Lesh, N., and Williamson, M. (1992). An approach to planning with incomplete information. In *KR-92*.
- Etzioni**, O., Banko, M., and Cafarella, M. J. (2006). Machine reading. In *AAAI-06*.
- Etzioni**, O., Cafarella, M. J., Downey, D., Popescu, A.-M., Shaked, T., Soderland, S., Weld, D. S., and Yates, A. (2005). Unsupervised named-entity extraction from the web: An experimental study. *AIJ*, 165(1), 91–134.
- Evans**, T. G. (1968). A program for the solution of a class of geometric-analogy intelligence-test questions. In Minsky, M. L. (Ed.), *Semantic Information Processing*. MIT Press.
- Fagin**, R., Halpern, J. Y., Moses, Y., and Vardi, M. Y. (1995). *Reasoning about Knowledge*. MIT Press.
- Fahlman**, S. E. (1974). A planning system for robot construction tasks. *AIJ*, 5, 1–49.
- Faugeras**, O. (1992). What can be seen in three dimensions with an uncalibrated stereo rig? In *ECCV*, Vol. 588 of *Lecture Notes in Computer Science*.
- Faugeras**, O., Luong, Q.-T., and Papadopoulou, T. (2001). *The Geometry of Multiple Images*. MIT Press.
- Fawcett**, T. and Provost, F. (1997). Adaptive fraud detection. *Data mining and knowledge discovery*, 1, 291–316.
- Fearing**, R. S. and Hollerbach, J. M. (1985). Basic solid mechanics for tactile sensing. *Int. J. Robotics Research*, 4, 40–54.
- Featherstone**, R. (1987). *Robot Dynamics Algorithms*. Kluwer Academic Publishers.
- Feigenbaum**, E. A. (1961). The simulation of verbal learning behavior. *Proc. Western Joint Computer Conference*, 19, 121–131.
- Feigenbaum**, E. A., Buchanan, B. G., and Ledderberg, J. (1971). On generality and problem solving: A case study using the DENDRAL program. In Meltzer, B. and Michie, D. (Eds.), *Machine Intelligence 6*. Edinburgh University Press.
- Feldman**, J. and Sprout, R. F. (1977). Decision theory and artificial intelligence II: The hungry monkey. Technical report, Computer Science Department, University of Rochester.
- Feldman**, J. and Yakimovsky, Y. (1974). Decision theory and artificial intelligence I: Semantics-based region analyzer. *AIJ*, 5, 349–371.
- Feldman**, M. (2017). Oak Ridge readies Summit supercomputer for 2018 debut. *Top500.org*, bit.ly/2ERRF9.
- Fellbaum**, C. (2001). *Wordnet: An Electronic Lexical Database*. MIT Press.
- Fellegi**, I. and Sunter, A. (1969). A theory for record linkage". *JASA*, 64, 1183–1210.
- Felner**, A., Korf, R. E., and Hanan, S. (2004). Additive pattern database heuristics. *JAIR*, 22, 279–318.
- Felner**, A. (2018). Position paper: Using early goal test in A\*. In *Eleventh Annual Symposium on Combinatorial Search*.
- Felner**, A., Korf, R. E., Meshulam, R., and Holte, R. C. (2007). Compressed pattern databases. *JAIR*, 30.
- Felner**, A., Zahavi, U., Holte, R. C., Schaeffer, J., Sturtevant, N. R., and Zhang, Z. (2011). Inconsistent heuristics in theory and practice. *AIJ*, 175, 1570–1603.
- Felzenszwalb**, P. and McAllester, D. A. (2007). The generalized A\* architecture. *JAIR*.
- Fenton**, N. and Neil, M. (2018). *Risk Assessment and Decision Analysis with Bayesian Networks* (2nd edition). Chapman and Hall.
- Ferguson**, T. (1992). Mate with knight and bishop in kriegspiel. *Theoretical Computer Science*, 96, 389–403.
- Ferguson**, T. (1995). Mate with the two bishops in kriegspiel. www.math.ucla.edu/~tom/papers.
- Ferguson**, T. (2001). *Optimal Stopping and Applications*. www.math.ucla.edu/~tom/Stopping/Contents.html.
- Ferguson**, T. (1973). Bayesian analysis of some non-parametric problems. *Annals of Statistics*, 1, 209–230.
- Fern**, A., Natarajan, S., Judah, K., and Tadepalli, P. (2014). A decision-theoretic model of assistance. *JAIR*, 50, 71–104.
- Fernandez**, J. M. F. and Mahlmann, T. (2018). The Dota 2 bot competition. *IEEE Transactions on Games*.
- Ferraris**, P. and Giunchiglia, E. (2000). Planning as satisfiability in nondeterministic domains. In *AAAI-00*.
- Ferriss**, T. (2007). *The 4-Hour Workweek*. Crown.
- Ferrucci**, D., Brown, E., Chu-Carroll, J., Fan, J., Gondek, D., Kalyanpur, A. A., Lally, A., Murdoch, J. W., Nyberg, E., Prager, J., Schlaefler, N., and Welty, C. (2010). Building Watson: An overview of the DeepQA project. *AI Magazine*, Fall.

- Fikes, R. E., Hart, P. E., and Nilsson, N. J.** (1972). Learning and executing generalized robot plans. *AIJ*, 3, 251–288.
- Fikes, R. E. and Nilsson, N. J.** (1971). STRIPS: A new approach to the application of theorem proving to problem solving. *AIJ*, 2, 189–208.
- Fikes, R. E. and Nilsson, N. J.** (1993). STRIPS, a retrospective. *AIJ*, 59, 227–232.
- Fine, S., Singer, Y., and Tishby, N.** (1998). The hierarchical hidden Markov model: Analysis and applications. *Machine Learning*, 32.
- Finn, C., Abbeel, P., and Levine, S.** (2017). Model-agnostic meta-learning for fast adaptation of deep networks. In *Proc. 34th International Conference on Machine Learning—Volume 70*.
- Finney, D. J.** (1947). *Probit analysis: A statistical treatment of the sigmoid response curve*. Cambridge University Press.
- Firoiu, V., Whitney, W. F., and Tenenbaum, J. B.** (2017). Beating the world's best at Super Smash Bros. with deep reinforcement learning. arXiv:1702.06230.
- Firth, J.** (1957). *Papers in Linguistics*. Oxford University Press.
- Fisher, R. A.** (1922). On the mathematical foundations of theoretical statistics. *Phil. Trans. Roy. Soc. A*, 222, 309–368.
- Fix, E. and Hodges, J. L.** (1951). Discriminatory analysis—Nonparametric discrimination: Consistency properties. Tech. rep., USAF School of Aviation Medicine.
- Floreano, D., Zufferey, J. C., Srinivasan, M. V., and Ellington, C.** (2009). *Flying Insects and Robots*. Springer.
- Floyd, R. W.** (1962). Algorithm 97: Shortest path. *CACM*, 5, 345.
- Fogel, D. B.** (2000). *Evolutionary Computation: Toward a New Philosophy of Machine Intelligence*. IEEE Press.
- Fogel, L. J., Owens, A. J., and Walsh, M. J.** (1966). *Artificial Intelligence through Simulated Evolution*. Wiley.
- Forbes, J., Huang, T., Kanazawa, K., and Russell, S. J.** (1995). The BATmobile: Towards a Bayesian automated taxi. In *IJCAI-95*.
- Forbus, K. D.** (1985). Qualitative process theory. In Bobrow, D. (Ed.), *Qualitative Reasoning About Physical Systems*. MIT Press.
- Forbus, K. D. and de Kleer, J.** (1993). *Building Problem Solvers*. MIT Press.
- Forbus, K. D., Hinrichs, T. R., De Kleer, J., and Usher, J. M.** (2010). FIRE: Infrastructure for experience-based systems with common sense. In *AAAI Fall Symposium: Commonsense Knowledge*.
- Ford, K. M. and Hayes, P. J.** (1995). Turing Test considered harmful. In *IJCAI-95*.
- Ford, L. R.** (1956). Network flow theory. Tech. rep., RAND Corporation.
- Ford, M.** (2015). *Rise of the Robots: Technology and the Threat of a Jobless Future*. Basic Books.
- Ford, M.** (2018). *Architects of Intelligence*. Packt.
- Forestier, J.-P. and Varaiya, P.** (1978). Multilayer control of large Markov chains. *IEEE Transactions on Automatic Control*, 23, 298–304.
- Forgy, C.** (1981). OPS5 user's manual. Technical report, Computer Science Department, Carnegie-Mellon University.
- Forgy, C.** (1982). A fast algorithm for the many patterns/many objects match problem. *AIJ*, 19, 17–37.
- Forster, E. M.** (1909). *The Machine Stops*. Sheba Blake.
- Forsyth, D. and Ponce, J.** (2002). *Computer Vision: A Modern Approach*. Prentice Hall.
- Fouhey, D., Kuo, W.-C., Efros, A., and Malik, J.** (2018). From lifestyle vlogs to everyday interactions. In *CVPR-18*.
- Fourier, J.** (1827). Analyse des travaux de l'Académie Royale des Sciences, pendant l'année 1824; partie mathématique. *Histoire de l'Académie Royale des Sciences de France*, 7, xlvi–lv.
- Fox, C. and Tversky, A.** (1995). Ambiguity aversion and comparative ignorance. *Quarterly Journal of Economics*, 110, 585–603.
- Fox, D., Burgard, W., Dellaert, F., and Thrun, S.** (1999). Monte Carlo localization: Efficient position estimation for mobile robots. In *AAAI-99*.
- Fox, M. S.** (1990). Constraint-guided scheduling: A short history of research at CMU. *Computers in Industry*, 14, 79–88.
- Fox, M. S., Allen, B., and Strohm, G.** (1982). Job shop scheduling: An investigation in constraint-directed reasoning. In *AAAI-82*.
- Franco, J. and Paull, M.** (1983). Probabilistic analysis of the Davis Putnam procedure for solving the satisfiability problem. *Discrete Applied Mathematics*, 5, 77–87.
- Francois-Lavet, V., Henderson, P., Islam, R., Bellemare, M. G., and Pineau, J.** (2018). An introduction to deep reinforcement learning. *Foundations and Trends in Machine Learning*, 11, 219–354.
- Frank, I., Basin, D. A., and Matsubara, H.** (1998). Finding optimal strategies for imperfect information games. In *AAAI-98*.
- Frank, R. H. and Cook, P. J.** (1996). *The Winner-Take-All Society*. Penguin.
- Frans, K., Ho, J., Chen, X., Abbeel, P., and Schulman, J.** (2018). Meta learning shared hierarchies. In *ICLR-18*.
- Franz, A. and Brants, T.** (2006). All our n-gram are belong to you. Google blog, ai.googleblog.com/2006/08/all-our-n-gram-are-belong-to-you.html.
- Frege, G.** (1879). *Begriffsschrift, eine der arithmetischen nachgebildete Formelsprache des reinen Denkens*. Halle, Berlin. English translation appears in van Heijenoort (1967).
- Freitag, D. and McCallum, A.** (2000). Information extraction with hmm structures learned by stochastic optimization. In *AAAI-00*.
- Freuder, E. C.** (1978). Synthesizing constraint expressions. *CACM*, 21, 958–966.
- Freuder, E. C.** (1982). A sufficient condition for backtrack-free search. *JACM*, 29, 24–32.
- Freuder, E. C.** (1985). A sufficient condition for backtrack-bounded search. *JACM*, 32, 755–761.
- Freund, Y. and Schapire, R. E.** (1996). Experiments with a new boosting algorithm. In *ICML-96*.
- Freund, Y. and Schapire, R. E.** (1999). Large margin classification using the perceptron algorithm. *Machine Learning*, 37, 277–296.
- Frey, B. J.** (1998). *Graphical models for machine learning and digital communication*. MIT Press.
- Frey, C. B. and Osborne, M. A.** (2017). The future of employment: How susceptible are jobs to computerisation? *Technological forecasting and social change*, 114, 254–280.
- Friedberg, R. M.** (1958). A learning machine: Part I. *IBM Journal of Research and Development*, 2, 2–13.
- Friedberg, R. M., Dunham, B., and North, T.** (1959). A learning machine: Part II. *IBM Journal of Research and Development*, 3, 282–287.
- Friedman, G. J.** (1959). Digital simulation of an evolutionary process. *General Systems Yearbook*, 4, 171–184.
- Friedman, J., Hastie, T., and Tibshirani, R.** (2000). Additive logistic regression: A statistical view of boosting. *Annals of Statistics*, 28, 337–374.
- Friedman, J.** (2001). Greedy function approximation: A gradient boosting machine. *Annals of statistics*, 29, 1189–1232.
- Friedman, N.** (1998). The Bayesian structural EM algorithm. In *UAI-98*.
- Friedman, N. and Goldszmidt, M.** (1996). Learning Bayesian networks with local structure. In *UAI-96*.
- Friedman, N. and Koller, D.** (2003). Being Bayesian about Bayesian network structure: A Bayesian approach to structure discovery in Bayesian networks. *Machine Learning*, 50, 95–125.
- Friedman, N., Murphy, K., and Russell, S. J.** (1998). Learning the structure of dynamic probabilistic networks. In *UAI-98*.
- Friedman, N.** (2004). Inferring cellular networks using probabilistic graphical models. *Science*, 303.
- Fruhwirth, T. and Abdenadher, S.** (2003). *Essentials of constraint programming*. Cambridge University Press.
- Fuchs, J. J., Gasquet, A., Olalanty, B., and Currie, K. W.** (1990). PlanERS-1: An expert planning system for generating spacecraft mission plans. In *First International Conference on Expert Planning Systems*. Institute of Electrical Engineers.
- Fudenberg, D. and Tirole, J.** (1991). *Game theory*. MIT Press.
- Fukunaga, A. S., Rabideau, G., Chien, S., and Yan, D.** (1997). ASPEN: A framework for automated planning and scheduling of spacecraft control and operations. In *Proc. International Symposium on AI, Robotics and Automation in Space*.
- Fukushima, K.** (1980). Neocognitron: A self-organizing neural network model for a mechanism of pattern recognition unaffected by shift in position. *Biological Cybernetics*, 36, 193–202.
- Fukushima, K. and Miyake, S.** (1982). Neocognitron: A self-organizing neural network model for a mechanism of visual pattern recognition. In *Competition and cooperation in neural nets*. Springer.
- Fuller, S. B., Straw, A. D., Peek, M. Y., Murray, R. M., and Dickinson, M. H.** (2014). Flying Drosophila stabilize their vision-based velocity controller by sensing wind with their antennae. *Proc. National Academy of Sciences of the United States of America*, 111 13, E1182–91.
- Fung, C., Yoon, C. J. M., and Beschastnikh, I.** (2018). Mitigating sybils in federated learning poisoning. arXiv:1808.04866.
- Fung, R. and Chang, K. C.** (1989). Weighting and integrating evidence for stochastic simulation in Bayesian networks. In *UAI-5*.
- Gaddum, J. H.** (1933). Reports on biological standard III: Methods of biological assay depending on a quantal response. Special report series of the medical research council, Medical Research Council.
- Gaifman, H.** (1964a). Concerning measures in first order calculi. *Israel J. Mathematics*, 2, 1–18.
- Gaifman, H.** (1964b). Concerning measures on Boolean algebras. *Pacific J. Mathematics*, 14, 61–73.
- Gallaire, H. and Minker, J.** (Eds.). (1978). *Logic and Databases*. Plenum.
- Gallier, J. H.** (1986). *Logic for Computer Science: Foundations of Automatic Theorem Proving*. Harper and Row.
- Galton, F.** (1886). Regression towards mediocrity in hereditary stature. *J. Anthropological Institute of Great Britain and Ireland*, 15, 246–263.

- Gamba**, A., Gamberini, L., Palmieri, G., and Sanna, R. (1961). Further experiments with PAPA. *Nuovo Cimento Supplemento*, 20, 221–231.
- Gandomi**, A. and Haider, M. (2015). Beyond the hype: Big data concepts, methods, and analytics. *International journal of information management*, 35, 137–144.
- Gao**, J. (2014). Machine learning applications for data center optimization. Google Research.
- García**, J. and Fernández, F. (2015). A comprehensive survey on safe reinforcement learning. *JMLR*, 16, 1437–1480.
- Gardner**, M. (1968). *Logic Machines, Diagrams and Boolean Algebra*. Dover.
- Garey**, M. R. and Johnson, D. S. (1979). *Computers and Intractability*. W. H. Freeman.
- Gaschnig**, J. (1977). A general backtrack algorithm that eliminates most redundant tests. In *IJCAI-77*.
- Gaschnig**, J. (1979). Performance measurement and analysis of certain search algorithms. Technical report, Computer Science Department, Carnegie-Mellon University.
- Gasser**, R. (1995). *Efficiently harnessing computational resources for exhaustive search*. Ph.D. thesis, ETH Zürich.
- Gat**, E. (1998). Three-layered architectures. In Kortenkamp, D., Bonasso, R. P., and Murphy, R. (Eds.), *AI-based Mobile Robots: Case Studies of Successful Robot Systems*. MIT Press.
- Gatys**, L. A., Ecker, A. S., and Bethge, M. (2016). Image style transfer using convolutional neural networks. In *CVPR-16*.
- Gauci**, J., Conti, E., Liang, Y., Virochsiri, K., He, Y., Kaden, Z., Narayanan, V., and Ye, X. (2018). Horizon: Facebook’s open source applied reinforcement learning platform. arXiv:1811.00260.
- Gauss**, C. F. (1809). *Theoria Motus Corporum Coelestium in Sectionibus Conicis Solem Ambientium*. Sumtibus F. Perthes et I. H. Besser, Hamburg.
- Gauss**, C. F. (1829). Beiträge zur theorie der algebraischen gleichungen. *Werke*, 3, 71–102.
- Gazzaniga**, M. (2018). *The Consciousness Instinct*. Farrar, Straus and Giroux.
- Gebru**, T., Morgenstern, J., Vecchione, B., Vaughan, J. W., Wallach, H. M., III, H. D., and Crawford, K. (2018). Datasheets for datasets. arXiv:1803.09010.
- Geiger**, D., Verma, T., and Pearl, J. (1990a). d-separation: From theorems to algorithms. In Henrion, M., Shachter, R. D., Kanal, L. N., and Lemmer, J. F. (Eds.), *UAI-90*. Elsevier.
- Geiger**, D., Verma, T., and Pearl, J. (1990b). Identifying independence in Bayesian networks. *Networks*, 20, 507–534.
- Gelb**, A. (1974). *Applied Optimal Estimation*. MIT Press.
- Gelernter**, H. (1959). Realization of a geometry-theorem proving machine. In *Proc. an International Conference on Information Processing*. UNESCO House.
- Gelfond**, M. and Lifschitz, V. (1988). Compiling circumscriptive theories into logic programs. In *Non-Monotonic Reasoning: 2nd International Workshop Proceedings*.
- Gelfond**, M. (2008). Answer sets. In van Harmelen, F., Lifschitz, V., and Porter, B. (Eds.), *Handbook of Knowledge Representation*. Elsevier.
- Gelman**, A. (2004). Exploratory data analysis for complex models. *Journal of Computational and Graphical Statistics*, 13, 755–779.
- Gelman**, A., Carlin, J. B., Stern, H. S., and Rubin, D. (1995). *Bayesian Data Analysis*. Chapman & Hall.
- Geman**, S. and Geman, D. (1984). Stochastic relaxation, Gibbs distributions, and Bayesian restoration of images. *PAMI*, 6, 721–741.
- Gene Ontology Consortium**, The. (2008). The gene ontology project in 2008. *Nucleic Acids Research*, 36(D440–D444).
- Genesereth**, M. R. (1984). The use of design descriptions in automated diagnosis. *AIJ*, 24, 411–436.
- Genesereth**, M. R. and Nilsson, N. J. (1987). *Logical Foundations of Artificial Intelligence*. Morgan Kaufmann.
- Genesereth**, M. R. and Nourbakhsh, I. (1993). Time-saving tips for problem solving with incomplete information. In *AAAI-93*.
- Genesereth**, M. R. and Smith, D. E. (1981). Meta-level architecture. Memo, Computer Science Department, Stanford University.
- Gent**, I., Petrie, K., and Puget, J.-F. (2006). Symmetry in constraint programming. In Rossi, F., van Beek, P., and Walsh, T. (Eds.), *Handbook of Constraint Programming*. Elsevier.
- Géron**, A. (2019). *Hands-On Machine Learning with Scikit-Learn, Keras, and TensorFlow: Concepts, Tools, and Techniques to Build Intelligent Systems*. O’Reilly.
- Gers**, F. A., Schmidhuber, J., and Cummins, F. (2000). Learning to forget: Continual prediction with LSTM. *Neural Computation*, 12, 2451–2471.
- Getoor**, L. and Taskar, B. (Eds.). (2007). *Introduction to Statistical Relational Learning*. MIT Press.
- Ghaheri**, A., Shoar, S., Naderan, M., and Hoseini, S. S. (2015). The applications of genetic algorithms in medicine. *Oman medical journal*, 30, 406–416.
- Ghahramani**, Z. (1998). Learning dynamic Bayesian networks. In *Adaptive Processing of Sequences and Data Structures*.
- Ghahramani**, Z. (2005). Tutorial on nonparametric Bayesian methods. Given at the UAI-05 Conference.
- Ghallab**, M., Howe, A., Knoblock, C. A., and McDermott, D. (1998). PDDL—The planning domain definition language. Tech. rep., Yale Center for Computational Vision and Control.
- Ghallab**, M. and Laruelle, H. (1994). Representation and control in IxTeT, a temporal planner. In *AIPS-94*.
- Ghallab**, M., Nau, D. S., and Traverso, P. (2004). *Automated Planning: Theory and practice*. Morgan Kaufmann.
- Ghallab**, M., Nau, D. S., and Traverso, P. (2016). *Automated Planning and aTing*. Cambridge University Press.
- Gibbs**, R. W. (2006). Metaphor interpretation as embodied simulation. *Mind*, 21, 434–458.
- Gibson**, J. J. (1950). *The Perception of the Visual World*. Houghton Mifflin.
- Gibson**, J. J. (1979). *The Ecological Approach to Visual Perception*. Houghton Mifflin.
- Gibson**, J. J. (1955). Parallax and perspective during aircraft landings. *American Journal of Psychology*, 68, 372–385.
- Gilks**, W. R., Richardson, S., and Spiegelhalter, D. J. (Eds.). (1996). *Markov chain Monte Carlo in practice*. Chapman and Hall.
- Gilks**, W. R., Thomas, A., and Spiegelhalter, D. J. (1994). A language and program for complex Bayesian modelling. *The Statistician*, 43, 169–178.
- Gilks**, W. R. and Berzuini, C. (2001). Following a moving target—Monte Carlo inference for dynamic Bayesian models. *J. Royal Statistical Society, Series B*, 63, 127–146.
- Gilks**, W. R. and Wild, P. P. (1992). Adaptive rejection sampling for Gibbs sampling. *Applied Statistics*, 41, 337–348.
- Gillies**, D. B. (1959). Solutions to general non-zero-sum games. In Tucker, A. W. and Luce, L. D. (Eds.), *Contributions to the Theory of Games, volume IV*. Princeton University Press.
- Gilmore**, P. C. (1960). A proof method for quantification theory: Its justification and realization. *IBM Journal of Research and Development*, 4, 28–35.
- Gilpin**, A., Sandholm, T., and Sorenson, T. (2008). A heads-up no-limit Texas Hold’em poker player: Discretized betting models and automatically generated equilibrium-finding programs. In *AAMAS-08*.
- Ginsberg**, M. L. (1993). *Essentials of Artificial Intelligence*. Morgan Kaufmann.
- Ginsberg**, M. L. (2001). GIB: Imperfect information in a computationally challenging game. *JAIR*, 14, 303–358.
- Gionis**, A., Indyk, P., and Motwani, R. (1999). Similarity search in high dimensions via hashing. In *Proc. 25th Very Large Database (VLDB) Conference*.
- Girshick**, R., Donahue, J., Darrell, T., and Malik, J. (2016). Region-based convolutional networks for accurate object detection and segmentation. *PAMI*, 38, 142–58.
- Gittins**, J. C. (1989). *Multi-Armed Bandit Allocation Indices*. Wiley.
- Gittins**, J. C. and Jones, D. M. (1974). A dynamic allocation index for the sequential design of experiments. In Gani, J. (Ed.), *Progress in Statistics*. North-Holland.
- Glanc**, A. (1978). On the etymology of the word “robot”. *SIGART Newsletter*, 67, 12.
- Glickman**, M. E. (1999). Parameter estimation in large dynamic paired comparison experiments. *Applied Statistics*, 48, 377–394.
- Glorot**, X., Bordes, A., and Bengio, Y. (2011). Deep sparse rectifier neural networks. In *AISTATS’11*.
- Glover**, F. and Laguna, M. (Eds.). (1997). *Tabu search*. Kluwer.
- Gluss**, B. (1959). An optimum policy for detecting a fault in a complex system. *Operations Research*, 7, 468–477.
- Godefroid**, P. (1990). Using partial orders to improve automatic verification methods. In *Proc. 2nd Int’l Workshop on Computer Aided Verification*.
- Gödel**, K. (1930). Über die Vollständigkeit des Logikkalküls. Ph.D. thesis, University of Vienna.
- Gödel**, K. (1931). Über formal unentscheidbare Sätze der Principia mathematica und verwandter Systeme I. *Monatshefte für Mathematik und Physik*, 38, 173–198.
- Goebel**, J., Volk, K., Walker, H., and Gerbault, F. (1989). Automatic classification of spectra from the infrared astronomical satellite (IRAS). *Astronomy and Astrophysics*, 222, L5–L8.
- Goertzel**, B. and Pennachin, C. (2007). *Artificial General Intelligence*. Springer.
- Gogate**, V. and Domingos, P. (2011). Approximation by quantization. In *UAI-11*.
- Gold**, E. M. (1967). Language identification in the limit. *Information and Control*, 10, 447–474.
- Goldberg**, A. V., Kaplan, H., and Werneck, R. F. (2006). Reach for A\*: Efficient point-to-point shortest path algorithms. In *Workshop on algorithm engineering and experiments*.
- Goldberg**, Y. (2017). Neural network methods for natural language processing. *Synthesis Lectures on Human Language Technologies*, 10.
- Goldberg**, Y., Zhao, K., and Huang, L. (2013). Efficient implementation of beam-search incremental parsers. In *ACL-13*.

- Goldman**, R. and Boddy, M. (1996). Expressive planning and explicit knowledge. In *AIPS-96*.
- Goldszmidt**, M. and Pearl, J. (1996). Qualitative probabilities for default reasoning, belief revision, and causal modeling. *AIJ*, 84, 57–112.
- Golomb**, S. and Baumert, L. (1965). Backtrack programming. *JACM*, 14, 516–524.
- Golub**, G., Heath, M., and Wahba, G. (1979). Generalized cross-validation as a method for choosing a good ridge parameter. *Technometrics*, 21.
- Gomes**, C., Selman, B., Crato, N., and Kautz, H. (2000). Heavy-tailed phenomena in satisfiability and constraint processing. *JAR*, 24, 67–100.
- Gomes**, C., Kautz, H., Sabharwal, A., and Selman, B. (2008). Satisfiability solvers. In van Harmelen, F., Lifschitz, V., and Porter, B. (Eds.), *Handbook of Knowledge Representation*. Elsevier.
- Gomes**, C. and Selman, B. (2001). Algorithm portfolios. *AIJ*, 126, 43–62.
- Gomes**, C., Selman, B., and Kautz, H. (1998). Boosting combinatorial search through randomization. In *AAAI-98*.
- Gonthier**, G. (2008). Formal proof—The four-color theorem. *Notices of the AMS*, 55, 1382–1393.
- Good**, I. J. (1961). A causal calculus. *British Journal of the Philosophy of Science*, 11, 305–318.
- Good**, I. J. (1965a). The mystery of Go. *New Scientist*, 427, 172–174.
- Good**, I. J. (1965b). Speculations concerning the first ultraintelligent machine. In Alt, F. L. and Rubinoff, M. (Eds.), *Advances in Computers*, Vol. 6. Academic Press.
- Good**, I. J. (1983). *Good Thinking: The Foundations of Probability and Its Applications*. University of Minnesota Press.
- Goodfellow**, I., Bengio, Y., and Courville, A. (2016). *Deep Learning*. MIT Press.
- Goodfellow**, I., Bulatov, Y., Ibarz, J., Arnoud, S., and Shet, V. (2014). Multi-digit number recognition from Street View imagery using deep convolutional neural networks. In *International Conference on Learning Representations*.
- Goodfellow**, I., Pouget-Abadie, J., Mirza, M., Xu, B., Warde-Farley, D., Ozair, S., Courville, A., and Bengio, Y. (2015a). Generative adversarial nets. In *NeurIPS 27*.
- Goodfellow**, I., Vinyals, O., and Saxe, A. M. (2015b). Qualitatively characterizing neural network optimization problems. In *International Conference on Learning Representations*.
- Goodman**, J. (2001). A bit of progress in language modeling. Tech. rep., Microsoft Research.
- Goodman**, N. D., Mansinghka, V. K., Roy, D., Bonawitz, K., and Tenenbaum, J. B. (2008). Church: A language for generative models. In *UAI-08*.
- Goodman**, N. (1977). *The Structure of Appearance* (3rd edition). D. Reidel.
- Gopnik**, A. and Glymour, C. (2002). Causal maps and Bayes nets: A cognitive and computational account of theory-formation. In Caruthers, P., Stich, S., and Siegal, M. (Eds.), *The Cognitive Basis of Science*. Cambridge University Press.
- Gordon**, A. D., Graepel, T., Rolland, N., Russo, C., Borgström, J., and Guiver, J. (2014). Tabular: A schema-driven probabilistic programming language. In *POPL-14*.
- Gordon**, A. S. and Hobbs, J. R. (2017). *A Formal Theory of Commonsense Psychology: How People Think People Think*. Cambridge University Press.
- Gordon**, M. J., Milner, A. J., and Wadsworth, C. P. (1979). *Edinburgh LCF*. Springer-Verlag.
- Gordon**, N. (1994). *Bayesian methods for tracking*. Ph.D. thesis, Imperial College.
- Gordon**, N., Salmond, D. J., and Smith, A. F. M. (1993). Novel approach to nonlinear/non-Gaussian Bayesian state estimation. *IEE Proceedings F (Radar and Signal Processing)*, 140, 107–113.
- Gordon**, S. A. (1994). A faster Scrabble move generation algorithm. *Software Practice and Experience*, 24, 219–232.
- Gorry**, G. A. (1968). Strategies for computer-aided diagnosis. *Math. Biosciences*, 2, 293–318.
- Gorry**, G. A., Kassirer, J. P., Essig, A., and Schwartz, W. B. (1973). Decision analysis as the basis for computer-aided management of acute renal failure. *American Journal of Medicine*, 55, 473–484.
- Gottlob**, G., Leone, N., and Scarcello, F. (1999a). A comparison of structural CSP decomposition methods. In *IJCAI-99*.
- Gottlob**, G., Leone, N., and Scarcello, F. (1999b). Hypertree decompositions and tractable queries. In *PODS-99*.
- Goyal**, Y., Khot, T., Summers-Stay, D., Batra, D., and Parikh, D. (2017). Making the V in VQA matter: Elevating the role of image understanding in visual question answering. In *CVPR-17*.
- Grace**, K., Salvatier, J., Dafoe, A., Zhang, B., and Evans, O. (2017). When will AI exceed human performance? Evidence from AI experts. arXiv:1705.08807.
- Graham**, S. L., Harrison, M. A., and Ruzzo, W. L. (1980). An improved context-free recognizer. *ACM Transactions on Programming Languages and Systems*, 2, 415–462.
- Grassmann**, H. (1861). *Lehrbuch der Arithmetik*. Th. Chr. Fr. Enslin, Berlin.
- Grayson**, C. J. (1960). Decisions under uncertainty: Drilling decisions by oil and gas operators. Tech. rep., Harvard Business School.
- Green**, B., Wolf, A., Chomsky, C., and Laugherty, K. (1961). BASEBALL: An automatic question answerer. In *Proc. Western Joint Computer Conference*.
- Green**, C. (1969a). Application of theorem proving to problem solving. In *IJCAI-69*.
- Green**, C. (1969b). Theorem-proving by resolution as a basis for question-answering systems. In Meltzer, B., Michie, D., and Swann, M. (Eds.), *Machine Intelligence 4*. Edinburgh University Press.
- Green**, C. and Raphael, B. (1968). The use of theorem-proving techniques in question-answering systems. In *Proc. 23rd ACM National Conference*.
- Gribkoff**, E., Van den Broeck, G., and Suciu, D. (2014). Understanding the complexity of lifted inference and asymmetric weighted model counting. In *UAI-14*.
- Griffiths**, T. L., Kemp, C., and Tenenbaum, J. B. (2008). Bayesian models of cognition. In Sun, R. (Ed.), *The Cambridge handbook of computational cognitive modeling*. Cambridge University Press.
- Grinstead**, C. and Snell, J. (1997). *Introduction to Probability*. American Mathematical Society.
- Grosz**, B. J. and Stone, P. (2018). A century long commitment to assessing artificial intelligence and its impact on society. *Communications of the ACM*, 61.
- Grove**, W. and Meehl, P. (1996). Comparative efficiency of informal (subjective, impressionistic) and formal (mechanical, algorithmic) prediction procedures: The clinical statistical controversy. *Psychology, Public Policy, and Law*, 2, 293–323.
- Gruber**, T. (2004). Interview of Tom Gruber. *AIS SIGSEMIS Bulletin*, 1.
- Gu**, J. (1989). *Parallel Algorithms and Architectures for Very Fast AI Search*. Ph.D. thesis, Univ. of Utah.
- Guard**, J., Oglesby, F., Bennett, J., and Settle, L. (1969). Semi-automated mathematics. *JACM*, 16, 49–62.
- Guestrin**, C., Koller, D., Gearhart, C., and Kanodia, N. (2003a). Generalizing plans to new environments in relational MDPs. In *IJCAI-03*.
- Guestrin**, C., Koller, D., Parr, R., and Venkataraman, S. (2003b). Efficient solution algorithms for factored MDPs. *JAIR*, 19, 399–468.
- Guestrin**, C., Lagoudakis, M. G., and Parr, R. (2002). Coordinated reinforcement learning. In *ICML-02*.
- Guibas**, L. J., Knuth, D. E., and Sharir, M. (1992). Randomized incremental construction of Delaunay and Voronoi diagrams. *Algorithmica*, 7, 381–413.
- Gulshan**, V., Peng, L., Coram, M., Stumpe, M. C., Wu, D., Narayanaswamy, A., Venugopalan, S., Widner, K., Madams, T., Cuadros, J., et al. (2016). Development and validation of a deep learning algorithm for detection of diabetic retinopathy in retinal fundus photographs. *Jama*, 316, 2402–2410.
- Gunkel**, D. J. (2018). *Robot Rights*. MIT Press.
- Gunning**, D. (2016). Explainable artificial intelligence (xai). Tech. rep., DARPA.
- Guo**, C., Goldstein, T., Hannun, A., and van der Maaten, L. (2019). Certified data removal from machine learning models. arXiv:1911.03030.
- Gururangan**, S., Swayamdipta, S., Levy, O., Schwartz, R., Bowman, S., and Smith, N. A. (2018). Annotation artifacts in natural language inference data. arXiv:1803.02324.
- Guyon**, I., Bennett, K., Cawley, G. C., Escalante, H. J., Escalera, S., Ho, T. K., Macia, N., Ray, B., Saeed, M., Statnikov, A. R., and Viegas, E. (2015). Design of the 2015 ChaLearn AutoML challenge. In *IJCNN-15*.
- Guyon**, I. and Elisseeff, A. (2003). An introduction to variable and feature selection. *JMLR*, 3, 1157–1182.
- Hacking**, I. (1975). *The Emergence of Probability*. Cambridge University Press.
- Hadfield-Menell**, D., Dragan, A. D., Abbeel, P., and Russell, S. J. (2017a). Cooperative inverse reinforcement learning. In *NeurIPS 29*.
- Hadfield-Menell**, D., Dragan, A. D., Abbeel, P., and Russell, S. J. (2017b). The off-switch game. In *IJCAI-17*.
- Hadfield-Menell**, D. and Russell, S. J. (2015). Multitasking: Efficient optimal planning for bandit super-processes. In *UAI-15*.
- Hailperin**, T. (1984). Probability logic. *Notre Dame Journal of Formal Logic*, 25, 198–212.
- Hald**, A. (1990). *A History of Probability and Statistics and Their Applications before 1750*. Wiley.
- Hales**, T. (2005). A proof of the Kepler conjecture. *Annals of mathematics*, 162, 1065–1185.
- Hales**, T., Adams, M., Bauer, G., Gang, T. D., Harrison, J., Le Truong, H., Kaliszyk, C., Magron, V., McLaughlin, S., Nguyen, T. T., et al. (2017). A formal proof of the Kepler conjecture. In *Forum of Mathematics, Pi*.
- Halevy**, A. (2007). Dataspaces: A new paradigm for data integration. In *Brazilian Symposium on Databases*.
- Halevy**, A., Norvig, P., and Pereira, F. (2009). The unreasonable effectiveness of data. *IEEE Intelligent Systems, March/April*, 8–12.
- Halpern**, J. Y. (1990). An analysis of first-order logics of probability. *AIJ*, 46, 311–350.
- Halpern**, J. Y. (1999). Technical addendum, Cox's theorem revisited. *JAIR*, 11, 429–435.
- Halpern**, J. Y. and Weissman, V. (2008). Using first-order logic to reason about policies. *ACM Transactions on Information and System Security*, 11, 1–41.

- Hammersley, J. M.** and Handscomb, D. C. (1964). *Monte Carlo Methods*. Methuen.
- Han, J.**, Pei, J., and Kamber, M. (2011). *Data Mining: Concepts and Techniques*. Elsevier.
- Han, X.** and Boyden, E. (2007). Multiple-color optical activation, silencing, and desynchronization of neural activity, with single-spike temporal resolution. *PLoS One*, e299.
- Handschin, J. E.** and Mayne, D. Q. (1969). Monte Carlo techniques to estimate the conditional expectation in multi-stage nonlinear filtering. *Int. J. Control.*, 9, 547–559.
- Hans, A.**, Schneegäß, D., Schäfer, A. M., and Udluft, S. (2008). Safe exploration for reinforcement learning. In *ESANN*.
- Hansen, E.** (1998). Solving POMDPs by searching in policy space. In *UAI-98*.
- Hansen, E.** and Zilberman, S. (2001). LAO\*: a heuristic search algorithm that finds solutions with loops. *AIJ*, 129, 35–62.
- Hansen, P.** and Jaumard, B. (1990). Algorithms for the maximum satisfiability problem. *Computing*, 44, 279–303.
- Hanski, I.** and Cambefort, Y. (Eds.). (1991). *Dung Beetle Ecology*. Princeton University Press.
- Hansson, O.** and Mayer, A. (1989). Heuristic search as evidential reasoning. In *UAI-5*.
- Haralick, R. M.** and Elliott, G. L. (1980). Increasing tree search efficiency for constraint satisfaction problems. *AIJ*, 14, 263–313.
- Hardin, G.** (1968). The tragedy of the commons. *Science*, 162, 1243–1248.
- Hardt, M.**, Price, E., Srebro, N., et al. (2017). Equality of opportunity in supervised learning. In *NeurIPS 29*.
- Harris, T.** (2016). How technology is hijacking your mind—From a magician and Google design ethicist. medium.com/thrive-global/how-technology-hijacks-peoples-minds-from-a-magician-and-google-s-design-ethicist-56d62ef5edf3.
- Harris, Z.** (1954). Distributional structure. *Word*, 10.
- Harrison, J.** and March, J. G. (1984). Decision making and postdecision surprises. *Administrative Science Quarterly*, 29, 26–42.
- Harrow, A. W.**, Hassidim, A., and Lloyd, S. (2009). Quantum algorithm for linear systems of equations. *Physical Review Letters*, 103 15, 150502.
- Harsanyi, J.** (1967). Games with incomplete information played by Bayesian players. *Management Science*, 14, 159–182.
- Hart, P. E.**, Nilsson, N. J., and Raphael, B. (1968). A formal basis for the heuristic determination of minimum cost paths. *IEEE Transactions on Systems Science and Cybernetics*, SSC-4(2), 100–107.
- Hart, T. P.** and Edwards, D. J. (1961). The tree prune (TP) algorithm. Artificial intelligence project memo, MIT.
- Hartley, H.** (1958). Maximum likelihood estimation from incomplete data. *Biometrika*, 14, 174–194.
- Hartley, R.** and Zisserman, A. (2000). *Multiple view geometry in computer vision*. Cambridge University Press.
- Hashimoto, K.**, Xiong, C., Tsuruoka, Y., and Socher, R. (2016). A joint many-task model: Growing a neural network for multiple NLP tasks. arXiv:1611.01587.
- Haslum, P.**, Botea, A., Helmert, M., Bonet, B., and Koenig, S. (2007). Domain-independent construction of pattern database heuristics for cost-optimal planning. In *AAAI-07*.
- Haslum, P.** and Geffner, H. (2001). Heuristic planning with time and resources. In *Proc. IJCAI-01 Workshop on Planning with Resources*.
- Haslum, P.** (2006). Improving heuristics through relaxed search – An analysis of TP4 and HSP\* in the 2004 planning competition. *JAIR*, 25, 233–267.
- Hastie, T.** and Tibshirani, R. (1996). Discriminant adaptive nearest neighbor classification and regression. In *NeurIPS 8*.
- Hastie, T.**, Tibshirani, R., and Friedman, J. (2009). *The Elements of Statistical Learning: Data Mining, Inference and Prediction* (2nd edition). Springer-Verlag.
- Hastings, W. K.** (1970). Monte Carlo sampling methods using Markov chains and their applications. *Biometrika*, 57, 97–109.
- Hatem, M.** and Ruml, W. (2014). Simpler bounded suboptimal search. In *AAAI-14*.
- Haugeland, J.** (1985). *Artificial Intelligence: The Very Idea*. MIT Press.
- Havelund, K.**, Lowry, M., Park, S., Pecheur, C., Penix, J., Visser, W., and White, J. L. (2000). Formal analysis of the remote agent before and after flight. In *Proc. 5th NASA Langley Formal Methods Workshop*.
- Havenstein, H.** (2005). Spring comes to AI winter. *Computer World*, Fe. 14.
- Hawkins, J.** (1961). Self-organizing systems: A review and commentary. *Proc. IRE*, 49, 31–48.
- Hay, N.**, Russell, S. J., Shimony, S. E., and Tolpin, D. (2012). Selecting computations: Theory and applications. In *UAI-12*.
- Hayes, P. J.** (1978). The naive physics manifesto. In Michie, D. (Ed.), *Expert Systems in the Microelectronic Age*. Edinburgh University Press.
- Hayes, P. J.** (1979). The logic of frames. In Metzing, D. (Ed.), *Frame Conceptions and Text Understanding*. de Gruyter.
- Hayes, P. J.** (1985a). Naive physics I: Ontology for liquids. In Hobbs, J. R. and Moore, R. C. (Eds.), *Formal Theories of the Commonsense World*, chap. 3. Ablex.
- Hayes, P. J.** (1985b). The second naive physics manifesto. In Hobbs, J. R. and Moore, R. C. (Eds.), *Formal Theories of the Commonsense World*, chap. 1. Ablex.
- Hays, J.** and Efros, A. (2007). Scene completion Using millions of photographs. *ACM Transactions on Graphics (SIGGRAPH)*, 26.
- He, H.**, Bai, Y., Garcia, E. A., and Li, S. (2008). ADASYN: Adaptive synthetic sampling approach for imbalanced learning. In *2008 IEEE International Joint Conference on Neural Networks (IEEE World Congress on Computational Intelligence)*.
- He, K.**, Zhang, X., Ren, S., and Sun, J. (2016). Deep residual learning for image recognition. In *CVPR-16*.
- Heawood, P. J.** (1890). Map colouring theorem. *Quarterly Journal of Mathematics*, 24, 332–338.
- Hebb, D. O.** (1949). *The Organization of Behavior*. Wiley.
- Heckerman, D.** (1986). Probabilistic interpretation for MYCIN's certainty factors. In Kanal, L. N. and Lemmer, J. F. (Eds.), *UAI 2*. Elsevier.
- Heckerman, D.** (1991). *Probabilistic Similarity Networks*. MIT Press.
- Heckerman, D.** (1998). A tutorial on learning with Bayesian networks. In Jordan, M. I. (Ed.), *Learning in graphical models*. Kluwer.
- Heckerman, D.**, Geiger, D., and Chickering, D. M. (1994). Learning Bayesian networks: The combination of knowledge and statistical data. Technical report, Microsoft Research.
- Heess, N.**, Wayne, G., Silver, D., Lillicrap, T., Erez, T., and Tassa, Y. (2016). Learning continuous control policies by stochastic value gradients. In *NeurIPS 28*.
- Heidegger, M.** (1927). *Being and Time*. SCM Press.
- Heinlein, R. A.** (1973). *Time Enough for Love*. Putnam.
- Held, M.** and Karp, R. M. (1970). The traveling salesman problem and minimum spanning trees. *Operations Research*, 18, 1138–1162.
- Helmert, M.** (2001). On the complexity of planning in transportation domains. In *ECP-01*.
- Helmert, M.** (2006). The fast downward planning system. *JAIR*, 26, 191–246.
- Helmert, M.** and Röger, G. (2008). How good is almost perfect? In *AAAI-08*.
- Helmert, M.**, Röger, G., and Karpas, E. (2011). Fast downward stone soup: A baseline for building planner portfolios. In *ICAPS*.
- Hendeby, G.**, Karlsson, R., and Gustafsson, F. (2010). Particle filtering: The need for speed. *EURASIP J. Adv. Sig. Proc.*, June.
- Henrion, M.** (1988). Propagation of uncertainty in Bayesian networks by probabilistic logic sampling. In Lemmer, J. F. and Kanal, L. N. (Eds.), *UAI 2*. Elsevier.
- Henzinger, T. A.** and Sastry, S. (Eds.). (1998). *Hybrid Systems: Computation and Control*. Springer-Verlag.
- Herbrand, J.** (1930). *Recherches sur la Théorie de la Démonstration*. Ph.D. thesis, University of Paris.
- Herbrich, R.**, Minka, T., and Graepel, T. (2007). TrueSkill: A Bayesian skill rating system. In *NeurIPS 19*.
- Hernández-Orallo, J.** (2016). Evaluation in artificial intelligence: From task-oriented to ability-oriented measurement. *Artificial Intelligence Review*, 48, 397–447.
- Hess, C.** and Ostrom, E. (2007). *Understanding Knowledge as a Commons*. MIT Press.
- Hewitt, C.** (1977). Viewing control structures as patterns of passing messages. *AIJ*, 8, 323–364.
- Hewitt, C.** (1969). PLANNER: a language for proving theorems in robots. In *IJCAI-69*.
- Hezaveh, Y. D.**, Levasseur, L. P., and Marshall, P. J. (2017). Fast automated analysis of strong gravitational lenses with convolutional neural networks. *Nature*, 548, 555–557.
- Hierholzer, C.** (1873). Über die Möglichkeit, einen Linienzug ohne Wiederholung und ohne Unterbrechung zu umfahren. *Mathematische Annalen*, 6, 30–32.
- Hilbert, M.** and Lopez, P. (2011). The world's technological capacity to store, communicate, and compute information. *Science*, 332, 60–65.
- Hilgard, E. R.** and Bower, G. H. (1975). *Theories of Learning* (4th edition). Prentice-Hall.
- Hind, M.**, Mehta, S., Mojsilovic, A., Nair, R., Ramanamurthy, K. N., Olteanu, A., and Varshney, K. R. (2018). Increasing trust in AI services through supplier's declarations of conformity. arXiv:1808.07261.
- Hintikka, J.** (1962). *Knowledge and Belief*. Cornell University Press.
- Hinton, G. E.** and Anderson, J. A. (1981). *Parallel Models of Associative Memory*. Lawrence Erlbaum.
- Hinton, G. E.** and Nowlan, S. J. (1987). How learning can guide evolution. *Complex Systems*, 1, 495–502.
- Hinton, G. E.** and Sejnowski, T. (1983). Optimal perceptual inference. In *CVPR-83*.
- Hinton, G. E.** and Sejnowski, T. (1986). Learning and relearning in Boltzmann machines. In Rumelhart, D. E. and McClelland, J. L. (Eds.), *Parallel Distributed Processing*. MIT Press.
- Hinton, G. E.** (1987). Learning translation invariant recognition in a massively parallel network. In Goos, G. and Hartmanis, J. (Eds.), *PARLE: Parallel Architectures and Languages Europe*. Springer-Verlag.

- Hinton, G. E., Deng, L., Yu, D., Dahl, G., Mohamed, A. R., Jaitly, N., Senior, A., Vanhoucke, V., Nguyen, P., Sainath, T., and Kingsbury, B.** (2012). Deep neural networks for acoustic modeling in speech recognition. *Signal Processing Magazine*, 29, 82 – 97.
- Hinton, G. E., Osindero, S., and Teh, Y. W.** (2006). A fast learning algorithm for deep belief nets. *Neural Computation*, 18, 1527–1554.
- Hirth, M., Hößfeld, T., and Tran-Gia, P.** (2013). Analyzing costs and accuracy of validation mechanisms for crowdsourcing platforms. *Mathematical and Computer Modelling*, 57, 2918–2932.
- Ho, M. K., Littman, M. L., MacGlashan, J., Cushman, F., and Austerweil, J. L.** (2017). Showing versus doing: Teaching by demonstration. In *NeurIPS* 29.
- Ho, T. K.** (1995). Random decision forests. In *Proc. 3rd Int'l Conf. on Document Analysis and Recognition*.
- Hobbs, J. R.** (1990). *Literature and Cognition*. CSLI Press.
- Hobbs, J. R. and Moore, R. C. (Eds.)** (1985). *Formal Theories of the Commonsense World*. Ablex.
- Hochreiter, S.** (1991). Untersuchungen zu dynamischen neuronalen Netzen. Diploma thesis, Technische Universität München.
- Hochreiter, S. and Schmidhuber, J.** (1997). Long short-term memory. *Neural Computation*, 9, 1735–1780.
- Hoffmann, M., Bach, F. R., and Blei, D. M.** (2011). Online learning for latent Dirichlet allocation. In *NeurIPS* 23.
- Hoffmann, J.** (2001). FF: The fast-forward planning system. *AIMag*, 22, 57–62.
- Hoffmann, J. and Brafman, R. I.** (2006). Conformant planning via heuristic forward search: A new approach. *AIJ*, 170, 507–541.
- Hoffmann, J. and Brafman, R. I.** (2005). Contingent planning via heuristic forward search with implicit belief states. In *ICAPS-05*.
- Hoffmann, J.** (2005). Where “ignoring delete lists” works: Local search topology in planning benchmarks. *JAIR*, 24, 685–758.
- Hoffmann, J. and Nebel, B.** (2001). The FF planning system: Fast plan generation through heuristic search. *JAIR*, 14, 253–302.
- Hoffmann, J., Sabharwal, A., and Domshlak, C.** (2006). Friends or foes? An AI planning perspective on abstraction and search. In *ICAPS-06*.
- Hofleitner, A., Herring, R., Abbeel, P., and Bayen, A. M.** (2012). Learning the dynamics of arterial traffic from probe data using a dynamic Bayesian network. *IEEE Transactions on Intelligent Transportation Systems*, 13, 1679–1693.
- Hogan, N.** (1985). Impedance control: An approach to manipulation. Parts I, II, and III. *J. Dynamic Systems, Measurement, and Control*, 107, 1–24.
- Hoiem, D., Efros, A., and Hebert, M.** (2007). Recovering surface layout from an image. *IJCV*, 75, 151–172.
- Holland, J. H.** (1975). *Adaption in Natural and Artificial Systems*. University of Michigan Press.
- Holland, J. H.** (1995). *Hidden Order: How Adaptation Builds Complexity*. Addison-Wesley.
- Holte, R. C., Felner, A., Sharon, G., and Sturtevant, N. R.** (2016). Bidirectional search that is guaranteed to meet in the middle. In *AAAI-16*.
- Holzmann, G. J.** (1997). The Spin model checker. *IEEE Transactions on Software Engineering*, 23, 279–295.
- Hood, A.** (1824). Case 4th—28 July 1824 (Mr. Hood's cases of injuries of the brain). *Phrenological Journal and Miscellany*, 2, 82–94.
- Hooker, J.** (1995). Testing heuristics: We have it all wrong. *J. Heuristics*, 1, 33–42.
- Hoos, H. H. and Stützle, T.** (2004). *Stochastic Local Search: Foundations and Applications*. Morgan Kaufmann.
- Hoos, H. H. and Tsang, E.** (2006). Local search methods. In Rossi, F., van Beek, P., and Walsh, T. (Eds.), *Handbook of Constraint Processing*. Elsevier.
- Hopfield, J. J.** (1982). Neural networks and physical systems with emergent collective computational abilities. *PNAS*, 79, 2554–2558.
- Horn, A.** (1951). On sentences which are true of direct unions of algebras. *JSL*, 16, 14–21.
- Horn, B. K. P.** (1970). Shape from shading: A method for obtaining the shape of a smooth opaque object from one view. Technical report, MIT Artificial Intelligence Laboratory.
- Horn, B. K. P. and Brooks, M. J.** (1989). *Shape from Shading*. MIT Press.
- Horn, K. V.** (2003). Constructing a logic of plausible inference: A guide to Cox's theorem. *IJAR*, 34, 3–24.
- Horning, J. J.** (1969). *A Study of Grammatical Inference*. Ph.D. thesis, Stanford University.
- Horswill, I.** (2000). Functional programming of behavior-based systems. *Autonomous Robots*, 9, 83–93.
- Horvitz, E. J.** (1987). Problem-solving design: Reasoning about computational value, trade-offs, and resources. In *Proc. Second Annual NASA Research Forum*.
- Horvitz, E. J. and Barry, M.** (1995). Display of information for time-critical decision making. In *UAI-95*.
- Horvitz, E. J., Breese, J. S., Heckerman, D., and Hovel, D.** (1998). The Lumière project: Bayesian user modeling for inferring the goals and needs of software users. In *UAI-98*.
- Horvitz, E. J., Breese, J. S., and Henrion, M.** (1988). Decision theory in expert systems and artificial intelligence. *IJAR*, 2, 247–302.
- Horvitz, E. J. and Breese, J. S.** (1996). Ideal partition of resources for metareasoning. In *AAAI-96*.
- Hotelling, H.** (1933). Analysis of a complex of statistical variables into principal components. *J. Ed. Psych.*, 24, 417–441.
- Howard, J. and Gugger, S.** (2020). *Deep Learning for Coders with fastai and PyTorch*. O'Reilly.
- Howard, J. and Ruder, S.** (2018). Fine-tuned language models for text classification. arXiv:1801.06146.
- Howard, R. A.** (1960). *Dynamic Programming and Markov Processes*. MIT Press.
- Howard, R. A.** (1966). Information value theory. *IEEE Transactions on Systems Science and Cybernetics, SSC-2*, 22–26.
- Howard, R. A.** (1989). Microrisks for medical decision analysis. *Int. J. Technology Assessment in Health Care*, 5, 357–370.
- Howard, R. A. and Matheson, J. E.** (1984). Influence diagrams. In Howard, R. A. and Matheson, J. E. (Eds.), *Readings on the Principles and Applications of Decision Analysis*. Strategic Decisions Group.
- Howe, D.** (1987). The computational behaviour of Giard's paradox. In *LICS-87*.
- Howson, C.** (2003). Probability and logic. *J. Applied Logic*, 1, 151–165.
- Hsiao, K., Kaelbling, L. P., and Lozano-Perez, T.** (2007). Grasping POMDPs. In *ICRA-07*.
- Hsu, F.-H.** (2004). *Behind Deep Blue: Building the Computer that Defeated the World Chess Champion*. Princeton University Press.
- Hsu, F.-H., Anantharaman, T. S., Campbell, M. S., and Nowatzky, A.** (1990). A grandmaster chess machine. *Scientific American*, 263, 44–50.
- Hu, J. and Wellman, M. P.** (1998). Multiagent reinforcement learning: Theoretical framework and an algorithm. In *ICML-98*.
- Hu, J. and Wellman, M. P.** (2003). Nash Q-learning for general-sum stochastic games. *JMLR*, 4, 1039–1069.
- Huang, T., Koller, D., Malik, J., Ogasawara, G., Rao, B., Russell, S. J., and Weber, J.** (1994). Automatic symbolic traffic scene analysis using belief networks. In *AAAI-94*.
- Huang, T. and Russell, S. J.** (1998). Object identification: A Bayesian analysis with application to traffic surveillance. *AIJ*, 103, 1–17.
- Hubel, D. H. and Wiesel, T. N.** (1962). Receptive fields, binocular interaction and functional architecture in the cat's visual cortex. *J. Physiology*, 160, 106–154.
- Hubel, D. H. and Wiesel, T. N.** (1968). Receptive fields and functional architecture of monkey striate cortex. *J. Physiology*, 195, 215–243.
- Hubel, D. H.** (1988). *Eye, Brain, and Vision*. W. H. Freeman.
- Hubel, D. H. and Wiesel, T. N.** (1959). Receptive fields of single neurons in the cat's striate cortex. *Journal of Physiology*, 148, 574–591.
- Huddleston, R. D. and Pullum, G. K.** (2002). *The Cambridge Grammar of the English Language*. Cambridge University Press.
- Huffman, D. A.** (1971). Impossible objects as nonsense sentences. In Meltzer, B. and Michie, D. (Eds.), *Machine Intelligence 6*. Edinburgh University Press.
- Hughes, B. D.** (1995). *Random Walks and Random Environments, Vol. 1: Random Walks*. Oxford University Press.
- Hughes, G. E. and Cresswell, M. J.** (1996). *A New Introduction to Modal Logic*. Routledge.
- Huhns, M. N. and Singh, M. (Eds.)** (1998). *Readings in Agents*. Morgan Kaufmann.
- Hume, D.** (1739). *A Treatise of Human Nature* (2nd edition). Republished by Oxford University Press, 1978, Oxford.
- Humphrys, M.** (2008). How my program passed the Turing test. In Epstein, R., Roberts, G., and Beber, G. (Eds.), *Parsing the Turing Test*. Springer.
- Hunsberger, L. and Grosz, B. J.** (2000). A combinatorial auction for collaborative planning. In *Int. Conference on Multi-Agent Systems*.
- Hunt, W. and Brock, B.** (1992). A formal HDL and its use in the FM9001 verification. *Phil. Trans. Roy. Soc.*, 339.
- Hunter, L. and States, D. J.** (1992). Bayesian classification of protein structure. *IEEE Expert*, 7, 67–75.
- Hur, C.-K., Nori, A. V., Rajamani, S. K., and Samuel, S.** (2014). Slicing probabilistic programs. In *PLDI-14*.
- Hurst, M.** (2000). *The Interpretation of Text in Tables*. Ph.D. thesis, Edinburgh.
- Hurwicz, L.** (1973). The design of mechanisms for resource allocation. *American Economic Review Papers and Proceedings*, 63, 1–30.
- Huth, M. and Ryan, M.** (2004). *Logic in Computer Science: Modelling and Reasoning About Systems* (2nd edition). Cambridge University Press.
- Huttenlocher, D. and Ullman, S.** (1990). Recognizing solid objects by alignment with an image. *IJCV*, 5, 195–212.
- Hutter, F., Kotthoff, L., and Vanschoren, J.** (2019). *Automated Machine Learning*. Springer.

- Huygens**, C. (1657). De ratiociniis in ludo aleae. In van Schooten, F. (Ed.), *Exercitionum Mathematicorum*. Elsevierii, Amsterdam. Translated into English by John Arbuthnot (1692).
- Huyn**, N., Dechter, R., and Pearl, J. (1980). Probabilistic analysis of the complexity of A\*. *AII*, 15, 241–254.
- Huynh**, V. A. and Roy, N. (2009). icLQG: Combining local and global optimization for control in information space. In *ICRA-09*.
- Hwa**, R. (1998). An empirical evaluation of probabilistic lexicalized tree insertion grammars. In *ACL-98*.
- Hwang**, C. H. and Schubert, L. K. (1993). EL: A formal, yet natural, comprehensive knowledge representation. In *AAAI-93*.
- Hyafil**, L. and Rivest, R. (1976). Constructing optimal binary decision trees is NP-complete. *Information Processing Letters*, 5, 15–17.
- Ieong**, S. and Shoham, Y. (2005). Marginal contribution nets: A compact representation scheme for coalitional games. In *Proc. Sixth ACM Conference on Electronic Commerce (EC'05)*.
- Ingerman**, P. Z. (1967). Panini–Backus form suggested. *CACM*, 10, 137.
- Intille**, S. and Bobick, A. (1999). A framework for recognizing multi-agent action from visual evidence. In *AAAI-99*.
- Ioffe**, S. and Szegedy, C. (2015). Batch normalization: Accelerating deep network training by reducing internal covariate shift. arXiv:1502.03167.
- Irpan**, A. (2018). Deep reinforcement learning doesn't work yet. [www.alexirpan.com/2018/02/14/il-hard.html](http://www.alexirpan.com/2018/02/14/il-hard.html).
- Isard**, M. and Blake, A. (1996). Contour tracking by stochastic propagation of conditional density. In *ECCV-96*.
- Isola**, P., Zhu, J.-Y., Zhou, T., and Efros, A. (2017). Image-to-image translation with conditional adversarial networks. In *CVPR-17*.
- Jaakkola**, T. and Jordan, M. I. (1996). Computing upper and lower bounds on likelihoods in intractable networks. In *UAI-96*.
- Jacobson**, D. H. and Mayne, D. Q. (1970). *Differential Dynamic Programming*. North-Holland.
- Jaderberg**, M., Czarnecki, W. M., Dunning, I., Marrs, L., Lever, G., Castaneda, A. G., Beattie, C., Rabinowitz, N. C., Morcos, A. S., Ruderman, A., et al. (2019). Human-level performance in 3D multiplayer games with population-based reinforcement learning. *Science*, 364, 859–865.
- Jaderberg**, M., Dalibard, V., Osindero, S., Czarnecki, W. M., Donahue, J., Razavi, A., Vinyals, O., Green, T., Dunning, I., Simonyan, K., Fernando, C., and Kavukcuoglu, K. (2017). Population based training of neural networks. arXiv:1711.09846.
- Jaffar**, J. and Lassez, J.-L. (1987). Constraint logic programming. In *Proc. Fourteenth ACM POPL Conference*. Association for Computing Machinery.
- Jaffar**, J., Michaylov, S., Stuckey, P. J., and Yap, R. H. C. (1992). The CLP(R) language and system. *ACM Transactions on Programming Languages and Systems*, 14, 339–395.
- Jain**, D., Barthels, A., and Beetz, M. (2010). Adaptive Markov logic networks: Learning statistical relational models with dynamic parameters. In *ECAI-10*.
- Jain**, D., Kirchlechner, B., and Beetz, M. (2007). Extending Markov logic to model probability distributions in relational domains. In *30th Annual German Conference on AI (KI)*.
- James**, G., Witten, D., Hastie, T., and Tibshirani, R. (2013). *An Introduction to Statistical Learning with Applications in R*. Springer-Verlag.
- Jarrett**, K., Kavukcuoglu, K., Ranzato, M., and LeCun, Y. (2009). What is the best multi-stage architecture for object recognition? In *ICCV-09*.
- Jaynes**, E. T. (2003). *Probability Theory: The Logic of Science*. Cambridge Univ. Press.
- Jeffrey**, R. C. (1983). *The Logic of Decision* (2nd edition). University of Chicago Press.
- Jeffreys**, H. (1948). *Theory of Probability*. Oxford.
- Jelinek**, F. (1976). Continuous speech recognition by statistical methods. *Proc. IEEE*, 64, 532–556.
- Jelinek**, F. and Mercer, R. L. (1980). Interpolated estimation of Markov source parameters from sparse data. In *Proc. Workshop on Pattern Recognition in Practice*.
- Jennings**, H. S. (1906). *Behavior of the Lower Organisms*. Columbia University Press.
- Jenniskens**, P., Betlem, H., Betlem, J., and Barifaijo, E. (1994). The Mbale meteorite shower. *Meteoritics*, 29, 246–254.
- Jensen**, F. V. (2007). *Bayesian Networks and Decision Graphs*. Springer-Verlag.
- Ji**, Z., Lipton, Z. C., and Elkan, C. (2014). Differential program and machine learning: A survey and review. arXiv:1412.7584.
- Jiang**, H. and Nachum, O. (2019). Identifying and correcting label bias in machine learning. arXiv:1901.04966.
- Jimenez**, P. and Torras, C. (2000). An efficient algorithm for searching implicit AND/OR graphs with cycles. *AII*, 124, 1–30.
- Joachims**, T. (2001). A statistical learning model of text classification with support vector machines. In *SIGIR-01*.
- Johnson**, M. (1998). PCFG models of linguistic tree representations. *Comput. Linguist.*, 24, 613–632.
- Johnson**, W. W. and Story, W. E. (1879). Notes on the “15” puzzle. *American Journal of Mathematics*, 2, 397–404.
- Johnston**, M. D. and Adorf, H.-M. (1992). Scheduling with neural networks: The case of the Hubble space telescope. *Computers and Operations Research*, 19, 209–240.
- Jonathan**, P. J. Y., Fung, C. C., and Wong, K. W. (2009). Devious chatbots-interactive malware with a plot. In *FIRA RoboWorld Congress*.
- Jones**, M. and Love, B. C. (2011). Bayesian fundamentalism or enlightenment? On the explanatory status and theoretical contributions of Bayesian models of cognition. *BBS*, 34, 169–231.
- Jones**, R. M., Laird, J., and Nielsen, P. E. (1998). Automated intelligent pilots for combat flight simulation. In *AAAI-98*.
- Jones**, R., McCallum, A., Nigam, K., and Riloff, E. (1999). Bootstrapping for text learning tasks. In *Proc. IJCAI-99 Workshop on Text Mining: Foundations, Techniques, and Applications*.
- Jones**, T. (2007). *Artificial Intelligence: A Systems Approach*. Infinity Science Press.
- Jonsson**, A., Morris, P., Muscettola, N., Rajan, K., and Smith, B. (2000). Planning in interplanetary space: Theory and practice. In *AIPS-00*.
- Jordan**, M. I. (2005). Dirichlet processes, Chinese restaurant processes and all that. Tutorial presentation at the NeurIPS Conference.
- Jordan**, M. I. (1986). Serial order: A parallel distributed processing approach. Tech. rep., UCSD Institute for Cognitive Science.
- Jordan**, M. I., Ghahramani, Z., Jaakkola, T., and Saul, L. K. (1999). An introduction to variational methods for graphical models. *Machine Learning*, 37, 183–233.
- Jouannaud**, J.-P. and Kirchner, C. (1991). Solving equations in abstract algebras: A rule-based survey of unification. In Lassez, J.-L. and Plotkin, G. (Eds.), *Computational Logic*. MIT Press.
- Joulin**, A., Grave, E., Bojanowski, P., and Mikolov, T. (2016). Bag of tricks for efficient text classification. arXiv:1607.01759.
- Jouppi**, N. P., Young, C., Patil, N., Patterson, D. A., et al. (2017). In-datacenter performance analysis of a tensor processing unit. In *ACM/IEEE 44th International Symposium on Computer Architecture*.
- Joy**, B. (2000). Why the future doesn't need us. *Wired*, 8.
- Jozefowicz**, R., Vinyals, O., Schuster, M., Shazeer, N., and Wu, Y. (2016). Exploring the limits of language modeling. arXiv:1602.02410.
- Jozefowicz**, R., Zaremba, W., and Sutskever, I. (2015). An empirical exploration of recurrent network architectures. In *ICML-15*.
- Juels**, A. and Wattenberg, M. (1996). Stochastic hill-climbing as a baseline method for evaluating genetic algorithms. In *NeurIPS 8*.
- Julesz**, B. (1971). *Foundations of Cyclopean Perception*. University of Chicago Press.
- Julian**, K. D., Kochenderfer, M. J., and Owen, M. P. (2018). Deep neural network compression for aircraft collision avoidance systems. arXiv:1810.04240.
- Juliani**, A., Berges, V., Vckay, E., Gao, Y., Henry, H., Mattar, M., and Lange, D. (2018). Unity: A general platform for intelligent agents. arXiv:1809.02627.
- Junker**, U. (2003). The logic of ilog (j)configurator: Combining constraint programming with a description logic. In *Proc. IJCAI-03 Configuration Workshop*.
- Jurafsky**, D. and Martin, J. H. (2020). *Speech and Language Processing: An Introduction to Natural Language Processing, Computational Linguistics, and Speech Recognition* (3rd edition). Prentice-Hall.
- Kadane**, J. B. and Simon, H. A. (1977). Optimal strategies for a class of constrained sequential problems. *Annals of Statistics*, 5, 237–255.
- Kadane**, J. B. and Larkey, P. D. (1982). Subjective probability and the theory of games. *Management Science*, 28, 113–120.
- Kaelbling**, L. P., Littman, M. L., and Cassandra, A. R. (1998). Planning and acting in partially observable stochastic domains. *AII*, 101, 99–134.
- Kaelbling**, L. P. and Rosenschein, S. J. (1990). Action and planning in embedded agents. *Robotics and Autonomous Systems*, 6, 35–48.
- Kager**, R. (1999). *Optimality Theory*. Cambridge University Press.
- Kahn**, H. and Marshall, A. W. (1953). Methods of reducing sample size in Monte Carlo computations. *Operations Research*, 1, 263–278.
- Kahn**, H. (1950a). Random sampling (Monte Carlo) techniques in neutron attenuation problems—I. *Nucleonics*, 6, 27–passim.
- Kahn**, H. (1950b). Random sampling (Monte Carlo) techniques in neutron attenuation problems—II. *Nucleonics*, 6, 60–65.
- Kahneman**, D. (2011). *Thinking, Fast and Slow*. Farar, Straus and Giroux.
- Kahneman**, D., Slovic, P., and Tversky, A. (Eds.). (1982). *Judgment under Uncertainty: Heuristics and Biases*. Cambridge University Press.
- Kahneman**, D. and Tversky, A. (1979). Prospect theory: An analysis of decision under risk. *Econometrica*, 47, 263–291.
- Kaindl**, H. and Khorsand, A. (1994). Memory-bounded bidirectional search. In *AAAI-94*.

- Kalman**, R. (1960). A new approach to linear filtering and prediction problems. *J. Basic Engineering*, 82, 35–46.
- Kambhampati**, S. (1994). Exploiting causal structure to control retrieval and refitting during plan reuse. *Computational Intelligence*, 10, 213–244.
- Kanade**, T., Thorpe, C., and Whittaker, W. (1986). Autonomous land vehicle project at CMU. In *ACM Fourteenth Annual Conference on Computer Science*.
- Kanal**, E. (2017). Machine learning in cybersecurity. CMU SEI Blog, insights.sei.cmu.edu/sei\_blog/2017/06/machine-learning-in-cybersecurity.html.
- Kanazawa**, A., Black, M., Jacobs, D., and Malik, J. (2018a). End-to-end recovery of human shape and pose. In *CVPR-18*.
- Kanazawa**, A., Tulsiani, M., Efros, A., and Malik, J. (2018b). Learning category-specific mesh reconstruction from image collections. In *ECCV-18*.
- Kanazawa**, K., Koller, D., and Russell, S. J. (1995). Stochastic simulation algorithms for dynamic probabilistic networks. In *UAI-95*.
- Kang**, S. M. and Wildes, R. P. (2016). Review of action recognition and detection methods. arXiv:1610.06906.
- Kanter**, J. M. and Veeramachaneni, K. (2015). Deep feature synthesis: Towards automating data science endeavors. In *Proc. IEEE Int'l Conf. on Data Science and Advanced Analytics*.
- Kantorovich**, L. V. (1939). Mathematical methods of organizing and planning production. Published in translation in *Management Science*, 6(4), 366–422, 1960.
- Kaplan**, D. and Montague, R. (1960). A paradox regained. *Notre Dame Formal Logic*, 1, 79–90.
- Karaboga**, D. and Basturk, B. (2007). A powerful and efficient algorithm for numerical function optimization: Artificial bee colony (ABC) algorithm. *Journal of global optimization*, 39, 459–471.
- Karamchandani**, A., Bjerager, P., and Cornell, C. A. (1989). Adaptive importance sampling. In *Proc. Fifth International Conference on Structural Safety and Reliability*.
- Karmarkar**, N. (1984). A new polynomial-time algorithm for linear programming. *Combinatorica*, 4, 373–395.
- Karp**, R. M. (1972). Reducibility among combinatorial problems. In Miller, R. E. and Thatcher, J. W. (Eds.), *Complexity of Computer Computations*. Plenum.
- Karpathy**, A. (2015). The unreasonable effectiveness of recurrent neural networks. Andrej Karpathy blog, karpathy.github.io/2015/05/21/rnn-effectiveness/.
- Karpathy**, A. and Fei-Fei, L. (2015). Deep visual-semantic alignments for generating image descriptions. In *CVPR-15*.
- Karras**, T., Aila, T., Laine, S., and Lehtinen, J. (2017). Progressive growing of GANs for improved quality, stability, and variation. arXiv:1710.10196.
- Karsch**, K., Hedau, V., Forsyth, D., and Hoiem, D. (2011). Rendering synthetic objects into legacy photographs. In *SIGGRAPH Asia*.
- Kartam**, N. A. and Levitt, R. E. (1990). A constraint-based approach to construction planning of multi-story buildings. In *Expert Planning Systems*. Institute of Electrical Engineers.
- Kasami**, T. (1965). An efficient recognition and syntax analysis algorithm for context-free languages. Tech. rep., Air Force Cambridge Research Laboratory.
- Katehakis**, M. N. and Veinott, A. F. (1987). The multi-armed bandit problem: Decomposition and computation. *Mathematics of Operations Research*, 12, 185–376.
- Katz**, B. (1997). Annotating the world wide web using natural language. In *RIAO '97*.
- Kaufmann**, M., Manolios, P., and Moore, J. S. (2000). *Computer-Aided Reasoning: An Approach*. Kluwer.
- Kautz**, H. (2006). Deconstructing planning as satisfiability. In *AAAI-06*.
- Kautz**, H., McAllester, D. A., and Selman, B. (1996). Encoding plans in propositional logic. In *KR-96*.
- Kautz**, H. and Selman, B. (1992). Planning as satisfiability. In *ECAI-92*.
- Kautz**, H. and Selman, B. (1998). BLACKBOX: A new approach to the application of theorem proving to problem solving. Working Notes of the AIPS-98 Workshop on Planning as Combinatorial Search.
- Kavraki**, L., Svestka, P., Latombe, J.-C., and Overmars, M. (1996). Probabilistic roadmaps for path planning in high-dimensional configuration spaces. *IEEE Transactions on Robotics and Automation*, 12, 566–580.
- Kazemi**, S. M., Kimmig, A., Van den Broeck, G., and Poole, D. (2017). New liftable classes for first-order probabilistic inference. In *NeurIPS 29*.
- Kearns**, M. (1990). *The Computational Complexity of Machine Learning*. MIT Press.
- Kearns**, M., Mansour, Y., and Ng, A. Y. (2000). Approximate planning in large POMDPs via reusable trajectories. In *NeurIPS 12*.
- Kearns**, M. and Singh, S. (1998). Near-optimal reinforcement learning in polynomial time. In *ICML-98*.
- Kearns**, M. and Vazirani, U. (1994). *An Introduction to Computational Learning Theory*. MIT Press.
- Kearns**, M. (1988). Thoughts on hypothesis boosting.
- Kearns**, M., Mansour, Y., and Ng, A. Y. (2002). A sparse sampling algorithm for near-optimal planning in large Markov decision processes. *Machine Learning*, 49, 193–208.
- Kebeasy**, R. M., Hussein, A. I., and Dahy, S. A. (1998). Discrimination between natural earthquakes and nuclear explosions using the Aswan Seismic Network. *Annali di Geofisica*, 41, 127–140.
- Keeney**, R. L. (1974). Multiplicative utility functions. *Operations Research*, 22, 22–34.
- Keeney**, R. L. and Raiffa, H. (1976). *Decisions with Multiple Objectives: Preferences and Value Tradeoffs*. Wiley.
- Kelley**, H. J. (1960). Gradient theory of optimal flight paths. *ARS Journal*, 30, 947–954.
- Kemp**, M. (Ed.). (1989). *Leonardo on Painting: An Anthology of Writings*. Yale University Press.
- Kempe**, A. B. (1879). On the geographical problem of the four-colors. *American Journal of Mathematics*, 2, 193–200.
- Kephart**, J. O. and Chess, D. M. (2003). The vision of autonomic computing. *IEEE Computer*, 36, 41–50.
- Kersting**, K., Raedt, L. D., and Kramer, S. (2000). Interpreting Bayesian logic programs. In *Proc. AAAI-00 Workshop on Learning Statistical Models from Relational Data*.
- Keskar**, N. S., McCann, B., Varshney, L., Xiong, C., and Socher, R. (2019). CTRL: A conditional transformer language model for controllable generation. arXiv:1909.
- Keynes**, J. M. (1921). *A Treatise on Probability*. Macmillan.
- Khare**, R. (2006). Microformats: The next (small) thing on the semantic web. *IEEE Internet Computing*, 10, 68–75.
- Khatib**, O. (1986). Real-time obstacle avoidance for robot manipulator and mobile robots. *Int. J. Robotics Research*, 5, 90–98.
- Kim**, B., Khanna, R., and Koyejo, O. O. (2017). Examples are not enough, learn to criticize! Criticism for interpretability. In *NeurIPS 29*.
- Kim**, J. H. (1983). *CONVINCE: A Conversational Inference Consolidation Engine*. Ph.D. thesis, Department of Computer Science, UCLA.
- Kim**, J. H. and Pearl, J. (1983). A computational model for combined causal and diagnostic reasoning in inference systems. In *IJCAI-83*.
- Kim**, J.-H., Lee, C.-H., Lee, K.-H., and Kuppuswamy, N. (2007). Evolving personality of a genetic robot in ubiquitous environment. In *Proc. 16th IEEE International Symposium on Robot and Human Interactive Communication*.
- Kim**, T. W. (2018). Explainable artificial intelligence (XAI), the goodness criteria and the grasp-ability test. arXiv:1810.09598.
- Kingma**, D. P. and Welling, M. (2013). Auto-encoding variational Bayes. arXiv:1312.6114.
- Kirk**, D. E. (2004). *Optimal Control Theory: An Introduction*. Dover.
- Kirkpatrick**, S., Gelatt, C. D., and Vecchi, M. P. (1983). Optimization by simulated annealing. *Science*, 220, 671–680.
- Kisynski**, J. and Poole, D. (2009). Lifted aggregation in directed first-order probabilistic models. In *IJCAI-09*.
- Kitaev**, N., Kaiser, L., and Levskaya, A. (2020). Reformer: The efficient transformer. arXiv:2001.04451.
- Kitaev**, N. and Klein, D. (2018). Constituency parsing with a self-attentive encoder. arXiv:1805.01052.
- Kitani**, K. M., abd James Andrew Bagnell, B. D. Z., and Hebert, M. (2012). Activity forecasting. In *ECCV-12*.
- Kitano**, H., Asada, M., Kuniyoshi, Y., Noda, I., and Osawa, E. (1997). RoboCup: The robot world cup initiative. In *Proc. First International Conference on Autonomous Agents*.
- Kjaerulff**, U. (1992). A computational scheme for reasoning in dynamic probabilistic networks. In *UAI-92*.
- Klarman**, H. E., Francis, J., and Rosenthal, G. D. (1968). Cost effectiveness analysis applied to the treatment of chronic renal disease. *Medical Care*, 6, 48–54.
- Klein**, D. and Manning, C. (2001). Parsing with treebank grammars: Empirical bounds, theoretical models, and the structure of the Penn treebank. In *ACL-01*.
- Klein**, D. and Manning, C. (2003). A\* parsing: Fast exact Viterbi parse selection. In *HLT-NAACL-03*.
- Kleinberg**, J. M., Mullaithathan, S., and Raghavan, M. (2016). Inherent trade-offs in the fair determination of risk scores. arXiv:1609.05807.
- Klemperer**, P. (2002). What really matters in auction design. *J. Economic Perspectives*, 16.
- Kneser**, R. and Ney, H. (1995). Improved backoff for M-gram language modeling. In *ICASSP-95*.
- Knoblock**, C. A. (1991). Search reduction in hierarchical problem solving. In *AAAI-91*.
- Knuth**, D. E. (1964). Representing numbers using only one 4. *Mathematics Magazine*, 37, 308–310.
- Knuth**, D. E. (1975). An analysis of alpha–beta pruning. *AIJ*, 6, 293–326.
- Knuth**, D. E. (2015). *The Art of Computer Programming*, Vol. 4, Fascicle 6: Satisfiability. Addison-Wesley.
- Knuth**, D. E. and Bendix, P. B. (1970). Simple word problems in universal algebras. In Leech, J. (Ed.), *Computational Problems in Abstract Algebra*. Pergamon.
- Kober**, J., Bagnell, J. A., and Peters, J. (2013). Reinforcement learning in robotics: A survey. *International Journal of Robotics Research*, 32, 1238–1274.
- Koch**, C. (2019). *The Feeling of Life Itself*. MIT Press.

- Kochenderfer, M. J.** (2015). *Decision Making Under Uncertainty: Theory and Application*. MIT Press.
- Kocsis, L.** and Szepesvari, C. (2006). Bandit-based Monte-Carlo planning. In *ECML-06*.
- Koditschek, D.** (1987). Exact robot navigation by means of potential functions: Some topological considerations. In *ICRA-87*.
- Koehn, P.** (2009). *Statistical Machine Translation*. Cambridge University Press.
- Koelsch, S.** and Siebel, W. A. (2005). Towards a neural basis of music perception. *Trends in Cognitive Sciences*, 9, 578–584.
- Koenderink, J. J.** (1990). *Solid Shape*. MIT Press.
- Koenderink, J. J.** and van Doorn, A. J. (1991). Affine structure from motion. *J. Optical Society of America A*, 8, 377–385.
- Koenig, S.** (1991). Optimal probabilistic and decision-theoretic planning using Markovian decision theory. Master's report, Computer Science Division, University of California, Berkeley.
- Koenig, S.** (2000). Exploring unknown environments with real-time search or reinforcement learning. In *NeurIPS-12*.
- Koenig, S.** (2001). Agent-centered search. *AI Mag*, 22, 109–131.
- Koenig, S.** and Likhachev, M. (2002). D\* Lite. *AAAI-15, 15*.
- Koenig, S.**, Likhachev, M., and Furcy, D. (2004). Lifelong planning A\*. *AIJ*, 155, 93–146.
- Kolesky, D. B.**, Truby, R. L., Gladman, A. S., Busbee, T. A., Homan, K. A., and Lewis, J. A. (2014). 3D bioprinting of vascularized, heterogeneous cell-laden tissue constructs. *Advanced Materials*, 26, 3124–3130.
- Koller, D.**, Meggido, N., and von Stengel, B. (1996). Efficient computation of equilibria for extensive two-person games. *Games and Economic Behaviour*, 14, 247–259.
- Koller, D.** and Pfeffer, A. (1997). Representations and solutions for game-theoretic problems. *AIJ*, 94, 167–215.
- Koller, D.** and Pfeffer, A. (1998). Probabilistic frame-based systems. In *AAAI-98*.
- Koller, D.** and Friedman, N. (2009). *Probabilistic Graphical Models: Principles and Techniques*. MIT Press.
- Koller, D.**, McAllester, D. A., and Pfeffer, A. (1997). Effective Bayesian inference for stochastic programs. In *AAAI-97*.
- Koller, D.** and Parr, R. (2000). Policy iteration for factored MDPs. In *UAI-00*.
- Koller, D.** and Sahami, M. (1997). Hierarchically classifying documents using very few words. In *ICML-97*.
- Kolmogorov, A. N.** (1941). Interpolation und extrapolation von stationären zufälligen folgen. *Bulletin of the Academy of Sciences of the USSR, Ser. Math.* 5, 3–14.
- Kolmogorov, A. N.** (1950). *Foundations of the Theory of Probability*. Chelsea.
- Kolmogorov, A. N.** (1963). On tables of random numbers. *Sankhya, the Indian Journal of Statistics: Series A*, 25(4), 369–376.
- Kolmogorov, A. N.** (1965). Three approaches to the quantitative definition of information. *Problems in Information Transmission*, 1, 1–7.
- Kolter, J. Z.**, Abbeel, P., and Ng, A. Y. (2008). Hierarchical apprenticeship learning, with application to quadruped locomotion. In *NeurIPS-20*.
- Kondrak, G.** and van Beek, P. (1997). A theoretical evaluation of selected backtracking algorithms. *AIJ*, 89, 365–387.
- Konečný, J.**, McMahan, H. B., Yu, F. X., Richtárik, P., Suresh, A. T., and Bacon, D. (2016). Federated learning: Strategies for improving communication efficiency. arXiv:1610.05492.
- Konolige, K.** (1997). COLBERT: A language for reactive control in Saphira. In *Künstliche Intelligenz: Advances in Artificial Intelligence*, LNAI.
- Konolige, K.** (2004). Large-scale map-making. In *AAAI-04*.
- Konolige, K.** (1982). A first order formalization of knowledge and action for a multi-agent planning system. In Hayes, J. E., Michie, D., and Pao, Y.-H. (Eds.), *Machine Intelligence 10*. Ellis Horwood.
- Konolige, K.** (1994). Easy to be hard: Difficult problems for greedy algorithms. In *KR-94*.
- Koopmans, T. C.** (1972). Representation of preference orderings over time. In McGuire, C. B. and Radner, R. (Eds.), *Decision and Organization*. Elsevier.
- Korb, K. B.** and Nicholson, A. (2010). *Bayesian Artificial Intelligence*. CRC Press.
- Korf, R. E.** (1985a). Depth-first iterative-deepening: an optimal admissible tree search. *AIJ*, 27, 97–109.
- Korf, R. E.** (1985b). Iterative-deepening A\*: An optimal admissible tree search. In *ICAI-85*.
- Korf, R. E.** (1987). Planning as search: A quantitative approach. *AIJ*, 33, 65–88.
- Korf, R. E.** (1990). Real-time heuristic search. *AIJ*, 42, 189–212.
- Korf, R. E.** (1993). Linear-space best-first search. *AIJ*, 62, 41–78.
- Korf, R. E.** and Chickering, D. M. (1996). Best-first minimax search. *AIJ*, 84, 299–337.
- Korf, R. E.** and Felner, A. (2002). Disjoint pattern database heuristics. *AIJ*, 134, 9–22.
- Korf, R. E.** and Zhang, W. (2000). Divide-and-conquer frontier search applied to optimal sequence alignment. In *AAAI-00*.
- Korf, R. E.** (1997). Finding optimal solutions to Rubik's Cube using pattern databases. In *AAAI-97*.
- Korf, R. E.** and Reid, M. (1998). Complexity analysis of admissible heuristic search. In *AAAI-98*.
- Koutsoupias, E.** and Papadimitriou, C. H. (1992). On the greedy algorithm for satisfiability. *Information Processing Letters*, 43, 53–55.
- Kovacs, D. L.** (2011). BNF definition of PDDL3.1. Unpublished manuscript from the IPC-2011 website.
- Kowalski, R.** (1974). Predicate logic as a programming language. In *Proc. IFIP Congress*.
- Kowalski, R.** (1979). *Logic for Problem Solving*. Elsevier.
- Kowalski, R.** (1988). The early years of logic programming. *CACM*, 31, 38–43.
- Kowalski, R.** and Sergot, M. (1986). A logic-based calculus of events. *New Generation Computing*, 4, 67–95.
- Koza, J. R.** (1992). *Genetic Programming: On the Programming of Computers by Means of Natural Selection*. MIT Press.
- Koza, J. R.** (1994). *Genetic Programming II: Automatic Discovery of Reusable Programs*. MIT Press.
- Koza, J. R.**, Bennett, F. H., Andre, D., and Keane, M. A. (1999). *Genetic Programming III: Darwinian Invention and Problem Solving*. Morgan Kaufmann.
- Krakovna, V.** (2018). Specification gaming examples in AI.
- Kraska, T.**, Beutel, A., Chi, E. H., Dean, J., and Polyzotis, N. (2017). The case for learned index structures. arXiv:1712.01208.
- Kraus, S.** (2001). *Strategic Negotiation in Multiagent Environments*. MIT Press.
- Kraus, S.**, Ephrat, E., and Lehmann, D. (1991). Negotiation in a non-cooperative environment. *AIJ*, 3, 255–281.
- Krause, A.** and Guestrin, C. (2005). Optimal nonmyopic value of information in graphical models: Efficient algorithms and theoretical limits. In *IJCAI-05*.
- Krause, A.** and Guestrin, C. (2009). Optimal value of information in graphical models. *JAIR*, 35, 557–591.
- Krause, A.**, McMahan, B., Guestrin, C., and Gupta, A. (2008). Robust submodular observation selection. *JMLR*, 9, 2761–2801.
- Kripke, S. A.** (1963). Semantical considerations on modal logic. *Acta Philosophica Fennica*, 16, 83–94.
- Krishna, V.** (2002). *Auction Theory*. Academic Press.
- Krishnamurthy, V.** (2016). *Partially Observed Markov Decision Processes: From Filtering to Controlled Sensing*. Cambridge University Press.
- Krishnanand, K.** and Ghose, D. (2009). Glowworm swarm optimisation: A new method for optimising multi-modal functions. *International Journal of Computational Intelligence Studies*, 1, 93–119.
- Krizhevsky, A.**, Sutskever, I., and Hinton, G. E. (2013). ImageNet classification with deep convolutional neural networks. In *NeurIPS-25*.
- Krogh, A.**, Brown, M., Mian, I. S., Sjolander, K., and Haussler, D. (1994). Hidden Markov models in computational biology: Applications to protein modeling. *J. Molecular Biology*, 235, 1501–1531.
- Krogh, A.** and Hertz, J. A. (1992). A simple weight decay can improve generalization. In *NeurIPS-4*.
- Kruppa, E.** (1913). Zur Ermittlung eines Objektes aus zwei Perspektiven mit innerer Orientierung. *Sitz.-Ber. Akad. Wiss., Wien, Math. Naturw., Kl. Abt. IIa*, 122, 1939–1948.
- Kübler, S.**, McDonald, R., and Nivre, J. (2009). *Dependency Parsing*. Morgan & Claypool.
- Kuffner, J. J.** and LaValle, S. (2000). RRT-connect: An efficient approach to single-query path planning. In *ICRA-00*.
- Kuhn, H. W.** (1953). Extensive games and the problem of information. In Kuhn, H. W. and Tucker, A. W. (Eds.), *Contributions to the Theory of Games II*. Princeton University Press.
- Kuhn, H. W.** (1955). The Hungarian method for the assignment problem. *Naval Research Logistics Quarterly*, 2, 83–97.
- Kuipers, B. J.** (1985). Qualitative simulation. In Bobrow, D. (Ed.), *Qualitative Reasoning About Physical Systems*. MIT Press.
- Kuipers, B. J.** and Levitt, T. S. (1988). Navigation and mapping in large-scale space. *AI Mag*, 9, 25–43.
- Kuipers, B. J.** (2001). Qualitative simulation. In Meyers, R. A. (Ed.), *Encyclopedia of Physical Science and Technology*. Academic Press.
- Kulkarni, T.**, Kohli, P., Tenenbaum, J. B., and Manning, K. V. (2015). Picture: A probabilistic programming language for scene perception. In *CVPR-15*.
- Kumar, P. R.** and Varaiya, P. (1986). *Stochastic Systems: Estimation, Identification, and Adaptive Control*. Prentice-Hall.
- Kumar, S.** (2017). A survey of deep learning methods for relation extraction. arXiv:1705.03645.
- Kumar, V.** and Kanal, L. N. (1988). The CDP: A unifying formulation for heuristic search, dynamic programming, and branch-and-bound. In Kanal, L. N. and Kumar, V. (Eds.), *Search in Artificial Intelligence*. Springer-Verlag.
- Kurien, J.**, Nayak, P., and Smith, D. E. (2002). Fragment-based conformant planning. In *AIPS-02*.

- Kurth**, T., Treichler, S., Romero, J., Mudigonda, M., Luehr, N., Phillips, E. H., Mahesh, A., Matheson, M., Deslippe, J., Fatica, M., Prabhat, and Houston, M. (2018). Exascale deep learning for climate analytics. arXiv:1810.01993.
- Kurzweil**, R. (2005). *The Singularity is Near*. Viking.
- Kwok**, C., Etzioni, O., and Weld, D. S. (2001). Scaling question answering to the web. In *Proc. 10th International Conference on the World Wide Web*.
- La Mettrie**, J. O. (1748). *L'homme machine*. E. Luzac, Leyde, France.
- La Mura**, P. and Shoham, Y. (1999). Expected utility networks. In *UAI-99*.
- Laborie**, P. (2003). Algorithms for propagating resource constraints in AI planning and scheduling. *AIJ*, 143, 151–188.
- Ladkin**, P. (1986a). Primitives and units for time specification. In *AAAI-86*.
- Ladkin**, P. (1986b). Time representation: a taxonomy of interval relations. In *AAAI-86*.
- Lafferty**, J., McCallum, A., and Pereira, F. (2001). Conditional random fields: Probabilistic models for segmenting and labeling sequence data. In *ICML-01*.
- Lai**, T. L. and Robbins, H. (1985). Asymptotically efficient adaptive allocation rules. *Advances in Applied Mathematics*, 6, 4–22.
- Laird**, J., Newell, A., and Rosenbloom, P. S. (1987). SOAR: An architecture for general intelligence. *AIJ*, 33, 1–64.
- Laird**, J., Rosenbloom, P. S., and Newell, A. (1986). Chunking in Soar: The anatomy of a general learning mechanism. *Machine Learning*, 1, 11–46.
- Laird**, J. (2008). Extending the Soar cognitive architecture. In *Artificial General Intelligence Conference*.
- Lake**, B., Salakhutdinov, R., and Tenenbaum, J. B. (2015). Human-level concept learning through probabilistic program induction. *Science*, 350, 1332–1338.
- Lakoff**, G. (1987). *Women, Fire, and Dangerous Things: What Categories Reveal About the Mind*. University of Chicago Press.
- Lakoff**, G. and Johnson, M. (1980). *Metaphors We Live By*. University of Chicago Press.
- Lakoff**, G. and Johnson, M. (1999). *Philosophy in the Flesh : The Embodied Mind and Its Challenge to Western Thought*. Basic Books.
- Lam**, J. and Greenspan, M. (2008). Eye-in-hand visual servoing for accurate shooting in pool robotics. In *5th Canadian Conference on Computer and Robot Vision*.
- Lamarek**, J. B. (1809). *Philosophie zoologique*. Chez Dentu et L'Auteur, Paris.
- Lample**, G. and Conneau, A. (2019). Cross-lingual language model pretraining. arXiv:1901.07291.
- Landhuis**, E. (2004). Lifelong debunker takes on arbiter of neural choices: Magician-turned-mathematician uncovers bias in a flip of a coin. *Stanford Report*, June 7.
- Langdon**, W. and Poli, R. (2002). *Foundations of Genetic Programming*. Springer.
- Langton**, C. (Ed.). (1995). *Artificial Life*. MIT Press.
- LaPaugh**, A. S. (2010). Algorithms and theory of computation handbook. In Atallah, M. J. and Blanton, M. (Eds.), *VLSI Layout Algorithms*. Chapman & Hall/CRC.
- Laplace**, P. (1816). *Essai philosophique sur les probabilités* (3rd edition). Courcier Imprimeur, Paris.
- Larochelle**, H. and Murray, I. (2011). The neural autoregressive distribution estimator. In *AISTATS-11*.
- Larson**, S. C. (1931). The shrinkage of the coefficient of multiple correlation. *J. Educational Psychology*, 22, 45–55.
- Laskey**, K. B. (1995). Sensitivity analysis for probability assessments in Bayesian networks. *IEEE Transactions on Systems, Man and Cybernetics*, 25, 901–909.
- Laskey**, K. B. (2008). MEBN: A language for first-order Bayesian knowledge bases. *AIJ*, 172, 140–178.
- Latombe**, J.-C. (1991). *Robot Motion Planning*. Kluwer.
- Lauritzen**, S. (1995). The EM algorithm for graphical association models with missing data. *Computational Statistics and Data Analysis*, 19, 191–201.
- Lauritzen**, S., Dawid, A. P., Larsen, B., and Leimer, H. (1990). Independence properties of directed Markov fields. *Networks*, 20, 491–505.
- Lauritzen**, S. and Spiegelhalter, D. J. (1988). Local computations with probabilities on graphical structures and their application to expert systems. *J. Royal Statistical Society, B* 50, 157–224.
- Lauritzen**, S. and Wermuth, N. (1989). Graphical models for associations between variables, some of which are qualitative and some quantitative. *Annals of Statistics*, 17, 31–57.
- LaValle**, S. (2006). *Planning Algorithms*. Cambridge University Press.
- Lawler**, E. L., Lenstra, J. K., Kan, A., and Shmoys, D. B. (1992). *The Travelling Salesman Problem*. Wiley Interscience.
- Lawler**, E. L., Lenstra, J. K., Kan, A., and Shmoys, D. B. (1993). Sequencing and scheduling: Algorithms and complexity. In Graves, S. C., Zipkin, P. H., and Kan, A. H. G. R. (Eds.), *Logistics of Production and Inventory: Handbooks in Operations Research and Management Science, Volume 4*. North-Holland.
- Lawler**, E. L. and Wood, D. E. (1966). Branch-and-bound methods: A survey. *Operations Research*, 14, 699–719.
- Lazanas**, A. and Latombe, J.-C. (1992). Landmark-based robot navigation. In *AAAI-92*.
- Le**, T. A., Baydin, A. G., and Wood, F. (2017). Inference compilation and universal probabilistic programming. In *AISTATS-17*.
- Lebedev**, M. A. and Nicolelis, M. A. (2006). Brain-machine interfaces: Past, present and future. *Trends in Neurosciences*, 29, 536–546.
- Lecoutre**, C. (2009). *Constraint Networks: Techniques and Algorithms*. Wiley-IEEE Press.
- LeCun**, Y., Denker, J., and Solla, S. (1990). Optimal brain damage. In *NeurIPS-2*.
- LeCun**, Y., Jackel, L., Bottou, L., Brunot, A., Cortes, C., Denker, J., Drucker, H., Guyon, I., Müller, U., Sackinger, E., Simard, P., and Vapnik, V. N. (1995). Comparison of learning algorithms for handwritten digit recognition. In *Int. Conference on Artificial Neural Networks*.
- LeCun**, Y., Bengio, Y., and Hinton, G. E. (2015). Deep learning. *Nature*, 521, 436–444.
- Lee**, D., Seo, H., and Jung, M. W. (2012). Neural basis of reinforcement learning and decision making. *Annual Review of Neuroscience*, 35, 287–308.
- Lee**, K.-F. (2018). *AI Superpowers: China, Silicon Valley, and the New World Order*. Houghton Mifflin.
- Leech**, G., Rayson, P., and Wilson, A. (2001). *Word Frequencies in Written and Spoken English: Based on the British National Corpus*. Longman.
- Legendre**, A. M. (1805). *Nouvelles méthodes pour la détermination des orbites des comètes*. Chez Firmin Didot, Paris.
- Lehmann**, J., Isele, R., Jakob, M., Jentzsch, A., Kontokostas, D., Mendes, P. N., Hellmann, S., Morse, M., van Kleef, P., Auer, S., and Bizer, C. (2015). DBpedia – A large-scale, multilingual knowledge base extracted from Wikipedia. *Semantic Web*, 6, 167–195.
- Lehrer**, J. (2009). *How We Decide*. Houghton Mifflin.
- Leike**, J., Martic, M., Krakovna, V., Ortega, P. A., Everett, T., Lefrancq, A., Orseau, L., and Legg, S. (2017). AI safety gridworlds. arXiv:1711.09883.
- Lelis**, L., Arfaei, S. J., Zilles, S., and Holte, R. C. (2012). Learning heuristic functions faster by using predicted solution costs. In *Proc. Fifth Annual Symposium on Combinatorial Search*.
- Lenat**, D. B. (1975). BEINGS: Knowledge as interacting experts. In *IJCAI-75*.
- Lenat**, D. B. and Guha, R. V. (1990). *Building Large Knowledge-Based Systems: Representation and Inference in the CYC Project*. Addison-Wesley.
- Leonard**, H. S. and Goodman, N. (1940). The calculus of individuals and its uses. *JSL*, 5, 45–55.
- Leonard**, J. and Durrant-Whyte, H. (1992). *Directed Sonar Sensing for Mobile Robot Navigation*. Kluwer.
- Lepage**, G. P. (1978). A new algorithm for adaptive multidimensional integration. *Journal of Computational Physics*, 27, 192–203.
- Lerner**, U. (2002). *Hybrid Bayesian Networks for Reasoning About Complex Systems*. Ph.D. thesis, Stanford University.
- Leśniewski**, S. (1916). Podstawy ogólnej teorii mnogości. Popławski.
- Lesser**, V. R. and Corkill, D. D. (1988). The distributed vehicle monitoring testbed: A tool for investigating distributed problem solving networks. In Engelmore, R. and Morgan, T. (Eds.), *Blackboard Systems*. Addison-Wesley.
- Letz**, R., Schumann, J., Bayerl, S., and Bibel, W. (1992). SETHEO: A high-performance theorem prover. *JAR*, 8, 183–212.
- Levesque**, H. J. and Brachman, R. J. (1987). Expressiveness and tractability in knowledge representation and reasoning. *Computational Intelligence*, 3, 78–93.
- Levin**, D. A., Peres, Y., and Wilmer, E. L. (2008). *Markov Chains and Mixing Times*. American Mathematical Society.
- Levine**, S., Finn, C., Darrell, T., and Abbeel, P. (2016). End-to-end training of deep visuomotor policies. *JMLR*, 17, 1334–1373.
- Levine**, S., Pastor, P., Krizhevsky, A., Ibarz, J., and Quillen, D. (2018). Learning hand-eye coordination for robotic grasping with deep learning and large-scale data collection. *International Journal of Robotics Research*, 37, 421–436.
- Levy**, D. (1989). The million pound bridge program. In Levy, D. and Beal, D. (Eds.), *Heuristic Programming in Artificial Intelligence*. Ellis Horwood.
- Levy**, D. (2008). *Love and Sex with Robots: The Evolution of Human—Robot Relationships*. Harper.
- Levy**, O. and Goldberg, Y. (2014). Linguistic regularities in sparse and explicit word representations. In *Proc. Eighteenth Conference on Computational Natural Language Learning*.
- Leyton-Brown**, K. and Shoham, Y. (2008). *Essentials of Game Theory: A Concise, Multidisciplinary Introduction*. Morgan & Claypool.
- Li**, C. M. and Anbulagan (1997). Heuristics based on unit propagation for satisfiability problems. In *IJCAI-97*.
- Li**, K. and Malik, J. (2018a). Implicit maximum likelihood estimation. arXiv:1809.09087.
- Li**, K. and Malik, J. (2018b). On the implicit assumptions of GANs. arXiv:1811.12402.

- Li, M., Vitányi, P., et al.** (2008). *An Introduction to Kolmogorov Complexity and Its Applications* (3rd edition). Springer-Verlag.
- Li, T.-M., Gharbi, M., Adams, A., Durand, F., and Ragan-Kelley, J.** (2018). Differentiable programming for image processing and deep learning in Halide. *ACM Transactions on Graphics*, 37, 139.
- Li, W. and Todorov, E.** (2004). Iterative linear quadratic regulator design for nonlinear biological movement systems. In *Proc. 1st International Conference on Informatics in Control, Automation and Robotics*.
- Li, X. and Yao, X.** (2012). Cooperatively coevolving particle swarms for large scale optimization. *IEEE Trans. Evolutionary Computation*, 16, 210–224.
- Li, Z., Li, P., Krishnan, A., and Liu, J.** (2011). Large-scale dynamic gene regulatory network inference combining differential equation models with local dynamic Bayesian network analysis. *Bioinformatics*, 27 19, 2686–91.
- Liang, P., Jordan, M. I., and Klein, D.** (2011). Learning dependency-based compositional semantics. arXiv:1109.6841.
- Liang, P. and Potts, C.** (2015). Bringing machine learning and compositional semantics together. *Annual Review of Linguistics*, 1, 355–376.
- Liberatore, P.** (1997). The complexity of the language *A*. *Electronic Transactions on Artificial Intelligence*, 1, 13–38.
- Lifschitz, V.** (2001). Answer set programming and plan generation. *AII*, 138, 39–54.
- Lighthill, J.** (1973). Artificial intelligence: A general survey. In Lighthill, J., Sutherland, N. S., Needham, R. M., Longuet-Higgins, H. C., and Michie, D. (Eds.), *Artificial Intelligence: A Paper Symposium*. Science Research Council of Great Britain.
- Lillicrap, T., Hunt, J. J., Pritzel, A., Heess, N., Erez, T., Tassa, Y., Silver, D., and Wierstra, D.** (2015). Continuous control with deep reinforcement learning. arXiv:1509.02971.
- Lin, S.** (1965). Computer solutions of the travelling salesman problem. *Bell Systems Technical Journal*, 44(10), 2245–2269.
- Lin, S. and Kernighan, B. W.** (1973). An effective heuristic algorithm for the travelling-salesman problem. *Operations Research*, 21, 498–516.
- Lindley, D. V.** (1956). On a measure of the information provided by an experiment. *Annals of Mathematical Statistics*, 27, 986–1005.
- Lindsay, R. K., Buchanan, B. G., Feigenbaum, E. A., and Lederberg, J.** (1980). *Applications of Artificial Intelligence for Organic Chemistry: The DENDRAL Project*. McGraw-Hill.
- Lindsten, F., Jordan, M. I., and Schön, T. B.** (2014). Particle Gibbs with ancestor sampling. *JMLR*, 15, 2145–2184.
- Littman, M. L.** (1994). Markov games as a framework for multi-agent reinforcement learning. In *ICML-94*.
- Littman, M. L., Cassandra, A. R., and Kaelbling, L. P.** (1995). Learning policies for partially observable environments: Scaling up. In *ICML-95*.
- Littman, M. L.** (2015). Reinforcement learning improves behaviour from evaluative feedback. *Nature*, 521, 445–451.
- Littman, M. L., Topeu, U., Fu, J., Isbell, C., Wen, M., and MacGlashan, J.** (2017). Environment-independent task specifications via GLTL. arXiv:1704.04341.
- Liu, B., Gemp, I., Ghavamzadeh, M., Liu, J., Mahadevan, S., and Petrik, M.** (2018). Proximal gradient temporal difference learning: Stable reinforcement learning with polynomial sample complexity. *JAIR*, 63, 461–494.
- Liu, H., Simonyan, K., Vinyals, O., Fernando, C., and Kavukcuoglu, K.** (2017). Hierarchical representations for efficient architecture search. arXiv:1711.00436.
- Liu, H., Simonyan, K., and Yang, Y.** (2019). DARTS: Differentiable architecture search. In *ICLR-19*.
- Liu, J. and Chen, R.** (1998). Sequential Monte Carlo methods for dynamic systems. *JASA*, 93, 1022–1031.
- Liu, J. and West, M.** (2001). Combined parameter and state estimation in simulation-based filtering. In Doucet, A., de Freitas, J. F. G., and Gordon, N. (Eds.), *Sequential Monte Carlo Methods in Practice*. Springer.
- Liu, L. T., Dean, S., Rolf, E., Simchowitz, M., and Hardt, M.** (2018a). Delayed impact of fair machine learning. arXiv:1803.04383.
- Liu, M.-Y., Breuel, T., and Kautz, J.** (2018b). Unsupervised image-to-image translation networks. In *NeurIPS 30*.
- Liu, X., Faes, L., Kale, A. U., Wagner, S. K., Fu, D. J., Bruynseels, A., Mahendiran, T., Moraes, G., Shamdas, M., Kern, C., Ledsam, J. R., Schmid, M., Balaskas, K., Topol, E., Bachmann, L. M., Keane, P. A., and Dennisson, A. K.** (2019a). A comparison of deep learning performance against health-care professionals in detecting diseases from medical imaging: A systematic review and meta-analysis. *The Lancet Digital Health*.
- Liu, Y., Ott, M., Goyal, N., Du, J., Joshi, M., Chen, D., Levy, O., Lewis, M., Zettlemoyer, L., and Stoyanov, V.** (2019b). RoBERTa: A robustly optimized BERT pre-training approach. arXiv:1907.11692.
- Liu, Y., Jain, A., Eng, C., Way, D. H., Lee, K., Bui, P., Kanada, K., de Oliveira Marinho, G., Gallegos, J., Gabriele, S., Gupta, V., Singh, N., Natarajan, V., Hofmann-Wellenhof, R., Corrado, G., Peng, L., Webster, D. R., Ai, D., Huang, S., Liu, Y., Dunn, R. C., and Coz, D.** (2019c). A deep learning system for differential diagnosis of skin diseases. arXiv:1909.
- Liu, Y., Gadepalli, K. K., Norouzi, M., Dahl, G., Kohlberger, T., Venugopalan, S., Boyko, A. S., Timofeev, A., Nelson, P. Q., Corrado, G., Hipp, J. D., Peng, L., and Stumpe, M. C.** (2017). Detecting cancer metastases on gigapixel pathology images. arXiv:1703.02442.
- Liu, Y., Kohlberger, T., Norouzi, M., Dahl, G., Smith, J. L., Mohtashamian, A., Olson, N., Peng, L., Hipp, J. D., and Stumpe, M. C.** (2018). Artificial intelligence-based breast cancer nodal metastasis detection: Insights into the black box for pathologists. *Archives of Pathology & Laboratory Medicine*, 143, 859–868.
- Livescu, K., Glass, J., and Bilmes, J.** (2003). Hidden feature modeling for speech recognition using dynamic Bayesian networks. In *EUROSPEECH-2003*.
- Lloyd, S.** (2000). Ultimate physical limits to computation. *Nature*, 406, 1047–1054.
- Lloyd, W. F.** (1833). *Two Lectures on the Checks to Population*. Oxford University.
- Lull, R.** (1305). *Ars Magna*. Published as Salzinger, I. et al. (Eds.), *Raymundi Lulli Opera omnia*, Mainz, 1721–1742.
- Loftus, E. and Palmer, J.** (1974). Reconstruction of automobile destruction: An example of the interaction between language and memory. *J. Verbal Learning and Verbal Behavior*, 13, 585–589.
- Lohn, J. D., Kraus, W. F., and Colombano, S. P.** (2001). Evolutionary optimization of yagi-uda antennas. In *Proc. Fourth International Conference on Evolvable Systems*.
- Longuet-Higgins, H. C.** (1981). A computer algorithm for reconstructing a scene from two projections. *Nature*, 293, 133–135.
- Loos, S., Irving, G., Szegedy, C., and Kaliszyk, C.** (2017). Deep network guided proof search. In *Proc. 21st Int'l Conf. on Logic for Programming, Artificial Intelligence and Reasoning*.
- Lopez de Segura, R.** (1561). *Libro de la invención liberal y arte del juego del axedrez*. Andres de Angulo.
- Lorentz, R.** (2015). Early playout termination in MCTS. In Plaat, A., van den Herik, J., and Kosters, W. (Eds.), *Advances in Computer Games*. Springer-Verlag.
- Love, N., Hinrichs, T., and Genesereth, M. R.** (2006). General game playing: Game description language specification. Tech. rep., Stanford University Computer Science Dept.
- Lovejoy, W. S.** (1991). A survey of algorithmic methods for partially observed Markov decision processes. *Annals of Operations Research*, 28, 47–66.
- Lovelace, A.** (1843). Sketch of the analytical engine invented by Charles Babbage. Notes appended to Lovelace's translation of an article of the above title written by L. F. Menabrea based on lectures by Charles Babbage in 1840. The translation appeared in R. Taylor (Ed.), *Scientific Memoirs*, vol. III. R. and J. E. Taylor, London.
- Loveland, D.** (1970). A linear format for resolution. In *Proc. IRIA Symposium on Automatic Demonstration*.
- Lowe, D.** (1987). Three-dimensional object recognition from single two-dimensional images. *AIJ*, 31, 355–395.
- Lowe, D.** (2004). Distinctive image features from scale-invariant keypoints. *IJCV*, 60, 91–110.
- Löwenheim, L.** (1915). Über möglichkeiten im Relativkalkül. *Mathematische Annalen*, 76, 447–470.
- Lowerre, B. T.** (1976). *The HARPY Speech Recognition System*. Ph.D. thesis, Computer Science Department, Carnegie-Mellon University.
- Lowry, M.** (2008). Intelligent software engineering tools for NASA's crew exploration vehicle. In *ISMIS-08*.
- Loyd, S.** (1959). *Mathematical Puzzles of Sam Loyd: Selected and Edited by Martin Gardner*. Dover.
- Lozano-Perez, T.** (1983). Spatial planning: A configuration space approach. *IEEE Transactions on Computers*, C-32, 108–120.
- Lozano-Perez, T., Mason, M., and Taylor, R.** (1984). Automatic synthesis of fine-motion strategies for robots. *Int. J. Robotics Research*, 3, 3–24.
- Liu, F. and Milios, E.** (1997). Globally consistent range scan alignment for environment mapping. *Autonomous Robots*, 4, 333–349.
- Lubberts, A. and Miikkulainen, R.** (2001). Co-evolving a Go-playing neural network. In *GECCO-01*.
- Luby, M., Sinclair, A., and Zuckerman, D.** (1993). Optimal speedup of Las Vegas algorithms. *Information Processing Letters*, 47, 173–180.
- Lucas, J. R.** (1961). Minds, machines, and Gödel. *Philosophy*, 36.
- Lucas, J. R.** (1976). This Gödel is killing me: A rejoinder. *Philosophia*, 6, 145–148.
- Lucas, P., van der Gaag, L., and Abu-Hanna, A.** (2004). Bayesian networks in biomedicine and healthcare. *Artificial Intelligence in Medicine*.
- Luce, D. R. and Raiffa, H.** (1957). *Games and Decisions*. Wiley.
- Lukasiewicz, T.** (1998). Probabilistic logic programming. In *ECAI-98*.
- Lundberg, S. M. and Lee, S.-I.** (2018). A unified approach to interpreting model predictions. In *NeurIPS 30*.

- Lunn, D., Jackson, C., Best, N., Thomas, A., and Spiegelhalter, D. J.** (2013). *The BUGS Book: A Practical Introduction to Bayesian Analysis*. Chapman and Hall.
- Lunn, D., Thomas, A., Best, N., and Spiegelhalter, D. J.** (2000). WinBUGS—a Bayesian modelling framework: Concepts, structure, and extensibility. *Statistics and Computing*, 10, 325–337.
- Luo, S., Bimbo, J., Dahiya, R., and Liu, H.** (2017). Robotic tactile perception of object properties: A review. *Mechatronics*, 48, 54–67.
- Lyman, P. and Varian, H. R.** (2003). How much information? [www.sims.berkeley.edu/how-much-info-2003](http://www.sims.berkeley.edu/how-much-info-2003).
- Lynch, K. and Park, F. C.** (2017). *Modern Robotics*. Cambridge University Press.
- Machina, M.** (2005). Choice under uncertainty. In *Encyclopedia of Cognitive Science*. Wiley.
- MacKay, D. J. C.** (2002). *Information Theory, Inference and Learning Algorithms*. Cambridge University Press.
- MacKenzie, D.** (2004). *Mechanizing Proof*. MIT Press.
- Mackworth, A. K.** (1977). Consistency in networks of relations. *AIJ*, 8, 99–118.
- Mackworth, A. K. and Freuder, E. C.** (1985). The complexity of some polynomial network consistency algorithms for constraint satisfaction problems. *AIJ*, 25, 65–74.
- Madhavan, R. and Schlenoff, C. I.** (2003). Moving object prediction for off-road autonomous navigation. In *Unmanned Ground Vehicle Technology V*.
- Mailath, G. and Samuelson, L.** (2006). *Repeated Games and Reputations: Long-Run Relationships*. Oxford University Press.
- Majercik, S. M. and Littman, M. L.** (2003). Contingent planning under uncertainty via stochastic satisfiability. *AIJ*, 147, 119–162.
- Malhotra, P., Vig, L., Shroff, G., and Agarwal, P.** (2015). Long short term memory networks for anomaly detection in time series. In *ISANN-15*.
- Malik, D., Palaniappan, M., Fisac, J. F., Hadfield-Menell, D., Russell, S. J., and Dragan, A. D.** (2018). An efficient, generalized bellman update for cooperative inverse reinforcement learning. In *ICML-18*.
- Malone, T. W.** (2004). *The Future of Work*. Harvard Business Review Press.
- Maneva, E., Mossel, E., and Wainwright, M.** (2007). A new look at survey propagation and its generalizations. *arXiv:cs/0409012*.
- Manna, Z. and Waldinger, R.** (1971). Toward automatic program synthesis. *CACM*, 14, 151–165.
- Manna, Z. and Waldinger, R.** (1985). *The Logical Basis for Computer Programming: Volume 1: Deductive Reasoning*. Addison-Wesley.
- Manne, A. S.** (1960). Linear programming and sequential decisions. *Management Science*, 6, 259–267.
- Manning, C. and Schütze, H.** (1999). *Foundations of Statistical Natural Language Processing*. MIT Press.
- Manning, C., Raghavan, P., and Schütze, H.** (2008). *Introduction to Information Retrieval*. Cambridge University Press.
- Mannion, M.** (2002). Using first-order logic for product line model validation. In *Software Product Lines: Second International Conference*.
- Mansinghka, V. K., Selsam, D., and Perov, Y.** (2013). Venture: A higher-order probabilistic programming platform with programmable inference. *arXiv:1404.0099*.
- Marbach, P. and Tsitsiklis, J. N.** (1998). Simulation-based optimization of Markov reward processes. Technical report, Laboratory for Information and Decision Systems, MIT.
- Marcus, G.** (2009). *Kluge: The Haphazard Evolution of the Human Mind*. Mariner Books.
- Marcus, M. P., Santorini, B., and Marcinkiewicz, M. A.** (1993). Building a large annotated corpus of English: The Penn treebank. *Computational Linguistics*, 19, 313–330.
- Marinescu, R. and Dechter, R.** (2009). AND/OR branch-and-bound search for combinatorial optimization in graphical models. *AIJ*, 173, 1457–1491.
- Markov, A.** (1913). An example of statistical investigation in the text of “Eugène Onegin” illustrating coupling of “tests” in chains. *Proc. Academy of Sciences of St. Petersburg*, 7, 153–162.
- Marler, R. T. and Arora, J. S.** (2004). Survey of multi-objective optimization methods for engineering. *Structural and Multidisciplinary Optimization*, 26, 369–395.
- Maron, M. E.** (1961). Automatic indexing: An experimental inquiry. *JACM*, 8, 404–417.
- Márquez, L. and Rodríguez, H.** (1998). Part-of-speech tagging using decision trees. In *ECML-98*.
- Marr, D. and Poggio, T.** (1976). Cooperative computation of stereo disparity. *Science*, 194, 283–287.
- Marr, D.** (1982). *Vision: A Computational Investigation into the Human Representation and Processing of Visual Information*. W. H. Freeman.
- Marriott, K. and Stuckey, P. J.** (1998). *Programming with Constraints: An Introduction*. MIT Press.
- Marsland, S.** (2014). *Machine Learning: An Algorithmic Perspective* (2nd edition). CRC Press.
- Martelli, A. and Montanari, U.** (1973). Additive AND/OR graphs. In *IJCAI-73*.
- Martelli, A.** (1977). On the complexity of admissible search algorithms. *AIJ*, 8, 1–13.
- Marthi, B., Pasula, H., Russell, S. J., and Peres, Y.** (2002). Decayed MCMC filtering. In *UAI-02*.
- Marthi, B., Russell, S. J., Latham, D., and Guestrin, C.** (2005). Concurrent hierarchical reinforcement learning. In *IJCAI-05*.
- Marthi, B., Russell, S. J., and Wolfe, J.** (2007). An-gelic semantics for high-level actions. In *ICAPS-07*.
- Marthi, B., Russell, S. J., and Wolfe, J.** (2008). An-gelic hierarchical planning: Optimal and online algorithms. In *ICAPS-08*.
- Martin, D., Fowlkes, C., and Malik, J.** (2004). Learning to detect natural image boundaries using local brightness, color, and texture cues. *PAMI*, 26, 530–549.
- Martin, F. G.** (2012). Will massive open online courses change how we teach? *CACM*, 55, 26–28.
- Martin, J. H.** (1990). *A Computational Model of Metaphor Interpretation*. Academic Press.
- Mason, M.** (1993). Kicking the sensing habit. *AIMag*, 14, 58–59.
- Mason, M.** (2001). *Mechanics of Robotic Manipulation*. MIT Press.
- Mason, M. and Salisbury, J.** (1985). *Robot Hands and the Mechanics of Manipulation*. MIT Press.
- Matarie, M. J.** (1997). Reinforcement learning in the multi-robot domain. *Autonomous Robots*, 4, 73–83.
- Mates, B.** (1953). *Stoic Logic*. University of California Press.
- Matuszek, C., Cabral, J., Witbrock, M., and DeOliveira, J.** (2006). An introduction to the syntax and semantics of Cyc. In *Proc. AAAI Spring Symposium on Formalizing and Compiling Background Knowledge and Its Applications to Knowledge Representation and Question Answering*.
- Mausam and Kolobov, A.** (2012). *Planning with Markov Decision Processes: An AI Perspective*. Morgan & Claypool.
- Maxwell, J.** (1868). On governors. *Proc. Roy. Soc.*, 16, 270–283.
- Mayer, J., Khairy, K., and Howard, J.** (2010). Drawing an elephant with four complex parameters. *American Journal of Physics*, 78, 648–649.
- Mayor, A.** (2018). *Gods and Robots: Myths, Machines, and Ancient Dreams of Technology*. Princeton University Press.
- McAllester, D. A.** (1980). An outlook on truth maintenance. AI memo, MIT AI Laboratory.
- McAllester, D. A.** (1988). Conspiracy numbers for min-max search. *AIJ*, 35, 287–310.
- McAllester, D. A.** (1998). What is the most pressing issue facing AI and the AAAI today? Candidate statement, election for Councilor of the American Association for Artificial Intelligence.
- McAllester, D. A. and Rosenblitt, D.** (1991). Systematic nonlinear planning. In *AAAI-91*.
- McAllester, D. A.** (1990). Truth maintenance. In *AAAI-90*.
- McAllester, D. A., Milch, B., and Goodman, N. D.** (2008). Random-world semantics and syntactic independence for expressive languages. Technical report, MIT.
- McCallum, A.** (2003). Efficiently inducing features of conditional random fields. In *UAI-03*.
- McCallum, A., Schultz, K., and Singh, S.** (2009). FACTORIE: Probabilistic programming via imperatively defined factor graphs. In *NeurIPS 22*.
- McCarthy, J.** (1958). Programs with common sense. In *Proc. Symposium on Mechanisation of Thought Processes*.
- McCarthy, J.** (1963). Situations, actions, and causal laws. Memo, Stanford University Artificial Intelligence Project.
- McCarthy, J.** (1968). Programs with common sense. In Minsky, M. L. (Ed.), *Semantic Information Processing*. MIT Press.
- McCarthy, J.** (1980). Circumscription: A form of non-monotonic reasoning. *AIJ*, 13, 27–39.
- McCarthy, J.** (2007). From here to human-level AI. *AIJ*, 171.
- McCarthy, J. and Hayes, P. J.** (1969). Some philosophical problems from the standpoint of artificial intelligence. In Meltzer, B., Michie, D., and Swann, M. (Eds.), *Machine Intelligence 4*. Edinburgh University Press.
- McCawley, J. D.** (1988). *The Syntactic Phenomena of English*. University of Chicago Press.
- McCorduck, P.** (2004). *Machines Who Think: A Personal Inquiry Into the History and Prospects of Artificial Intelligence* (Revised edition). A K Peters.
- McCulloch, W. S. and Pitts, W.** (1943). A logical calculus of the ideas immanent in nervous activity. *Bulletin of Mathematical Biophysics*, 5, 115–137.
- McCune, W.** (1997). Solution of the Robbins problem. *JAR*, 19, 263–276.
- McCune, W.** (1990). Otter 2.0. In *International Conference on Automated Deduction*.
- McDermott, D.** (1976). Artificial intelligence meets natural stupidity. *SIGART Newsletter*, 57, 4–9.

- McDermott**, D. (1978a). Planning and acting. *Cognitive Science*, 2, 71–109.
- McDermott**, D. (1978b). Tarskian semantics, or no notation without denotation! *Cognitive Science*, 2, 277–282.
- McDermott**, D. (1985). Reasoning about plans. In Hobbs, J. and Moore, R. (Eds.), *Formal theories of the commonsense world*. Ablex.
- McDermott**, D. (1987). A critique of pure reason. *Computational Intelligence*, 3, 151–237.
- McDermott**, D. (1996). A heuristic estimator for means-ends analysis in planning. In *ICAPS-96*.
- McDermott**, D. and Doyle, J. (1980). Non-monotonic logic: i. *AIJ*, 13, 41–72.
- McDermott**, J. (1982). R1: A rule-based configurer of computer systems. *AIJ*, 19, 39–88.
- McEliece**, R. J., MacKay, D. J. C., and Cheng, J.-F. (1998). Turbo decoding as an instance of Pearl's "belief propagation" algorithm. *IEEE Journal on Selected Areas in Communications*, 16, 140–152.
- McGregor**, J. J. (1979). Relational consistency algorithms and their application in finding subgraph and graph isomorphisms. *Information Sciences*, 19, 229–250.
- McIlraith**, S. and Zeng, H. (2001). Semantic web services. *IEEE Intelligent Systems*, 16, 46–53.
- McKinney**, W. (2012). *Python for Data Analysis: Data Wrangling with Pandas*. O'Reilly.
- McLachlan**, G. J. and Krishnan, T. (1997). *The EM Algorithm and Extensions*. Wiley.
- McMahan**, H. B. and Andrew, G. (2018). A general approach to adding differential privacy to iterative training procedures. arXiv:1812.06210.
- McMillan**, K. L. (1993). *Symbolic Model Checking*. Kluwer.
- McWhorter**, J. H. (2014). *The Language Hoax: Why the World Looks the Same in Any Language*. Oxford University Press.
- Meehl**, P. (1955). *Clinical vs. Statistical Prediction*. University of Minnesota Press.
- Mehrabi**, N., Morstatter, F., Saxena, N., Lerman, K., and Galstyan, A. (2019). A survey on bias and fairness in machine learning. arXiv:1908.09635.
- Mendel**, G. (1866). Versuche über pflanzen-hybriden. *Verhandlungen des Naturforschenden Vereins, Abhandlungen*, Brünn, 4, 3–47. Translated into English by C. T. Drury, published by Bateson (1902).
- Mercer**, J. (1909). Functions of positive and negative type and their connection with the theory of integral equations. *Phil. Trans. Roy. Soc., A*, 209, 415–446.
- Merleau-Ponty**, M. (1945). *Phenomenology of Perception*. Routledge.
- Metropolis**, N., Rosenbluth, A., Rosenbluth, M., Teller, A., and Teller, E. (1953). Equations of state calculations by fast computing machines. *J. Chemical Physics*, 21, 1087–1091.
- Metropolis**, N. and Ulam, S. (1949). The beginning of the Monte Carlo method. *Journal of the American Statistical Association*, 44, 335–341.
- Mézard**, M., Parisi, G., and Virasoro, M. (1987). *Spin Glass Theory and Beyond: An Introduction to the Replica Method and Its Applications*. World Scientific.
- Michie**, D. (1966). Game-playing and game-learning automata. In Fox, L. (Ed.), *Advances in Programming and Non-Numerical Computation*. Pergamon.
- Michie**, D. (1972). Machine intelligence at Edinburgh. *Management Informatics*, 2, 7–12.
- Michie**, D. and Chambers, R. A. (1968). BOXES: An experiment in adaptive control. In Dale, E. and Michie, D. (Eds.), *Machine Intelligence 2*. Elsevier.
- Michie**, D. (1963). Experiments on the mechanization of game-learning Part I. Characterization of the model and its parameters. *The Computer Journal*, 6, 232–236.
- Miiikkulainen**, R., Liang, J., Meyerson, E., Rawal, A., Fink, D., Francon, O., Raju, B., Shahrad, H., Navruzyan, A., Duffy, N., et al. (2019). Evolving deep neural networks. In *Artificial Intelligence in the Age of Neural Networks and Brain Computing*. Elsevier.
- Mikolov**, T., Chen, K., Corrado, G., and Dean, J. (2013). Efficient estimation of word representations in vector space. arXiv:1301.3781.
- Mikolov**, T., Karafiat, M., Burget, L., Černocký, J., and Khudanpur, S. (2010). Recurrent neural network based language model. In *Eleventh Annual Conference of the International Speech Communication Association*.
- Mikolov**, T., Sutskever, I., Chen, K., Corrado, G., and Dean, J. (2014). Distributed representations of words and phrases and their compositionality. In *NIPS 26*.
- Milch**, B. (2006). *Probabilistic Models with Unknown Objects*. Ph.D. thesis, UC Berkeley.
- Milch**, B., Marthi, B., Sontag, D., Russell, S. J., Ong, D., and Kolobov, A. (2005). BLOG: Probabilistic models with unknown objects. In *IJCAI-05*.
- Milch**, B., Zettlemoyer, L., Kersting, K., Haimes, M., and Kaelbling, L. P. (2008). Lifted probabilistic inference with counting formulas. In *AAAI-08*.
- Milgrom**, P. (1997). Putting auction theory to work: The simultaneous ascending auction. Tech. rep., Stanford University Department of Economics.
- Mill**, J. S. (1863). *Utilitarianism*. Parker, Son and Bourne, London.
- Miller**, A. C., Merkhofer, M. M., Howard, R. A., Matheson, J. E., and Rice, T. R. (1976). Development of automated aids for decision analysis. Technical report, SRI International.
- Miller**, T., Howe, P., and Sonnenberg, L. (2017). Explainable AI: Beware of inmates running the asylum. In *Proc. IJCAI-17 Workshop on Explainable AI*.
- Minka**, T. (2010). Bayesian linear regression. Unpublished manuscript.
- Minka**, T., Cleven, R., and Zaykov, Y. (2018). TrueSkill 2: An improved Bayesian skill rating system. Tech. rep., Microsoft Research.
- Minker**, J. (2001). *Logic-Based Artificial Intelligence*. Kluwer.
- Minsky**, M. L. (1975). A framework for representing knowledge. In Winston, P. H. (Ed.), *The Psychology of Computer Vision*. McGraw-Hill.
- Minsky**, M. L. (1986). *The Society of Mind*. Simon and Schuster.
- Minsky**, M. L. (2007). *The Emotion Machine: Commonsense Thinking, Artificial Intelligence, and the Future of the Human Mind*. Simon and Schuster.
- Minsky**, M. L. and Papert, S. (1969). *Perceptrons: An Introduction to Computational Geometry*. MIT Press.
- Minsky**, M. L. and Papert, S. (1988). *Perceptrons: An Introduction to Computational Geometry* (Expanded edition). MIT Press.
- Minsky**, M. L., Singh, P., and Sloman, A. (2004). The St. Thomas common sense symposium: Designing architectures for human-level intelligence. *AIMag*, 25, 113–124.
- Minton**, S., Johnston, M. D., Philips, A. B., and Laird, P. (1992). Minimizing conflicts: A heuristic repair method for constraint satisfaction and scheduling problems. *AIJ*, 58, 161–205.
- Mirjalili**, S. M. and Lewis, A. (2014). Grey wolf optimizer. *Advances in Engineering Software*, 69, 46–61.
- Misak**, C. (2004). *The Cambridge Companion to Peirce*. Cambridge University Press.
- Mitchell**, M., Wu, S., Zaldivar, A., Barnes, P., Vasserman, L., Hutchinson, B., Spitzer, E., Raji, I. D., and Gebru, T. (2019). Model cards for model reporting. *Proc. of the Conference on Fairness, Accountability, and Transparency*.
- Mitchell**, M. (1996). *An Introduction to Genetic Algorithms*. MIT Press.
- Mitchell**, M. (2019). *Artificial Intelligence: A Guide for Thinking Humans*. Farrar, Straus and Giroux.
- Mitchell**, M., Holland, J. H., and Forrest, S. (1996). When will a genetic algorithm outperform hill climbing? In *NIPS 6*.
- Mitchell**, T. M. (1997). *Machine Learning*. McGraw-Hill.
- Mitchell**, T. M. (2005). Reading the web: A breakthrough goal for AI. *AIMag*, 26.
- Mitchell**, T. M. (2007). Learning, information extraction and the web. In *ECML-07*.
- Mitchell**, T. M., Cohen, W., Hruschka, E., Talukdar, P., Yang, B., Betteridge, J., Carlson, A., Dalvi, B., Gardner, M., Kisiel, B., et al. (2018). Never-ending learning. *CACM*, 61, 103–115.
- Mitchell**, T. M., Shinkareva, S. V., Carlson, A., Chang, K. M., Malave, V. L., Mason, R. A., and Just, M. A. (2008). Predicting human brain activity associated with the meanings of nouns. *Science*, 320, 1191–1195.
- Mittelstadt**, B. (2019). Principles alone cannot guarantee ethical AI. *Nature Machine Intelligence*, 1, 501–507.
- Mitten**, L. G. (1960). An analytic solution to the least cost testing sequence problem. *Journal of Industrial Engineering*, 11, 17.
- Miyato**, T., Kataoka, T., Koyama, M., and Yoshida, Y. (2018). Spectral normalization for generative adversarial networks. arXiv:1802.05957.
- Mnih**, V., Kavukcuoglu, K., Silver, D., Graves, A., Antonoglou, I., Wierstra, D., and Riedmiller, M. A. (2013). Playing Atari with deep reinforcement learning. arXiv:1312.5602.
- Mnih**, V., Kavukcuoglu, K., Silver, D., Rusu, A. A., Veness, J., Bellemare, M. G., Graves, A., Riedmiller, M. A., Fidjeland, A., Ostrovski, G., Petersen, S., Beattie, C., Sadik, A., Antonoglou, I., King, H., Kumaran, D., Wierstra, D., Legg, S., and Hassabis, D. (2015). Human-level control through deep reinforcement learning. *Nature*, 518, 529–533.
- Mohr**, R. and Henderson, T. C. (1986). Arc and path consistency revisited. *AIJ*, 28, 225–233.
- Montague**, R. (1970). English as a formal language. In Visentini, B. (Ed.), *Linguaggi nella Società e nella Tecnica*. Edizioni di Comunità.
- Montague**, R. (1973). The proper treatment of quantification in ordinary English. In Hintikka, K. J. J., Moravcsik, J. M. E., and Suppes, P. (Eds.), *Approaches to Natural Language*. D. Reidel.
- Montanari**, U. (1974). Networks of constraints: Fundamental properties and applications to picture processing. *Information Sciences*, 7, 95–132.
- Montemerlo**, M. and Thrun, S. (2004). Large-scale robotic 3-D mapping of urban structures. In *Proc. International Symposium on Experimental Robotics*.
- Montemerlo**, M., Thrun, S., Koller, D., and Wegbreit, B. (2002). FastSLAM: A factored solution to the simultaneous localization and mapping problem. In *AAAI-02*.
- Mooney**, R. (1999). Learning for semantic interpretation: Scaling up without dumbing down. In *Proc. 1st Workshop on Learning Language in Logic*.

- Moore, A. M.** and Wong, W.-K. (2003). Optimal reinsertion: A new search operator for accelerated and more accurate Bayesian network structure learning. In *ICML-03*.
- Moore, A. W.** and Atkeson, C. G. (1993). Prioritized sweeping—Reinforcement learning with less data and less time. *Machine Learning*, 13, 103–130.
- Moore, A. W.** and Lee, M. S. (1997). Cached sufficient statistics for efficient machine learning with large datasets. *JAIR*, 8, 67–91.
- Moore, E. F.** (1959). The shortest path through a maze. In *Proc. International Symposium on the Theory of Switching, Part II*. Harvard University Press.
- Moore, R. C.** (1980). Reasoning about knowledge and action. Artificial intelligence center technical note, SRI International.
- Moore, R. C.** (1985). A formal theory of knowledge and action. In Hobbs, J. R. and Moore, R. C. (Eds.), *Formal Theories of the Commonsense World*. Ablex.
- Moore, R. C.** and DeNero, J. (2011). L1 and L2 regularization for multiclass hinge loss models. In *Symposium on Machine Learning in Speech and Natural Language Processing*.
- Moravčík, M.**, Schmid, M., Burch, N., Lisý, V., Morhill, D., Bard, N., Davis, T., Waugh, K., Johanson, M., and Bowling, M. (2017). Deepstack: Expert-level artificial intelligence in no-limit poker. arXiv:1701.01724.
- Moravec, H. P.** (1983). The Stanford cart and the CMU rover. *Proc. IEEE*, 71, 872–884.
- Moravec, H. P.** and Elfes, A. (1985). High resolution maps from wide angle sonar. In *ICRA-85*.
- Moravec, H. P.** (2000). *Robot: Mere Machine to Transcendent Mind*. Oxford University Press.
- Morgan, C. L.** (1896). *Habit and Instinct*. Edward Arnold.
- Morgan, T. J. H.** and Griffiths, T. L. (2015). What the Baldwin Effect affects. In *COGSCI-15*.
- Morjaria, M. A.**, Rink, F. J., Smith, W. D., Klempner, G., Burns, C., and Stein, J. (1995). Elicitation of probabilities for belief networks: Combining qualitative and quantitative information. In *UAI-95*.
- Morrison, P.** and Morrison, E. (Eds.). (1961). *Charles Babbage and His Calculating Engines: Selected Writings by Charles Babbage and Others*. Dover.
- Moskewicz, M. W.**, Madigan, C. F., Zhao, Y., Zhang, L., and Malik, S. (2001). Chaff: Engineering an efficient SAT solver. In *Proc. 38th Design Automation Conference*.
- Mott, A.**, Job, J., Vlimant, J.-R., Lidar, D., and Spiropulu, M. (2017). Solving a Higgs optimization problem with quantum annealing for machine learning. *Nature*, 550, 375.
- Motzkin, T. S.** and Schoenberg, I. J. (1954). The relaxation method for linear inequalities. *Canadian Journal of Mathematics*, 6, 393–404.
- Moutarlier, P.** and Chatila, R. (1989). Stochastic multisensory data fusion for mobile robot location and environment modeling. In *ISRR-89*.
- Mueller, E. T.** (2006). *Commonsense Reasoning*. Morgan Kaufmann.
- Muggleton, S. H.** and De Raedt, L. (1994). Inductive logic programming: Theory and methods. *J. Logic Programming*, 19/20, 629–679.
- Müller, M.** (2002). Computer Go. *AIJ*, 134, 145–179.
- Mumford, D.** and Shah, J. (1989). Optimal approximations by piece-wise smooth functions and associated variational problems. *Commun. Pure Appl. Math.*, 42, 577–685.
- Mundy, J.** and Zisserman, A. (Eds.). (1992). *Geometric Invariance in Computer Vision*. MIT Press.
- Munos, R.**, Stepleton, T., Harutyunyan, A., and Bellamare, M. G. (2017). Safe and efficient off-policy reinforcement learning. In *NeurIPS 29*.
- Murphy, K.** (2002). *Dynamic Bayesian Networks: Representation, Inference and Learning*. Ph.D. thesis, UC Berkeley.
- Murphy, K.** (2012). *Machine Learning: A Probabilistic Perspective*. MIT Press.
- Murphy, K.** and Mian, I. S. (1999). Modelling gene expression data using Bayesian networks. Tech. rep., Computer Science Division, UC Berkeley.
- Murphy, K.** and Russell, S. J. (2001). Rao-Blackwellised particle filtering for dynamic Bayesian networks. In Doucet, A., de Freitas, J. F. G., and Gordon, N. J. (Eds.), *Sequential Monte Carlo Methods in Practice*. Springer-Verlag.
- Murphy, K.** and Weiss, Y. (2001). The factored frontier algorithm for approximate inference in DBNs. In *UAI-01*.
- Murphy, R.** (2000). *Introduction to AI Robotics*. MIT Press.
- Murray, L. M.** (2013). Bayesian state-space modelling on high-performance hardware using LibBi. arXiv:1306.3277.
- Murray, R. M.** (2017). *A Mathematical Introduction to Robotic Manipulation*. CRC Press.
- Murray-Rust, P.**, Rzepa, H. S., Williamson, J., and Willighagen, E. L. (2003). Chemical markup, XML and the world-wide web. 4. CML schema. *J. Chem. Inf. Comput. Sci.*, 43, 752–772.
- Murthy, C.** and Russell, J. R. (1990). A constructive proof of Higman's lemma. In *LICS-90*.
- Muscettola, N.** (2002). Computing the envelope for stepwise-constant resource allocations. In *CP-02*.
- Muscettola, N.**, Nayak, P., Pell, B., and Williams, B. (1998). Remote agent: To boldly go where no AI system has gone before. *AIJ*, 103, 5–48.
- Muslea, I.** (1999). Extraction patterns for information extraction tasks: A survey. In *Proc. AAAI-99 Workshop on Machine Learning for Information Extraction*.
- Muth, J. T.**, Vogt, D. M., Truby, R. L., Mengüç, Y., Kolesky, D. B., Wood, R. J., and Lewis, J. A. (2014). Embedded 3D printing of strain sensors within highly stretchable elastomers. *Advanced Materials*, 26, 6307–6312.
- Myerson, R.** (1981). Optimal auction design. *Mathematics of Operations Research*, 6, 58–73.
- Myerson, R.** (1986). Multistage games with communication. *Econometrica*, 54, 323–358.
- Myerson, R.** (1991). *Game Theory: Analysis of Conflict*. Harvard University Press.
- Nair, V.** and Hinton, G. E. (2010). Rectified linear units improve restricted Boltzmann machines. In *ICML-10*.
- Nalwa, V. S.** (1993). *A Guided Tour of Computer Vision*. Addison-Wesley.
- Narayanan, A.**, Shi, E., and Rubinstein, B. I. (2011). Link prediction by de-anonymization: How we won the Kaggle social network challenge. In *IJCNN-11*.
- Narayanan, A.** and Shmatikov, V. (2006). How to break anonymity of the Netflix prize dataset. arXiv:cs/0610105.
- Nash, J.** (1950). Equilibrium points in N-person games. *PNAS*, 36, 48–49.
- Nash, P.** (1973). *Optimal Allocation of Resources Between Research Projects*. Ph.D. thesis, University of Cambridge.
- Nayak, P.** and Williams, B. (1997). Fast context switching in real-time propositional reasoning. In *AAAI-97*.
- Neches, R.**, Swartout, W. R., and Moore, J. D. (1985). Enhanced maintenance and explanation of expert systems through explicit models of their development. *IEEE Transactions on Software Engineering*, SE-11, 1337–1351.
- Nemhauser, G. L.**, Wolsey, L. A., and Fisher, M. L. (1978). An analysis of approximations for maximizing submodular set functions I. *Mathematical Programming*, 14, 265–294.
- Nesterov, Y.** and Nemirovski, A. (1994). *Interior-Point Polynomial Methods in Convex Programming*. SIAM (Society for Industrial and Applied Mathematics).
- Newell, A.** (1982). The knowledge level. *AIJ*, 18, 82–127.
- Newell, A.** (1990). *Unified Theories of Cognition*. Harvard University Press.
- Newell, A.** and Ernst, G. (1965). The search for generality. In *Proc. IFIP Congress*.
- Newell, A.**, Shaw, J. C., and Simon, H. A. (1957). Empirical explorations with the logic theory machine. *Proc. Western Joint Computer Conference*, 15, 218–239. Reprinted in Feigenbaum and Feldman (1963).
- Newell, A.** and Simon, H. A. (1961). GPS, a program that simulates human thought. In Billig, H. (Ed.), *Lernende Automaten*. R. Oldenbourg.
- Newell, A.** and Simon, H. A. (1972). *Human Problem Solving*. Prentice-Hall.
- Newell, A.** and Simon, H. A. (1976). Computer science as empirical inquiry: Symbols and search. *CACM*, 19, 113–126.
- Newton, I.** (1664–1671). *Methodus fluxionum et serierum infinitarum*. Unpublished notes.
- Ng, A. Y.** (2004). Feature selection,  $L_1$  vs.  $L_2$  regularization, and rotational invariance. In *ICML-04*.
- Ng, A. Y.** (2019). *Machine Learning Yearning*. www.mlyarning.org.
- Ng, A. Y.**, Harada, D., and Russell, S. J. (1999). Policy invariance under reward transformations: Theory and application to reward shaping. In *ICML-99*.
- Ng, A. Y.** and Jordan, M. I. (2000). PEGASUS: A policy search method for large MDPs and POMDPs. In *UAI-00*.
- Ng, A. Y.** and Jordan, M. I. (2002). On discriminative vs. generative classifiers: A comparison of logistic regression and naive Bayes. In *NeurIPS 14*.
- Ng, A. Y.**, Kim, H. J., Jordan, M. I., and Sastry, S. (2003). Autonomous helicopter flight via reinforcement learning. In *NeurIPS 16*.
- Ng, A. Y.** and Russell, S. J. (2000). Algorithms for inverse reinforcement learning. In *ICML-00*.
- Nicholson, A.** and Brady, J. M. (1992). The data association problem when monitoring robot vehicles using dynamic belief networks. In *ECAI-92*.
- Nielsen, M. A.** (2015). *Neural Networks and Deep Learning*. Determination Press.
- Nielsen, T.** and Jensen, F. (2003). Sensitivity analysis in influence diagrams. *IEEE Transactions on Systems, Man and Cybernetics*, 33, 223–234.
- Niemelä, I.**, Simons, P., and Syrjänen, T. (2000). Smodels: A system for answer set programming. In *Proc. 8th International Workshop on Non-Monotonic Reasoning*.
- Nikolaidis, S.** and Shah, J. (2013). Human-robot cross-training: computational formulation, modeling and evaluation of a human team training strategy. In *HRI-13*.
- Niles, I.** and Pease, A. (2001). Towards a standard upper ontology. In *Proc. International Conference on Formal Ontology in Information Systems*.

- Nilsson, D.** and Lauritzen, S. (2000). Evaluating influence diagrams using LIMIDs. In *UAI-00*.
- Nilsson, N. J.** (1965). *Learning Machines: Foundations of Trainable Pattern-Classifying Systems*. McGraw-Hill.
- Nilsson, N. J.** (1971). *Problem-Solving Methods in Artificial Intelligence*. McGraw-Hill.
- Nilsson, N. J.** (1984). Shakey the robot. Technical note, SRI International.
- Nilsson, N. J.** (1986). Probabilistic logic. *AIJ*, 28, 71–87.
- Nilsson, N. J.** (1995). Eye on the prize. *AIMag*, 16, 9–17.
- Nilsson, N. J.** (2009). *The Quest for Artificial Intelligence: A History of Ideas and Achievements*. Cambridge University Press.
- Nisan, N.** (2007). Introduction to mechanism design (for computer scientists). In Nisan, N., Roughgarden, T., Tardos, E., and Vazirani, V. V. (Eds.), *Algorithmic Game Theory*. Cambridge University Press.
- Nisan, N.**, Roughgarden, T., Tardos, E., and Vazirani, V. (Eds.). (2007). *Algorithmic Game Theory*. Cambridge University Press.
- Niv, Y.** (2009). Reinforcement learning in the brain. *Journal of Mathematical Psychology*, 53, 139–154.
- Nivre, J.**, De Marneffe, M.-C., Ginter, F., Goldberg, Y., Hajic, J., Manning, C., McDonald, R., Petrov, S., et al. (2016). Universal dependencies v1: A multilingual treebank collection. In *Proc. International Conference on Language Resources and Evaluation*.
- Nodelman, U.**, Shelton, C., and Koller, D. (2002). Continuous time Bayesian networks. In *UAI-02*.
- Noe, A.** (2009). *Out of Our Heads: Why You Are Not Your Brain, and Other Lessons from the Biology of Consciousness*. Hill and Wang.
- Nordfors, D.**, Cerf, V., and Senges, M. (2018). *Disrupting Unemployment*. Amazon Digital Services.
- Norvig, P.** (1988). Multiple simultaneous interpretations of ambiguous sentences. In *COGSCI-88*.
- Norvig, P.** (1992). *Paradigms of Artificial Intelligence Programming: Case Studies in Common Lisp*. Morgan Kaufmann.
- Norvig, P.** (2009). Natural language corpus data. In Segaran, T. and Hammerbacher, J. (Eds.), *Beautiful Data*. O'Reilly.
- Nowick, S. M.**, Dean, M. E., Dill, D. L., and Horowitz, M. (1993). The design of a high-performance cache controller: A case study in asynchronous synthesis. *Integration: The VLSI Journal*, 15, 241–262.
- Och, F. J.** and Ney, H. (2003). A systematic comparison of various statistical alignment models. *Computational Linguistics*, 29, 19–51.
- Och, F. J.** and Ney, H. (2004). The alignment template approach to statistical machine translation. *Computational Linguistics*, 30, 417–449.
- Och, F. J.** and Ney, H. (2002). Discriminative training and maximum entropy models for statistical machine translation. In *COLING-02*.
- Ogawa, S.**, Lee, T.-M., Kay, A. R., and Tank, D. W. (1990). Brain magnetic resonance imaging with contrast dependent on blood oxygenation. *PNAS*, 87, 9868–9872.
- Oh, M.-S.** and Berger, J. O. (1992). Adaptive importance sampling in Monte Carlo integration. *Journal of Statistical Computation and Simulation*, 41, 143–168.
- Oh, S.**, Russell, S. J., and Sastry, S. (2009). Markov chain Monte Carlo data association for multi-target tracking. *IEEE Transactions on Automatic Control*, 54, 481–497.
- Oizumi, M.**, Albantakis, L., and Tononi, G. (2014). From the phenomenology to the mechanisms of consciousness: Integrated information theory 3.0. *PLoS Computational Biology*, 10, e1003588.
- Olesen, K. G.** (1993). Causal probabilistic networks with both discrete and continuous variables. *PAMI*, 15, 275–279.
- Oliver, N.**, Garg, A., and Horvitz, E. J. (2004). Layered representations for learning and inferring office activity from multiple sensory channels. *Computer Vision and Image Understanding*, 96, 163–180.
- Oliver, R. M.** and Smith, J. Q. (Eds.). (1990). *Influence Diagrams, Belief Nets and Decision Analysis*. Wiley.
- Omohundro, S.** (2008). The basic AI drives. In *AGI-08 Workshop on the Sociocultural, Ethical and Futureological Implications of Artificial Intelligence*.
- O'Neil, C.** (2017). *Weapons of Math Destruction: How Big Data Increases Inequality and Threatens Democracy*. Broadway Books.
- O'Neil, C.** and Schutt, R. (2013). *Doing Data Science: Straight Talk from the Frontline*. O'Reilly.
- O'Reilly, U.-M.** and Oppacher, F. (1994). Program search with a hierarchical variable length representation: Genetic programming, simulated annealing and hill climbing. In *Proc. Third Conference on Parallel Problem Solving from Nature*.
- Osborne, M. J.** (2004). *An Introduction to Game Theory*. Oxford University Press.
- Osborne, M. J.** and Rubinstein, A. (1994). *A Course in Game Theory*. MIT Press.
- Osherson, D. N.**, Stob, M., and Weinstein, S. (1986). *Systems That Learn: An Introduction to Learning Theory for Cognitive and Computer Scientists*. MIT Press.
- Ostrom, E.** (1990). *Governing the Commons*. Cambridge University Press.
- Padgham, L.** and Winikoff, M. (2004). *Developing Intelligent Agent Systems: A Practical Guide*. Wiley.
- Paige, B.** and Wood, F. (2014). A compilation target for probabilistic programming languages. In *ICML-14*.
- Paige, B.**, Wood, F., Doucet, A., and Teh, Y. W. (2015). Asynchronous anytime sequential Monte Carlo. In *NeurIPS 27*.
- Palacios, H.** and Geffner, H. (2007). From conformant into classical planning: Efficient translations that may be complete too. In *ICAPS-07*.
- Palmer, S.** (1999). *Vision Science: Photons to Phenomenology*. MIT Press.
- Patapadimitriou, C. H.** (1994). *Computational Complexity*. Addison-Wesley.
- Patapadimitriou, C. H.** and Tsitsiklis, J. N. (1987). The complexity of Markov decision processes. *Mathematics of Operations Research*, 12, 441–450.
- Patapadimitriou, C. H.** and Yannakakis, M. (1991). Shortest paths without a map. *Theoretical Computer Science*, 84, 127–150.
- Papavassiliou, V.** and Russell, S. J. (1999). Convergence of reinforcement learning with general function approximators. In *IJCAI-99*.
- Parisi, G.** (1988). *Statistical Field Theory*. Addison-Wesley.
- Parisi, M. M. G.** and Zecchina, R. (2002). Analytic and algorithmic solution of random satisfiability problems. *Science*, 297, 812–815.
- Park, J. D.** and Darwiche, A. (2004). Complexity results and approximation strategies for MAP explanations. *JAIR*, 21, 101–133.
- Parker, A.**, Nau, D. S., and Subrahmanian, V. S. (2005). Game-tree search with combinatorially large belief states. In *IJCAI-05*.
- Parker, D. B.** (1985). Learning logic. Technical report, Center for Computational Research in Economics and Management Science, MIT.
- Parker, L. E.** (1996). On the design of behavior-based multi-robot teams. *J. Advanced Robotics*, 10, 547–578.
- Parr, R.** and Russell, S. J. (1998). Reinforcement learning with hierarchies of machines. In *NeurIPS 10*.
- Parzen, E.** (1962). On estimation of a probability density function and mode. *Annals of Mathematical Statistics*, 33, 1065–1076.
- Pasca, M.** and Harabagiu, S. M. (2001). High performance question/answering. In *SIGIR-01*.
- Pasca, M.**, Lin, D., Bigham, J., Lifchits, A., and Jain, A. (2006). Organizing and searching the world wide web of facts—Step one: The one-million fact extraction challenge. In *AAAI-06*.
- Paskin, M.** (2002). Maximum entropy probabilistic logic. Tech. report, UC Berkeley.
- Passala, H.**, Marti, B., Milch, B., Russell, S. J., and Shpitser, I. (2003). Identity uncertainty and citation matching. In *NeurIPS 15*.
- Passala, H.**, Russell, S. J., Ostland, M., and Ritov, Y. (1999). Tracking many objects with many sensors. In *IJCAI-99*.
- Patel-Schneider, P.** (2014). Analyzing schema.org. In *Proc. International Semantic Web Conference*.
- Patrick, B. G.**, Almulla, M., and Newborn, M. (1992). An upper bound on the time complexity of iterative-deepening-A\*. *AIJ*, 5, 265–278.
- Paul, R. P.** (1981). *Robot Manipulators: Mathematics, Programming, and Control*. MIT Press.
- Pauls, A.** and Klein, D. (2009). K-best A\* parsing. In *ACL-09*.
- Peano, G.** (1889). *Arithmetices principia, nova methodo exposita*. Fratres Bocca, Turin.
- Pearce, J.**, Tambe, M., and Maheswaran, R. (2008). Solving multiagent networks using distributed constraint optimization. *AIMag*, 29, 47–62.
- Pearl, J.** (1982a). Reverend Bayes on inference engines: A distributed hierarchical approach. In *AAAI-82*.
- Pearl, J.** (1982b). The solution for the branching factor of the alpha-beta pruning algorithm and its optimality. *CACM*, 25, 559–564.
- Pearl, J.** (1984). *Heuristics: Intelligent Search Strategies for Computer Problem Solving*. Addison-Wesley.
- Pearl, J.** (1985). Bayesian networks: A model of self-activated memory for evidential reasoning. In *COGSCI-85*.
- Pearl, J.** (1986). Fusion, propagation, and structuring in belief networks. *AIJ*, 29, 241–288.
- Pearl, J.** (1987). Evidential reasoning using stochastic simulation of causal models. *AIJ*, 32, 247–257.
- Pearl, J.** (1988). *Probabilistic Reasoning in Intelligent Systems: Networks of Plausible Inference*. Morgan Kaufmann.
- Pearl, J.** (2000). *Causality: Models, Reasoning, and Inference*. Cambridge University Press.
- Pearl, J.** and McKenzie, D. (2018). *The Book of Why*. Basic Books.
- Pearl, J.** and Verma, T. (1991). A theory of inferred causation. In *KR-91*.
- Pearson, K.** (1895). Contributions to the mathematical theory of evolution, II: Skew variation in homogeneous material. *Phil. Trans. Roy. Soc.*, 186, 343–414.
- Pearson, K.** (1901). On lines and planes of closest fit to systems of points in space. *Philosophical Magazine*, 2, 559–572.

- Pease, A.** and Niles, I. (2002). IEEE standard upper ontology: A progress report. *Knowledge Engineering Review*, 17, 65–70.
- Pednault, E. P. D.** (1986). Formulating multiagent, dynamic-world problems in the classical planning framework. In *Reasoning About Actions and Plans: Proc. 1986 Workshop*.
- Pedregosa, F.**, Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., Blondel, M., Prettenhofer, P., Weiss, R., Dubourg, V., et al. (2011). Scikit-learn: Machine learning in Python. *JMLR*, 12, 2825–2830.
- Peirce, C. S.** (1870). Description of a notation for the logic of relatives, resulting from an amplification of the conceptions of Boole's calculus of logic. *Memoirs of the American Academy of Arts and Sciences*, 9, 317–378.
- Peirce, C. S.** (1883). A theory of probable inference. Note B. The logic of relatives. In Peirce, C. S. (Ed.), *Studies in Logic*. Little, Brown.
- Peirce, C. S.** (1909). Existential graphs. Unpublished manuscript; reprinted in (Buchler 1955).
- Peleg, B.** and Sudholter, P. (2002). *Introduction to the Theory of Cooperative Games* (2nd edition). Springer-Verlag.
- Pelikan, M.**, Goldberg, D. E., and Cantu-Paz, E. (1999). BOA: The Bayesian optimization algorithm. In *GECCO-99*.
- Pemberton, J. C.** and Korf, R. E. (1992). Incremental planning on graphs with cycles. In *AIPS-92*.
- Penberthy, J. S.** and Weld, D. S. (1992). UCPOP: A sound, complete, partial order planner for ADL. In *KR-92*.
- Peng, J.** and Williams, R. J. (1993). Efficient learning and planning within the Dyna framework. *Adaptive Behavior*, 2, 437–454.
- Pennington, J.**, Socher, R., and Manning, C. (2014). Glove: Global vectors for word representation. In *EMNLP-14*.
- Penrose, R.** (1989). *The Emperor's New Mind*. Oxford University Press.
- Penrose, R.** (1994). *Shadows of the Mind*. Oxford University Press.
- Peot, M.** and Smith, D. E. (1992). Conditional nonlinear planning. In *ICAPS-92*.
- Pereira, F.** and Schabes, Y. (1992). Inside-outside reestimation from partially bracketed corpora. In *ACL-92*.
- Pereira, F.** and Warren, D. H. D. (1980). Definite clause grammars for language analysis: A survey of the formalism and a comparison with augmented transition networks. *AIJ*, 13, 231–278.
- Peters, J.** and Schaal, S. (2008). Reinforcement learning of motor skills with policy gradients. *Neural Networks*, 21, 682–697.
- Peters, J.**, Janzing, D., and Schölkopf, B. (2017). *Elements of Causal Inference: Foundations and Learning Algorithms*. MIT press.
- Peters, M. E.**, Neumann, M., Iyyer, M., Gardner, M., Clark, C., Lee, K., and Zettlemoyer, L. (2018). Deep contextualized word representations. arXiv:1802.05365.
- Peterson, C.** and Anderson, J. R. (1987). A mean field theory learning algorithm for neural networks. *Complex Systems*, 1, 995–1019.
- Petosa, N.** and Balch, T. (2019). Multiplayer AlphaZero. arXiv:1910.13012.
- Pfeffer, A.** (2001). IBAL: A probabilistic rational programming language. In *IJCAI-01*.
- Pfeffer, A.**, Koller, D., Milch, B., and Takusagawa, K. T. (1999). SPOOK: A system for probabilistic object-oriented knowledge representation. In *UAI-99*.
- Pfeffer, A.** (2016). *Practical Probabilistic Programming*. Manning.
- Pfeffer, A.** (2000). *Probabilistic Reasoning for Complex Systems*. Ph.D. thesis, Stanford University.
- Pfeffer, A.** (2007). The design and implementation of IBAL: A general-purpose probabilistic language. In Getoor, L. and Taskar, B. (Eds.), *Introduction to Statistical Relational Learning*. MIT Press.
- Pfeifer, R.**, Bongard, J., Brooks, R. A., and Iwasawa, S. (2006). *How the Body Shapes the Way We Think: A New View of Intelligence*. Bradford.
- Pham, H.**, Guan, M. Y., Zoph, B., Le, Q. V., and Dean, J. (2018). Efficient neural architecture search via parameter sharing. arXiv:1802.03268.
- Pineau, J.**, Gordon, G., and Thrun, S. (2003). Point-based value iteration: An anytime algorithm for POMDPs. In *IJCAI-03*.
- Pinedo, M.** (2008). *Scheduling: Theory, Algorithms, and Systems*. Springer Verlag.
- Pinkas, G.** and Dechter, R. (1995). Improving connectionist energy minimization. *JAIR*, 3, 223–248.
- Pinker, S.** (1995). Language acquisition. In Gleitman, L. R., Liberman, M., and Osherson, D. N. (Eds.), *An Invitation to Cognitive Science* (2nd edition). MIT Press.
- Pinker, S.** (2003). *The Blank Slate: The Modern Denial of Human Nature*. Penguin.
- Pinto, D.**, McCallum, A., Wei, X., and Croft, W. B. (2003). Table extraction using conditional random fields. In *SIGIR-03*.
- Pinto, L.** and Gupta, A. (2016). Supersizing self-supervision: Learning to grasp from 50k tries and 700 robot hours. In *ICRA-16*.
- Platt, J.** (1999). Fast training of support vector machines using sequential minimal optimization. In *Advances in Kernel Methods: Support Vector Learning*. MIT Press.
- Plotkin, G.** (1972). Building-in equational theories. In Meltzer, B. and Michie, D. (Eds.), *Machine Intelligence 7*. Edinburgh University Press.
- Plummer, M.** (2003). JAGS: A program for analysis of Bayesian graphical models using Gibbs sampling. In *Proc. Third Int'l Workshop on Distributed Statistical Computing*.
- Pnueli, A.** (1977). The temporal logic of programs. In *FOCS-77*.
- Pohl, I.** (1971). Bi-directional search. In Meltzer, B. and Michie, D. (Eds.), *Machine Intelligence 6*. Edinburgh University Press.
- Pohl, I.** (1973). The avoidance of (relative) catastrophe, heuristic competence, genuine dynamic weighting and computational issues in heuristic problem solving. In *IJCAI-73*.
- Pohl, I.** (1977). Practical and theoretical considerations in heuristic search algorithms. In Elcock, E. W. and Michie, D. (Eds.), *Machine Intelligence 8*. Ellis Horwood.
- Pohl, I.** (1970). Heuristic search viewed as path finding in a graph. *AIJ*, 1, 193–204.
- Poli, R.**, Langdon, W., and McPhee, N. (2008). *A Field Guide to Genetic Programming*. Lulu.com.
- Pomerleau, D. A.** (1993). *Neural Network Perception for Mobile Robot Guidance*. Kluwer.
- Poole, B.**, Lahiri, S., Raghu, M., Sohl-Dickstein, J., and Ganguli, S. (2017). Exponential expressivity in deep neural networks through transient chaos. In *NeurIPS 29*.
- Poole, D.** (1993). Probabilistic Horn abduction and Bayesian networks. *AIJ*, 64, 81–129.
- Poole, D.** (2003). First-order probabilistic inference. In *IJCAI-03*.
- Poole, D.** and Mackworth, A. K. (2017). *Artificial Intelligence: Foundations of Computational Agents* (2 edition). Cambridge University Press.
- Poppe, R.** (2010). A survey on vision-based human action recognition. *Image Vision Comput.*, 28, 976–990.
- Popper, K. R.** (1959). *The Logic of Scientific Discovery*. Basic Books.
- Popper, K. R.** (1962). *Conjectures and Refutations: The Growth of Scientific Knowledge*. Basic Books.
- Portner, P.** and Partee, B. H. (2002). *Formal Semantics: The Essential Readings*. Wiley-Blackwell.
- Post, E. L.** (1921). Introduction to a general theory of elementary propositions. *American Journal of Mathematics*, 43, 163–185.
- Poulton, C.** and Watts, M. (2016). MIT and DARPA pack Lidar sensor onto single chip. *IEEE Spectrum*, August 4.
- Poundstone, W.** (1993). *Prisoner's Dilemma*. Anchor.
- Pourret, O.**, Naïm, P., and Marcot, B. (2008). *Bayesian Networks: A Practical Guide to Applications*. Wiley.
- Pradhan, M.**, Provan, G. M., Middleton, B., and Henrion, M. (1994). Knowledge engineering for large belief networks. In *UAI-94*.
- Prawitz, D.** (1960). An improved proof procedure. *Theoria*, 26, 102–139.
- Press, W. H.**, Teukolsky, S. A., Vetterling, W. T., and Flannery, B. P. (2007). *Numerical Recipes: The Art of Scientific Computing* (3rd edition). Cambridge University Press.
- Preston, J.** and Bishop, M. (2002). *Views into the Chinese Room: New Essays on Searle and Artificial Intelligence*. Oxford University Press.
- Prieditis, A. E.** (1993). Machine discovery of effective admissible heuristics. *Machine Learning*, 12, 117–141.
- Prosser, P.** (1993). Hybrid algorithms for constraint satisfaction problems. *Computational Intelligence*, 9, 268–299.
- Pullum, G. K.** (1991). *The Great Eskimo Vocabulary Hoax (and Other Irreverent Essays on the Study of Language)*. University of Chicago Press.
- Pullum, G. K.** (1996). Learnability, hyperlearning, and the poverty of the stimulus. In *22nd Annual Meeting of the Berkeley Linguistics Society*.
- Puterman, M. L.** (1994). *Markov Decision Processes: Discrete Stochastic Dynamic Programming*. Wiley.
- Puterman, M. L.** and Shin, M. C. (1978). Modified policy iteration algorithms for discounted Markov decision problems. *Management Science*, 24, 1127–1137.
- Putnam, H.** (1963). ‘Degree of confirmation’ and inductive logic. In Schilpp, P. A. (Ed.), *The Philosophy of Rudolf Carnap*. Open Court.
- Quillian, M. R.** (1961). A design for an understanding machine. Paper presented at a colloquium: Semantic Problems in Natural Language, King's College, Cambridge, England.
- Quine, W. V.** (1953). Two dogmas of empiricism. In *From a Logical Point of View*. Harper and Row.
- Quine, W. V.** (1960). *Word and Object*. MIT Press.
- Quine, W. V.** (1982). *Methods of Logic* (4th edition). Harvard University Press.
- Quinlan, J. R.** (1979). Discovering rules from large collections of examples: A case study. In Michie, D. (Ed.), *Expert Systems in the Microelectronic Age*. Edinburgh University Press.
- Quinlan, J. R.** (1986). Induction of decision trees. *Machine Learning*, 1, 81–106.
- Quinlan, J. R.** (1993). *C4.5: Programs for Machine Learning*. Morgan Kaufmann.

- Quinlan**, S. and Khatib, O. (1993). Elastic bands: Connecting path planning and control. In *ICRA-93*.
- Quirk**, R., Greenbaum, S., Leech, G., and Svartvik, J. (1985). *A Comprehensive Grammar of the English Language*. Longman.
- Rabani**, Y., Rabinovich, Y., and Sinclair, A. (1998). A computational view of population genetics. *Random Structures and Algorithms*, 12, 313–334.
- Rabiner**, L. R. and Juang, B.-H. (1993). *Fundamentals of Speech Recognition*. Prentice-Hall.
- Radford**, A., Metz, L., and Chintala, S. (2015). Unsupervised representation learning with deep convolutional generative adversarial networks. arXiv:1511.06434.
- Radford**, A., Wu, J., Child, R., Luan, D., Amodei, D., and Sutskever, I. (2019). Language models are unsupervised multitask learners. *OpenAI Blog*, 1.
- Raffel**, C., Shazeer, N., Roberts, A., Lee, K., Narang, S., Matena, M., Zhou, Y., Li, W., and Liu, P. J. (2019). Exploring the limits of transfer learning with a unified text-to-text transformer. arXiv:1910.10683.
- Rafferty**, A. N., Brunskill, E., Griffiths, T. L., and Shafto, P. (2016). Faster teaching via POMDP planning. *Cognitive Science*, 40, 1290–1332.
- Rahwan**, T., Michalak, T. P., Wooldridge, M., and Jennings, N. R. (2015). Coalition structure generation: A survey. *AII*, 229, 139–174.
- Raihert**, M., Blankespoor, K., Nelson, G., and Playter, R. (2008). Bigdog, the rough-terrain quadruped robot. *IFAC Proceedings Volumes*, 41, 10822–10825.
- Rajpurkar**, P., Zhang, J., Lopyrev, K., and Liang, P. (2016). Squad: 100,000+ questions for machine comprehension of text. In *EMNLP-16*.
- Ramsey**, F. P. (1931). Truth and probability. In Braithwaite, R. B. (Ed.), *The Foundations of Mathematics and Other Logical Essays*. Harcourt Brace Jovanovich.
- Ramsundar**, B. and Zadeh, R. B. (2018). *TensorFlow for Deep Learning: From Linear Regression to Reinforcement Learning*. O'Reilly.
- Rao**, D. A. S. and Verweij, G. (2017). Sizing the prize: What's the real value of AI for your business and how can you capitalise? PwC.
- Raphael**, B. (1976). *The Thinking Computer: Mind Inside Matter*. W. H. Freeman.
- Raphson**, J. (1690). *Analysis aequationum universalis*. Apud Abelem Swalle, London.
- Raschka**, S. (2015). *Python Machine Learning*. Packt.
- Rashevsky**, N. (1936). Physico-mathematical aspects of excitation and conduction in nerves. In *Cold Springs Harbor Symposia on Quantitative Biology. IV: Excitation Phenomena*.
- Rashevsky**, N. (1938). *Mathematical Biophysics: Physico-Mathematical Foundations of Biology*. University of Chicago Press.
- Rasmussen**, C. E. and Williams, C. K. I. (2006). *Gaussian Processes for Machine Learning*. MIT Press.
- Rassenti**, S., Smith, V., and Bulfin, R. (1982). A combinatorial auction mechanism for airport time slot allocation. *Bell Journal of Economics*, 13, 402–417.
- Ratliff**, N., Bagnell, J. A., and Zinkevich, M. (2006). Maximum margin planning. In *ICML-06*.
- Ratliff**, N., Zucker, M., Bagnell, J. A., and Srinivasa, S. (2009). CHOMP: Gradient optimization techniques for efficient motion planning. In *ICRA-09*.
- Ratnaparkhi**, A. (1996). A maximum entropy model for part-of-speech tagging. In *EMNLP-96*.
- Ratner**, D. and Warmuth, M. (1986). Finding a shortest solution for the  $n \times n$  extension of the 15-puzzle is intractable. In *AAAI-86*.
- Rauch**, H. E., Tung, F., and Striebel, C. T. (1965). Maximum likelihood estimates of linear dynamic systems. *AIAA Journal*, 3, 1445–1450.
- Rayward-Smith**, V., Osman, I., Reeves, C., and Smith, G. (Eds.). (1996). *Modern Heuristic Search Methods*. Wiley.
- Real**, E., Aggarwal, A., Huang, Y., and Le, Q. V. (2018). Regularized evolution for image classifier architecture search. arXiv:1802.01548.
- Rechenberg**, I. (1965). Cybernetic solution path of an experimental problem. Library translation, Royal Aircraft Establishment.
- Regin**, J. (1994). A filtering algorithm for constraints of difference in CSPs. In *AAAI-94*.
- Reid**, D. B. (1979). An algorithm for tracking multiple targets. *IEEE Trans. Automatic Control*, 24, 843–854.
- Reif**, J. (1979). Complexity of the mover's problem and generalizations. In *FOCS-79*.
- Reiter**, R. (1980). A logic for default reasoning. *AII*, 13, 81–132.
- Reiter**, R. (1991). The frame problem in the situation calculus: A simple solution (sometimes) and a completeness result for goal regression. In Lifschitz, V. (Ed.), *Artificial Intelligence and Mathematical Theory of Computation: Papers in Honor of John McCarthy*. Academic Press.
- Reiter**, R. (2001). *Knowledge in Action: Logical Foundations for Specifying and Implementing Dynamical Systems*. MIT Press.
- Renner**, G. and Ekart, A. (2003). Genetic algorithms in computer aided design. *Computer Aided Design*, 35, 709–726.
- Rényi**, A. (1970). *Probability Theory*. Elsevier.
- Resnick**, P. and Varian, H. R. (1997). Recommender systems. *CACM*, 40, 56–58.
- Resende**, D. J., Mohamed, S., and Wierstra, D. (2014). Stochastic backpropagation and approximate inference in deep generative models. In *ICML-14*.
- Riazanov**, A. and Voronkov, A. (2002). The design and implementation of VAMPIRE. *AI Communications*, 15, 91–110.
- Ribeiro**, M. T., Singh, S., and Guestrin, C. (2016). Why should I trust you?: Explaining the predictions of any classifier. In *KDD-16*.
- Richardson**, M. and Domingos, P. (2006). Markov logic networks. *Machine Learning*, 62, 107–136.
- Richter**, S. and Helmert, M. (2009). Preferred operators and deferred evaluation in satisficing planning. In *ICAPS-09*.
- Ridley**, M. (2004). *Evolution*. Oxford Reader.
- Riley**, J. and Samuelson, W. (1981). Optimal auctions. *American Economic Review*, 71, 381–392.
- Riley**, P. (2019). Three pitfalls to avoid in machine learning. *Nature*, 572, 27–29.
- Riloff**, E. (1993). Automatically constructing a dictionary for information extraction tasks. In *AAAI-93*.
- Ringgaard**, M., Gupta, R., and Pereira, F. (2017). SLING: A framework for frame semantic parsing. arXiv:1710.07032.
- Rintanen**, J. (1999). Improvements to the evaluation of quantified Boolean formulae. In *IJCAI-99*.
- Rintanen**, J. (2007). Asymptotically optimal encodings of conformant planning in QBF. In *AAAI-07*.
- Rintanen**, J. (2012). Planning as satisfiability: Heuristics. *AII*, 193, 45–86.
- Rintanen**, J. (2016). Computational complexity in automated planning and scheduling. In *ICAPS-16*.
- Ripley**, B. D. (1996). *Pattern Recognition and Neural Networks*. Cambridge University Press.
- Rissanen**, J. (1984). Universal coding, information, prediction, and estimation. *IEEE Transactions on Information Theory*, IT-30, 629–636.
- Rissanen**, J. (2007). *Information and Complexity in Statistical Modeling*. Springer.
- Rivest**, R. (1987). Learning decision lists. *Machine Learning*, 2, 229–246.
- Robbins**, H. (1952). Some aspects of the sequential design of experiments. *Bulletin of the American Mathematical Society*, 58, 527–535.
- Robbins**, H. and Monro, S. (1951). A stochastic approximation method. *Annals of Mathematical Statistics*, 22, 400–407.
- Roberts**, L. G. (1963). Machine perception of three-dimensional solids. Technical report, MIT Lincoln Laboratory.
- Robertson**, N. and Seymour, P. D. (1986). Graph minors. II. Algorithmic aspects of tree-width. *J. Algorithms*, 7, 309–322.
- Robertson**, S. E. and Sparck Jones, K. (1976). Relevance weighting of search terms. *J. American Society for Information Science*, 27, 129–146.
- Robins**, J. (1986). A new approach to causal inference in mortality studies with a sustained exposure period: Application to control of the healthy worker survivor effect. *Mathematical Modelling*, 7, 1393–1512.
- Robinson**, A. and Voronkov, A. (Eds.). (2001). *Handbook of Automated Reasoning*. Elsevier.
- Robinson**, J. A. (1965). A machine-oriented logic based on the resolution principle. *JACM*, 12, 23–41.
- Robinson**, S. (2002). Computer scientists find unexpected depths in airfare search problem. *SIAM News*, 35(6).
- Roche**, E. and Schabes, Y. (Eds.). (1997). *Finite-State Language Processing*. Bradford Books.
- Rock**, I. (1984). *Perception*. W. H. Freeman.
- Rokicki**, T., Kociemba, H., Davidson, M., and Dethridge, J. (2014). The diameter of the Rubik's Cube group is twenty. *SIAM Review*, 56, 645–670.
- Rolf**, D. (2006). Improved bound for the PPSZ/Schöning-algorithm for 3-SAT. *Journal on Satisfiability, Boolean Modeling and Computation*, 1, 111–122.
- Rolnick**, D., Donti, P. L., Kaack, L. H., et al. (2019). Tackling climate change with machine learning. arXiv:1906.05433.
- Rolnick**, D. and Tegmark, M. (2018). The power of deeper networks for expressing natural functions. In *ICLR-18*.
- Romanovskii**, I. (1962). Reduction of a game with complete memory to a matrix game. *Soviet Mathematics*, 3, 678–681.
- Ros**, G., Sellart, L., Materzynska, J., Vazquez, D., and Lopez, A. M. (2016). The SYNTHIA dataset: A large collection of synthetic images for semantic segmentation of urban scenes. In *CVPR-16*.
- Rosenblatt**, F. (1957). The perceptron: A perceiving and recognizing automaton. Report, Project PARA, Cornell Aeronautical Laboratory.
- Rosenblatt**, F. (1960). On the convergence of reinforcement procedures in simple perceptrons. Report, Cornell Aeronautical Laboratory.
- Rosenblatt**, F. (1962). *Principles of Neurodynamics: Perceptrons and the Theory of Brain Mechanisms*. Spartan.
- Rosenblatt**, M. (1956). Remarks on some nonparametric estimates of a density function. *Annals of Mathematical Statistics*, 27, 832–837.
- Rosenblueth**, A., Wiener, N., and Bigelow, J. (1943). Behavior, purpose, and teleology. *Philosophy of Science*, 10, 18–24.

- Rosenschein**, J. S. and Zlotkin, G. (1994). *Rules of Encounter*. MIT Press.
- Rosenschein**, S. J. (1985). Formal theories of knowledge in AI and robotics. *New Generation Computing*, 3, 345–357.
- Ross**, G. (2012). Fisher and the millionaire: The statistician and the calculator. *Significance*, 9, 46–48.
- Ross**, S. (2015). *A First Course in Probability* (9th edition). Pearson.
- Ross**, S., Gordon, G., and Bagnell, D. (2011). A reduction of imitation learning and structured prediction to no-regret online learning. In *AISTATS-11*.
- Rossi**, F., van Beek, P., and Walsh, T. (2006). *Handbook of Constraint Processing*. Elsevier.
- Roth**, D. (1996). On the hardness of approximate reasoning. *AIJ*, 82, 273–302.
- Roussel**, P. (1975). Prolog: Manual de référence et d'utilisation. Tech. rep., Groupe d'Intelligence Artificielle, Université d'Aix-Marseille.
- Rowat**, P. F. (1979). *Representing the Spatial Experience and Solving Spatial Problems in a Simulated Robot Environment*. Ph.D. thesis, University of British Columbia.
- Roweis**, S. T. and Ghahramani, Z. (1999). A unifying review of linear Gaussian models. *Neural Computation*, 11, 305–345.
- Rowley**, H., Baluja, S., and Kanade, T. (1998). Neural network-based face detection. *PAMI*, 20, 23–38.
- Roy**, N., Gordon, G., and Thrun, S. (2005). Finding approximate POMDP solutions through belief compression. *JAIR*, 23, 1–40.
- Rubin**, D. (1974). Estimating causal effects of treatments in randomized and nonrandomized studies. *Journal of Educational Psychology*, 66, 688–701.
- Rubin**, D. (1988). Using the SIR algorithm to simulate posterior distributions. In Bernardo, J. M., de Groot, M. H., Lindley, D. V., and Smith, A. F. M. (Eds.), *Bayesian Statistics 3*. Oxford University Press.
- Rubinstein**, A. (1982). Perfect equilibrium in a bargaining model. *Econometrica*, 50, 97–109.
- Rubinstein**, A. (2003). Economics and psychology? The case of hyperbolic discounting. *International Economic Review*, 44, 1207–1216.
- Ruder**, S. (2018). NLP's ImageNet moment has arrived. *The Gradient*, July 8.
- Ruder**, S., Peters, M. E., Swayamdipta, S., and Wolf, T. (2019). Transfer learning in natural language processing. In *COLING-19*.
- Rumelhart**, D. E., Hinton, G. E., and Williams, R. J. (1986). Learning representations by back-propagating errors. *Nature*, 323, 533–536.
- Rumelhart**, D. E. and McClelland, J. L. (Eds.). (1986). *Parallel Distributed Processing*. MIT Press.
- Rummery**, G. A. and Niranjan, M. (1994). On-line Q-learning using connectionist systems. Tech. rep., Cambridge University Engineering Department.
- Ruspini**, E. H., Lowrance, J. D., and Strat, T. M. (1992). Understanding evidential reasoning. *IJAR*, 6, 401–424.
- Russakovsky**, O., Deng, J., Su, H., Krause, J., Satheesh, S., Ma, S., Huang, Z., Karpathy, A., Khosla, A., Bernstein, M., Berg, A. C., and Fei-Fei, L. (2015). ImageNet large scale visual recognition challenge. *IJCV*, 115, 211–252.
- Russell**, J. G. B. (1990). Is screening for abdominal aortic aneurysm worthwhile? *Clinical Radiology*, 41, 182–184.
- Russell**, S. J. (1985). The compleat guide to MRS. Report, Computer Science Department, Stanford University.
- Russell**, S. J. (1992). Efficient memory-bounded search methods. In *ECAI-92*.
- Russell**, S. J. (1998). Learning agents for uncertain environments. In *COLT-98*.
- Russell**, S. J. (1999). Expressive probability models in science. In *Proc. Second International Conference on Discovery Science*.
- Russell**, S. J. (2019). *Human Compatible*. Penguin.
- Russell**, S. J., Binder, J., Koller, D., and Kanazawa, K. (1995). Local learning in probabilistic networks with hidden variables. In *IJCAI-95*.
- Russell**, S. J. and Norvig, P. (2003). *Artificial Intelligence: A Modern Approach* (2nd edition). Prentice-Hall.
- Russell**, S. J. and Subramanian, D. (1995). Provably bounded-optimal agents. *JAIR*, 3, 575–609.
- Russell**, S. J. and Wefald, E. H. (1989). On optimal game-tree search using rational meta-reasoning. In *IJCAI-89*.
- Russell**, S. J. and Wefald, E. H. (1991). *Do the Right Thing: Studies in Limited Rationality*. MIT Press.
- Russell**, S. J. and Wolfe, J. (2005). Efficient belief-state AND-OR search, with applications to Kriegspiel. In *IJCAI-05*.
- Russell**, S. J. and Zimdars, A. (2003). Q-decomposition of reinforcement learning agents. In *ICML-03*.
- Rustagi**, J. S. (1976). *Variational Methods in Statistics*. Academic Press.
- Saad**, F. and Mansinghka, V. K. (2017). A probabilistic programming approach to probabilistic data analysis. In *NeurIPS 29*.
- Sabin**, D. and Freuder, E. C. (1994). Contradicting conventional wisdom in constraint satisfaction. In *ECAI-94*.
- Sabri**, K. E. (2015). Automated verification of role-based access control policies constraints using Prover9. arXiv:1503.07645.
- Sacerdoti**, E. D. (1974). Planning in a hierarchy of abstraction spaces. *AIJ*, 5, 115–135.
- Sacerdoti**, E. D. (1975). The nonlinear nature of plans. In *IJCAI-75*.
- Sacerdoti**, E. D. (1977). *A Structure for Plans and Behavior*. Elsevier.
- Sadeghi**, F. and Levine, S. (2016). CAD2RL: Real single-image flight without a single real image. arXiv:1611.04201.
- Sadigh**, D., Sastry, S., Seshia, S. A., and Dragan, A. D. (2016). Planning for autonomous cars that leverage effects on human actions. In *Proc. Robotics: Science and Systems*.
- Sadler**, M. and Regan, N. (2019). *Game Changer*. New in Chess.
- Sadri**, F. and Kowalski, R. (1995). Variants of the event calculus. In *ICLP-95*.
- Sagae**, K. and Lavie, A. (2006). A best-first probabilistic shift-reduce parser. In *COLING-06*.
- Sahami**, M., Hearst, M. A., and Saund, E. (1996). Applying the multiple cause mixture model to text categorization. In *ICML-96*.
- Sahin**, N. T., Pinker, S., Cash, S. S., Schomer, D., and Halgren, E. (2009). Sequential processing of lexical, grammatical, and phonological information within Broca's area. *Science*, 326, 445–449.
- Sakuta**, M. and Iida, H. (2002). AND/OR-tree search for solving problems with uncertainty: A case study using screen-shogi problems. *Trans. Inf. Proc. Society of Japan*, 43, 1–10.
- Salomaa**, A. (1969). Probabilistic and weighted grammars. *Information and Control*, 15, 529–544.
- Samadi**, M., Felner, A., and Schaeffer, J. (2008). Learning from multiple heuristics. In *AAAI-08*.
- Samet**, H. (2006). *Foundations of Multidimensional and Metric Data Structures*. Morgan Kaufmann.
- Sammut**, C., Hurst, S., Kedzior, D., and Michie, D. (1992). Learning to fly. In *ICML-92*.
- Samuel**, A. (1959). Some studies in machine learning using the game of checkers. *IBM Journal of Research and Development*, 3, 210–229.
- Samuel**, A. (1967). Some studies in machine learning using the game of checkers II—Recent progress. *IBM Journal of Research and Development*, 11, 601–617.
- Sanchez-Lengeling**, B., Wei, J. N., Lee, B. K., Gerkin, R. C., Aspuru-Guzik, A., and Wiltschko, A. B. (2019). Machine learning for scent: Learning generalizable perceptual representations of small molecules. arXiv:1910.10685.
- Sandholm**, T. (1999). Distributed rational decision making. In Weiß, G. (Ed.), *Multiagent Systems*. MIT Press.
- Sandholm**, T., Larson, K., Andersson, M., Shehory, O., and Tohmé, F. (1999). Coalition structure generation with worst case guarantees. *AIJ*, 111, 209–238.
- Sandholm**, T. (1993). An implementation of the contract net protocol based on marginal cost calculations. In *AAAI-93*.
- Sang**, T., Beame, P., and Kautz, H. (2005). Performing Bayesian inference by weighted model counting. In *AAAI-05*.
- Sapir**, E. (1921). *Language: An Introduction to the Study of Speech*. Harcourt Brace Jovanovich.
- Sarawagi**, S. (2007). Information extraction. *Foundations and Trends in Databases*, 1, 261–377.
- Sargent**, T. J. (1978). Estimation of dynamic labor demand schedules under rational expectations. *J. Political Economy*, 86, 1009–1044.
- Sartre**, J.-P. (1960). *Critique de la Raison dialectique*. Editions Gallimard.
- Satia**, J. K. and Lave, R. E. (1973). Markovian decision processes with probabilistic observation of states. *Management Science*, 20, 1–13.
- Sato**, T. and Kameya, Y. (1997). PRISM: A symbolic-statistical modeling language. In *IJCAI-97*.
- Saul**, L. K., Jaakkola, T., and Jordan, M. I. (1996). Mean field theory for sigmoid belief networks. *AIJ*, 4, 61–76.
- Saunders**, W., Sastry, G., Stuhlmüller, A., and Evans, O. (2018). Trial without error: Towards safe reinforcement learning via human intervention. In *AAMAS-18*.
- Savage**, L. J. (1954). *The Foundations of Statistics*. Wiley.
- Savva**, M., Kadian, A., MakSYMets, O., Zhao, Y., Wijmans, E., Jain, B., Straub, J., Liu, J., Kolton, V., Malik, J., Parikh, D., and Batra, D. (2019). Habitat: A platform for embodied AI research. arXiv:1904.01201.
- Sayre**, K. (1993). Three more flaws in the computational model. Paper presented at the APA (Central Division) Annual Conference, Chicago, Illinois.
- Schaeffer**, J. (2008). *One Jump Ahead: Computer Perfection at Checkers*. Springer-Verlag.
- Schaeffer**, J., Burch, N., Björnsson, Y., Kishimoto, A., Müller, M., Lake, R., Lu, P., and Suphen, S. (2007). Checkers is solved. *Science*, 317, 1518–1522.
- Schank**, R. C. and Abelson, R. P. (1977). *Scripts, Plans, Goals, and Understanding*. Lawrence Erlbaum.
- Schank**, R. C. and Riesbeck, C. (1981). *Inside Computer Understanding: Five Programs Plus Miniatures*. Lawrence Erlbaum.
- Schapire**, R. E. and Singer, Y. (2000). Boostexter: A boosting-based system for text categorization. *Machine Learning*, 39, 135–168.

- Schapire, R. E.** (1990). The strength of weak learnability. *Machine Learning*, 5, 197–227.
- Schapire, R. E.** (2003). The boosting approach to machine learning: An overview. In Denison, D. D., Hansen, M. H., Holmes, C., Mallick, B., and Yu, B. (Eds.), *Nonlinear Estimation and Classification*. Springer.
- Scharre, P.** (2018). *Army of None*. W. W. Norton.
- Schmid, C.** and Mohr, R. (1996). Combining grey-value invariants with local constraints for object recognition. In *CVPR-96*.
- Schmidhuber, J.** (2015). Deep learning in neural networks: An overview. *Neural Networks*, 61, 85–117.
- Schofield, M.** and Thielscher, M. (2015). Lifting model sampling for general game playing to incomplete-information models. In *AAAI-15*.
- Schölkopf, B.** and Smola, A. J. (2002). *Learning with Kernels*. MIT Press.
- Schöning, T.** (1999). A probabilistic algorithm for k-SAT and constraint satisfaction problems. In *FOCS-99*.
- Schoppers, M. J.** (1989). In defense of reaction plans as caches. *AIMag*, 10, 51–60.
- Schraudolph, N. N.**, Dayan, P., and Sejnowski, T. (1994). Temporal difference learning of position evaluation in the game of Go. In *NeurIPS 6*.
- Schriftwieser, J.**, Antonoglou, I., Hubert, T., Simonyan, K., Sifre, L., Schmitt, S., Guez, A., Lockhart, E., Hassabis, D., Graepel, T., Lillicrap, T., and Silver, D. (2019). Mastering Atari, Go, chess and shogi by planning with a learned model. arXiv:1911.08265.
- Schröder, E.** (1877). *Der Operationskreis des Logikkalküls*. B. G. Teubner, Leipzig.
- Schulman, J.**, Ho, J., Lee, A. X., Awwal, I., Bradlow, H., and Abbeel, P. (2013). Finding locally optimal, collision-free trajectories with sequential convex optimization. In *Proc. Robotics: Science and Systems*.
- Schulman, J.**, Levine, S., Abbeel, P., Jordan, M. I., and Moritz, P. (2015a). Trust region policy optimization. In *ICML-15*.
- Schulman, J.**, Levine, S., Moritz, P., Jordan, M., and Abbeel, P. (2015b). Trust region policy optimization. In *ICML-15*.
- Schultz, W.**, Dayan, P., and Montague, P. R. (1997). A neural substrate of prediction and reward. *Science*, 275, 1593.
- Schulz, D.**, Burgard, W., Fox, D., and Cremers, A. B. (2003). People tracking with mobile robots using sample-based joint probabilistic data association filters. *Int. J. Robotics Research*, 22, 99–116.
- Schulz, S.** (2004). System Description: E 0.81. In *Proc. International Joint Conference on Automated Reasoning*, Vol. 3097 of *LNAI*.
- Schulz, S.** (2013). System description: E 1.8. In *Proc. Int. Conf. on Logic for Programming Artificial Intelligence and Reasoning*.
- Schütze, H.** (1995). *Ambiguity in Language Learning: Computational and Cognitive Models*. Ph.D. thesis, Stanford University. Also published by CSLI Press, 1997.
- Schwartz, J. T.**, Scharir, M., and Hopcroft, J. (1987). *Planning, Geometry and Complexity of Robot Motion*. Ablex.
- Schwartz, S. P.** (Ed.). (1977). *Naming, Necessity, and Natural Kinds*. Cornell University Press.
- Scott, D.** and Krauss, P. (1966). Assigning probabilities to logical formulas. In Hintikka, J. and Suppes, P. (Eds.), *Aspects of Inductive Logic*. North-Holland.
- Searle, J. R.** (1980). Minds, brains, and programs. *BBS*, 3, 417–457.
- Searle, J. R.** (1990). Is the brain's mind a computer program? *Scientific American*, 262, 26–31.
- Searle, J. R.** (1992). *The Rediscovery of the Mind*. MIT Press.
- Sedgewick, R.** and Wayne, K. (2011). *Algorithms*. Addison-Wesley.
- Sefidgar, Y. S.**, Agarwal, P., and Cakmak, M. (2017). Situated tangible robot programming. In *HRI-17*.
- Segaran, T.** (2007). *Programming Collective Intelligence: Building Smart Web 2.0 Applications*. O'Reilly.
- Seipp, J.** and Röger, G. (2018). Fast downward stone soup 2018. *IPC 2018 Classical Track*.
- Seipp, J.**, Sievers, S., Helmert, M., and Hutter, F. (2015). Automatic configuration of sequential planning portfolios. In *AAAI-15*.
- Selman, B.**, Kautz, H., and Cohen, B. (1996). Local search strategies for satisfiability testing. In Johnson, D. S. and Trick, M. A. (Eds.), *Cliques, Coloring, and Satisfiability*. American Mathematical Society.
- Selman, B.** and Levesque, H. J. (1993). The complexity of path-based defeasible inheritance. *AIJ*, 62, 303–339.
- Selman, B.**, Levesque, H. J., and Mitchell, D. (1992). A new method for solving hard satisfiability problems. In *AAAI-92*.
- Seni, G.** and Elder, J. F. (2010). Ensemble methods in data mining: Improving accuracy through combining predictions. *Synthesis Lectures on Data Mining and Knowledge Discovery*, 2, 1–126.
- Seo, M.**, Kembhavi, A., Farhadi, A., and Hajishirzi, H. (2017). Bidirectional attention flow for machine comprehension. In *ICLR-17*.
- Shachter, R. D.** (1986). Evaluating influence diagrams. *Operations Research*, 34, 871–882.
- Shachter, R. D.** (1998). Bayes-ball: The rational pastime (for determining irrelevance and requisite information in belief networks and influence diagrams). In *UAI-98*.
- Shachter, R. D.**, D'Ambrosio, B., and Del Favero, B. A. (1990). Symbolic probabilistic inference in belief networks. In *AAAI-90*.
- Shachter, R. D.** and Kenley, C. R. (1989). Gaussian influence diagrams. *Management Science*, 35, 527–550.
- Shachter, R. D.** and Peot, M. (1989). Simulation approaches to general probabilistic inference on belief networks. In *UAI-98*.
- Shafer, G.** (1976). *A Mathematical Theory of Evidence*. Princeton University Press.
- Shanahan, M.** (1997). *Solving the Frame Problem*. MIT Press.
- Shanahan, M.** (1999). The event calculus explained. In Wooldridge, M. J. and Veloso, M. (Eds.), *Artificial Intelligence Today*. Springer-Verlag.
- Shanahan, M.** (2015). *The Technological Singularity*. MIT Press.
- Shani, G.**, Pineau, J., and Kaplow, R. (2013). A survey of point-based POMDP solvers. *Autonomous Agents and Multi-Agent Systems*, 27, 1–51.
- Shankar, N.** (1986). *Proof-Checking Metamathematics*. Ph.D. thesis, Computer Science Department, University of Texas at Austin.
- Shannon, C. E.** and Weaver, W. (1949). *The Mathematical Theory of Communication*. University of Illinois Press.
- Shannon, C. E.** (1950). Programming a computer for playing chess. *Philosophical Magazine*, 41, 256–275.
- Shapley, L. S.** (1953a). A value for  $n$ -person games. In Kuhn, H. W. and Tucker, A. W. (Eds.), *Contributions to the Theory of Games*. Princeton University Press.
- Shapley, S.** (1953b). Stochastic games. *PNAS*, 39, 1095–1100.
- Sharan, R. V.** and Moir, T. J. (2016). An overview of applications and advancements in automatic sound recognition. *Neurocomputing*, 200, 22–34.
- Shatkay, H.** and Kaelbling, L. P. (1997). Learning topological maps with weak local odometric information. In *IJCAI-97*.
- Shazeer, N.**, Mirhoseini, A., Maziarz, K., Davis, A., Le, Q. V., Hinton, G. E., and Dean, J. (2017). Outrageously large neural networks: The sparsely-gated mixture-of-experts layer. arXiv:1701.06538.
- Shelley, M.** (1818). *Frankenstein: Or, the Modern Prometheus*. Pickering and Chatto.
- Sheppard, B.** (2002). World-championship-caliber scrabble. *AIJ*, 134, 241–275.
- Shi, J.** and Malik, J. (2000). Normalized cuts and image segmentation. *PAMI*, 22, 888–905.
- Shieber, S.** (1994). Lessons from a restricted Turing test. *CACM*, 37, 70–78.
- Shieber, S.** (Ed.). (2004). *The Turing Test*. MIT Press.
- Shimony, S. E.** (1994). Finding MAPs for belief networks is NP-hard. *AIJ*, 68, 399–410.
- Shoham, Y.** (1993). Agent-oriented programming. *AIJ*, 60, 51–92.
- Shoham, Y.** (1994). *Artificial Intelligence Techniques in Prolog*. Morgan Kaufmann.
- Shoham, Y.** and Leyton-Brown, K. (2009). *Multiagent Systems: Algorithmic, Game-Theoretic, and Logical Foundations*. Cambridge Univ. Press.
- Shoham, Y.**, Powers, R., and Grenager, T. (2004). If multi-agent learning is the answer, what is the question? In *Proc. AAAI Fall Symposium on Artificial Multi-Agent Learning*.
- Shortliffe, E. H.** (1976). *Computer-Based Medical Consultations: MYCIN*. Elsevier.
- Siciliano, B.** and Khatib, O. (Eds.). (2016). *Springer Handbook of Robotics* (2nd edition). Springer-Verlag.
- Sigaud, O.** and Buffet, O. (2010). *Markov Decision Processes in Artificial Intelligence*. Wiley.
- Sigmund, K.** (2017). *Exact Thinking in Demented Times*. Basic Books.
- Silberstein, M.**, Weissbrod, O., Otten, L., Tzemach, A., Anisena, A., Shtark, O., Tuberg, D., Galfrin, E., Gannon, I., Shalata, A., Borochowitz, Z. U., Dechter, R., Thompson, E., and Geiger, D. (2013). A system for exact and approximate genetic linkage analysis of SNP data in large pedigrees. *Bioinformatics*, 29, 197–205.
- Silva, R.**, Melo, F. S., and Veloso, M. (2015). Towards table tennis with a quadrotor autonomous learning robot and onboard vision. In *IROS-15*.
- Silver, D.** and Veness, J. (2011). Monte-Carlo planning in large POMDPs. In *NeurIPS 23*.
- Silver, D.**, Huang, A., Maddison, C. J., Guez, A., and Hassabis, D. (2016). Mastering the game of Go with deep neural networks and tree search. *Nature*, 529, 484–489.
- Silver, D.**, Hubert, T., Schriftwieser, J., Antonoglou, I., Lai, M., Guez, A., Lanctot, M., Sifre, L., Kumaran, D., Graepel, T., et al. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362, 1140–1144.
- Silver, D.**, Schriftwieser, J., Simonyan, K., Antonoglou, I., Huang, A., Guez, A., Hubert, T., Baker, L., Lai, M., Bolton, A., Chen, Y., Lillicrap, T., Hui, F., Sifre, L., van den Driessche, G., Graepel, T., and Hassabis, D. (2017). Mastering the game of Go without human knowledge. *Nature*, 550, 354–359.
- Silverman, B. W.** (1986). *Density Estimation for Statistics and Data Analysis*. Chapman and Hall.

- Silverstein**, C., Henzinger, M., Marais, H., and Moricz, M. (1998). Analysis of a very large AltaVista query log. Tech. rep., Digital Systems Research Center.
- Simmons**, R. and Koenig, S. (1995). Probabilistic robot navigation in partially observable environments. In *IJCAI-95*.
- Simon**, D. (2006). *Optimal State Estimation: Kalman, H Infinity, and Nonlinear Approaches*. Wiley.
- Simon**, H. A. (1947). *Administrative Behavior*. Macmillan.
- Simon**, H. A. (1963). Experiments with a heuristic compiler. *JACM*, 10, 493–506.
- Simon**, H. A. and Newell, A. (1958). Heuristic problem solving: The next advance in operations research. *Operations Research*, 6, 1–10.
- Simon**, J. C. and Dubois, O. (1989). Number of solutions to satisfiability instances—Applications to knowledge bases. *AII*, 3, 53–65.
- Simonis**, H. (2005). Sudoku as a constraint problem. In *CP-05 Workshop on Modeling and Reformulating Constraint Satisfaction Problems*.
- Singer**, P. W. (2009). *Wired for War*. Penguin Press.
- Singh**, P., Lin, T., Mueller, E. T., Lim, G., Perkins, T., and Zhu, W. L. (2002). Open mind common sense: Knowledge acquisition from the general public. In *Proc. First International Conference on Ontologies, Databases, and Applications of Semantics for Large Scale Information Systems*.
- Sisbot**, E. A., Marin-Urias, L. F., Alami, R., and Simeon, T. (2007). A human aware mobile robot motion planner. *IEEE Transactions on Robotics*, 23, 874–883.
- Siskind**, J. M. and Pearlmuter, B. A. (2016). Efficient implementation of a higher-order language with built-in AD. arXiv:1611.03416.
- Sistla**, A. P. and Godefroid, P. (2004). Symmetry and reduced symmetry in model checking. *ACM Trans. Program. Lang. Syst.*, 26, 702–734.
- Sittler**, R. W. (1964). An optimal data association problem in surveillance theory. *IEEE Transactions on Military Electronics*, 8, 125–139.
- Skolem**, T. (1920). Logisch-kombinatorische Untersuchungen über die Erfüllbarkeit oder Beweisbarkeit mathematischer Sätze nebst einem Theoreme über die dichten Mengen. *Videnskapssekskapets skrifter, I. Matematisk-naturvidenskabelig klasse*, 4, 1–36.
- Skolem**, T. (1928). Über die mathematische Logik. *Norsk matematisk tidsskrift*, 10, 125–142.
- Slagle**, J. R. (1963). A heuristic program that solves symbolic integration problems in freshman calculus. *JACM*, 10.
- Slate**, D. J. and Atkin, L. R. (1977). CHESS 4.5—Northwestern University chess program. In Frey, P. W. (Ed.), *Chess Skill in Man and Machine*. Springer-Verlag.
- Slater**, E. (1950). Statistics for the chess computer and the factor of mobility. In *Symposium on Information Theory*. Ministry of Supply.
- Slocum**, J. and Sonneveld, D. (2006). *The 15 Puzzle*. Slocum Puzzle Foundation.
- Smallwood**, R. D. and Sondik, E. J. (1973). The optimal control of partially observable Markov processes over a finite horizon. *Operations Research*, 21, 1071–1088.
- Smith**, B. (2004). Ontology. In Floridi, L. (Ed.), *The Blackwell Guide to the Philosophy of Computing and Information*. Wiley-Blackwell.
- Smith**, B., Ashburner, M., Rosse, C., et al. (2007). The OBO Foundry: Coordinated evolution of ontologies to support biomedical data integration. *Nature Biotechnology*, 25, 1251–1255.
- Smith**, D. E., Genesereth, M. R., and Ginsberg, M. L. (1986). Controlling recursive inference. *AII*, 30, 343–389.
- Smith**, D. A. and Eisner, J. (2008). Dependency parsing by belief propagation. In *EMNLP-08*.
- Smith**, D. E. and Weld, D. S. (1998). Conformant Graphplan. In *AAAI-98*.
- Smith**, J. Q. (1988). *Decision Analysis*. Chapman and Hall.
- Smith**, J. E. and Winkler, R. L. (2006). The optimizer's curse: Skepticism and postdecision surprise in decision analysis. *Management Science*, 52, 311–322.
- Smith**, J. M. (1982). *Evolution and the Theory of Games*. Cambridge University Press.
- Smith**, J. M. and Szathmáry, E. (1999). *The Origins of Life: From the Birth of Life to the Origin of Language*. Oxford University Press.
- Smith**, M. K., Welty, C., and McGuinness, D. (2004). OWL web ontology language guide. Tech. rep., W3C.
- Smith**, R. G. (1980). *A Framework for Distributed Problem Solving*. UMI Research Press.
- Smith**, R. C. and Cheeseman, P. (1986). On the representation and estimation of spatial uncertainty. *Int. J. Robotics Research*, 5, 56–68.
- Smith**, S. J. J., Nau, D. S., and Throop, T. A. (1998). Success in spades: Using AI planning techniques to win the world championship of computer bridge. In *AAAI-98*.
- Smith**, W. E. (1956). Various optimizers for single-stage production. *Naval Research Logistics Quarterly*, 3, 59–66.
- Smolensky**, P. (1988). On the proper treatment of connectionism. *BBS*, 2, 1–74.
- Smolensky**, P. and Prince, A. (1993). Optimality theory: Constraint interaction in generative grammar. Tech. rep., Department of Computer Science, University of Colorado at Boulder.
- Smulyan**, R. M. (1995). *First-Order Logic*. Dover.
- Smyth**, P., Heckerman, D., and Jordan, M. I. (1997). Probabilistic independence networks for hidden Markov probability models. *Neural Computation*, 9, 227–269.
- Snoek**, J., Larochelle, H., and Adams, R. P. (2013). Practical Bayesian optimization of machine learning algorithms. In *NeurIPS 25*.
- Solomonoff**, R. J. (1964). A formal theory of inductive inference. *Information and Control*, 7, 1–22, 224–254.
- Solomonoff**, R. J. (2009). Algorithmic probability-theory and applications. In Emmert-Streib, F. and Dehmer, M. (Eds.), *Information Theory and Statistical Learning*. Springer.
- Sondik**, E. J. (1971). *The Optimal Control of Partially Observable Markov Decision Processes*. Ph.D. thesis, Stanford University.
- Sosic**, R. and Gu, J. (1994). Efficient local search with conflict minimization: A case study of the n-queens problem. *IEEE Transactions on Knowledge and Data Engineering*, 6, 661–668.
- Sowa**, J. (1999). *Knowledge Representation: Logical, Philosophical, and Computational Foundations*. Blackwell.
- Spaan**, M. T. J. and Vlassis, N. (2005). Perseus: Randomized point-based value iteration for POMDPs. *JAIR*, 24, 195–220.
- Sparrow**, R. (2004). The Turing triage test. *Ethics and Information Technology*, 6, 203–213.
- Spiegelhalter**, D. J., Dawid, A. P., Lauritzen, S., and Cowell, R. (1993). Bayesian analysis in expert systems. *Statistical Science*, 8, 219–282.
- Spirites**, P., Glymour, C., and Scheines, R. (1993). *Causation, Prediction, and Search*. Springer-Verlag.
- Spitkovsky**, V. I., Alshawi, H., and Jurafsky, D. (2010a). From baby steps to leapfrog: How less is more in unsupervised dependency parsing. In *NAACL HLT*.
- Spitkovsky**, V. I., Jurafsky, D., and Alshawi, H. (2010b). Profiting from mark-up: Hyper-text annotations for guided parsing. In *ACL-10*.
- Srivastava**, M. and Bickford, M. (1990). Formal verification of a pipelined microprocessor. *IEEE Software*, 7, 52–64.
- Srivastava**, N., Hinton, G. E., Krizhevsky, A., Sutskever, I., and Salakhutdinov, R. (2014a). Dropout: A simple way to prevent neural networks from overfitting. *JMLR*, 15, 1929–1958.
- Srivastava**, S., Russell, S. J., and Ruan, P. (2014b). First-order open-universe POMDPs. In *UAI-14*.
- Staab**, S. (2004). *Handbook on Ontologies*. Springer.
- Stallman**, R. M. and Sussman, G. J. (1977). Forward reasoning and dependency-directed backtracking in a system for computer-aided circuit analysis. *AII*, 9, 135–196.
- Stanfill**, C. and Waltz, D. (1986). Toward memory-based reasoning. *ACM*, 29, 1213–1228.
- Stanisławska**, K., Krawiec, K., and Vihma, T. (2015). Genetic programming for estimation of heat flux between the atmosphere and sea ice in polar regions. In *GECCO-15*.
- Stefik**, M. (1995). *Introduction to Knowledge Systems*. Morgan Kaufmann.
- Steiner**, D. F., MacDonald, R., Liu, Y., Truszkowski, P., Hipp, J. D., Gammie, C., Thng, F., Peng, L., and Stumpe, M. C. (2018). Impact of deep learning assistance on the histopathologic review of lymph nodes for metastatic breast cancer. *Am. J. Surgical Pathology*, 42, 1636–1646.
- Steinruecken**, C., Smith, E., Janz, D., Lloyd, J., and Ghahramani, Z. (2019). The Automatic Statistician. In Hutter, F., Kotthoff, L., and Vanschoren, J. (Eds.), *Automated Machine Learning*. Springer.
- Stergiou**, K. and Walsh, T. (1999). The difference all-difference makes. In *IJCAI-99*.
- Stickel**, M. E. (1992). A Prolog technology theorem prover: a new exposition and implementation in Prolog. *Theoretical Computer Science*, 104, 109–128.
- Stiller**, L. (1992). KQNKR. *J. International Computer Chess Association*, 15, 16–18.
- Stiller**, L. (1996). Multilinear algebra and chess endgames. In Nowakowski, R. J. (Ed.), *Games of No Chance, MSRI*, 29, 1996. Mathematical Sciences Research Institute.
- Stockman**, G. (1979). A minimax algorithm better than alpha-beta? *AII*, 12, 179–196.
- Stoffel**, K., Taylor, M., and Hendler, J. (1997). Efficient management of very large ontologies. In *AAAI-97*.
- Stone**, M. (1974). Cross-validatory choice and assessment of statistical predictions. *J. Royal Statistical Society, Series B*, 36, 111–133.
- Stone**, P. (2000). *Layered Learning in Multi-Agent Systems: A Winning Approach to Robotic Soccer*. MIT Press.
- Stone**, P. (2003). Multiagent competitions and research: Lessons from RoboCup and TAC. In Lima, P. U. and Rojas, P. (Eds.), *RoboCup-2002: Robot Soccer World Cup VI*. Springer Verlag.
- Stone**, P. (2016). What's hot at RoboCup. In *AAAI-16*.

- Stone, P.**, Brooks, R. A., Brynjolfsson, E., Calo, R., Etzioni, O., Hager, G., Hirschberg, J., Kalyanakrishnan, S., Kamar, E., Kraus, S., *et al.* (2016). Artificial intelligence and life in 2030. Tech. rep., Stanford University One Hundred Year Study on Artificial Intelligence: Report of the 2015–2016 Study Panel.
- Stone, P.**, Kaminka, G., and Rosenschein, J. S. (2009). Leading a best-response teammate in an ad hoc team. In *AAMAS Workshop in Agent Mediated Electronic Commerce*.
- Stone, P.**, Sutton, R. S., and Kuhlmann, G. (2005). Reinforcement learning for robocup soccer keepaway. *Adaptive Behavior*, 13, 165–188.
- Storvik, G.** (2002). Particle filters for state-space models with the presence of unknown static parameters. *IEEE Transactions on Signal Processing*, 50, 281–289.
- Strachey, C.** (1952). Logical or non-mathematical programmes. In *Proc. 1952 ACM National Meeting*.
- Stratonovich, R. L.** (1959). Optimum nonlinear systems which bring about a separation of a signal with constant parameters from noise. *Radiofizika*, 2, 892–901.
- Stratonovich, R. L.** (1965). On value of information. *Izvestiya of USSR Academy of Sciences, Technical Cybernetics*, 5, 3–12.
- Sturtevant, N. R.** and Bulitko, V. (2016). Scrubbing during learning in real-time heuristic search. *JAIR*, 57, 307–343.
- Subramanian, D.** and Wang, E. (1994). Constraint-based kinematic synthesis. In *Proc. International Conference on Qualitative Reasoning*.
- Suk, H.-I.**, Sin, B.-K., and Lee, S.-W. (2010). Hand gesture recognition based on dynamic Bayesian network framework. *Pattern Recognition*, 43, 3059–3072.
- Sun, Y.**, Wang, S., Li, Y., Feng, S., Tian, H., Wu, H., and Wang, H. (2019). ERNIE 2.0: A continual pre-training framework for language understanding. arXiv:1907.12412.
- Sussman, G. J.** (1975). *A Computer Model of Skill Acquisition*. Elsevier.
- Sutcliffe, G.** (2016). The CADE ATP system competition - CASC. *AIMag*, 37, 99–101.
- Sutcliffe, G.** and Sutttner, C. (1998). The TPTP Problem Library: CNF Release v1.2.1. *JAR*, 21, 177–203.
- Sutcliffe, G.**, Schulz, S., Claessen, K., and Gelder, A. V. (2006). Using the TPTP language for writing derivations and finite interpretations. In *Proc. International Joint Conference on Automated Reasoning*.
- Sutherland, I.** (1963). Sketchpad: A man-machine graphical communication system. In *Proc. Spring Joint Computer Conference*.
- Sutskever, I.**, Vinyals, O., and Le, Q. V. (2015). Sequence to sequence learning with neural networks. In *NeurIPS 27*.
- Sutton, C.** and McCallum, A. (2007). An introduction to conditional random fields for relational learning. In Getoor, L. and Taskar, B. (Eds.), *Introduction to Statistical Relational Learning*. MIT Press.
- Sutton, R. S.** (1988). Learning to predict by the methods of temporal differences. *Machine Learning*, 3, 9–44.
- Sutton, R. S.**, McAllester, D. A., Singh, S., and Mansour, Y. (2000). Policy gradient methods for reinforcement learning with function approximation. In *NeurIPS 12*.
- Sutton, R. S.** (1990). Integrated architectures for learning, planning, and reacting based on approximating dynamic programming. In *ICML-90*.
- Sutton, R. S.** and Barto, A. G. (2018). *Reinforcement Learning: An Introduction* (2nd edition). MIT Press.
- Swade, D.** (2000). *Difference Engine: Charles Babbage And The Quest To Build The First Computer*. Diane Publishing Co.
- Sweeney, L.** (2000). Simple demographics often identify people uniquely. *Health (San Francisco)*, 671, 1–34.
- Sweeney, L.** (2002a). Achieving k-anonymity privacy protection using generalization and suppression. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, 10, 571–588.
- Sweeney, L.** (2002b). k-anonymity: A model for protecting privacy. *International Journal of Uncertainty, Fuzziness and Knowledge-Based Systems*, 10, 557–570.
- Swerling, P.** (1959). First order error propagation in a stagewise smoothing procedure for satellite observations. *J. Astronautical Sciences*, 6, 46–52.
- Swift, T.** and Warren, D. S. (1994). Analysis of SLG-WM evaluation of definite programs. In *Logic Programming: Proc. 1994 International Symposium*.
- Szegedy, C.**, Zaremba, W., Sutskever, I., Bruna, J., Erhan, D., Goodfellow, I., and Fergus, R. (2013). Intriguing properties of neural networks. arXiv:1312.6199.
- Szeliski, R.** (2011). *Computer Vision: Algorithms and Applications*. Springer-Verlag.
- Szepesvari, C.** (2010). Algorithms for reinforcement learning. *Synthesis Lectures on Artificial Intelligence and Machine Learning*, 4, 1–103.
- Tadepalli, P.**, Givan, R., and Driessens, K. (2004). Relational reinforcement learning: An overview. In *ICML-04*.
- Tait, P. G.** (1880). Note on the theory of the “15 puzzle”. *Proc. Royal Society of Edinburgh*, 10, 664–665.
- Tamaki, H.** and Sato, T. (1986). OLD resolution with tabulation. In *ICLP-86*.
- Tan, P.**, Steinbach, M., Karpatne, A., and Kumar, V. (2019). *Introduction to Data Mining* (2nd edition). Pearson.
- Tang, E.** (2018). A quantum-inspired classical algorithm for recommendation systems. arXiv:1807.04271.
- Tarski, A.** (1935). Die Wahrheitsbegriff in den formalisierten Sprachen. *Studia Philosophica*, 1, 261–405.
- Tarski, A.** (1941). *Introduction to Logic and to the Methodology of Deductive Sciences*. Dover.
- Tarski, A.** (1956). *Logic, Semantics, Metamathematics: Papers from 1923 to 1938*. Oxford University Press.
- Tash, J. K.** and Russell, S. J. (1994). Control strategies for a stochastic planner. In *AAAI-94*.
- Tassa, Y.**, Doron, Y., Muldal, A., Erez, T., Li, Y., Casas, D. d. L., Budden, D., Abdolmaleki, A., Merel, J., Lefrancq, A., *et al.* (2018). Deepmind control suite. arXiv:1801.00690.
- Tate, A.** (1975a). Interacting goals and their use. In *IJCAI-75*.
- Tate, A.** (1975b). *Using Goal Structure to Direct Search in a Problem Solver*. Ph.D. thesis, University of Edinburgh.
- Tate, A.** (1977). Generating project networks. In *IJCAI-77*.
- Tate, A.** and Whiter, A. M. (1984). Planning with multiple resource constraints and an application to a naval planning problem. In *Proc. First Conference on AI Applications*.
- Tatman, J. A.** and Shachter, R. D. (1990). Dynamic programming and influence diagrams. *IEEE Transactions on Systems, Man and Cybernetics*, 20, 365–379.
- Tattersall, C.** (1911). *A Thousand End-Games: A Collection of Chess Positions That Can Be Won or Drawn by the Best Play*. British Chess Magazine.
- Taylor, A. D.** and Zwicker, W. S. (1999). *Simple Games: Desirability Relations, Trading, Pseudoweightings*. Princeton University Press.
- Taylor, G.**, Stensrud, B., Eitelman, S., and Dunham, C. (2007). Towards automating airspace management. In *Proc. Computational Intelligence for Security and Defense Applications (CISDA) Conference*.
- Taylor, P.** (2009). *Text-to-Speech Synthesis*. Cambridge University Press.
- Tedrake, R.**, Zhang, T. W., and Seung, H. S. (2004). Stochastic policy gradient reinforcement learning on a simple 3D biped. In *IROS-04*.
- Tellex, S.**, Kollar, T., Dickerson, S., Walter, M. R., Banerjee, A., Teller, S., and Roy, N. (2011). Understanding natural language commands for robotic navigation and mobile manipulation. In *AAAI-11*.
- Tenenbaum, J. B.**, Griffiths, T. L., and Niyogi, S. (2007). Intuitive theories as grammars for causal inference. In Gopnik, A. and Schulz, L. (Eds.), *Causal Learning: Psychology, Philosophy, and Computation*. Oxford University Press.
- Tesauro, G.** (1990). Neurogammon: A neural-network backgammon program. In *IJCNN-90*.
- Tesauro, G.** (1992). Practical issues in temporal difference learning. *Machine Learning*, 8, 257–277.
- Tesauro, G.** (1995). Temporal difference learning and TD-Gammon. *CACM*, 38, 58–68.
- Tesauro, G.** and Galperin, G. R. (1997). On-line policy improvement using Monte-Carlo search. In *NeurIPS 9*.
- Tetlock, P. E.** (2017). *Expert Political Judgment: How Good Is It? How Can We Know?* Princeton University Press.
- Teyssier, M.** and Koller, D. (2005). Ordering-based search: A simple and effective algorithm for learning Bayesian networks. In *UAI-05*.
- Thaler, R.** (1992). *The Winner's Curse: Paradoxes and Anomalies of Economic Life*. Princeton University Press.
- Thaler, R.** and Sunstein, C. (2009). *Nudge: Improving Decisions About Health, Wealth, and Happiness*. Penguin.
- Thayer, J. T.**, Dionne, A., and Ruml, W. (2011). Learning inadmissible heuristics during search. In *ICAPS-11*.
- Theocharous, G.**, Murphy, K., and Kaelbling, L. P. (2004). Representing hierarchical POMDPs as DBNs for multi-scale robot localization. In *ICRA-04*.
- Thiele, T.** (1880). Om anvendelse af mindste kvadraters metode i nogle tilfælde, hvor en komplikation af visse slags uensartede tilfældige fejlkilder giver fejlene en ‘systematisk’ karakter. *Vidensk. Selsk. Skr. 5. Rk., naturvid. og mat. Afd.*, 12, 381–408.
- Thielscher, M.** (1999). From situation calculus to fluent calculus: State update axioms as a solution to the inferential frame problem. *AIJ*, 111, 277–299.
- Thomas, P. S.**, da Silva, B. C., Barto, A. G., and Brunskill, E. (2017). On ensuring that intelligent machines are well-behaved. arXiv:1708.05448.
- Thomaz, A.**, Hoffman, G., Cakmak, M., *et al.* (2016). Computational human-robot interaction. *Foundations and Trends in Robotics*, 4, 105–223.
- Thompson, K.** (1986). Retrograde analysis of certain endgames. *J. International Computer Chess Association*, 9, 131–139.
- Thompson, K.** (1996). 6-piece endgames. *J. International Computer Chess Association*, 19, 215–226.
- Thompson, W. R.** (1933). On the likelihood that one unknown probability exceeds another in view of the evidence of two samples. *Biometrika*, 25, 285–294.
- Thorndike, E.** (1911). *Animal Intelligence*. Macmillan.

- Thornton**, C., Hutter, F., Hoos, H. H., and Leyton-Brown, K. (2013). Auto-WEKA: Combined selection and hyperparameter optimization of classification algorithms. In *KDD-13*.
- Thrun**, S., Burgard, W., and Fox, D. (2005). *Probabilistic Robotics*. MIT Press.
- Thrun**, S., Fox, D., and Burgard, W. (1998). A probabilistic approach to concurrent mapping and localization for mobile robots. *Machine Learning*, 31, 29–53.
- Thrun**, S. (2006). Stanley, the robot that won the DARPA Grand Challenge. *J. Field Robotics*, 23, 661–692.
- Thrun**, S. and Pratt, L. (2012). *Learning to Learn*. Springer.
- Thorstone**, L. L. (1927). A law of comparative judgment. *Psychological Review*, 34, 273–286.
- Tian**, J., Paz, A., and Pearl, J. (1998). Finding a minimal  $d$ -separator. Tech. rep., UCLA Department of Computer Science.
- Tikhonov**, A. N. (1963). Solution of incorrectly formulated problems and the regularization method. *Soviet Math. Dokl.*, 5, 1035–1038.
- Tipping**, M. E. and Bishop, C. M. (1999). Probabilistic principal component analysis. *J. Royal Statistical Society*, 61, 611–622.
- Titterington**, D. M., Smith, A. F. M., and Makov, U. E. (1985). *Statistical Analysis of Finite Mixture Distributions*. Wiley.
- Toma**, P. (1977). SYSTRAN as a multilingual machine translation system. In *Proc. Third European Congress on Information Systems and Networks: Overcoming the Language Barrier*.
- Tomasi**, C. and Kanade, T. (1992). Shape and motion from image streams under orthography: A factorization method. *IJCV*, 9, 137–154.
- Topol**, E. (2019). *Deep Medicine: How Artificial Intelligence Can Make Healthcare Human Again*. Basic Books.
- Torralba**, A., Fergus, R., and Weiss, Y. (2008). Small codes and large image databases for recognition. In *CVPR*.
- Torralba**, A., Linares López, C., and Borrajo, D. (2016). Abstraction heuristics for symbolic bidirectional search. In *IJCAI-16*.
- Tramèr**, F., Zhang, F., Juels, A., Reiter, M. K., and Risdenpart, T. (2016). Stealing machine learning models via prediction APIs. In *USENIX Security Symposium*.
- Tran**, D., Hoffman, M., Saurois, R. A., Brevdo, E., Murphy, K., and Blei, D. M. (2017). Deep probabilistic programming. In *ICLR-17*.
- Trappenberg**, T. (2010). *Fundamentals of Computational Neuroscience* (2nd edition). Oxford University Press.
- Tsang**, E. (1993). *Foundations of Constraint Satisfaction*. Academic Press.
- Tshitoyan**, V., Dagdelen, J., Weston, L., Dunn, A., Rong, Z., Kononova, O., Persson, K. A., Ceder, G., and Jain, A. (2019). Unsupervised word embeddings capture latent knowledge from materials science literature. *Nature*, 571, 95.
- Tsitsiklis**, J. N. and Van Roy, B. (1997). An analysis of temporal-difference learning with function approximation. *IEEE Transactions on Automatic Control*, 42, 674–690.
- Tukey**, J. W. (1977). *Exploratory Data Analysis*. Addison-Wesley.
- Tumer**, K. and Wolpert, D. (2000). Collective intelligence and Braess' paradox. In *AAAI-00*.
- Turian**, J., Ratinov, L., and Bengio, Y. (2010). Word representations: a simple and general method for semi-supervised learning. In *ACL-10*.
- Turing**, A. (1936). On computable numbers, with an application to the Entscheidungsproblem. *Proc. London Mathematical Society, 2nd series*, 42, 230–265.
- Turing**, A. (1948). Intelligent machinery. Tech. rep., National Physical Laboratory. reprinted in (Ince, 1992).
- Turing**, A. (1950). Computing machinery and intelligence. *Mind*, 59, 433–460.
- Turing**, A., Strachey, C., Bates, M. A., and Bowden, B. V. (1953). Digital computers applied to games. In Bowden, B. V. (Ed.), *Faster than Thought*. Pitman.
- Turing**, A. (1947). Lecture to the London Mathematical Society on 20 February 1947.
- Turing**, A. (1996). Intelligent machinery, a heretical theory. *Philosophia Mathematica*, 4, 256–260. Originally written c. 1951.
- Tversky**, A. and Kahneman, D. (1982). Causal schemata in judgements under uncertainty. In Kahneman, D., Slovic, P., and Tversky, A. (Eds.), *Judgement Under Uncertainty: Heuristics and Biases*. Cambridge University Press.
- Tygar**, J. D. (2011). Adversarial machine learning. *IEEE Internet Computing*, 15, 4–6.
- Ullman**, J. D. (1985). Implementation of logical query languages for databases. *ACM Transactions on Database Systems*, 10, 289–321.
- Ullman**, S. (1979). *The Interpretation of Visual Motion*. MIT Press.
- Urmon**, C. and Whittaker, W. (2008). Self-driving cars and the Urban Challenge. *IEEE Intelligent Systems*, 23, 66–68.
- Valiant**, L. (1984). A theory of the learnable. *CACM*, 27, 1134–1142.
- Vallati**, M., Chrpa, L., and Kitchin, D. E. (2015). Portfolio-based planning: State of the art, common practice and open challenges. *AI Commun.*, 28(4), 717–733.
- van Beek**, P. (2006). Backtracking search algorithms. In Rossi, F., van Beek, P., and Walsh, T. (Eds.), *Handbook of Constraint Programming*. Elsevier.
- van Beek**, P. and Chen, X. (1999). CPlan: A constraint programming approach to planning. In *AAAI-99*.
- van Beek**, P. and Manchak, D. (1996). The design and experimental analysis of algorithms for temporal reasoning. *JAIR*, 4, 1–18.
- van Benthem**, J. and ter Meulen, A. (1997). *Handbook of Logic and Language*. MIT Press.
- van den Oord**, A., Dieleman, S., and Schrauwen, B. (2014). Deep content-based music recommendation. In *NeurIPS 26*.
- van den Oord**, A., Dieleman, S., Zen, H., Simonyan, K., Vinyals, O., Graves, A., Kalchbrenner, N., Senior, A., and Kavukcuoglu, K. (2016a). WaveNet: A generative model for raw audio. arXiv:1609.03499.
- van den Oord**, A., Kalchbrenner, N., and Kavukcuoglu, K. (2016b). Pixel recurrent neural networks. arXiv:1601.06759.
- van Harmelen**, F., Lifschitz, V., and Porter, B. (2007). *The Handbook of Knowledge Representation*. Elsevier.
- van Heijenoort**, J. (Ed.). (1967). *From Frege to Gödel: A Source Book in Mathematical Logic, 1879–1931*. Harvard University Press.
- Van Hentenryck**, P., Saraswat, V., and Deville, Y. (1998). Design, implementation, and evaluation of the constraint language cc(FD). *J. Logic Programming*, 37, 139–164.
- Van Hoeve**, W.-J. (2001). The alldifferent constraint: a survey. In *6th Annual Workshop of the ERCIM Working Group on Constraints*.
- van Hoeve**, W.-J. and Katriel, I. (2006). Global constraints. In Rossi, F., van Beek, P., and Walsh, T. (Eds.), *Handbook of Constraint Processing*. Elsevier.
- van Lambalgen**, M. and Hamm, F. (2005). *The Proper Treatment of Events*. Wiley-Blackwell.
- van Nunen**, J. A. E. E. (1976). A set of successive approximation methods for discounted Markovian decision problems. *Zeitschrift für Operations Research, Serie A*, 20, 203–208.
- Van Roy**, P. L. (1990). Can logic programming execute as fast as imperative programming? Report, Computer Science Division, UC Berkeley.
- Vapnik**, V. N. (1998). *Statistical Learning Theory*. Wiley.
- Vapnik**, V. N. and Chervonenkis, A. Y. (1971). On the uniform convergence of relative frequencies of events to their probabilities. *Theory of Probability and Its Applications*, 16, 264–280.
- Vardi**, M. Y. (1996). An automata-theoretic approach to linear temporal logic. In Moller, F. and Birtwistle, G. (Eds.), *Logics for Concurrency*. Springer.
- Varian**, H. R. (1995). Economic mechanism design for computerized agents. In *USENIX Workshop on Electronic Commerce*.
- Vasilache**, N., Johnson, J., Mathieu, M., Chintala, S., Piantino, S., and LeCun, Y. (2014). Fast convolutional nets with fbfft: A GPU performance evaluation. arXiv:1412.7580.
- Vaswani**, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., Kaiser, L., and Polosukhin, I. (2018). Attention is all you need. In *NeurIPS 30*.
- Veach**, E. and Guibas, L. J. (1995). Optimally combining sampling techniques for Monte Carlo rendering. In *Proc. 22rd Annual Conference on Computer Graphics and Interactive Techniques (SIGGRAPH)*.
- Venkatesh**, S. (2012). *The Theory of Probability: Explorations and Applications*. Cambridge University Press.
- Vere**, S. A. (1983). Planning in time: Windows and durations for activities and goals. *PAMI*, 5, 246–267.
- Verma**, S. and Rubin, J. (2018). Fairness definitions explained. In *2018 IEEE/ACM International Workshop on Software Fairness*.
- Verma**, V., Gordon, G., Simmons, R., and Thrun, S. (2004). Particle filters for rover fault diagnosis. *IEEE Robotics and Automation Magazine*, June.
- Vinge**, V. (1993). The coming technological singularity: How to survive in the post-human era. In *Proc. Vision-21: Interdisciplinary Science and Engineering in the Era of Cyberspace*. NASA.
- Vinyals**, O., Babuschkin, I., Czarnecki, W. M., Mathieu, M., Dudzik, A., Chung, J., Choi, D. H., Powell, R., Ewalds, T., Georgiev, P., Hassabis, D., Apps, C., and Silver, D. (2019). Grandmaster level in StarCraft II using multi-agent reinforcement learning. *Nature*, 575, 350–354.
- Vinyals**, O., Ewalds, T., Bartunov, S., and Georgiev, P. (2017a). StarCraft II: A new challenge for reinforcement learning. arXiv:1708.04782.
- Vinyals**, O., Toshev, A., Bengio, S., and Erhan, D. (2017b). Show and tell: Lessons learned from the 2015 MSCOCO image captioning challenge. *PAMI*, 39, 652–663.
- Viola**, P. and Jones, M. (2004). Robust real-time face detection. *IJCV*, 57, 137–154.
- Visser**, U., Ribeiro, F., Ohashi, T., and Dellaert, F. (Eds.). (2008). *RoboCup 2007: Robot Soccer World Cup XI*. Springer.
- Viterbi**, A. J. (1967). Error bounds for convolutional codes and an asymptotically optimum decoding algorithm. *IEEE Transactions on Information Theory*, 13, 260–269.

- Vlassis**, N. (2008). *A Concise Introduction to Multi-agent Systems and Distributed Artificial Intelligence*. Morgan & Claypool.
- von Mises**, R. (1928). *Wahrscheinlichkeit, Statistik und Wahrheit*. J. Springer.
- von Neumann**, J. (1928). Zur Theorie der Gesellschaftsspiele. *Mathematische Annalen*, 100, 295–320.
- von Neumann**, J. and Morgenstern, O. (1944). *Theory of Games and Economic Behavior* (first edition). Princeton University Press.
- von Winterfeldt**, D. and Edwards, W. (1986). *Decision Analysis and Behavioral Research*. Cambridge University Press.
- Vossen**, T., Ball, M., Lotem, A., and Nau, D. S. (2001). Applying integer programming to AI planning. *Knowledge Engineering Review*, 16, 85–100.
- Wainwright**, M. and Jordan, M. I. (2008). Graphical models, exponential families, and variational inference. *Foundations and Trends in Machine Learning*, 1, 1–305.
- Walker**, G. (1931). On periodicity in series of related terms. *Proc. Roy. Soc. A*, 131, 518–532.
- Walker**, R. J. (1960). An enumerative technique for a class of combinatorial problems. In *Proc. Sympos. Appl. Math.*, Vol. 10.
- Wallace**, A. R. (1858). On the tendency of varieties to depart indefinitely from the original type. *Proc. Linnean Society of London*, 3, 53–62.
- Walpole**, R. E., Myers, R. H., Myers, S. L., and Ye, K. E. (2016). *Probability and Statistics for Engineers and Scientists* (9th edition). Pearson.
- Walsh**, T. (2015). Turing's red flag. arXiv:1510.09033.
- Waltz**, D. (1975). Understanding line drawings of scenes with shadows. In Winston, P. H. (Ed.), *The Psychology of Computer Vision*. McGraw-Hill.
- Wang**, A., Prukaschatkun, Y., Nangia, N., Singh, A., Michael, J., Hill, F., Levy, O., and Bowman, S. R. (2019). SuperGLUE: A stickier benchmark for general-purpose language understanding systems. arXiv:1905.00537.
- Wang**, A., Singh, A., Michael, J., Hill, F., Levy, O., and Bowman, S. (2018a). GLUE: A multi-task benchmark and analysis platform for natural language understanding. arXiv:1804.07461.
- Wang**, J., Zhu, T., Li, H., Hsueh, C.-H., and Wu, I.-C. (2018b). Belief-state Monte Carlo tree search for phantom Go. *IEEE Transactions on Games*, 10, 139–154.
- Wanner**, E. (1974). *On Remembering, Forgetting and Understanding Sentences*. Mouton.
- Warren**, D. H. D. (1974). WARPLAN: A System for Generating Plans. Department of Computational Logic Memo, University of Edinburgh.
- Warren**, D. H. D. (1983). An abstract Prolog instruction set. Technical note, SRI International.
- Wasserman**, L. (2004). *All of Statistics*. Springer.
- Watkins**, C. J. (1989). *Models of Delayed Reinforcement Learning*. Ph.D. thesis, Psychology Department, Cambridge University.
- Watson**, J. D. and Crick, F. (1953). A structure for deoxyribose nucleic acid. *Nature*, 171, 737.
- Wattenberg**, M., Viégas, F., and Johnson, I. (2016). How to use t-SNE effectively. *Distill*, 1.
- Waugh**, K., Schnizlein, D., Bowling, M., and Szafron, D. (2009). Abstraction pathologies in extensive games. In *AAMAS-09*.
- Weibull**, J. (1995). *Evolutionary Game Theory*. MIT Press.
- Weidenbach**, C. (2001). SPASS: Combining superposition, sorts and splitting. In Robinson, A. and Voronkov, A. (Eds.), *Handbook of Automated Reasoning*. MIT Press.
- Weiss**, G. (2000a). *Multiagent Systems*. MIT Press.
- Weiss**, Y. (2000b). Correctness of local probability propagation in graphical models with loops. *Neural Computation*, 12, 1–41.
- Weiss**, Y. and Freeman, W. (2001). Correctness of belief propagation in Gaussian graphical models of arbitrary topology. *Neural Computation*, 13, 2173–2200.
- Weizenbaum**, J. (1976). *Computer Power and Human Reason*. W. H. Freeman.
- Weld**, D. S. (1994). An introduction to least commitment planning. *AIMag*, 15, 27–61.
- Weld**, D. S. (1999). Recent advances in AI planning. *AIMag*, 20, 93–122.
- Weld**, D. S., Anderson, C. R., and Smith, D. E. (1998). Extending Graphplan to handle uncertainty and sensing actions. In *AAAI-98*.
- Weld**, D. S. and de Kleer, J. (1990). *Readings in Qualitative Reasoning about Physical Systems*. Morgan Kaufmann.
- Weld**, D. S. and Etzioni, O. (1994). The first law of robotics: A call to arms. In *AAAI-94*.
- Wellman**, M. P. (1985). Reasoning about preference models. Technical report, Laboratory for Computer Science, MIT.
- Wellman**, M. P. (1988). *Formulation of Tradeoffs in Planning under Uncertainty*. Ph.D. thesis, MIT.
- Wellman**, M. P. (1990a). Fundamental concepts of qualitative probabilistic networks. *AIJ*, 44, 257–303.
- Wellman**, M. P. (1990b). The STRIPS assumption for planning under uncertainty. In *AAAI-90*.
- Wellman**, M. P., Breese, J. S., and Goldman, R. (1992). From knowledge bases to decision models. *Knowledge Engineering Review*, 7, 35–53.
- Wellman**, M. P. and Doyle, J. (1992). Modular utility representation for decision-theoretic planning. In *ICAPS-92*.
- Wellman**, M. P., Wurman, P., O’Malley, K., Bangera, R., Lin, S., Reeves, D., and Walsh, W. (2001). Designing the market game for a trading agent competition. *IEEE Internet Computing*, 5, 43–51.
- Werbos**, P. (1974). *Beyond Regression: New Tools for Prediction and Analysis in the Behavioral Sciences*. Ph.D. thesis, Harvard University.
- Werbos**, P. (1990). Backpropagation through time: What it does and how to do it. *Proc. IEEE*, 78, 1550–1560.
- Werbos**, P. (1992). Approximate dynamic programming for real-time control and neural modeling. In White, D. A. and Sofge, D. A. (Eds.), *Handbook of Intelligent Control: Neural, Fuzzy, and Adaptive Approaches*. Van Nostrand Reinhold.
- Werbos**, P. (1977). Advanced forecasting methods for global crisis warning and models of intelligence. *General Systems Yearbook*, 22, 25–38.
- Wesley**, M. A. and Lozano-Perez, T. (1979). An algorithm for planning collision-free paths among polyhedral objects. *CACM*, 22, 560–570.
- West**, D. M. (2018). *The Future of Work: Robots, AI, and Automation*. Brookings Institution Press.
- West**, S. M., Whittaker, M., and Crawford, K. (2019). Discriminating systems: Gender, race and power in AI. Tech. rep., AI Now Institute.
- Wexler**, Y. and Meek, C. (2009). MAS: A multiplicative approximation scheme for probabilistic inference. In *NeurIPS 21*.
- Wheatstone**, C. (1838). On some remarkable, and hitherto unresolved, phenomena of binocular vision. *Phil. Trans. Roy. Soc.*, 2, 371–394.
- White**, C., Neiswanger, W., and Savani, Y. (2019). BANANAS: Bayesian optimization with neural architectures for neural architecture search. arXiv:1910.11858.
- Whitehead**, A. N. and Russell, B. (1910). *Principia Mathematica*. Cambridge University Press.
- Whittle**, P. (1979). Discussion of Dr Gittins’ paper. *J. Royal Statistical Society*, 41, 165.
- Whorf**, B. (1956). *Language, Thought, and Reality*. MIT Press.
- Widrow**, B. (1962). Generalization and information storage in networks of ADALINE “neurons”. In Yovits, M. C., Jacobi, G. T., and Goldstein, G. D. (Eds.), *Self-Organizing Systems*. Spartan.
- Widrow**, B. and Hoff, M. E. (1960). Adaptive switching circuits. In *IRE WESCON Convention Record*.
- Wiedijk**, F. (2003). Comparing mathematical provers. In *Proc. 2nd Int. Conf. on Mathematical Knowledge Management*.
- Wiegley**, J., Goldberg, K., Peshkin, M., and Brokowski, M. (1996). A complete algorithm for designing passive fences to orient parts. In *ICRA-96*.
- Wiener**, N. (1942). The extrapolation, interpolation, and smoothing of stationary time series. Tech. rep., Research Project DIC-6037, MIT.
- Wiener**, N. (1948). *Cybernetics*. Wiley.
- Wiener**, N. (1950). *The Human Use of Human Beings*. Houghton Mifflin.
- Wiener**, N. (1960). Some moral and technical consequences of automation. *Science*, 131, 1355–1358.
- Wiener**, N. (1964). *God & Golem, Inc: A Comment on Certain Points Where Cybernetics Impinges on Religion*. MIT Press.
- Wilensky**, R. (1978). *Understanding Goal-Based Stories*. Ph.D. thesis, Yale University.
- Wilkins**, D. E. (1988). *Practical Planning: Extending the AI Planning Paradigm*. Morgan Kaufmann.
- Wilkins**, D. E. (1990). Can AI planners solve practical problems? *Computational Intelligence*, 6, 232–246.
- Wilks**, Y. (2010). *Close Engagements With Artificial Companions: Key Social, Psychological, Ethical and Design Issues*. John Benjamins.
- Wilks**, Y. (2019). *Artificial Intelligence: Modern Magic or Dangerous Future*. Icon.
- Williams**, A., Nangia, N., and Bowman, S. (2018). A broad-coverage challenge corpus for sentence understanding through inference. In *NAACL HLT*.
- Williams**, B., Ingham, M., Chung, S., and Elliott, P. (2003). Model-based programming of intelligent embedded systems and robotic space explorers. *Proc. IEEE*, 91(212–237).
- Williams**, R. J. (1992). Simple statistical gradient-following algorithms for connectionist reinforcement learning. *Machine Learning*, 8, 229–256.
- Williams**, R. J. and Zipser, D. (1989). A learning algorithm for continually running fully recurrent neural networks. *Neural Computation*, 1, 270–280.
- Williams**, R. J. and Baird, L. C. I. (1993). Tight performance bounds on greedy policies based on imperfect value functions. Tech. rep., College of Computer Science, Northeastern University.
- Wilson**, D. H. (2011). *Robopocalypse*. Doubleday.
- Wilson**, R. A. and Keil, F. C. (Eds.). (1999). *The MIT Encyclopedia of the Cognitive Sciences*. MIT Press.
- Wilson**, R. (2004). *Four Colors Suffice*. Princeton University Press.

- Wilt**, C. M. and Ruml, W. (2014). Speedy versus greedy search. In *Seventh Annual Symposium on Combinatorial Search*.
- Wilt**, C. M. and Ruml, W. (2016). Effective heuristics for suboptimal best-first search. *JAIR*, 57, 273–306.
- Wingate**, D. and Seppi, K. D. (2005). Prioritization methods for accelerating MDP solvers. *JMLR*, 6, 851–881.
- Wingate**, D., Stuhlmüller, A., and Goodman, N. D. (2011). Lightweight implementations of probabilistic programming languages via transformational compilation. In *AISTATS-11*.
- Winograd**, S. and Cowan, J. D. (1963). *Reliable Computation in the Presence of Noise*. MIT Press.
- Winograd**, T. (1972). Understanding natural language. *Cognitive Psychology*, 3, 1–191.
- Winston**, P. H. (1970). Learning structural descriptions from examples. Technical report, Department of Electrical Engineering and Computer Science, MIT.
- Wintermute**, S., Xu, J., and Laird, J. (2007). SORTS: A human-level approach to real-time strategy AI. In *Proc. Third Artificial Intelligence and Interactive Digital Entertainment Conference*.
- Winternitz**, L. (2017). Autonomous navigation above the GNSS constellations and beyond: GPS navigation for the magnetospheric multiscale mission and SEXTANT pulsar navigation demonstration. Tech. rep., NASA Goddard Space Flight Center.
- Witten**, I. H. (1977). An adaptive optimal controller for discrete-time Markov environments. *Information and Control*, 34, 286–295.
- Witten**, I. H. and Bell, T. C. (1991). The zero-frequency problem: Estimating the probabilities of novel events in adaptive text compression. *IEEE Transactions on Information Theory*, 37, 1085–1094.
- Witten**, I. H. and Frank, E. (2016). *Data Mining: Practical Machine Learning Tools and Techniques* (4th edition). Morgan Kaufmann.
- Witten**, I. H., Moffat, A., and Bell, T. C. (1999). *Managing Gigabytes: Compressing and Indexing Documents and Images* (2nd edition). Morgan Kaufmann.
- Wittgenstein**, L. (1922). *Tractatus Logico-Philosophicus* (2nd edition). Routledge and Kegan Paul. Reprinted 1971, edited by D. F. Pears and B. F. McGuinness.
- Wittgenstein**, L. (1953). *Philosophical Investigations*. Macmillan.
- Wojciechowski**, W. S. and Wojcik, A. S. (1983). Automated design of multiple-valued logic circuits by automated theorem proving techniques. *IEEE Transactions on Computers*, C-32, 785–798.
- Wolfe**, J. and Russell, S. J. (2007). Exploiting belief state structure in graph search. In *ICAPS Workshop on Planning in Games*.
- Wolpert**, D. (2013). Ubiquity symposium: Evolutionary computation and the processes of life: what the no free lunch theorems really mean: how to improve search algorithms. *Ubiquity*, December, 1–15.
- Wolpert**, D. and Macready, W. G. (1997). No free lunch theorems for optimization. *IEEE Trans. Evolutionary Computation*, 1(1), 67–82.
- Wong**, C., Houlsby, N., Lu, Y., and Gesmundo, A. (2019). Transfer learning with neural AutoML. In *NeurIPS 31*.
- Woods**, W. A. (1973). Progress in natural language understanding: An application to lunar geology. In *AFIPS Conference Proceedings*.
- Woods**, W. A. (1975). What's in a link? Foundations for semantic networks. In Bobrow, D. G. and Collins, A. M. (Eds.), *Representation and Understanding: Studies in Cognitive Science*. Academic Press.
- Wooldridge**, M. (2009). *An Introduction to Multi-Agent Systems* (2nd edition). Wiley.
- Wooldridge**, M. and Rao, A. (Eds.). (1999). *Foundations of Rational Agency*. Kluwer.
- Wos**, L., Carson, D., and Robinson, G. (1964). The unit preference strategy in theorem proving. In *Proc. Fall Joint Computer Conference*.
- Wos**, L., Carson, D., and Robinson, G. (1965). Efficiency and completeness of the set-of-support strategy in theorem proving. *JACM*, 12, 536–541.
- Wos**, L., Overbeek, R., Lusk, E., and Boyle, J. (1992). *Automated Reasoning: Introduction and Applications* (2nd edition). McGraw-Hill.
- Wos**, L. and Robinson, G. (1968). Paramodulation and set of support. In *Proc. IRIA Symposium on Automatic Demonstration*.
- Wos**, L., Robinson, G., Carson, D., and Shalla, L. (1967). The concept of demodulation in theorem proving. *JACM*, 14, 698–704.
- Wos**, L. and Winker, S. (1983). Open questions solved with the assistance of AURA. In Bledsoe, W. W. and Loveland, D. (Eds.), *Automated Theorem Proving: After 25 Years*. American Mathematical Society.
- Wos**, L. and Pieper, G. (2003). *Automated Reasoning and the Discovery of Missing and Elegant Proofs*. Rinton Press.
- Wray**, R. E. and Jones, R. M. (2005). An introduction to Soar as an agent architecture. In Sun, R. (Ed.), *Cognition and Multi-Agent Interaction: From Cognitive Modeling to Social Simulation*. Cambridge University Press.
- Wright**, S. (1921). Correlation and causation. *J. Agricultural Research*, 20, 557–585.
- Wright**, S. (1931). Evolution in Mendelian populations. *Genetics*, 16, 97–159.
- Wright**, S. (1934). The method of path coefficients. *Annals of Mathematical Statistics*, 5, 161–215.
- Wu**, F. and Weld, D. S. (2008). Automatically refining the Wikipedia infobox ontology. In *17th World Wide Web Conference (WWW2008)*.
- Wu**, Y., Li, L., and Russell, S. J. (2016a). SWIFT: Compiled inference for probabilistic programming languages. In *IJCAI-16*.
- Wu**, Y., Schuster, M., Chen, Z., Le, Q. V., Norouzi, M., Macherey, W., Krikun, M., Cao, Y., Gao, Q., Macherey, K., et al. (2016b). Google's neural machine translation system: Bridging the gap between human and machine translation. arXiv:1609.08144.
- Wu**, Y. and He, K. (2018). Group normalization. arXiv:1803.08494.
- Xiong**, W., Wu, L., Alleva, F., Droppo, J., Huang, X., and Stolcke, A. (2017). The Microsoft 2017 conversational speech recognition system. arXiv:1708.06073.
- Yampolskiy**, R. V. (2018). *Artificial Intelligence Safety and Security*. Chapman and Hall/CRC.
- Yang**, G., Lin, Y., and Bhattacharya, P. (2010). A driver fatigue recognition model based on information fusion and dynamic Bayesian network. *Inf. Sci.*, 180, 1942–1954.
- Yang**, X.-S. (2009). Firefly algorithms for multimodal optimization. In *International Symposium on Stochastic Algorithms*.
- Yang**, X.-S. and Deb, S. (2014). Cuckoo search: Recent advances and applications. *Neural Computing and Applications*, 24, 169–174.
- Yang**, Z., Dai, Z., Yang, Y., Carbonell, J. G., Salakhutdinov, R., and Le, Q. V. (2019). XLNet: Generalized autoregressive pretraining for language understanding. arXiv:1906.08237.
- Yarowsky**, D. (1995). Unsupervised word sense disambiguation rivaling supervised methods. In *ACL-95*.
- Ye**, Y. (2011). The simplex and policy-iteration methods are strongly polynomial for the Markov decision problem with a fixed discount rate. *Mathematics of Operations Research*, 36, 593–784.
- Yedidia**, J., Freeman, W., and Weiss, Y. (2005). Constructing free-energy approximations and generalized belief propagation algorithms. *IEEE Transactions on Information Theory*, 51, 2282–2312.
- Yeo**, H.-S., Minami, R., Rodriguez, K., Shaker, G., and Quigley, A. (2018). Exploring tangible interactions with radar sensing. *Proc. ACM on Interactive, Mobile, Wearable and Ubiquitous Technologies*, 2, 1–25.
- Ying**, C., Kumar, S., Chen, D., Wang, T., and Cheng, Y. (2018). Image classification at supercomputer scale. arXiv:1811.06992.
- Yip**, K. M.-K. (1991). *KAM: A System for Intelligent Guiding Numerical Experimentation by Computer*. MIT Press.
- Yngve**, V. (1955). A model and an hypothesis for language structure. In Locke, W. N. and Booth, A. D. (Eds.), *Machine Translation of Languages*. MIT Press.
- Yob**, G. (1975). Hunt the wumpus! *Creative Computing*, Sep/Oct.
- Yoshikawa**, T. (1990). *Foundations of Robotics: Analysis and Control*. MIT Press.
- You**, Y., Pan, X., Wang, Z., and Lu, C. (2017). Virtual to real reinforcement learning for autonomous driving. arXiv:1704.03952.
- Young**, H. P. (2004). *Strategic Learning and Its Limits*. Oxford University Press.
- Young**, S., Gašić, M., Thompson, B., and Williams, J. (2013). POMDP-based statistical spoken dialog systems: A review. *Proc. IEEE*, 101, 1160–1179.
- Younger**, D. H. (1967). Recognition and parsing of context-free languages in time  $n^3$ . *Information and Control*, 10, 189–208.
- Yu**, D. and Deng, L. (2016). *Automatic Speech Recognition*. Springer-Verlag.
- Yu**, H.-F., Lo, H.-Y., Hsieh, H.-P., and Lou, J.-K. (2011). Feature engineering and classifier ensemble for KDD Cup 2010. In *Proc. KDD Cup 2010 Workshop*.
- Yu**, K., Sciuco, C., Jaggi, M., Musat, C., and Salzmann, M. (2019). Evaluating the search phase of neural architecture search. arXiv:1902.08142.
- Yudkowsky**, E. (2008). Artificial intelligence as a positive and negative factor in global risk. In Bostrom, N. and Cirkovic, M. (Eds.), *Global Catastrophic Risk*. Oxford University Press.
- Yule**, G. U. (1927). On a method of investigating periodicities in disturbed series, with special reference to Wolfer's sunspot numbers. *Phil. Trans. Roy. Soc. A*, 226, 267–298.
- Zadeh**, L. A. (1965). Fuzzy sets. *Information and Control*, 8, 338–353.
- Zadeh**, L. A. (1978). Fuzzy sets as a basis for a theory of possibility. *Fuzzy Sets and Systems*, 1, 3–28.
- Zaritskii**, V. S., Svetnik, V. B., and Shimelevich, L. I. (1975). Monte-Carlo technique in problems of optimal information processing. *Automation and Remote Control*, 36, 2015–22.
- Zeckhauser**, R. and Shepard, D. (1976). Where now for saving lives? *Law and Contemporary Problems*, 40, 5–45.
- Zeeberg**, A. (2017). D.I.Y. artificial intelligence comes to a Japanese family farm. *New Yorker*, August 10.
- Zelle**, J. and Mooney, R. (1996). Learning to parse database queries using inductive logic programming. In *AAAI-96*.
- Zemel**, R., Wu, Y., Swersky, K., Pitassi, T., and Dwork, C. (2013). Learning fair representations. In *ICML-13*.

- Zemelman**, B. V., Lee, G. A., Ng, M., and Miesenböck, G. (2002). Selective photostimulation of genetically chARGed neurons. *Neuron*, 33, 15–22.
- Zermelo**, E. (1913). Über Eine Anwendung der Mengenlehre auf die Theorie des Schachspiels. In *Proc. Fifth International Congress of Mathematicians*.
- Zermelo**, E. (1976). An application of set theory to the theory of chess-playing. *Fibibus News*, 6, 37–42. English translation of (Zermelo 1913).
- Zettlemoyer**, L. and Collins, M. (2005). Learning to map sentences to logical form: Structured classification with probabilistic categorial grammars. In *UAI-05*.
- Zhang**, C., Bengio, S., Hardt, M., Recht, B., and Vinyals, O. (2016). Understanding deep learning requires rethinking generalization. arXiv:1611.03530.
- Zhang**, H. and Stickel, M. E. (1996). An efficient algorithm for unit-propagation. In *Proc. Fourth International Symposium on Artificial Intelligence and Mathematics*.
- Zhang**, L., Pavlovic, V., Cantor, C. R., and Kasif, S. (2003). Human-mouse gene identification by comparative evidence integration and evolutionary analysis. *Genome Research*, 13, 1190–1202.
- Zhang**, N. L. and Poole, D. (1994). A simple approach to Bayesian network computations. In *Proc. 10th Canadian Conference on Artificial Intelligence*.
- Zhang**, S., Yao, L., and Sun, A. (2017). Deep learning based recommender system: A survey and new perspectives. arXiv:1707.07435.
- Zhang**, X., Zhao, J., and LeCun, Y. (2016). Character-level convolutional networks for text classification. In *NeurIPS 28*.
- Zhang**, Y., Pezeshki, M., Brakel, P., Zhang, S., Laurent, C., Bengio, Y., and Courville, A. (2017). Towards end-to-end speech recognition with deep convolutional neural networks. arXiv:1701.02720.
- Zhao**, K. and Huang, L. (2015). Type-driven incremental semantic parsing with polymorphism. In *NAAACL HLT*.
- Zhou**, K., Doyle, J., and Glover, K. (1995). *Robust and Optimal Control*. Pearson.
- Zhou**, R. and Hansen, E. (2002). Memory-bounded A\* graph search. In *Proc. 15th International FLAIRS Conference*.
- Zhou**, R. and Hansen, E. (2006). Breadth-first heuristic search. *AIJ*, 170, 385–408.
- Zhu**, B., Jiao, J., and Tse, D. (2019). Deconstructing generative adversarial networks. arXiv:1901.09465.
- Zhu**, D. J. and Latombe, J.-C. (1991). New heuristic algorithms for efficient hierarchical path planning. *IEEE Transactions on Robotics and Automation*, 7, 9–20.
- Zhu**, J.-Y., Park, T., Isola, P., and Efros, A. (2017). Unpaired image-to-image translation using cycle-consistent adversarial networks. In *ICCV-17*.
- Zhu**, M., Zhang, Y., Chen, W., Zhang, M., and Zhu, J. (2013). Fast and accurate shift-reduce constituent parsing. In *ACL-13*.
- Ziebart**, B. D., Maas, A. L., Dey, A. K., and Bagnell, J. A. (2008). Navigate like a cabbie: Probabilistic reasoning from observed context-aware behavior. In *Proc. 10th Int. Conf. on Ubiquitous Computing*.
- Ziebart**, B. D., Ratliff, N., Gallagher, G., Mertz, C., Peterson, K., Bagnell, J. A., Hebert, M., Dey, A. K., and Srinivasa, S. (2009). Planning-based prediction for pedestrians. In *IROS-09*.
- Zimmermann**, H.-J. (Ed.) (1999). *Practical Applications of Fuzzy Technologies*. Kluwer.
- Zimmermann**, H.-J. (2001). *Fuzzy Set Theory—And Its Applications* (4th edition). Kluwer.
- Zinkevich**, M., Johanson, M., Bowling, M., and Piccione, C. (2008). Regret minimization in games with incomplete information. In *NeurIPS 20*.
- Zipf**, G. (1935). *The Psychobiology of Language*. Houghton Mifflin.
- Zipf**, G. (1949). *Human Behavior and the Principle of Least Effort*. Addison-Wesley.
- Zobrist**, A. L. (1970). *Feature Extraction and Representation for Pattern Recognition and the Game of Go*. Ph.D. thesis, University of Wisconsin.
- Zollmann**, A., Venugopal, A., Och, F. J., and Ponte, J. (2008). A systematic comparison of phrase-based, hierarchical and syntax-augmented statistical MT. In *COLING-08*.
- Zoph**, B. and Le, Q. V. (2016). Neural architecture search with reinforcement learning. arXiv:1611.01578.
- Zuse**, K. (1945). The Plankalkül. Report, Gesellschaft für Mathematik und Datenverarbeitung.
- Zweig**, G. and Russell, S. J. (1998). Speech recognition with dynamic Bayesian networks. In *AAAI-98*.

# Index

Page numbers in **bold** refer to definitions of terms and algorithms.

Page numbers in *italics* refer to items in the bibliography.

## Symbols

$\alpha$  (alpha) learning rate, 816  
 $\alpha$  (alpha) normalization constant, 418  
 $\wedge$  (and), 235  
 $\chi^2$  (chi squared), 682  
 $\vdash$  (derives), 234  
 $\models$  (entailment), 232  
 $\varepsilon$  (epsilon)-ball, **691**  
 $\exists$  (there exists), 280  
 $\forall$  (for all), 279  
 $\gamma$  (gamma) discount factor, 632  
 $\sqsubset$  (gap) in sentence, 897  
 $|$  (given), 407  
 $\Leftrightarrow$  (if and only if), 235  
 $\Rightarrow$  (implies), 235  
 $\sim$  (indifferent), 520  
 $\lambda$  (lambda)-expression, **277**  
 $\nabla$  (nabla) gradient, 138  
 $\neg$  (not), 235  
 $\vee$  (or), 235  
 $\succ$  (preferred), 520  
 $\sigma$  (sigma) standard deviation, 1079  
 $T$  (matrix transpose), 1077

## A

$A(s)$  (actions in a state), 552  
A\* search, 103–108  
Aaronson, S., 1035, *1085*  
Aarts, E., 190, *1085*  
Aarup, M., 402, *1085*  
Abbas, A., 549, 550, *1085*  
Abbeel, P., 80, 516, 551, 638, 737, 841,  
    868, 872, 873, 985, 986, 1055,  
    1085, 1091, 1095, 1097, 1099,  
    1102, 1103, *1112*  
Abbott, L. F., 839, 873, *1092*  
ABC computer, 32  
Abdennadher, S., 191, *1095*  
abd James Andrew Bagnell, B. D. Z., 868,  
    *1101*  
Abdolmaleki, A., 873, *1114*  
Abelson, R. P., 41, *1111*  
Abney, S., 889, *1085*  
Aboian, M. S., 48, *1093*  
Abramson, B., 222, 587, *1085*  
Abreu, D., 637, *1085*  
absolute independence, 416, 419  
absorbing state, **850**  
abstraction, **84**, 218, 612  
abstraction hierarchy, **400**  
ABSTRIPS (planning system), 400

Abu-Hanna, A., 428, *1104*  
AC-3, **171**  
accessibility relation, **345**  
accountability, **729**  
accusative case, 892  
Acharya, A., 127, *1090*  
Achlioptas, D., 267, *1085*  
Ackerman, E., 47, *1085*  
Ackerman, N., 667, *1085*  
Ackley, D. H., 161, *1085*  
acoustic model  
    in disambiguation, 900  
ACT (cognitive architecture), 310  
acting rationally, 21  
action, 54, **83**, 123  
    egocentric, 85  
    high-level, **375**  
    irreversible, **154**, 850  
    joint, **593**  
    monitoring, 390, 391  
    primitive, **375**  
    rational, 25, 52  
action-utility function, **535**, 558  
action-utility learning, **841**  
action cost function, **83**, 123  
Action Description Language (ADL), 398  
action exclusion axiom, **263**, 594  
action monitoring, 390, 391  
action schema, **363**  
action sequence, 71, 84  
activation function, **803**  
active sensing, 988  
actor, **591**  
actor model, 636  
actuator, **54**, 61, **936**  
    electric, 936  
AD-tree, 798  
ADA BOOST, **720**  
adaline, 39  
Adams, A., 1067, *1104*  
Adams, J., 343  
Adams, M., 331, *1097*  
Adams, R. P., 690, *1113*  
adaptive control theory, 870  
adaptive dynamic programming (ADP),  
    **844**, 869  
adaptive perception, 945  
ADASYN (data generation system), 725,  
    1046  
add-one smoothing, 878  
additive decomposition (of value  
    functions), **861**  
add list, **363**  
Adida, B., 357, 1085  
adjustment formula, **471**  
ADL (Action Description Language), 398  
admissibility, 104  
admissible heuristic, **104**, 371  
Adolph, K.E., 852, *1085*  
Adorf, H.-M., 402, *1100*  
ADP (adaptive dynamic programming),  
    **844**, 869  
adversarial example, **821**, 838  
adversarial search, **192**  
adversarial training, 975  
adversary argument, **154**  
Advice Taker, 37  
AFSM (augmented finite state machine),  
    976  
Agarwal, P., 974, 986, 1041, *1105*, *1112*  
agent, 21, 54, 78  
    active learning, 848  
    architecture of, **65**, 1069  
    autonomous, 228  
    benevolent, 589  
    components, 1063–1069  
    decision-theoretic, 406, 518  
    function, **54**, 55, 65, 554  
    goal-based, 71–72, 78, 79  
    greedy, **848**  
    hybrid, **259**  
    impatient, 632  
    intelligent, 52  
    knowledge-based, 31, 226–228  
    learning, 74–76, 80  
    logical, 255–264, 297  
    model-based, **70**, 69–71  
    online planning, 397  
    personal, **1066**  
    problem-solving, **81**, 81–84  
    program, **55**, **65**, 66, 78  
    rational, **22**, 21–24, 54, **57**, 57–58, 73,  
        78, 79, 547  
    reflex, **67**, 67–69, 78, 554, 841  
    situated, 1033  
    software agent, **61**  
    taxi-driving, 75, 1070  
    utility-based, 72–74, 78  
vacuum, 57  
wumpus, 230, 288, 988  
Agerbeck, C., 189, *1085*

- Aggarwal, A., 838, **1110**  
 Aggarwal, G., 627, **1085**  
 aggregate querying, **1042**  
 aggregation, **394**  
 Agha, G., 636, **1085**  
 AGI (artificial general intelligence), **50**  
 Agichtein, E., 906, **1085**  
 Agmon, S., 836, **1085**  
 Agostinelli, F., 124, **1085**  
 Agrawal, P., 986, **1085**  
 Agrawal, R., 1060, **1086**  
 Agre, P. E., 401, **1085**  
 Ai, D., 48, **1104**  
 AI2 ARC (science test questions), 901  
 AI4People, 1059  
 AI FAIRNESS, 378, 1047  
 AI for Humanitarian Action, 1037  
 AI for Social Good, 1037  
 AI Habitat (simulated environment), 873  
 AI Index, **45**  
 Aila, T., 831, **1101**  
 AI Now Institute, 1046, 1059  
 Airborne Collision Avoidance System X (ACAS X), 588  
 aircraft carrier scheduling, 401  
 airport, driving to, 403  
 airport siting, 530, 535  
 AI safety, 1061  
 AI Safety Gridworlds, 873  
 AISB (Society for Artificial Intelligence and Simulation of Behaviour), 53  
 Aitken, S., 799, **1092**  
 AI winter, 42, 45  
 Aizerman, M., 735, **1085**  
 Akametalu, A. K., 872, **1085**  
 Akgun, B., 986, **1085**  
 al-Khwarizmi, M., 27  
 Alami, R., 986, **1113**  
 Albantakis, L., 1058, **1108**  
 Alberti, C., 904, **1085**  
 Alberti, L., 1027  
 Aldous, D., 160, **1085**  
 ALE (Arcade Learning Environment), 873  
 Alemi, A. A., 330, **1085**  
 Alexandria, 33  
 AlexNet (neural network system), 833  
 Algol-58, 903  
 algorithm, **27**  
 algorithmic complexity, 734  
 Algorithmic Justice League, 1046  
 Alhazen, 1027  
 Alibaba, 1068  
 Allais, M., 528, 550, **1085**  
 Allais paradox, 528, 550  
 Allan, J., 901, **1085**  
*Alldiff* constraint, 168  
 Allen, B., 401, **1095**  
 Allen, C., 550, **1091**  
 Allen, J. F., 342, 358, 401, 402, **1085**  
 Allen, P., 1061, **1085**  
 Allen-Zhu, Z., 837, **1085**  
 Alleva, F., 47, **1117**  
 alliance (in multiplayer games), **197**  
 Almulla, M., 126, **1108**  
 Alperin Resnick, L., 350, **1088**  
 alpha-beta pruning, **198**  
 alpha-beta search, 198–201, 220, 221  
 ALPHA-BETA-SEARCH, 200  
 ALPHAGO (Go program), 37, 45, 48, 222, 223, 867  
 ALPHASTAR (game-playing program), 218, 225  
 ALPHAZERO (game-playing program), 48, 218, 220, 223  
 Alshawi, H., 891, **1113**  
 Alterman, R., 400, **1085**  
 alternating offers bargaining model, **631**  
 Altman, A., 225, **1085**  
 altruism, 405  
 Alvey report, 41  
 ALVINN (autonomous vehicle), 974  
 Amarel, S., 162, 356, **1085**  
 Amazon, 47, 1068  
 ambient light, **993**  
 ambiguity, 270, **898**
  - lexical, **898**
  - semantic, **898**
  - syntactic, **898**, 904
 ambiguity aversion, **529**, 550  
 Amir, E., 267, 668, **1085**, 1092  
 Amit, Y., 736, **1085**  
 Amodei, D., 33, 930, 1061, 1069, **1085**, 1110  
 analogical reasoning, **768**  
 ANALOGY, 38  
 analysis of algorithms, **1074**  
 Analytical Engine, 33  
 analytical generalization, **768**  
 Anantharaman, T. S., 222, **1099**  
 Anbulagan, 266, **1103**  
 anchor box, 1007  
 anchoring effect, **529**  
 And-Elimination, **241**  
 AND-OR-SEARCH, **143**  
 AND-OR graph, 248  
 AND-OR tree, **141**  
 Andersen, S. K., 473, 474, **1085**  
 Anderson, C. R., 401, **1116**  
 Anderson, J. A., 836, **1098**  
 Anderson, J.R., 32, 310, 476, **1085**, 1109  
 Anderson, K., 126, 873, **1085**, 1087  
 Andersson, M., 638, **1111**  
 AND node, **141**  
 Andoni, A., 735, **1085**  
 Andor, D., 904, **1085**  
 Andre, D., 161, 586, 871, 873, **1085**, 1092, **1102**  
 Andraeae, P., 873, **1085**  
 Andrew, G., 1060, **1106**  
 Andrieu, C., 517, **1085**  
 Andrychowicz, M., 966, 986, **1085**  
 Aneja, J., 1016, **1085**  
 Angeli, G., 931, **1088**  
 ANGELIC-SEARCH, **382**  
 angelic semantics, 397  
 animatronics, **980**  
 Anisenia, A., 474, **1112**  
 answer set programming, 330  
 antecedent, 235  
 anthropomorphic robot, **933**, 937  
 Antonoglou, I., 45, 48, 201, 220, 223, 224, 835, 841, 871, 873, **1106**, 1112  
 anytime algorithm, **1070**  
 Aoki, M., 587, **1085**  
 aperture, **990**  
 apparent motion, 1000  
 Appel, K., 188, **1085**  
 Appelt, D., 905, **1085**  
 Apple, 1059  
 applicable, **83**, 363  
 apprenticeship learning, **864**, 1054, 1067  
 approval voting, **630**  
 approximate near-neighbors, **707**  
 Apps, C., 48, 225, **1115**  
 Apt, K. R., 189, 191, **1085**  
 Apté, C., 903, **1085**  
 Arbuthnot, J., 426, **1085**  
 Arcade Learning Environment (ALE), 873  
 arc consistency, **170**  
 Archibald, C., 225, **1085**  
 Architecture
  - agent, **65**, 1069
  - AI, 1069
  - cognitive, **52**, **310**
  - computer, 670
  - for speech recognition, 43
  - network, 819, 821, 838, 1029
  - reflective, **1070**
  - RNN, 912
  - rule-based, 310
  - subsumption, 976
  - transformer, 919, 931
 Arentoft, M. M., 402, **1085**  
 Arfaee, S.J., 127, **1085**, **1103**  
 Argall, B. D., 986, **1085**  
 argmax, 1077  
 argument
  - from disability, 1033–1034
  - from informality, 1032–1033
 Ariely, D., 528, 550, **1085**  
 ARISTO (question-answering system), 926, 927, 931  
 Aristotle, ix, 21, 24, 25, 29, 78, 79, 265, 296, 356, 357, 359, 733, 1027, 983

- arity, 275, 306  
 Arkin, R., 987, 1040, 1085  
 Armando, A., 268, 1085  
 Armstrong, S., 1053, 1085  
 Arnauld, A., 25, 28, 547, 1085  
 Arnould, S., 820, 1097  
 Arora, J. S., 137, 1105  
 Arora, N. S., 653, 667, 1085  
 Arora, S., 125, 1085  
 Arous, G. B., 837, 1090  
 Arpit, D., 734, 1086  
 Arrow's theorem, 630, 639  
 Arrow, K. J., 630, 639, 1086  
 artificial flight, 20  
 artificial general intelligence (AGI), 50  
 artificial intelligence, 1–1073
  - applications of, 45–48
  - conferences, 53
  - ethics, 1037–1052
  - foundations, 23–35, 870
  - future of, 49–52, 1063–1073
  - goals of, 1071–1072
  - history of, 35–45
  - journals, 53
  - philosophy of, 1032–1062
  - possibility of, 1032–1035
  - programming language, 37
  - provably beneficial, 23
  - real-time, 1070
  - risks, 49–52, 1038–1047
  - safety, 1052–1056
  - societies, 53
  - strong, 1032, 1056, 1057
  - weak, 1032, 1056, 1057
 artificial intelligence (AI), 1  
 artificial life, 161  
 artificial superintelligence (ASI), 51  
 Arulampalam, M. S., 517, 1086  
 Arulkumaran, K., 871, 1086  
 Arunachalam, R., 639, 1086  
 arXiv.org, 45, 839, 1069  
 Asada, M., 983, 1101  
 asbestos removal, 523  
 Ashburner, M., 334, 1113  
 Ashby, W. R., 34, 1086  
 Asimov, I., 982, 1058, 1086  
 ASKMSR (question-answering system), 901  
 Aspuru-Guzik, A., 1027, 1111  
 assertion (logical), 283  
 assertion (probabilistic), 406  
 assignment (in a CSP), 165  
 assistance, 551  
 assistance game, *see* game, assistance  
 assumption, 355  
 Astrom, K. J., 162, 587, 1086  
 asymmetry, 911  
 asymptotic analysis, 1075, 1074–1075  
 Atanasoff, J., 32
- Atari video game, 867  
 Athalye, A., 838, 1090  
 Atkeson, C.G., 586, 871, 986, 1086, 1107  
 Atkin, L. R., 125, 1113  
 Atlas (robot), 47atom, 278  
 atomic representation, 77, 81  
 atomic sentence, 235, 278, 278, 282  
 attention (neural net), 916, 917, 931  
 attentional sequence-to-sequence model, 917  
 attribute, 77  
 attribute-based learning algorithm, 759, 760  
 AUC (area under ROC curve), 728  
 auction, 624
  - ascending-bid, 624
  - English, 624
  - first-price, 626
  - sealed-bid, 626
    - second-price, 626
  - truth-revealing, 625
  - Vickrey, 626
 Audi, R., 1058, 1086  
 Auer, P., 587, 1086  
 Auer, S., 334, 357, 1088, 1103  
 augmentation, 902  
 augmented finite state machine (AFSM), 976  
 augmented grammar, 892  
 Aumann, R., 637, 1086  
 AURA (theorem prover), 327, 331  
 Auray, J. P., 549, 1087  
 Austerweil, J. L., 872, 1099  
 Australia, 165, 166, 177  
 author attribution, 877  
 AUTOCLASS (unsupervised learning algorithm), 799  
 autoencoder, 829
  - variational, 829
 automata, 1057  
 automated machine learning (AutoML), 737  
 automated reasoning, 20  
 automated taxi, 60, 61, 75, 228, 403, 1070  
 automatic differentiation, 807
  - reverse mode, 807
 Automatic Statistician, 737  
 AutoML, 737  
 automobile insurance, 529  
 Auton, L. D., 267, 1091  
 autonomatronics, 980  
 autonomic computing, 79  
 autonomous underwater vehicle (AUV), 934  
 autonomy, 60, 1031, 978  
 autoregressive model, 830, 838
  - deep, 830
 AUV (autonomous underwater vehicle), 934
- average pooling, 914  
 average reward, 557  
 Awwal, I., 985, 1112  
 Axelrod, R., 637, 1086  
 axiom, 227, 285
  - action exclusion, 263, 594
  - decomposability, 521
  - domain-specific, 334
  - effect axiom, 257
  - frame axiom, 257
  - Kolmogorov's, 411
  - of number theory, 286
  - of probability, 412
  - Peano, 286, 296, 307
  - precondition, 263
  - of probability, 411, 1078
  - of set theory, 287
  - successor-state, 258, 268
  - of utility theory, 521
  - wumpus world, 288
 axon, 30
- 
- B**
- b\* (branching factor), 116  
 B\* search, 221  
 Ba, J. L., 837, 1086  
 Baader, F., 330, 359, 1086  
 Babbage, C., 33, 221  
 Babuschkin, I., 48, 225, 1115  
 Bacchiani, M., 900, 1090  
 Bacchus, F., 188, 191, 428, 474, 559, 665, 1086  
 Bach, F. R., 903, 1099  
 bachelor, 336  
 Bachmann, L. M., 48, 1104  
 Bachmann, P. G. H., 1080, 1086  
 back-propagation, 40, 42, 208, 209, 806, 817, 836
  - through time, 825
 backgammon, 210, 224, 855, 866  
 background knowledge, 227, 320  
 backing up (in a search tree), 111, 195  
 backjumping, 179, 189  
 backmarking, 190  
 backoff model, 878  
**BACKTRACK**, 176  
 backtracking
  - chronological, 179
  - dependency-directed, 189
  - dynamic, 190
  - intelligent, 179–181, 252**BACKTRACKING-SEARCH**, 176  
 backtracking search, 98, 179–181, 183, 187  
 Backus, J. W., 903, 1086  
 Backus-Naur form (BNF), 1081  
 backward chaining, 248, 249–250, 265, 311–316, 329  
 backward induction, 604

- backward message, 487  
 backward search for planning, 368–369  
 Bacon, D., 1043, **1102**  
 Bacon, F., 24, 49, **1086**  
 Baeza-Yates, R., 901, **1086**  
 bag-of-words model, **875**, 877, 883, 902  
 Bagdasaryan, E., 1060, **1086**  
 bagging, **715**, 736  
 Bagnell, D., 986, **1111**  
 Bagnell, J. A., 868, 958, 970, 973, 985,  
     986, **1086**, **1101**, **1110**, **1118**  
 Bahdanau, D., 931, **1086**  
 Bahubalendruni, M. R., 125, **1086**  
 Bai, A., 873, **1086**  
 Bai, H., 588, **1086**  
 Bai, Y., 725, 985, 1046, **1088**, **1098**  
 Baidu, 901, 1031  
 Baird, L. C. I., 586, **1116**  
 Bajcsy, A., 974, **1086**  
 Baker, B., 966, 986, **1085**  
 Baker, C. L., 872, **1086**  
 Baker, J., 903, 905, **1086**  
 Baker, L., 45, 48, **1112**  
 Bakkes, S., 222, **1090**  
 Balaskas, K., 48, **1104**  
 Balch, T., 224, **1109**  
 Baldi, P., 124, 516, **1085**, **1086**  
 Baldwin, J.M., 136, 161, **1086**  
 Baldwin effect, **136**, 161  
 Ball, M., 399, **1116**  
 Ballard, B. W., 221, **1086**  
 Ballas, N., 734, **1086**  
 Baluja, S., 161, 1029, **1086**, **1111**  
 BANANAS (neural net architecture  
     search), 838  
 Bancilhon, F., 329, **1086**  
 bandit  
     Bernoulli, **574**  
     one-armed, **572**  
     problem, **571**, 587, 849  
     superprocess (BSP), **576**  
 Banerjee, A., 986, 1041, **1090**, **1114**  
 bang-bang control, **867**  
 Bangera, R., 639, **1116**  
 Banko, M., 44, 334, 357, 737, 901, 905,  
     906, **1086**, **1094**  
 Bansal, K., 327, **1086**  
 Bapna, A., 901, **1090**  
 Bar-Hillel, Y., 904, **1086**  
 Bar-Shalom, Y., 79, 515, 667, **1086**  
 Barber, D., 800, **1086**  
 Bard, N., 48, 224, **1107**  
 Barifaijo, E., 390, **1100**  
 Barnes, P., 1046, **1106**  
 Barr, A., 125, **1086**  
 Barreiro, J., 47, **1086**  
 Barreno, M., 1061, **1086**  
 Barrett, S., 225, **1086**  
 Barry, M., 473, **1099**  
 Barták, R., 189, 191, **1086**  
 Barthels, A., 666, **1100**  
 Bartholdi, J. J., 639, **1086**  
 Bartlett, F., 31  
 Bartlett, P., 872, **1086**  
 Barto, A. G., 163, 587, 588, 871–873,  
     **1086**, **1114**  
 Bartunov, S., 225, 873, **1115**  
 Barwise, J., 297, **1086**  
 baseline, 930, **1009**  
 base model, **714**  
 Basin, D. A., 222, **1095**  
 Basturk, B., 160, **1101**  
 Basye, K., 984, **1092**  
 batch normalization, **819**, 1030  
 Bates, E., 905, **1094**  
 Bates, M.A., 32, 221, **1115**  
 Batra, D., 873, 1017, **1097**, **1111**  
 Bauer, G., 331, **1097**  
 Baum, E., 134, 221, **1086**  
 Baum, L. E., 515, 799, **1086**  
 Baumert, L., 189, **1097**  
 Baxter, J., 872, **1086**  
 Bayardo, R. J., 190, 191, 266, 1060, **1086**  
 Baydin, A. G., 475, 668, **1103**  
 Bayen, A. M., 516, **1099**  
 Bayerl, S., 330, **1103**  
 Bayes' rule, 26, **417**, 417–418, 426  
 Bayes, T., 417, 428, 798, **1086**  
 Bayes-Nash equilibrium, **613**  
 Bayesian, 427  
 Bayesian classifier, 420  
 Bayesian learning, 719, **773**, 773–774,  
     797  
 Bayesian network, 43, **430**, 430–478, 799  
     continuous-time, 516  
     dynamic, **503**  
     hybrid, **440**, 472  
     inference in, 445–453  
     learning hidden variables in, 796–797  
     learning in, 785–786  
     multi-entity, 667  
     semantics, 432  
 Bayesian optimization, **690**  
 Bayesian reinforcement learning, **851**  
 BDD (binary decision diagram), 400  
 Beal, J., 50, **1087**  
 Beame, P., 474, **1111**  
 beam search, **110**, 124, **133**, 205, 882,  
     887, 919  
     local, **133**  
 Beardon, A. F., 548, **1087**  
 Beattie, C., 48, 225, 871, 873, **1087**,  
     **1100**, **1106**  
 Beber, G., 1057, **1094**  
 Bechhofer, R., 587, **1087**  
 Beck, J. C., 190, **1087**  
 Beckert, B., 330, **1087**  
 beer factory scheduling, 401  
 Beeri, C., 190, **1087**  
 beetle, dung, 59, 80, 392, 977  
 Beetz, M., 666, **1100**  
 behavioral cloning, 872, **973**  
 behaviorism, **31**, 34  
 BEINGS (multiagent system), 636  
 Bekey, G., 987, **1087**  
 belief, 344  
     degree of, 403, **404**, 412  
     desires and, 518–519  
     propagation, **476**  
         loopy, 476  
         revision, **353**  
         update, 353  
     belief network, *see* Bayesian network  
     belief state, **140**, 259, 383, 403, 406  
         in game theory, 610  
         probabilistic, 479, **483**  
         wiggly, 261  
 Belkin, M., 734, **1087**  
 Bell, C., 377, 401, **1087**  
 Bell, D. A., 798, **1090**  
 Bell, J. L., 297, **1087**  
 Bell, T.C., 902, 903, **1117**  
 Bellamy, E., **1087**  
 Bellamy, R. K. E., 1047, 1060, **1087**  
 BELLE (chess program), 222  
 Bellemare, M. G., 871–873, **1087**, **1095**,  
     **1106**, **1107**  
 Bellman, R. E., 28, 124, 125, 222, 223,  
     558, 586, 735, **1087**  
 Bellman equation, **558**  
 Bellman operator, 564  
 Bellman update, **563**  
 Ben-Tal, A., 161, **1087**  
 benchmark, 42, 930, 1029, **1074**  
 Bendix, P. B., 330, **1101**  
 benevolent agent assumption, **589**  
 Bengio, E., 734, **1086**  
 Bengio, S., 48, 734, 838, 930, **1087**,  
     **1115**, **1118**  
 Bengio, Y., 35, 734, 736, 737, 837–839,  
     900, 929, 931, 1067, **1086**, **1087**,  
     **1092**, **1096**, **1097**, **1103**, **1115**,  
     **1118**  
 Benjamin, M., 1040, **1087**  
 Bennett, B., 360, **1091**  
 Bennett, F.H., 161, **1102**  
 Bennett, J., 331, **1097**  
 Bennett, K., 737, **1097**  
 Bentham, J., 26, 548, **1087**  
 Benzmüller, C., 331, **1087**  
 Beresniak, A., 549, **1087**  
 Berg, A. C., 837, **1111**  
 Berger, H., 29  
 Berger, J. O., 475, 799, **1087**, **1108**  
 Berges, V., 873, **1100**  
 Bergstra, J., 737, **1087**  
 Berk, R., 1060, **1087**

- Berkeley, 696  
 Berkeley Parser, 904  
 Berkson, J., 473, *1087*  
 Berleur, J., 1059, *1087*  
 Berlin, K., 735, *1087*  
 Berliner, H. J., 221, 224, *1087*  
 Bermúdez-Chacón, R., 737, *1087*  
 Bernardo, J. M., 782, *1087*  
 Berners-Lee, T., 357, *1087*  
 Bernoulli, D., 25, 28, 525, 547, *1087*  
 Bernoulli, J., 26, 427  
 Bernstein, M., 837, *1111*  
 Bernstein, P. L., 429, *1087*  
 Berrada, L., 734, *1087*  
 Berrou, C., 476, *1087*  
 Berry, C., 32  
 Berry, D. A., 587, *1087*  
 BERT (natural language system), 930, 1072  
 Bertele, U., 474, *1087*  
 Bertoli, P., 400, 401, *1087*  
 Bertot, Y., 330, *1087*  
 Bertsekas, D., 79, 428, 588, 873, 1080, 1087  
 Bertsimas, D., 88, 734, *1087*  
 Berzuini, C., 517, *1096*  
 Beschastnikh, I., 1060, *1095*  
 Bessen, J., 1051, *1087*  
 Bessière, C., 189, *1087*  
 Best, N., 666, 798, *1105*  
 BEST-FIRST-SEARCH, **91**  
 best-first search, **91**, 123  
 best-fit function, 672  
 best possible prize, 523  
 best response, **598**  
 beta distribution, 505, **781**  
 Betancourt, M., 476, 668, 798, *1090*  
 Bethge, M., 1034, *1096*  
 Betlem, H., 390, *1100*  
 Betlem, J., 390, *1100*  
 Betteridge, J., 901, *1106*  
 betting game, 412  
 Beutel, A., 329, 1060, *1087*, *1102*  
 bfloat16, 33  
 Bhar, R., 516, *1087*  
 Bharath, A. A., 871, *1086*  
 Bhattacharya, P., 516, *1117*  
 bias  
     societal, **1043**  
 bias (statistical), **672**  
 bias (unfairness in outcomes), 724, 1043–1047  
 bias-variance tradeoff, **673**  
 Bibel, W., 330, 331, *1087*, *1103*  
 BIBF-SEARCH, **101**  
 Bible, 883  
 Bickford, M., 327, *1113*  
 biconditional, **235**  
 bicycle, 1061  
 bid, **623**  
 bidder, **624**  
 bidirectional RNN, **914**  
 bidirectional search, **100**, 114–115, 127  
 Bidlack, C., 983, *1091*  
 Bien, J., 737, *1087*  
 Biere, A., 267, *1087*  
 Bies, A., 903, *1087*  
 big computation, 1029  
 big data, **44**, 1029, 1066  
 BigDog, 47  
 Bigelow, J., 34, *1110*  
 Bigham, J., 905, *1108*  
 billiards, 225  
 Billings, D., 637, *1087*  
 Billingsley, P., 429, *1087*  
 Bilmes, J., 516, *1104*  
 Bimbo, J., 1027, *1105*  
 binary CSP, **168**  
 binary decision diagram (BDD), 400  
 binary resolution, **318**  
 Binder, J., 515, 516, 799, *1087*, *1111*  
 binding list, **284**  
 Bingham, E., 667, *1087*  
 Binmore, K., 637, *1087*  
 binocular stereopsis, **1009**, 1009–1010, 1028  
 binomial nomenclature, 357  
 bioinformatics, 903  
 biological naturalism, **1036**  
 Biran, O., 1060, *1088*  
 Birattari, M., 160, *1093*  
 Birbeck, M., 357, *1085*  
 Bischof, J., 1060, *1087*  
 Bishop, C. M., 160, 473, 735, 738, 799, 838, *1088*, *1115*  
 Bishop, M., 1058, *1109*  
 Bishop, R.H., 79, *1093*  
 Bisson, T., 1036, *1088*  
 Bistarelli, S., 188, *1088*  
 Biswal, B.B., 125, *1086*  
 Bitner, J. R., 189, *1088*  
 Bizer, C., 334, 357, *1088*, *1103*  
 Bjerager, P., 475, *1101*  
 Björnsson, Y., 223, *1111*  
 BKG (backgammon program), 224  
 Black, M., 1014, *1101*  
 BLACKBOX (planning system), 399  
 Blake, A., 516, *1100*  
 Blankspoor, K., 47, *1110*  
 Blau, H. M., 48, *1094*  
 Blazewicz, J., 402, *1088*  
 Blei, D. M., 667, 903, *1088*, *1099*, *1115*  
 Bliss, C. I., 473, *1088*  
 Blizzard, 873  
 Block, H.D., 39, *1088*  
 Block, N., 1058, *1088*  
 block sampling, **465**  
 blocks world, **38**, 360, 364  
 BLOG (probabilistic programming language), 667  
 Blondel, M., 738, *1109*  
 bluff, **217**  
 Blum, A. L., 399, 722, 736, 738, 906, 1088  
 Blum, C., 160, *1093*  
 Blumer, A., 735, *1088*  
 BNF (Backus-Naur form), **1081**  
 BO (bounded optimality), **1071**  
 Bobick, A., 516, *1100*  
 Bobrow, D. G., 38, *1088*  
 Bod, R., 891, *1088*  
 Boddington, P., 1059, *1088*  
 Boddy, M., 162, 400, 1070, *1092*, *1097*  
 Boden, M. A., 1058, *1088*  
 body (of Horn clause), **248**  
 Bojanowski, P., 903, *1100*  
 Bokeh (data analysis software), 727  
 Bolognesi, A., 222, *1088*  
 Bolton, A., 45, 48, 873, *1087*, *1112*  
 Bolton, R. J., 1041, *1088*  
 Boltzmann machine, **839**  
 Boltzmann rationality, **865**  
 Bonawitz, K., 667, 1043, *1088*, *1097*  
 Bond, A. H., 636, *1088*  
 Boneh, D., 134, *1086*  
 Bonet, B., 162, 398–401, 588, *1088*, *1098*  
 Bongard, J., 1057, *1109*  
 Boole, G., 26, 265, 427, *1088*  
 Boolean classification, 675  
 boosting, **717**  
 Booth, T. L., 903, *1088*  
 bootstrap, 715  
 Borda count, **630**  
 Bordes, A., 837, *1096*  
 Borel, E., 637, *1088*  
 Borenstein, J., 984, 985, *1088*  
 Borgida, A., 350, *1088*  
 Borgström, J., 667, 668, *1091*, *1097*  
 Borochowitz, Z. U., 474, *1112*  
 Boroditsky, L., 271, *1088*  
 Borrajo, D., 399, *1115*  
 Boser, B., 736, 1029, *1088*, *1103*  
 BOSS (autonomous vehicle), 979, 980, 984  
 Bosse, M., 984, *1088*  
 Boston Dynamics, ix  
 Bostrom, N., 51, 1061, *1088*  
 Botea, A., 399, *1098*  
 Bottou, L., 44, 735, 837, 1029, *1088*, *1103*  
 Boué, L., 734, *1088*  
 bounded-cost search, **110**  
 bounded optimality (BO), **1071**  
 bounded suboptimal search, **110**  
 bounding box, **1006**  
 bounds-consistent, **173**  
 bounds propagation, **173**

- Bousmalis, K., 985, *1088*  
 Bousquet, O., 735, 837, *1088*  
 Boutilier, C., 473, 549, 551, 587, 636, *1088*  
 Bouzy, B., 223, *1088*  
 Bowden, B. V., 32, 221, *1115*  
 Bower, G. H., 870, *1098*  
 Bowling, M., 48, 224, 637, 638, 873, *1087, 1088, 1107, 1116, 1118*  
 Bowman, D., 1036  
 Bowman, S., 928, 930, 931, *1088, 1097, 1116*  
 Bowman, S. R., 931, *1116*  
 Bowyer, K. W., 725, 1046, *1090*  
 Box, G. E. P., 161, 515, 798, 838, *1088*  
**BOXES** (reinforcement learning algorithm), *867*  
 Boyan, J. A., 160, *1088*  
 Boyce, M., 47, *1086*  
 Boyd, D., 1060, *1090*  
 Boyd, S., 161, *1088*  
 Boyden, E., 29, *1098*  
 Boyen, X., 517, *1088*  
 Boyen-Koller algorithm, *517*  
 Boyer, R.S., 327, 330, 331, *1088*  
 Boyer-Moore theorem prover, 330, 331  
 Boyko, A.S., 48, *1104*  
 Boyle, J., 331, *1117*  
 Boys, D., 225  
 Brachman, R. J., 350, 359, 361, *1088, 1103*  
 Bradlow, H., 985, *1112*  
 Bradt, R.N., 571, 587, *1088*  
 Bradtko, S.J., 163, 587, 871, *1086*  
 Brady, J. M., 516, *1107*  
 Brafman, O., 550, *1088*  
 Brafman, R., 550, *1088*  
 Brafman, R.I., 400, 401, 549, 636, 871, *1087, 1099*  
 Brahmagupta, 188  
 brain, 35, 801
  - computational power, 31
  - electronic super, 27
  - human, 29
  - imaging, **186**
 brain-machine interface, **29**, *978*  
 Braitenberg, V., 986, *1088*  
 Brakel, P., 900, *1118*  
 branch-and-bound, **126**, 395, 587  
 branching factor, **94**
  - effective, **116**, 125, 201
 Brandenburger, A., 637, *1086*  
 Brandt, F., 639, *1088*  
 Brants, T., 903, 931, *1088, 1095*  
 Bratko, I., 127, 330, *1088*  
 Bratman, M. E., 79, *1089*  
 Braverman, E., 735, *1085*  
**BREADTH-FIRST-SEARCH**, **95**  
 breadth-first search, **94**, 94–95, 123, 376  
 Breck, E., 731, 737, *1089*  
 Breese, J.S., 80, 473, 549, 666, 668, 1070, *1089, 1099, 1116*  
 Breiman, L., 716, 698, 736, *1089*  
 Brelaz, D., 189, *1089*  
 Brendel, W., 838, *1090*  
 Brent, R. P., 160, *1089*  
 Bresnan, J., 904, *1089*  
 Breuel, T., 930, *1104*  
 Brevdo, E., 667, *1115*  
 Brewka, G., 360, *1089*  
 Brickley, D., 357, *1089*  
 bridge (card game), 224  
**BRIDGE BARON**, 224  
 Briggs, R., 356, *1089*  
 brightness, 993  
 Brill, E., 44, 737, 901, 903, 905, *1086, 1089*  
 Brin, D., 906, *1089*  
 Brin, S., 905, *1089*  
 Bringsjord, S., 1057, *1089*  
 Brioschi, F., 474, *1087*  
 Britain, 40, 41  
 Broadbent, D. E., 32, *1089*  
 Broadhead, M., 901, 906, *1086*  
 Broca, P., 29  
 Brock, B., 330, *1099*  
 Brockman, G., 873, *1089*  
 Brokowski, M., 162, *1116*  
 Brooks, M. J., 1030, *1099*  
 Brooks, R. A., 45, 79, 267, 401, 976, 984, 986, 1057, 1061, *1089, 1109, 1114*  
 Brooks, S., 476, *1089*  
 Brouwer, P. S., 871, *1086*  
 Brown, C., 191, *1089*  
 Brown, E., 48, *1094*  
 Brown, J. S., 360, *1092*  
 Brown, K.C., 549, *1089*  
 Brown, M., 516, *1102*  
 Brown, N., 48, 224, 638, 841, *1089*  
 Brown, P.F., 929, 931, *1089*  
 Browne, C., 222, *1089*  
 Browning, B., 986, *1085*  
 Brownston, L., 329, *1089*  
 Brubaker, M., 476, 668, 798, *1090*  
 Bruce, V., 1031, *1089*  
 Brügmann, B., 223, 587, *1089*  
 Bruna, J., 838, *1114*  
 Brundage, M., 871, *1086*  
 Brunelleschi, F., 1027  
 Brunnstein, K., 1059, *1087*  
 Brunot, A., 44, 837, 1029, *1103*  
 Brunskill, E., 588, 872, *1110, 1114*  
 Bruynseels, A., 48, *1104*  
 Bryce, D., 162, 399, 401, *1089*  
 Brynjolfsson, E., 45, 1049, 1056, 1062, *1087, 1114*  
 Bryson, A. E., 40, 836, *1089*  
 Bryson, J. J., 1048, 1058, *1089*  
**BSP** (bandit superprocess), *576*  
 Buchanan, B. G., 40, 41, 80, 356, 477, *1089, 1094, 1104*  
 Buck, C., 903, *1089*  
 Budden, D., 873, *1114*  
 Buehler, M., 984, *1089*  
 Buffet, O., 588, *1112*  
 Buffon, G., 475, *1089*  
**BUGS** (probabilistic reasoning system), *476, 666*  
 Bui, P., 48, *1104*  
**BUILD** (planning system), *360*  
 Bulatov, Y., 820, *1097*  
 Bulfin, R., 639, *1110*  
 Bulitko, V., 163, *1114*  
 bunch, **337**  
 Bunt, H. C., 358, *1089*  
 Buolamwini, J., 1046, *1089*  
 Burch, N., 48, 223, 224, 637, *1087, 1088, 1107, 1111*  
 Burgard, W., 79, 667, 984, 985, 987, *1089, 1090, 1095, 1112, 1115*  
 Burget, L., 929, 930, *1106*  
 burglar alarm, 431–432  
 Burkov, A., 738, *1089*  
 Burns, C., 473, *1107*  
 Burns, E., 126, *1089*  
 Buro, M., 205, 224, *1089*  
 Burstein, J., 356, *1089*  
 Burton, R., 550, *1089*  
 Busbee, T. A., 1064, *1102*  
 business process automation, **1050**  
 Buss, D. M., 550, *1089*  
 Butler, S., 51, 1061, *1089*  
 Bylander, T., 402, *1089*  
 Byrd, R. H., 735, *1089*
- 
- C**
- c* (action cost), *83*  
**C-space**, **946**  
 C-space obstacle, **947**  
**C4** (Colossal Clean Crawled Corpus), *928, 930*  
**C4.5** (decision tree learning algorithm), *733*  
 Cabeza, R., 29, *1089*  
 Cabral, J., 357, *1105*  
 caching, 126, **259**, *648*  
 Cafarella, M. J., 901, 906, *1086, 1089, 1094*  
**CAFFE** (machine learning software), *1072*  
 Cai, S., 588, 731, 737, *1086, 1089*  
 Cain, A., 873, *1087*  
 Cajal, S., 29  
 Cakmak, M., 974, 986, *1085, 1112, 1114*  
 calculus, 138, 696, 805  
 calculus of variations, 160  
 Calo, R., 45, *1114*

- Calvanese, D., 359, 1086, 1089  
 Camacho, R., 872, 1089  
 Cambefort, Y., 80, 1098  
 Cambridge, 31  
 camera  
     for robots, 934  
     stereo, 934  
     surveillance, 1041  
     time-of-flight, 935  
 Campbell, D. E., 639, 1089  
 Campbell, M. S., 222, 1089, 1099  
 Campbell, W., 549, 1090  
 Candeal, J. C., 548, 1087  
 candidate elimination algorithm, 744  
 Cannings, C., 474, 1089  
 Canny, J., 985, 986, 1028, 1089  
 Canny edge detection, 1028  
 canonical distribution, 438  
 Cant, M., 873, 1087  
 Cantor, C. R., 474, 1118  
 Cantu-Paz, E., 161, 1109  
 Cao, Y., 47, 834, 901, 916, 1117  
 Čapek, K., 982, 1052  
 Capen, E., 549, 1090  
 Carbone, R., 268, 1085  
 Carbonell, J. G., 400, 930, 1090, 1117  
 Carbonnel, C., 191, 1090  
 Cardano, G., 26, 224, 426, 1090  
 card games, 217  
 Carlin, J. B., 799, 1096  
 Carlini, N., 838, 1090  
 Carlson, A., 271, 901, 1106  
 CARMEL (mobile robot), 983  
 Carnap, R., 25, 412, 427, 428, 1090  
 Carnegie Mellon University, 36  
 Carpenter, B., 476, 668, 798, 1090  
 Carpenter, M., 402, 1092  
 Carroll, S., 161, 1090  
 Carson, D., 330, 1117  
 CART, 683, 734  
 cart-pole problem, 867  
 Casas, D. d. L., 873, 1114  
 Casati, R., 358, 1090  
 case  
     accusative, 892  
     dative, 892  
     nominative, 892  
     objective, 892  
     subjective, 892  
 case-based reasoning, 769  
 Cash, S. S., 271, 1111  
 Cassandra, A. R., 586, 587, 1090, 1100, 1104  
 Cassandras, C. G., 79, 1090  
 Castaneda, A.G., 48, 225, 1100  
 Casteran, P., 330, 1087  
 Castro, R., 473, 1090  
 catastrophic forgetting, 856  
 categorical distribution, 409  
 category, 317–340, 335, 347  
 Cauchy, A., 735, 1090  
 causal network, 430, 467–471, 799  
 causal probability, 417  
 causal rule, 436  
 causation, 237, 419  
 Cawley, G. C., 737, 1097  
 Cazenave, T., 223, 1088  
 Ceder, G., 923, 1115  
 cell decomposition, 952  
 cell layout, 88  
 center (in mechanism design), 622  
 Center for Human-Compatible AI, 1059  
 Center for Humane Technology, 1059, 1066  
 central limit theorem, 1079  
 Centre for the Study of Existential Risk, 1059  
 cerebral cortex, 30  
 Cerf, V., 1062, 1108  
 Černocký, J., 929, 930, 1106  
 certainty effect, 528  
 certainty equivalent, 525  
 certainty factor, 41, 477  
 certification, 1047  
 Cesa-Bianchi, N., 587, 736, 1086, 1090  
 CGP (Conformant Graphplan), 400  
 CHAFF (logical reasoning system), 266  
 chain rule  
     for differentiation, 696, 703, 805  
     for probabilities, 434  
 Chajewska, U., 551, 1090  
 Chakrabarti, P.P., 127, 163, 1090, 1092  
 Chalkiadakis, G., 638, 1090  
 Chalmers, D. J., 1058, 1090  
 Chambers, R. A., 867, 1106  
 chance node (decision network), 535  
 chance node (game tree), 211  
 chance of winning, 202  
 Chandola, V., 1041, 1090  
 Chandra, A.K., 329, 1090  
 Chang, C.-L., 331, 1090  
 Chang, H. S., 587, 1090  
 Chang, K.-M., 271, 1106  
 Chang, K. C., 475, 1095  
 Chang, M.-W., 930, 1093  
 channel (in neural networks), 815  
 channel routing, 88  
 Chao, W.-L., 928, 1090  
 Chapman, D., 398, 401, 1085, 1090  
 Chapman, N., 124  
 character-level model, 877, 911  
 characteristic function, 616  
 Charniak, E., 41, 329, 473, 515, 666, 839, 903, 904, 1090  
 chart parser, 886, 902  
 Chaslot, G., 222, 1090  
 chatbot, 1035  
 Chater, N., 550, 1090  
 Chatfield, C., 515, 1090  
 Chatila, R., 984, 1107  
 Chauvin, Y., 516, 1086  
 Chavira, M., 474, 1090  
 Chawla, N. V., 725, 1046, 1090  
 checkers, 37, 80, 223, 870, 871  
 checkmate, 215  
 Cheeseman, P., 190, 472, 473, 799, 984, 1087, 1113  
 Chekaluk, R., 984, 1092  
 chemistry, 40, 356, 1034  
 Chen, D., 33, 904, 927, 930, 1069, 1090, 1104, 1117  
 Chen, J., 114, 126, 667, 1060, 1087, 1090, 1094  
 Chen, K., 909, 929, 1106  
 Chen, M. X., 901, 1090  
 Chen, R., 516, 1104  
 Chen, S. F., 903, 1090  
 Chen, T., 736, 1090  
 Chen, W., 904, 1118  
 Chen, X., 399, 873, 1095, 1115  
 Chen, Y., 45, 48, 1112  
 Chen, Z., 47, 834, 900, 901, 916, 1090, 1117  
 Cheng, J., 475, 798, 1090  
 Cheng, J.-F., 476, 1106  
 Cheng, Y., 33, 1069, 1117  
 Chernova, S., 986, 1085  
 Chervonenkis, A. Y., 735, 1115  
 chess, 23, 32, 39, 48, 64, 125, 193, 201–204, 222  
 Chess, D. M., 79, 1101  
 CHESS 4.5, 125  
 Cheung, V., 873, 1089  
 Chi, E.H., 329, 1060, 1087, 1102  
 $\chi^2$  pruning, 682  
 Chickering, D. M., 221, 798, 1098, 1102  
 Chien, S., 401, 1095  
 Child, R., 930, 1110  
 child node, 90  
 Chin, C.-S., 735, 1087  
 Chinese room, 1036, 1058  
 CHINOOK (checkers program), 223  
 Chintala, S., 828, 1069, 1110, 1115  
 Chiu, C., 900, 1090  
 Chklovski, T., 334, 1090  
 Cho, K., 837, 931, 1086, 1092  
 Chociej, M., 966, 986, 1085  
 Choi, D. H., 48, 225, 1115  
 Chollet, F., 330, 738, 1058, 1085, 1090  
 Chomsky, C., 904, 1097  
 Chomsky, N., 32, 34, 902, 903, 905, 906, 1090  
 Chomsky Normal Form, 887, 902  
 Choromanska, A., 837, 1090  
 Chorowski, J., 900, 1090  
 Choset, H., 79, 985, 987, 1090  
 Chouldechova, A., 1060, 1090

- Christian, B., 1057, 1090  
 Christiano, P., 1061, 1085  
 Christin, A., 1060, 1090  
 chronicle, 358  
 chronological backtracking, 179  
 Chrpa, L., 399, 1115  
 Chu-Carroll, J., 48, 1094  
 cHUGIN, 473  
 Chung, J., 48, 225, 1115  
 Chung, K. L., 1080, 1090  
 Chung, S., 267, 1116  
 chunking, 768  
 Church, A., 27, 297, 300, 328, 1090  
 Church, K., 902–904, 906, 1090  
 Church-Turing thesis, 27  
 Churchland, P. M., 1058, 1090  
 Ciancarini, P., 79, 222, 1088, 1091  
 Ciccolini, J., 1060, 1094  
 CIGOL program, 769  
 Cimatti, A., 399–401, 1087, 1091  
 circuit, 88, 161, 289
  - Boolean, 67, 70, 802
  - domain, 291–295
  - integrated, 128
  - verification, 266, 294
 circumscription, 351, 356, 359
  - prioritized, 352
 CiteSeer, 652, 653  
 city-block distance, 116  
 Claessen, K., 331, 1114  
 clairvoyance, 217  
 Clamp, S. E., 428, 1092  
 Clapp, R., 549, 1090  
 Clapp, T., 517, 1086  
 Claret, G., 668, 1091  
 Clark, A., 1033, 1057, 1091  
 Clark, C., 930, 1109  
 Clark, K. L., 360, 1091  
 Clark, P., 901, 927, 931, 1091  
 Clark, S., 904, 984, 1091, 1093  
 Clark completion, 360  
 Clarke, A. C., 472, 1091  
 Clarke, E., 399, 1091  
 CLASSIC (description logic), 350  
 classical planning, 362  
 classification (in description logic), 349  
 classification (in learning), 670  
 clause, 244  
 Clearwater, S. H., 639, 1091  
 Clerc, M., 160, 1093  
 Cleven, R., 667, 1106  
 Clocksin, W.F., 330, 1091  
 closed-loop, 82, 958  
 closed-world assumption, 282, 315, 356, 385, 643  
 closed class, 886  
 closed list, 90  
 CLP (constraint logic programming), 189, 316
  - CLP(R) (constraint logic programming system), 330
  - clustering (in Bayesian networks), 452, 452–453, 474
  - clustering (unsupervised learning), 671, 789
  - clutter (in data association), 657
  - CNF (conjunctive normal form), 244, 244–245, 265, 317–318
  - CNLP (conditional nonlinear planning), 401
  - CNN (convolutional neural network), 811, 1003
  - co-NP, 1076
  - co-NP-complete, 240, 1076
  - coalition, 616
  - coalition structure, 616
  - coalition structure graph, 621
  - coarse-to-fine search, 126
  - Coase, R. H., 639, 1091
  - coastal navigation, 964
  - Coates, A., 868, 872, 1055, 1091
  - Coates, M., 473, 1090
  - Cobham, A., 27, 1091
  - Cocke, A., 886
  - Cocke, J., 931, 1089
  - COCO (image data set), 43, 832, 1016
  - codes of ethics, 1059
  - coercion, 384
  - cognitive architecture, 52, 310
  - cognitive modeling, 20–21
  - cognitive psychology, 31, 874
  - cognitive science, 21, 667
  - Cohen, B., 267, 1112
  - Cohen, C., 983, 1091
  - Cohen, P. R., 42, 636, 1091
  - Cohen, W., 901, 1106
  - Cohn, A. G., 360, 1091
  - COLBERT (robot control language), 987
  - collaboration, 968, 971
  - Collin, Z., 191, 1091
  - Collins, M., 736, 896, 904, 905, 1085, 1087, 1118
  - collision checker, 953
  - collusion, 624
  - Colmerauer, A., 296, 329, 1091
  - Colombano, S.P., 161, 1104
  - color, 995
  - Colossal Clean Crawled Corpus (C4), 928, 930
  - Colossus, 32
  - Colton, S., 222, 1089
  - column player, 596
  - combinatorial auction, 628
  - combinatorial explosion, 39
  - commitment
    - epistemological, 273, 295, 404
    - ontological, 272, 295, 404
  - Common Crawl, 903, 922
  - common goal, 590
  - common sense, 426
  - common value, 624
  - communication, 595, 874
  - commutativity (in search problems), 175
  - Compagna, L., 268, 1085
  - COMPAS (expert system), 1044, 1060
  - competitive environment, 192
  - competitive ratio, 153
  - complementary literals, 244
  - complete-state formulation, 129
  - complete assignment, 165
  - complete data, 775
  - completeness
    - of a proof procedure, 234, 240, 264
    - of resolution, 244, 321–324
    - of a search algorithm, 93, 123
  - completing the square, 499
  - completion (of a data base), 315
  - complexity, 1074–1076
    - analysis, 1075
    - sample, 692
    - space, 93, 123
    - time, 93, 123
  - complex sentence, 235, 278
  - component (of mixture distribution), 790
  - composite decision process, 126
  - composite object, 336
  - compositionality, 269
  - compositional semantics, 894
  - computability, 27
  - computational learning theory, 690, 691
  - computational linguistics, 34, 904
  - computation graph, 805
  - computed torque control, 961
  - computer engineering, 32–33
  - computer vision, 30, 38, 186, 188, 989–1026
  - concession, 634
  - conclusion (of an implication), 235
  - concurrency, 591
  - concurrent action constraint, 594
  - condensation, 516
  - condition-action rule, 549
  - conditional distribution, 409
  - conditional distributions, 438–442
  - conditional effect, 386
  - conditional Gaussian, 441
  - conditional independence, 419, 424, 426, 433, 436–445, 472, 487
  - conditional plan, 128, 140, 141, 580
  - conditional probability, 407, 414, 417, 425, 434
  - conditional probability table (CPT), 431
  - conditioning, 414
  - conditioning case, 432
  - Condon, J.H., 222, 1091
  - Condorcet's Paradox, 629
  - configuration space, 946

- confirmation theory, **25**, 427  
 conflict-directed backjumping, **180**, 187  
 conflict clause learning, 252  
 conflict deal, **631**  
 conflict set, **179**  
 conformant planning, 383, 385–388, 397, 400  
 confusion matrix, **728**  
 Congdon, C. B., 983, *1091*  
 Conitzer, V., 639, *1088*  
 conjugate prior, **782**  
 conjunct, 235  
 conjunction (logic), **235**  
 conjunctive normal form (CNF), **244**, 244–245, 265, 317–318  
 conjunct ordering, **308**  
 Conlisk, J., 550, *1091*  
 Conneau, A., 930, *1103*  
 connected component, **183**  
 connectionism, **42**, 836  
 connective, logical, 35, 235, 264, 278  
 Connell, J., 987, *1091*  
 consciousness, 29, **1036**, 1058  
 consequent, 235  
 consequentialism, 26, **57**  
 conservative approximation, **261**, 387  
 consistency, **349**
  - arc, **170**
  - condition, 125
  - of a CSP assignment, **165**
  - of a heuristic, **106**
  - path, **172**, 188
 consistent estimation, **455**  
 consistent hypothesis, **671**  
 conspiracy number, 221  
 constant symbol, **275**, 277  
 constrained optimization problem, 139, **169**  
 constraint
  - binary, 168
  - global, **168**, 172
  - nonlinear, **167**
  - preference constraint, **169**
  - propagation, **169**, 169–175, 178–179
  - resource constraint, **173**
  - symmetry-breaking, **187**
  - unary, 168
 constraint-based generalization, **768**  
 constraint graph, **165**, 184  
 constraint hypergraph, **168**  
 constraint learning, **180**, 187, **190**  
 constraint logic programming, 316, 330  
 constraint logic programming (CLP), **189**, 316  
 constraint satisfaction problem (CSP), 38, **164**, 164–169, 370
  - finite-domain, **167**, 316
 constraint weighting, **182**  
 constructive induction algorithms, **760**
- consumable resource, **393**  
 context, 911  
 context (in computer vision), **1006**  
 context-free grammar, 884, 902, 903, 1081  
 context-specific independence, 645  
 Conti, E., 873, *1096*  
 Conti-Cook, C., 1060, *1094*  
 contingency planning, 383, 388–389, 397  
 continuity (of preferences), **520**  
 continuous domains, **167**  
 continuous values, 440  
 contour (of a state space), **107**  
 contraction mapping, **564**  
 contract net protocol, **622**  
 contradiction, **241**  
 control, **938**
  - bang-bang, **867**
  - lateral, 1024
  - longitudinal, 1024
 control law, **959**  
 controller, 22, 79, 833, 868, 959  
 control theory, **34**, 33–34, 79, 160, 398, 836, 867, **957**
  - adaptive, 870
  - robust, **852**
 control uncertainty, 964  
 convention, **595**  
 conversion to normal form, 317–318  
 convexity, 140  
 convex optimization, **140**, 159  
 CONVINCE (Bayesian expert system), 472  
 convolution, **997**  
 convolution (in neural networks), **811**  
 convolutional neural network (CNN), **811**, 1003  
 Conway, D., 738, *1091*  
 Cook, P. J., 1051, *1095*  
 Cook, S. A., 27, 266, 267, 1080, *1091*  
 Cooper, G., 475, 798, *1091*  
 Cooper, M. C., 191, *1090*  
 cooperation, 594–595  
 cooperative distributed problem solving, 636  
 cooperative game, **591**, 635  
 coordination problem, **590**, 598, 968  
 Copeland, J., 358, 1058, *1091*  
 Copeland, T. P., 48, *1093*  
 CoQ (theorem prover), 188, 330  
 Coram, M., 48, *1097*  
 Corbett-Davies, S., 1060, *1091*  
 core, **617**, 635, 638  
 Corkill, D. D., 636, *1103*  
 Cormen, T. H., 125, 1080, *1091*  
 Cornell, C. A., 475, *1101*  
 corpus, **876**  
 Corrado, G., 48, 909, 929, *1104*, *1106*  
 correspondence problem, **975**
- Cortes, C., 44, 736, 837, 1029, *1091*, *1103*  
 cost function, **34**  
 cost optimality, 123  
 cost to go, 962  
 cotraining, 906  
 Cotton, C., 1060, *1088*  
 counterparts, **590**  
 count noun, **340**  
 Cournot, A., 637, *1091*  
 Cournot competition, **612**  
 Courville, A., 734, 838, 839, 900, *1086*, *1097*, *1118*  
 covariance, **1080**  
 covariance matrix, 1079, **1080**  
 Cover, T., 738, *1091*  
 Cowan, J. D., 38, 836, *1091*, *1117*  
 Cowell, R., 549, 798, *1091*, *1113*  
 Cowhey, I., 901, *1091*  
 Cowling, P. I., 222, *1089*  
 Cox, I., 667, 983, *1091*  
 Cox, R. T., 412, 427, 428, *1091*  
 Coz, D., 48, *1104*  
 Cozmo (entertainment robot), 980  
 CPCS (medical diagnosis system), 439, 473  
 CPLAN (planning system), 399  
 CPT (conditional probability table), **431**  
 Craig, J., 985, *1091*  
 Craik, K., 31, *1091*  
 Crammer, K., 736, *1093*  
 Cramton, P., 639, *1091*  
 Crato, N., 190, *1097*  
 Craven, M., 905, *1091*  
 Crawford, J. M., 267, *1091*  
 Crawford, K., 1046, *1096*, *1116*  
 creativity, 34  
 credible threat, **609**  
 credit assignment, **858**  
 Cremers, A. B., 667, 984, *1089*, *1112*  
 Cresswell, M. J., 358, *1099*  
 Crick, F., 29, 136, 1058, *1091*, *1116*  
 Crisan, D., 517, *1091*  
 Cristianini, N., 736, *1091*  
 critic (in learning), **75**  
 critical path, **394**  
 Crockett, L., 267, *1091*  
 Croft, W. B., 901, 905, 906, *1091*, *1109*  
 Cross, S. E., 47, *1091*  
 cross-entropy, **809**  
 CROSS-VALIDATION, **685**  
 cross-validation, **684**, 734  
 crossover, 159  
 crossover point, **134**  
 crossword puzzle, 64  
 crowdsourcing, 723  
 Cruise, 1031  
 Cruse, A., 904, *1091*  
 cryptarithmetic, **168**

CSI, *see* independence, context-specific  
 Csorba, M., 984, 1093  
 CSP (constraint satisfaction problem), 38, 164, 164–169, 370  
 CTRL (language model), 884  
 Cuadros, J., 48, 1097  
 Cuellar, J., 268, 1085  
 Culberson, J., 127, 1091  
 culling, 134  
 cult of computationalism, 1033  
 Cummins, D., 550, 1091  
 Cummins, F., 838, 1096  
 cumulative distribution, 531, 1078  
 cumulative learning, 767  
 Curie, M., 19  
 curiosity, 850  
 Curran, J. R., 904, 1091  
 current-best-hypothesis approach, 768  
 current-best-hypothesis search, 741–743  
 curriculum learning, 891  
 Currie, K. W., 402, 1095  
 curse  
   of dimensionality, 706, 735, 952  
   optimizer's, 527, 549  
   winner's, 549  
 Curtis, F. E., 735, 1088  
 Cushing, W., 402, 1091  
 Cushman, F., 872, 1099  
 Cusumano-Towner, M. F., 667, 668, 1091  
 Cutler, A., 716  
 cutoff test, 202  
 cutset conditioning, 185, 188, 474  
 Cybenko, G., 836, 1091  
 CYBERLOVER (chatbot), 1035  
 cybernetics, 34, 33–34  
 cybersecurity, 50, 1041  
 CYC (knowledge base), 334, 358  
 cycle, 92  
 cycle constraint, 1021  
 cycle cutset, 185  
 cycle of net negative cost, 83  
 cyclic solution, 143  
 Cyert, R., 551, 1091  
 Cyganiak, R., 334, 357, 1088  
 CYK-PARSE, 888  
 CYK algorithm, 886, 902  
 Czarnecki, W. M., 48, 225, 838, 1100, 1115

**D**

D'Ambrosio, B., 474, 1112  
 d-separation, 437  
 Dafoe, A., 46, 1097  
 DAG (directed acyclic graph), 430, 472  
 Dagan, I., 931, 1091  
 Daganzo, C., 473, 1091  
 Dagdelen, J., 923, 1115  
 DAGGER (imitation learning system), 974, 986

Dagum, P., 475, 1092  
 Dahiya, R., 1027, 1105  
 Dahl, G., 48, 900, 905, 1099, 1104  
 Dahy, S. A., 700, 702, 1101  
 Dai, A. M., 930, 1092  
 Dai, Z., 930, 1117  
 Dalal, N., 1029, 1092  
 Dalibard, V., 838, 1100  
 Dalmao, S., 474, 1086  
 Dalvi, B., 901, 1106  
 Dalvi, N. N., 665, 1092  
 Daly, R., 799, 1092  
 Damasio, A. R., 1058, 1092  
 Damerau, F., 903, 1085  
 Danaher, J., 1057, 1092  
 Dang, T. D., 331, 1097  
 Danish, 877  
 Dantzig, G. B., 161, 1092  
 DARPA, 47, 1061  
 DARPA Grand Challenge, 46, 979, 984, 1071  
 DARPA Urban Challenge, 979  
 Darrell, T., 841, 986, 1030, 1096, 1103  
 Dartmouth workshop, 36  
 Darwiche, A., 474, 475, 478, 1090, 1092, 1108  
 Darwin, C., 136, 1092  
 Dasgupta, P., 163, 639, 1092  
 data-driven, 249  
 data-oriented parsing, 891  
 data association, 655, 942  
 data augmentation, 725  
 database, 77, 282  
   probabilistic, 665  
 database semantics, 283, 315, 363, 643  
 data complexity, 308  
 dataflow graph, 805  
 Datalog, 305, 328, 329  
 data matrix, 698  
 data provenance, 724  
 data science, 716, 717, 720, 738, 1066  
 data set augmentation, 1005  
 data sheet, 1046  
 dative case, 892  
 Daun, B., 402, 1092  
 Dauphin, Y., 837, 1092  
 Davidson, A., 637, 1087  
 Davidson, D., 358, 875, 1092  
 Davidson, M., 124, 1110  
 Davis, A., 736, 1072, 1112  
 Davis, E., 358–361, 1092  
 Davis, G., 402, 1092  
 Davis, M., 251, 266, 321, 328, 668, 1092  
 Davis, T., 48, 224, 1107  
 Davis-Putnam algorithm, 251  
 Dawid, A. P., 474, 549, 798, 1091, 1103, 1113  
 Dayan, P., 223, 839, 871–873, 1092, 1112  
 da Silva, B. C., 872, 1114  
 da Vinci, L., 24, 1027  
 DBN (dynamic Bayesian network), 479, 503–516  
 DBPEDIA (knowledge base), 334, 357  
 DDN (dynamic decision network), 560, 585  
 de-identification, 1042  
 Deacon, T. W., 42, 1092  
 dead end, 153  
 Deale, M., 402, 1092  
 Dean, J., 329, 670, 736, 838, 903, 909, 929, 931, 1069, 1072, 1088, 1092, 1102, 1106, 1109, 1112  
 Dean, M. E., 268, 1108  
 Dean, S., 1060, 1104  
 Dean, T., 401, 473, 516, 586, 587, 984, 985, 1070, 1092  
 Dearden, R., 587, 871, 1088, 1092  
 Deb, S., 160, 1117  
 Debevec, P., 1030, 1092  
 de Borda, J.-C., 630  
 Debreu, G., 533, 1092  
 debugging, 291  
 DEC (Digital Equipment Corporation), 41, 310  
 Dechter, A., 189, 1092  
 Dechter, R., 125, 188–191, 474, 475, 478, 1091, 1092, 1100, 1105, 1109, 1112  
 decision  
   rational, 403, 518  
   robust, 543  
   sequential, 537, 552  
 DECISION-LIST-LEARNING, 693  
 decision analysis, 548  
 decision boundary, 700  
 decision list, 692  
 decision maker, 548  
 decision network, 472, 518, 534, 534–537, 547  
   dynamic, 560, 585  
   evaluation of, 536  
 decision node, 535  
 decision stump, 718  
 decision theory, 28, 43, 405, 547  
 decision tree, 675, 733  
   expressiveness, 675  
   pruning, 681  
 declarative, 269  
 declarative bias, 757  
 declarativism, 228, 265  
 decoder (in autoencoders), 829  
 decoding, 918  
 decoding (in MT), 918  
   greedy, 918  
 decomposability (of lotteries), 521  
 DECOMPOSE, 382  
 decomposition, 374  
 DeCoste, D., 736, 1092

- Dedekind, R., 296, 1092  
 deduction theorem, 240  
 deductive database, 310, 328, 329  
 Deep Blue, 222  
 Deep Blue (chess program), viii, 48, 222  
 deepfake, 1022  
 DEEPHOL (theorem prover), 327  
 deep learning, 44, 716, 801–839  
   for NLP, 907–931  
   for robotics, 965–975  
   for vision, 1001–1025  
 DEEPMATH (theorem prover), 330  
 DeepMind, 49, 225, 830, 835, 867, 873, 1059  
 deep Q-network (DQN), 835, 867, 873  
 deep reinforcement learning, 857, 986  
 Deep Space One, 373, 402  
 DeepStack (poker program), 224, 612  
 DEEP THOUGHT (chess program), 222  
 Deerwester, S. C., 903, 929, 1093  
 default logic, 352, 356, 359  
 default reasoning, 351–353, 477  
 default value, 349  
 definite clause, 247, 304–305  
 definition (logical), 285  
 degree heuristic, 177, 189, 252  
 degree of belief, 403, 404, 411, 412  
 degree of truth, 273  
 degree of usefulness, 405  
 degrees of freedom (DOF), 947  
 DeGroot, M. H., 428, 799, 1093  
 Dehaene, S., 1058, 1093  
 Deisenroth, M. P., 871, 1086  
 Delarue, A., 88, 1087  
 delete list, 363  
 Delgrande, J., 359, 1093  
 deliberative, 975  
 Dellaert, F., 225, 984, 1095, 1115  
 Della Pietra, S. A., 931, 1089  
 Della Pietra, V. J., 931, 1089  
 Delling, D., 126, 1093  
 delta rule, 855  
 Del Favero, B. A., 474, 1112  
 Del Moral, P., 517, 1093  
 demodulation, 325, 330  
 demographic parity, 1044  
 Dempster, A. P., 477, 515, 799, 1093  
 Dempster-Shafer theory, 477  
 Denardo, E. V., 586, 1093  
 DENDRAL (expert system), 40, 41, 356  
 dendrite, 30  
 DeNero, J., 735, 1107  
 Deng, J., 44, 837, 1093, 1111  
 Deng, L., 839, 900, 905, 1093, 1099, 1117  
 Deng, X., 162, 638, 1093  
 Denker, J., 44, 837, 838, 1029, 1103  
 Denney, E., 330, 1093  
 Denniston, A. K., 48, 1104  
 density estimation, 775  
   nonparametric, 787  
 DeOliveira, J., 357, 1105  
 deontological ethics, 26  
 dependency grammar, 889  
 D'Épenoux, F., 586, 1093  
 depth-first search, 96, 96–98, 123, 376  
 DEPTH-LIMITED-SEARCH, 99  
 depth limit, 204  
 depth of field, 992  
 derivation, 234  
 derivational analogy, 769  
 Dervovic, D., 1069, 1093  
 Descartes, R., 24, 1027, 1057, 1093  
 descendant (in Bayesian networks), 436  
 Descotte, Y., 401, 1093  
 description logic, 331–351, 347, 349, 355, 359  
 descriptive theory, 528  
 Deshpande, A., 1016, 1085  
 Deshpande, I., 1023, 1093  
 Deslippe, J., 48, 1103  
 Desouza, P. V., 929, 1089  
 detailed balance, 462  
 detection failure (in data association), 657  
 deterministic environment, 63  
 deterministic node, 438  
 deterministic parser, 889  
 Dethridge, J., 124, 1110  
 detour index, 109  
 Deutscher, G., 296, 1093  
 development set, 684  
 Deville, Y., 189, 1115  
 DEVISER (planning system), 401  
 Devlin, J., 930, 1093  
 Devlin, K., 1057, 1093  
 Devol, G., 983  
 Devroye, L., 799, 1093  
 Dewey Decimal system, 335  
 Dey, A. K., 868, 970, 986, 1118  
 Dey, K., 1047, 1060, 1087  
 de Condorcet, M., 736, 1092  
 de Dombal, F. T., 428, 1092  
 de Farias, D. P., 586, 1092  
 de Finetti's theorem, 412  
 de Finetti, B., 412, 427, 1092  
 de Freitas, J. F. G., 516, 517, 1092, 1093  
 de Ghellinck, G., 586, 1092  
 de Groot, M., 551, 1091  
 De Kleer, J., 357, 1095  
 de Kleer, J., 190, 329, 360, 1092, 1095, 1116  
 de Marcken, C., 905, 1092  
 De Marneffe, M.-C., 890, 1108  
 De Morgan, A., 188, 296, 1092  
 De Morgan rules, 281  
 de Oliveira Marinho, G., 48, 1104  
 De Raedt, L., 905, 1107  
 de Salvo Braz, R., 668, 1092  
 de Sarkar, S. C., 127, 163, 1090, 1092  
 De Wever, A., 549, 1087  
 diabetic retinopathy, 716  
 Diaconis, P., 528  
 diagnosis, 404, 417, 418  
   dental, 404  
   medical, 41, 428, 436, 537  
 diagnostic rule, 436  
 dialysis, 524  
 diameter (of a graph), 98  
 Dias, M. B., 160, 1093  
 Dickerson, S., 986, 1114  
 Dickinson, M. H., 1064, 1095  
 Dickmanns, E. D., 46, 1030, 984, 1093  
 dictionary, 879  
 Dieleman, S., 47, 830, 838, 900, 1027, 1094, 1115  
 Dietterich, T., 873, 1093  
 Difference Engine, 33  
 differentiable programming, 1067  
 differential equation, stochastic, 480  
 differential GPS, 936  
 differential heuristic, 121  
 differential privacy, 1042, 1060  
 Digital Equipment Corporation (DEC), 41, 310  
 Dijkstra's algorithm, 95, 125  
 Dijkstra, E. W., 125, 1035, 1093  
 Dill, D. L., 268, 1108  
 Ding, Y., 48, 1093  
 Dinh, H., 125, 1093  
 Dionne, A., 127, 1114  
 Diophantine equations, 188  
 Diophantus, 188  
 DiPasquo, D., 905, 1091  
 Diplomacy, 197  
 direct collocation, 962  
 directed acyclic graph (DAG), 430, 472  
 directional arc consistency, 183  
 direct utility estimation, 869  
 Dirichlet distribution, 782  
 Dirichlet process, 799  
 disambiguation, 897–900  
 discontinuity (depth), 996  
 discount factor, 555, 586, 632, 843  
 discovery systems, 769  
 discretization, 138, 440  
 discriminative model, 779, 883  
 discriminator (in GANs), 831  
 disjoint sets, 336  
 disjunct, 235  
 disjunction, 235, 411  
 disjunctive constraint, 167  
 disparity, 1009  
 Dissanayake, G., 984, 1093  
 distribute  $\vee$  over  $\wedge$ , 245, 300  
 distributed constraint satisfaction, 191  
 distributed representation, 78  
 distribution

- beta, 505, **781**  
 categorical, 409  
 conditional, nonparametric, **440**  
 cumulative, 531, 1078  
 heavy-tailed, **160**  
 mixture, **790**
- Dittmer, S., 550, *1093*  
 dividing a pie, 631  
 Dix, J., 360, *1089*  
 DLV (logic programming system), 360  
 DNA, 134  
 Do, M., 47, 399, *1086, 1093*  
 Do, M. B., 401, *1093*  
 do-calculus, **470**  
 Doctorow, C., 30, 358, *1093*  
 Dodd, L., 891, *1093*  
 domain  
   continuous, **167**  
   element of, **274**  
   infinite, **167**  
   in first-order logic, **274**  
   in knowledge representation, **283**  
 domain closure, **282**, 643  
 domain randomization, **967**  
 dominance  
   stochastic, **531**, 547  
   strict, **530**  
 dominant strategy, **597**, 625  
 dominant strategy equilibrium, **597**  
 dominated plan (in POMDP), **582**  
 domination (of heuristics), **117**  
 Domingos, P., 53, 428, 474, 666, 726,  
   736, 738, 798, *1093, 1096, 1110*  
 Domshlak, C., 399, 549, 636, *1088, 1099*  
 Donahue, J., 838, 1030, *1096, 1100*  
 Dong, W., 44, *1093*  
 Dong, X., 334, *1093*  
 Donti, P. L., 48, *1110*  
 Doorenbos, R., 329, *1093*  
 Doran, J., 125, 126, *1093*  
 Dorf, R. C., 79, *1093*  
 Dorigo, M., 160, *1093*  
 Doron, Y., 873, *1114*  
 Doshi, T., 1060, *1087*  
 Doshi-Velez, F., 737, *1093*  
 Dota 2, 46, 48  
 do the right thing, 175, **22**, 26, 57, 519  
 Doucet, A., 516, 517, *1085, 1091, 1093,*  
   *1108*  
 Downey, D., 906, *1094*  
 Downs, L., 985, *1088*  
 downsampling, **813**  
 downward refinement property, **378**  
 Dowty, D., 904, *1093*  
 Doyle, J., 79, 190, 359, 360, 549, 550,  
   *1093, 1106, 1116, 1118*  
 DPPL, **252**, 266, 415  
 DPPL-SATISFIABLE?, **252**  
 DQN (deep Q-network), 835, **867**, 873  
 Drabble, B., 402, *1093*  
 Dragan, A. D., 80, 551, 638, 872, 971,  
   974, 986, *1086, 1093, 1097, 1105,*  
   *1111*  
 DRAGON (speech recognition system),  
   905  
 Drake, J. P., 735, *1087*  
 Draper, D., 401, *1094*  
 Drebbel, C., 34  
 Drechsler, R., 126, *1094*  
 Dredze, M., 736, *1093*  
 Dressel, J., 1044, 1060, *1093*  
 Dreyfus, H. L., 267, 1032, 1033, *1093*  
 Dreyfus, S. E., 125, 586, 836, 1033, *1087,*  
   *1093*  
 Driessens, K., 871, *1114*  
 drilling rights, 537  
 driver assist, **980**  
 dropout (in neural networks), **823**, 838  
 Droppo, J., 47, *1117*  
 Drucker, H., 44, 837, 1029, *1103*  
 Druzdzel, M. J., 475, *1090*  
 DT-AGENT, **406**  
 Du, J., 927, 930, *1104*  
 Du, S. S., 837, *1093*  
 dual graph, **168**  
 dualism, **24**  
 dual use, **1041**  
 Dubois, D., 478, *1093*  
 Dubois, O., 267, *1113*  
 Dubourg, V., 738, *1109*  
 Ducharme, R., 929, *1087*  
 duck, mechanical, 983  
 Duda, R. O., 428, 738, 798, 800, *1093,*  
   *1094*  
 Dudek, G., 987, *1094*  
 Dudzik, A., 48, 225, *1115*  
 Duffy, D., 331, *1094*  
 Duffy, K., 736, *1091*  
 Duffy, N., 137, 161, *1106*  
 Dumais, S. T., 901, 903, 905, 929, *1086,*  
   *1093*  
 dummy player, **619**  
 dung beetle, 59, 80, 392, 977  
 Dunham, B., 39, *1095*  
 Dunham, C., 329, *1114*  
 Dunn, A., 923, *1115*  
 Dunn, H. L., 666, *1094*  
 Dunn, J., 734, *1087*  
 Dunn, R. C., 48, *1104*  
 Dunning, I., 48, 225, 838, *1100*  
 DuPont, 41  
 Dupont, D., 549, *1087*  
 Durand, F., 1067, *1104*  
 duration, **393**  
 Dürer, A., 1027  
 Durfee, E. H., 639, *1094*  
 Durme, B. V., 906, *1094*  
 Durrant-Whyte, H., 984, *1093, 1103*  
 Duru, G., 549, *1087*  
 Dwork, C., 1046, 1060, *1094, 1117*  
 DYNA (reinforcement learning agent),  
   871  
 dynamical system, 515  
   quadratic, 161  
 dynamic backtracking, 190  
 dynamic Bayesian network (DBN), 479,  
   **503**, 503–516  
 approximate inference in, 509  
 exact inference in, 507  
 dynamic decision network (DDN), **560**,  
   585  
 dynamic environment, **63**  
 dynamic programming, 79, 119, 125, **315**,  
   488, **553**, 586, 886  
   adaptive (ADP), **844**, 844  
   nonserial, **474**  
 dynamics model, **958**  
 dynamic state, **958**  
 Dyson, F., 734, *1094*  
 Dyson, G., 1061, *1094*  
 dystopia, 1062, 1073
- 
- E**
- E (theorem prover), 330, 331  
 $\mathcal{E}_0$  (English fragment), 884  
 Earley, J., 904, *1094*  
 early playout termination, **210**  
 early stopping, **682**  
 earthquake, 431  
 Ebendt, R., 126, *1094*  
 EBL (explanation-based learning), 400  
 Eck, D., 1027, *1094*  
 Ecker, A. S., 1034, *1096*  
 Ecker, K., 402, *1088*  
 Eckerle, J., 114, 126, *1094*  
 Eckert, J., 32  
 Eckhouse, L., 1060, *1094*  
 economics, 27–28, 79, 524  
 economy, **192**  
 Edelkamp, S., 127, 399, *1094*  
 edge (in a scene), **996**  
 edge detection, 996–999  
 Edinburgh, 983  
 Edmonds, D., 35  
 Edmonds, J., 27, *1094*  
 Edward (probabilistic programming  
   language), 667  
 Edwards, D. J., 221, *1098*  
 Edwards, P., 1058, *1094*  
 Edwards, W., 548, *1116*  
 EEG, 29  
 Een, N., 330, *1085*  
 effect, **363**  
   missing, **390**  
 effective depth, **116**  
 effector, **932**  
 efficient auction, **624**

- Efros, A., 44, 930, 1015, 1020–1022, 1030, 1095, 1098–1101, 1118  
 egalitarian social welfare, **600**  
 egocentric action, 85  
 Ehrenfeucht, A., 735, 1088  
 8-puzzle, **86**, 115, 118, 124  
 8-queens problem, 128, 130, 134, 181  
 Einstein, A., 1  
 Eisner, J., 904, 1113  
 Eitelman, S., 329, 1114  
 Eiter, T., 360, 1094  
 Ekart, A., 161, 1110  
 EKF (extended Kalman filter), **501**, 942  
 ELBO (evidence lower bound), **829**  
 Elder, J. F., 736, 1112  
 Elementary Perceiver And Memorizer (EPAM), 733  
 Elfes, A., 984, 1107  
 ELIMINATION-ASK, **450**  
 Elio, R., 550, 1094  
 Eliot, T.S., 824  
 Elisseeff, A., 737, 1097  
 elitism, **134**  
 ELIZA (chatbot), 1035  
 Elkan, C., 798, 1060, 1094, 1100  
 Elkind, E., 638, 1090  
 Ellington, C., 1064, 1095  
 Elliott, G. L., 189, 1098  
 Elliott, P., 267, 1116  
 Ellsberg, D., 550, 1094  
 Ellsberg paradox, 528, 550  
 Elman, J. L., 837, 905, 1094  
 ELMO (natural language system), 930  
 Elo, A. E., 667, 1094  
 Elskens, T., 838, 1094  
 EM algorithm, 484, 788–797  
     structural, **797**  
 embodied cognition, **1033**  
 empirical gradient, **138**, 862  
 empirical loss, **688**  
 empiricism, **24**, 906  
 Empson, W., 904, 1094  
 EMV (expected monetary value), **524**  
 ENAS (Efficient Neural Architecture Search), 838  
 encoder (in autoencoders), 829  
 end-to-end learning, 967  
 Enderton, H. B., 297, 328, 1094  
 Endriss, U., 639, 1088  
 Eng, C., 48, 1104  
 Engel, J., 1027, 1094  
 Engelbart, D., 32  
 Engelberger, J., 983  
 ENIAC, 32  
 ensemble learning, **714**, 714–720  
 ensemble model, **714**  
 entailment, **232**, 264  
 entailment constraints, **767**  
 entropy (H), **679**, 680  
 ENUMERATE-ALL, **447**  
 ENUMERATION-ASK, **447**  
 environment, **54**, 60–65  
     class, **65**  
     competitive, **63**  
     continuous, **64**  
     cooperative, **63**  
     deterministic, **63**  
     discrete, **64**  
     dynamic, **63**  
     episodic, **63**  
     history, 553  
     known, **64**  
     multiagent, **62**, 589  
     nondeterministic, **63**, 128  
     observable, **61**  
     one-shot, **63**  
     partially observable, **61**  
     properties, **61**  
     semidynamic, **64**  
     sequential, **63**  
     single-agent, **62**  
     static, **63**  
     stochastic, **63**, 552  
     taxi, 60, 61  
     unknown, **64**  
     unobservable, **62**  
     virtual, **61**  
 EPAM (Elementary Perceiver And Memorizer), 733  
 Ephrati, E., 639, 1102  
 epistemological commitment, **273**, 295, 404  
 epoch, **697**  
 Epstein, R., 1057, 1094  
 EQP (theorem prover), 331  
 equality (in logic), 282, 324  
 equality symbol, **282**  
 equality test, 683  
 equilibrium, 217  
     Bayes-Nash, **613**  
     dominant strategy, **597**  
     maximin, **603**  
     Nash, **598**, 635  
     subgame perfect, **609**  
 equivalence (logical), **240**  
 Erdmann, M. A., 162, 1094  
 Erez, T., 873, 986, 1098, 1104, 1114  
 ergodic, **462**  
 Erhan, D., 48, 838, 930, 1114, 1115  
 ERNIE (NLP system), 930  
 Ernst, G., 125, 1107  
 Ernst, H. A., 983, 1094  
 Ernst, M., 399, 1094  
 Erol, K., 400, 1094  
 Erol, Y., 517, 1094  
 error (of a hypothesis), **684**, 691  
 error function, 1079  
 error rate, **684**  
 Escalante, H. J., 737, 1097  
 Escalera, S., 737, 1097  
 Essig, A., 428, 1097  
 Esteva, A., 48, 1094  
 Estrin, D., 1060, 1086  
 Etchemendy, J., 297, 1086  
 ethics, 1037–1056  
 Etzioni, A., 1059, 1060, 1094  
 Etzioni, O., 45, 334, 357, 401, 901, 905, 906, 927, 931, 1059, 1071, 1086, 1091, 1094, 1103, 1114, 1116  
 Euclid, 27, 1027  
 Eugene Goostman, 1035  
 Euler-Lagrange equation, **956**  
 EUROPA (planning system), 47  
 Europe, 41  
 European Space Agency, 402  
 evaluation function, **91**, 123, 192, 202–204, 854  
     linear, 122  
 Evans, O., 46, 872, 1097, 1111  
 Evans, T. G., 38, 1094  
 event, 340–343  
     exogenous, **390**  
     in probability, **407**, **445**  
 event calculus, 340–342, **341**, 358, 897  
 Everett, B., 984, 1088  
 Everitt, T., 873, 1054, 1103  
 evidence, **407**, 772  
     reversal, **516**  
 evidence lower bound (ELBO), **829**  
 evidence variable, 445  
 evolution, 136  
     machine, **39**  
     of machines, **39**  
 evolutionary algorithm, **133**, 159  
 evolutionary psychology, **529**  
 evolution strategies, **134**, 161  
 Ewalds, T., 48, 225, 873, 1115  
 exception, 333, 349  
 exclusive or, 237  
 execution monitoring, 389–392, **390**, 401  
 exhaustive decomposition, **336**  
 existence uncertainty, **648**  
 existential graph, **347**  
 Existential Instantiation, **299**  
 existential quantifier, **280**  
 expansion (of nodes), **89**  
 expectation, **1079**  
 expectation maximization, 797  
 expected monetary value (EMV), **524**  
 expected utility, **73**, 80, 405, 518, **519**, 524  
 expected value (in a game tree), **202**, 211  
 expectiminimax, 212, 220, 221  
     complexity of, 213  
     value, **211**  
 experience replay, **857**  
 expert system, **41**, 356, 548

commercial, 41, 310  
 medical, 477  
 Prolog-based, 312  
 explit model, **442**  
 explainability, **729**  
 explainable AI (XAI), 737, **1048**  
 explanation, **354**  
     most probable, **475**  
 explanation-based learning (EBL), 400, **749**, **767**  
     branching factor, **753**  
     definition, 750  
     efficiency, 752–754  
     general rules, 751–752  
     logic programming implementation, 751  
     memoization, **751**  
     operationality, **754**  
     prune, **752**  
 exploitation, **207**, 571, 849  
 exploration, 58, 59, 152–159, **207**, 581, 842, 848, 849  
     bonus, **575**  
     function, **850**, 853  
     safe, **154**  
 exploratory data analysis, **671**, 726  
 expressiveness, **77**  
 extended Kalman filter (EKF), **501**, 942  
 extension (of default theory), **353**  
 extensive form, **607**  
 externalities, **627**, 1053  
 extremely randomized trees (ExtraTrees), 716  
 extrinsic property, **340**  
 eyes, 988, 991, 992, 1027

---

**F**

Facebook, 47, 873, 1059  
 fact, **248**  
 factor (in variable elimination), **448**  
 factored frontier, **517**  
 factored representation, **77**, 81, 164, 362, 408, 560, 670  
 factoring, **244**, 319  
 Faes, L., 48, **1104**  
 Fagin, R., 190, 359, **1087**, **1094**  
 Fahlman, S. E., 38, 360, **1094**  
 failure model, 506  
 failure modes and effect analysis (FMEA), **1052**  
 fair division, 618  
 fairness, 724, 729, 1043–1047, 1060  
 fall in love, 1033  
 false alarm (in data association), **657**  
 false positive, **728**  
 Fan, J., 48, **1094**  
 Farhadi, A., 931, **1112**  
 Farid, H., 1044, 1060, **1093**  
 FARMVILLE (video game), 1050

Farrell, R., 329, **1089**  
 FASTDOWNWARD (planning system), 398  
 Fast Downward Stone Soup, 399  
 Faster RCNN (computer vision system), 1007  
 FASTFORWARD (planning system), 374  
 FASTTEXT (word embedding), 908  
 Fatica, M., 48, **1103**  
 Faugeras, O., 1030, **1094**  
 fault tree analysis (FTA), **1052**  
 Favini, G. P., 222, **1091**  
 Fawcett, T., 1041, **1094**  
 FDSS (planning system), 399  
 Fearing, R. S., 985, **1094**  
 Featherstone, R., 985, **1094**  
 feature (of a state), **122**, 202  
 feature expectation, **865**  
 feature extraction, 988  
 feature map, **815**  
 feature matching, **865**  
 feature selection, **689**, 876  
 federated learning, 724, **1043**, 1060  
 feedback, 33, 669, **671**  
 feedforward network, **802**  
 Fei-Fei, L., 44, 837, 930, **1093**, **1101**, **1111**  
 Feigenbaum, E. A., 35, 40, 41, 125, 356, 733, **1086**, **1089**, **1094**, **1104**  
 Feiten, W., 984, **1088**  
 Feldman, J., 80, 548, **1094**  
 Feldman, M., 31, **1094**  
 Fellbaum, C., 903, **1094**  
 Fellegi, I., 666, **1094**  
 Feller, A., 1060, **1091**  
 Felner, A., 106, 126, 127, 399, **1094**, **1099**, **1102**, **1111**  
 Felzenszwalb, P., 162, **1094**  
 Feng, L., 984, **1088**  
 Feng, S., 930, **1114**  
 Feng, T. K., 190, **1087**  
 Fenton, N., 548, **1094**  
 Fergus, R., 708, 838, **1114**, **1115**  
 Ferguson, T., 222, 587, 799, **1094**  
 Fermat, P., 26, 426  
 Fern, A., 551, **1094**  
 Fernández, F., 872, **1096**  
 Fernandez, J. M. F., 48, 225, **1094**  
 Fernando, C., 838, **1100**, **1104**  
 FERPA, 1041  
 Ferraris, P., 400, **1094**  
 Ferriss, T., 1051, **1094**  
 Ferrucci, D., 48, **1094**  
 FF (planning system), 374, 398  
 Fidjeland, A., 871, 873, **1106**  
 FIFO queue, **92**  
 15-puzzle, **86**, 124  
 Fifth Generation project, 41  
 figure of speech, 898, 899  
 Fikes, R.E., 79, 162, 296, 398, 400, 401, 983, **1095**  
 filtering, **150**, 353, 484–485, 514, 578, 795, 938  
     assumed-density, **517**  
 Fine, S., 516, **1095**  
 finite state machine, 604  
 Fink, D., 137, 161, **1106**  
 Finkelstein, L., 191, **1089**  
 Finn, C., 737, 841, 986, **1095**, **1103**  
 Finney, D. J., 473, **1095**  
 Firat, O., 901, **1090**  
 Firby, R. J., 401, **1092**  
 FIRE (theorem prover), 357  
 Firoiu, V., 225, **1095**  
 first-order logic, **269**, 269–297  
 first mover, 632  
 Firth, J., 907, **1095**  
 Fisac, J. F., 648, 872, **1085**, **1105**  
 Fischer, A., 734, **1086**  
 Fischer, B., 330, **1093**  
 Fischer, P., 587, **1086**  
 Fischer, R., 214  
 Fisher, M. L., 551, **1107**  
 Fisher, R. A., 27, 427, **1095**  
 fitness landscape, 161  
 Fix, E., 735, **1095**  
 fixation, **1010**  
 FIXED-LAG-SMOOTHING, **493**  
 fixed-lag smoothing, **489**  
 Flannery, B. P., 160, **1109**  
 Floreano, D., 1064, **1095**  
 Floyd, R. W., 125, **1095**  
 fluent, **256**, 265, 343  
     missing, **390**  
 fly eyes, 1012, 1024  
 FMEA (failure modes and effect analysis), **1052**  
 focal length, **990**  
 focal plane, **992**  
 focal point (in game theory), **598**  
 focus of expansion, **1011**  
 Fogel, D.B., 162, **1095**  
 Fogel, L. J., 161, **1095**  
 fog of war, 214  
 FOL-BC-AND, **311**  
 FOL-BC-ASK, **311**  
 FOL-BC-OR, **311**  
 FOL-FC-Ask, **306**  
 folk psychology, 361  
 Fong, R., 986, **1085**  
 FOPC, *see* logic, first-order  
 Forbes, J., 588, **1095**  
 FORBIN (planning system), 401, 402  
 Forbus, K.D., 329, 357, 360, **1095**  
 force sensor, **936**  
 Ford, K. M., 1057, **1095**  
 Ford, L. R., 125, **1095**  
 Ford, M., 46, 53, 1062, **1095**

- foreshortening, 989  
 Forestier, J.-P., 873, 1095  
 forget gate (in LSTM), 826  
 Forgy, C., 329, 1095  
 formal logic, 26  
 Forrest, S., 161, 1106  
 Forster, E. M., 1062, 1095  
 Forsyth, D., 1021, 1023, 1031, 1093, 1095, 1101  
 Fortmann, T. E., 515, 667, 1086  
 Fortran, 807  
 forward-backward, 488, 795  
**FORWARD-BACKWARD, 488**  
 forward chaining, 248, 248–249, 265, 304–311, 329  
 forward checking, 178, 178  
 forward kinematics, 947  
 forward message, 487  
 forward pruning, 205  
 forward search for planning, 366–368  
 Foster, G., 901, 1090  
 Fouhey, D., 1015, 1095  
 four-color map problem, 188, 1034  
 Fourier, J., 188, 1095  
 Fowlkes, C., 1028, 1105  
 Fox, C., 560, 1095  
 Fox, D., 667, 984, 987, 1089, 1095, 1112, 1115  
 Fox, M. S., 401, 1095  
 FPGA, 45  
 frame, 41, 359  
 FrameNet (lexical database), 357  
 frame problem, 257, 267, 268  
     representational, 257  
 framing effect, 529  
 Francis, J., 549, 1101  
 Franco, J., 266, 1095  
 Francois-Lavet, V., 871, 1095  
 Francon, O., 137, 161, 1106  
 Frank, E., 738, 1117  
 Frank, I., 222, 1095  
 Frank, J., 47, 1086  
 Frank, R. H., 1051, 1095  
 Frankenstein, 1052  
 Frans, K., 873, 1095  
 Franz, A., 903, 1095  
 Frasconi, P., 837, 1087  
**FREDDY (robot)**, 125, 162, 983  
 Fredkin Prize, 222  
 Freeman, W., 476, 1116, 1117  
 Freer, C., 667, 1085  
 free space, 947  
 free will, 24  
 Frege, G., 26, 266, 296, 328, 1095  
 Freitag, D., 905, 906, 1091, 1095  
 frequentism, 426  
 Freuder, E. C., 189, 190, 1095, 1105, 1111  
 Freund, Y., 718, 736, 1095  
 Frey, B. J., 838, 1095  
 Frey, C. B., 1050, 1095  
 Friedberg, R. M., 39, 161, 1095  
 Friedman, G. J., 161, 1095  
 Friedman, J., 734, 736, 738, 800, 1089, 1095, 1098  
 Friedman, N., 473, 474, 478, 516, 517, 586, 798, 799, 871, 1085, 1088, 1092, 1095, 1102  
 Friendly AI, 1061  
 frisbee, 48, 1033  
 Fristedt, B., 587, 1087  
 front-to-end, 114  
 front-to-front, 114  
 frontier, 90  
 Frost, D., 189, 191, 1092  
 Fruhwirth, T., 191, 1095  
**FTA (fault tree analysis)**, 1052  
 Fu, D. J., 48, 1104  
 Fu, J., 359, 1065, 1104  
 Fu, M. C., 587, 1090  
 Fuchs, J. J., 402, 1095  
 Fudenberg, D., 638, 1095  
 Fukunaga, A. S., 401, 1095  
 Fukushima, K., 837, 1029, 1095  
 Fuller, S. B., 1064, 1095  
 full joint distribution, 413, 415  
 fully connected (neural network), 805  
 fully observable, 578  
 function, 272  
     total, 275  
 functional, 956  
 functional magnetic resonance imaging (fMRI), 29, 271  
 function approximation, 854  
 function symbol, 275, 277  
 Fung, C., 1060, 1095  
 Fung, C. C., 1035, 1100  
 Fung, R., 475, 1095  
 Furcy, D., 163, 1102  
 Furnas, G. W., 903, 929, 1093  
 Furst, M., 399, 1088  
 futility pruning, 222  
 Future of Humanity Institute, 1059  
 Future of Life Institute, 1059  
 future of work, 1049–1051, 1062  
 fuzzy control, 477  
 fuzzy logic, 232, 273, 477  
 fuzzy set, 477
- 
- G**
- g* (path cost), 91  
 Gabriele, S., 48, 1104  
 Gabrilovich, E., 334, 1093  
 Gaddum, J. H., 473, 1095  
 Gadepalli, K. K., 48, 1104  
 Gaffney, S., 873, 1087  
 Gaifman, H., 665, 666, 1095  
 gain factor, 959  
 gait, 47, 976, 1041  
 Gale, W.A., 902, 1090  
 Galfrin, E., 474, 1112  
 Galileo, G., 175, 75  
 Gallagher, G., 970, 986, 1118  
 Gallaire, H., 329, 1095  
 Gallegos, J., 48, 1104  
 Gallier, J. H., 297, 1095  
 Galperin, G. R., 222, 1114  
 Galstyan, A., 1060, 1106  
 Galton, F., 1095  
 Gamba, A., 836, 1096  
 Gamba perceptrons, 836  
 Gamberini, L., 836, 1096  
 gambling, 26, 521  
 game, 28  
     assistance, 52, 613, 866  
     backgammon, 210, 224, 855, 866  
     billiards, 225  
     of chance, 212–214  
     checkers, 37, 80, 223, 870, 871  
     chess, 23, 32, 39, 48, 64, 125, 193, 201–204, 222  
     cooperative, 616  
     dice, 217  
     Diplomacy, 197  
     Go, 207  
     of imperfect information, 192  
     incomplete information, 551  
     inspection game, 596  
     Kriegspiel, 214  
     multiplayer, 197–198  
     normal form, 595  
     optimal decisions in, 194–201  
     Othello, 224  
     partially observable, 214–218  
     of perfect information, 193  
     physical, 225  
     poker, 224, 638  
     principal-agent, 638  
     repeated, 598, 604  
     Reversi, 224  
     Scrabble, 225  
     Settlers of Catan, 197  
     stochastic, 210  
     Tetris, 562, 571  
     Yahtzee, 214  
     zero-sum, 193, 600  
 game playing, 192–221  
 game show, 524  
 game theory, 28, 590, 635  
     cooperative, 616  
     non-cooperative, 595–615  
 Gammie, C., 48, 1113  
 GAN (generative adversarial network), 831, 838, 1022  
 Ganchev, K., 904, 1085  
 Gandomi, A., 737, 1096  
 Ganguli, S., 837, 838, 1092, 1109

- Gannon, I., 474, **1112**  
 Gao, J., 670, **1096**  
 Gao, Q., 47, 834, 901, 916, **1117**  
 Gao, Y., 873, **1100**  
 Garcia, E. A., 725, 1046, **1098**  
 García, J., 872, **1096**  
 Gardner, M., 266, 901, 930, **1096**, **1106**,  
     *1109*  
 Garey, M. R., 1076, 1080, **1096**  
 Garg, A., 516, **1108**  
 GARI (planning system), 401  
 Garrett, C., 134, **1086**  
 Gaschnig, J., 125, 189, 190, **1096**  
 Gašić, M., 588, **1117**  
 Gasquet, A., 402, **1095**  
 Gasser, L., 636, **1088**  
 Gasser, R., 127, **1096**  
 Gat, E., 986, **1096**  
 gate (logic), 292  
 Gates, B., 51  
 gating unit (in LSTM), **826**  
 Gatys, L. A., 1034, **1096**  
 Gauci, J., 873, **1096**  
 Gauss, C. F., 188, 515, 735, **1096**  
 Gaussian distribution, **1078**  
     multivariate, 497, 1079  
 Gaussian error model, **505**  
 Gaussian filter, **997**  
 Gaussian noise, 497  
 Gaussian process, 690, **799**  
 Gazzaniga, M., 1058, **1096**  
 GBM (gradient boosting machine), 719  
 GBRT (gradient boosted regression tree),  
     719  
 GDPR, 1041, **1048**  
 GDPR (General Data Protection  
     Regulation), 730  
 Gearhart, C., 587, **1097**  
 Gebru, T., 1046, 1089, **1096**, **1106**  
 Gee, A. H., 516, **1092**  
 Geffner, H., 162, 398, 400, 401, **1088**,  
     *1098*, **1108**  
 Geiger, D., 473, 474, 798, **1096**, **1098**,  
     *1112*  
 Gelatt, C. D., 160, 190, **1101**  
 Gelb, A., 515, **1096**  
 Gelder, A. V., 331, **1114**  
 Gelernter, H., 37, 330, **1096**  
 Gelfond, M., 330, 360, **1096**  
 Gelman, A., 476, 668, 737, 798, 799,  
     *1089*, **1090**, **1096**  
 Geman, D., 475, 736, 1028, **1085**, **1096**  
 Geman, S., 475, 1028, **1096**  
 Gemp, I., 871, **1104**  
 generalization, **973**  
 generalization loss, **688**  
 generalizing fields, 1042  
 General Problem Solver, 20, 25, 37, 398  
 GENERATE-IMAGE, **660**  
 GENERATE-LETTERS, **660**  
 GENERATE-MARKOV-LETTERS, **662**  
 generation (of nodes), **90**  
 generative adversarial network (GAN),  
     **831**, 838, 1022  
 generative model, **778**, 882  
 generator, **1082**  
 generator network (in GANs), **831**  
 Genesereth, M. R., 79, 162, 225, 296,  
     297, 316, 321, 329, 330, **1096**,  
     *1104*, **1113**  
 GENETIC-ALGORITHM, **137**  
 genetic algorithm, **134**, 133–137,  
     161–162  
 genetic programming, 39, **134**, 161  
 Gene Ontology Consortium, The., 358,  
     *1096*  
 Gent, I., 191, **1096**  
 Geometry Theorem Prover, 37  
 Georgeson, M., 1031, **1089**  
 Georgiev, P., 48, 225, 873, **1115**  
 Gerbault, F., 799, **1096**  
 Gerkin, R. C., 1027, **1111**  
 Géron, A., 738, **1096**  
 Gers, F. A., 838, **1096**  
 Gesmundo, A., 737, **1117**  
 Gestalt school, 1028  
 Getoor, L., 668, **1096**  
 Ghaheri, A., 161, **1096**  
 Ghahramani, Z., 473, 516, 737, 799, 838,  
     *1096*, **1100**, **1111**, **1113**  
 Ghallab, M., 362, 398, 401, 402, **1096**  
 Gharbi, M., 1067, **1104**  
 Ghavamzadeh, M., 871, **1104**  
 Ghose, D., 160, **1102**  
 Ghose, S., 127, **1090**  
 GIB (bridge program), 224  
 Gibbard-Satterthwaite Theorem, **631**  
 Gibbs, R. W., 904, **1096**  
 GIBBS-ASK, **461**  
 Gibbs sampling, **460**, 463, 475  
 Gibson, J. J., 1028, 1030, **1096**  
 Gil, Y., 334, **1090**  
 Gilks, W. R., 475, 476, 517, 666, 798,  
     *1096*  
 Gillies, D. B., 648, **1096**  
 Gillula, J. H., 872, **1085**  
 Gilmore, P. C., 328, **1096**  
 Gilpin, A., 637, **1096**  
 Gini coefficient, **600**  
 Ginsberg, M. L., 224, 190, 329, 478,  
     *1092*, **1096**, **1113**  
 Ginter, F., 890, **1108**  
 Gionis, A., 735, **1096**  
 Girshick, R., 1030, **1096**  
 Gittins, J. C., 587, **1096**  
 Gittins index, **573**  
 Giunchiglia, E., 400, **1094**  
 Givan, R., 871, **1114**  
 Gladman, A. S., 1064, **1102**  
 Glanc, A., 982, **1096**  
 Glass, J., 516, **1104**  
 Glavieux, A., 476, **1087**  
 Glickman, M. E., 667, **1096**  
 Glickman, O., 931, **1091**  
 GLIE (greedy in the limit of infinite  
     exploration), **849**  
 global constraint, **168**, 172  
 Global Positioning System, **935**  
 GLONASS (Russian GPS), 936  
 Glorot, X., 837, **1096**  
 GloVe (word embedding software), 908,  
     909, 923, 926  
 Glover, F., 160, **1096**  
 Glover, K., 550, **1118**  
 GLUE (General Language Understanding  
     Evaluation), 930  
 Gluss, B., 560, **1096**  
 Glymour, C., 296, 798, **1097**, **1113**  
 Go (game), 37, 45, 48, 207, 223, 835, 867  
 goal, **71**, **81**, 363  
     clause, **248**  
     common, 590  
     formulation of, **81**  
     inferential, 284  
     monitoring, 390  
     state, **83**, 123  
 goal-based action selection, 71  
 goal-based agent, 71–72, 78, 79  
 goal-directed reasoning, **250**  
 goal test  
     early, **94**  
     late, **94**  
 God, existence of, 427  
 Godefroid, P., 399, **1096**, **1113**  
 Gödel, K., 27, 328, 1034, **1096**  
 Gödel number, 323  
 Goebel, J., 799, **1096**  
 Goel, A., 627, **1085**  
 Goel, S., 1060, **1091**  
 Goertzel, B., 51, **1096**  
 GOFAI (Good Old-Fashioned AI), **1033**  
 Gogate, V., 474, **1096**  
 gold, 228  
 Gold, E. M., 734, 905, **1096**  
 Goldberg, A. V., 126, **1096**  
 Goldberg, D. E., 161, **1109**  
 Goldberg, K., 162, **1116**  
 Goldberg, Y., 890, 900, 904, 929, **1096**,  
     *1103*, **1108**  
 Goldman, R., 162, 400, 666, **1090**, **1097**,  
     *1116*  
 Goldstein, T., 1060, **1097**  
 Goldszmidt, M., 473, 478, 597, 798,  
     *1088*, **1095**, **1097**  
 Golgi, C., 29  
 Golomb, S., 189, **1097**  
 Golub, G., 734, **1097**

- Gomes, C., 160, 190, 266, 399, **1097**  
 Gomez, A. N., 901, 919, 931, **1115**  
 Gondek, D., 48, **1094**  
 Gonina, K., 900, **1090**  
 Gonnet, G. H., 737, **1087**  
 Gonthier, G., 188, **1097**  
 Good, I. J., 223, 427, 472, 473, 1055, **1097**  
 Good-Turing smoothing, 902  
 good and evil, 547  
 Gooday, J. M., 360, **1091**  
 Goodfellow, I., 820, 837–839, **1090**, **1096**, **1114**  
 Goodman, J., 903, **1090**, **1097**  
 Goodman, N., 358, **1097**, **1103**  
 Goodman, N. D., 667, 668, **1087**, **1097**, **1105**, **1117**  
 Good Old-Fashioned AI (GOFAI), **1033**  
 Goodrich, B., 476, 668, 798, **1090**  
 Google, 47, 49, 670, 814, 901, 903, 904, 1031, 1037, 1059, 1068  
 Google Duplex, 47  
 Google Knowledge Graph, 334  
 Google Scholar, 652  
 Gopnik, A., 296, **1097**  
 Gordon, A. D., 667, 668, **1091**, **1097**  
 Gordon, A. S., 359, **1097**  
 Gordon, G., 516, 598, 986, **1109**, **1111**, **1115**  
 Gordon, M. J., 296, **1097**  
 Gordon, N., 516, 517, **1086**, **1093**, **1097**  
 Gordon, S. A., 225, **1097**  
 gorilla problem, **51**  
 Gorry, G. A., 428, **1097**  
 Gottlob, G., 191, **1097**  
 Gotts, N., 360, **1091**  
 Goyal, N., 927, 930, **1104**  
 Goyal, Y., 1017, **1097**  
 GP-CSP (planning system), 399  
 GPT-2 (language model), 883, 884, 927, 930, 1072  
 GPU (graphics processing unit), 33, 45  
 Grace, K., 46, **1097**  
 graceful degradation, 584  
 gradient, **138**, 719
  - empirical, **138**, 862
  - exploding, **825**
  - vanishing, **807**, 825
 gradient boosted regression tree (GBRT), 719  
 gradient boosting, 716, **719**  
 gradient boosting machine (GBM), 719  
 gradient descent, 132, **695**
  - batch, **697**
  - stochastic, **697**, 816
 Graepel, T., 45, 48, 201, 220, 223, 224, 667, 871, **1097**, **1098**, **1112**  
 Graham, S. L., 904, **1097**  
 Gramfort, A., 738, **1109**  
 grammar, 874, 883, 884, 1081
  - augmented, **892**
  - categorial, 904
  - context-free, 884, 902, 903, **1081**
    - lexicalized, **892**
    - probabilistic, 833–886, **884**, 903
  - dependency, 904
  - induction of, 905
  - lexical-functional (LFG), 904
  - phrase structure, 902
 grand coalition, **616**  
 graph, **83**
  - coloring, 188
  - Eulerian, **162**
 graphical model, 430, 478  
 graphics processing unit (GPU), 33, 45  
 Graphplan (planning system), 370, 399  
 graph search, **92**  
 grasping, 985  
 Grassmann, H., 296, **1097**  
 Graunt, J., 26  
 Gravano, L., 906, **1085**  
 Grave, E., 903, **1100**  
 Graves, A., 830, 835, 838, 841, 871, 873, 900, **1106**, **1115**  
 Grayson, C. J., 525, **1097**  
 Greaves, M., 1061, **1085**  
 Greece, 265, 358  
 Green, B., 904, **1097**  
 Green, C., 38, 296, 327, 329, **1097**  
 Green, P., 1031, **1089**  
 Green, S., 873, **1087**  
 Green, T., 838, **1100**  
 Greenbaum, S., 904, **1110**  
 Greenspan, M., 225, **1103**  
 Greiner, R., 798, **1090**  
 Grenager, T., 638, **1112**  
 Gribkoff, E., 668, **1097**  
 grid search, **689**  
 Griffiths, T. L., 136, 161, 296, 550, 588, **1097**, **1107**, **1110**, **1114**  
 Grinstead, C., 428, **1097**  
 Grisel, O., 738, **1109**  
 GRL (robot control language), 987  
 Grosz, B. J., 45, 627, 639, **1097**, **1099**  
 grounding, **234**, **647**  
 ground resolution theorem, **246**, 321  
 ground term, **279**, 298  
 ground truth, **671**  
 Grove, A., 428, 549, **1086**  
 Grove, W., 356, **1097**  
 Gruber, T., 334, 358, **1097**  
 Grumberg, O., 399, **1091**  
 GSAT (satisfiability algorithm), 267  
 Gu, J., 190, 266, **1097**, **1113**  
 Guan, M. Y., 838, **1109**  
 Guard, J., 331, **1097**  
 Guestrin, C., 551, 587, 637, 736, 737, 873, **1090**, **1097**, **1102**, **1105**, **1110**  
 Guez, A., 37, 45, 48, 201, 220, 222–224, 871, **1112**  
 Gugger, S., 738, **1099**  
 Guha, R. V., 334, 357, **1089**, **1103**  
 Guibas, L.J., 475, 985, **1097**, **1115**  
 guided missile, 1038  
 Guiver, J., 667, **1097**  
 Guizzo, E., 47, **1085**  
 Gulcehre, C., 837, **1092**  
 Gulshan, V., 48, **1097**  
 Gunkel, D. J., 1062, **1097**  
 Gunning, D., 1061, **1097**  
 Guo, C., 1060, **1097**  
 Guo, J., 476, 668, 798, **1090**  
 Gupta, A., 561, 986, **1102**, **1109**  
 Gupta, R., 928, **1110**  
 Gupta, V., 48, **1104**  
 Gururangan, S., 928, **1097**  
 Gustafsson, F., 516, **1098**  
 Guterres, A., 1039  
 Guthrie, F., 188  
 Guugu Yimithirr, 271  
 Guyon, I., 44, 736, 737, 837, 1029, **1088**, **1097**, **1103**
- 
- H**
- $\mathcal{H}$  (hypothesis space), 671  
 H (entropy), **679**, 680  
 h (heuristic function), 102  
 $h_{MAP}$  (MAP hypothesis), 774  
 HACKER (planning system), 398  
 Hacking, I., 429, **1097**  
 Hadfield-Menell, D., 80, 551, 587, 638, 872, **1097**, **1105**  
 Hager, G., 45, **1114**  
 Hahn, M., 736, **1091**  
 Hahnel, D., 984, **1089**  
 Haider, M., 737, **1096**  
 Hailperin, T., 665, **1097**  
 Haimes, M., 668, **1106**  
 Hajic, J., 890, **1108**  
 Hajishirzi, H., 931, **1112**  
 Haken, W., 188, **1085**  
 HAL 9000 computer, 472, 1036, 1058  
 Hald, A., 429, **1097**  
 Hales, T., 331, **1097**  
 Halevy, A., 44, 329, 358, 737, 901, 906, **1089**, **1097**  
 Halgren, E., 271, **1111**  
 Hall, L. O., 725, 1046, **1090**  
 Halpern, J. Y., 296, 359, 428, 665, **1086**, **1094**, **1097**  
 halting problem, 300  
 Hamilton, A., 901  
 Hamiltonian Monte Carlo, 668  
 Hamm, F., 358, **1115**  
 Hammersley, J. M., 872, **1098**  
 Hamming distance, **706**

- Hamori, S., 516, 1087  
 ham sandwich, 899  
 Han, J., 738, 1098  
 Han, X., 29, 1098  
 Hanan, S., 399, 1094  
 Hand, D. J., 1041, 1088  
 hand-tuning, 689  
 Handschin, J. E., 516, 1098  
 Handscomb, D.C., 872, 1098  
 Hanks, S., 401, 1094  
 Hannun, A., 1060, 1097  
 Hans, A., 872, 1098  
 Hansen, E., 127, 162, 389, 401, 587, 1098, 1118  
 Hansen, M. O., 189, 1085  
 Hansen, P., 266, 1098  
 Hanski, I., 80, 1098  
 Hansson, O., 127, 1098  
 happy graph, 679  
 haptics, 985  
 Harabagiu, S. M., 905, 1108  
 Harada, D., 586, 1107  
 Haralick, R. M., 189, 1098  
 Hardin, G., 639, 1053, 1098  
 Hardt, M., 734, 1060, 1094, 1098, 1104, 1118  
 Harel, D., 329, 1090  
 Harman, D., 901, 1085  
 Harnish, R., 48, 1093  
 HARPY (speech recognition system), 160  
 Harris, T., 1066, 1098  
 Harris, Z., 902, 1098  
 Harrison, J., 331, 549, 1097, 1098  
 Harrison, M. A., 904, 1097  
 Harrow, A. W., 1069, 1098  
 Harsanyi, J., 551, 637, 1098  
 Harshman, R. A., 903, 929, 1093  
 Hart, P. E., 125, 162, 400, 401, 428, 738, 798, 800, 1093–1095, 1098  
 Hart, T. P., 221, 1098  
 Hartley, H., 799, 1098  
 Hartley, R., 1030, 1031, 1098  
 Harutyunyan, A., 872, 1107  
 Harvard, 529  
 Harvey Mudd University, 1045  
 Hashimoto, K., 1072, 1098  
 Haslum, P., 398, 399, 401, 1098  
 Hassabis, D., 37, 45, 48, 49, 222–225, 871, 873, 1087, 1106, 1112, 1115  
 Hassidim, A., 1069, 1098  
 Hastie, T., 735, 736, 738, 800, 1095, 1098, 1100  
 Hastings, W. K., 475, 1098  
 Hatem, M., 126, 1089, 1098  
 Haugeland, J., 1058, 1098  
 Haussler, D., 516, 735, 1088, 1102  
 Havelund, K., 327, 1098  
 Havenstein, H., 44, 1098  
 Hawking, S., 51  
 Hawkins, J., 836, 1098  
 Hay, N., 587, 1070, 1098  
 Hayes, P. J., 267, 358–360, 1057, 1095, 1098, 1105  
 Hays, J., 44, 1098  
 He, H., 725, 1046, 1098  
 He, K., 837, 1098, 1111  
 He, Y., 873, 1096  
 head (in NLP), 892  
 head (of Horn clause), 248  
 Heafield, K., 903, 1089  
 Hearst, M. A., 903, 1111  
 Heath, M., 734, 1097  
 Heath Robinson, 32  
 heavy-tailed distribution, 160  
 Heawood, P. J., 1034, 1098  
 Hebb, D. O., 35, 39, 870, 1098  
 Hebbian learning, 35  
 Hebert, M., 868, 970, 986, 1030, 1099, 1101, 1118  
 Heckerman, D., 473, 477, 516, 798, 1089, 1098, 1099, 1113  
 Hedau, V., 1021, 1101  
 hedonic calculus, 548  
 Heess, N., 986, 1098, 1104  
 Heidari, H., 1060, 1087  
 Heidegger, M., 1057, 1098  
 Heinlein, R. A., 1071, 1098  
 Heitz, G., 334, 1093  
 Held, M., 127, 1098  
 Hellmann, S., 357, 1103  
 Helmert, M., 125, 398, 399, 1098, 1110, 1112  
 Helmholtz, H., 30, 1027  
 Hempel, C., 25  
 Henaff, M., 837, 1090  
 Hendeby, G., 516, 1098  
 Henderson, P., 871, 1095  
 Henderson, T. C., 189, 1106  
 Hendler, J., 357, 400, 402, 1085, 1087, 1094, 1113  
 Henrion, M., 80, 439, 473, 475, 549, 1098, 1099, 1109  
 Henry, H., 873, 1100  
 Henzinger, M., 905, 1113  
 Henzinger, T. A., 79, 1098  
 Hephaistos, 982  
 Herbrand's theorem, 322, 328  
 Herbrard, J., 300, 322, 328, 1098  
 Herbrand base, 322  
 Herbrand universe, 321, 328  
 Herbrich, R., 667, 1098  
 Herbster, M., 1069, 1093  
 Herden, G., 548, 1087  
 Hernandez, D., 33, 1069, 1085  
 Hernández-Orallo, J., 1058, 1098  
 Herring, R., 516, 1099  
 Herskovits, E., 798, 1091  
 Hertz, J. A., 838, 1102  
 Hess, C., 1061, 1098  
 Hessian, 139  
 Hestness, J., 906, 1090  
 Heule, M., 267, 1087  
 heuristic, 123  
 admissible, 104, 371  
 composite, 118  
 degree, 177, 189, 252  
 for planning, 371–374  
 function, 102, 115–122  
 inadmissible, 107  
 least-constraining-value, 177  
 Manhattan, 116  
 min-conflicts, 181  
 minimum remaining values, 177, 189, 308, 396  
 null move, 222  
 search, 125  
 straight-line, 103  
 heuristic function, 102  
 Heuristic Programming Project (HPP), 41  
 Hewitt, C., 329, 636, 1098  
 hexapod robot, 975  
 Hezaveh, Y. D., 670, 1098  
 hidden Markov model (HMM), 43, 479, 491, 491–496, 503, 515, 795, 881  
 hidden variable, 443, 788  
 HIERARCHICAL-SEARCH, 377  
 hierarchical decomposition, 375  
 hierarchical look-ahead, 383  
 hierarchical reinforcement learning, 858, 1065  
 hierarchical structure, 1065  
 hierarchical task network (HTN), 375, 397, 858  
 Hierholzer, C., 162, 1098  
 high-level action, 375  
 higher-order logic, 273  
 Hilbert, M., 737, 1098  
 Hilgard, E. R., 870, 1098  
 Hill, F., 930, 931, 1116  
 HILL-CLIMBING, 129  
 hill climbing, 129, 159  
 first-choice, 131  
 random-restart, 131  
 stochastic, 131  
 Hind, M., 1046, 1047, 1060, 1061, 1087, 1098  
 Hingorani, S. L., 667, 1091  
 Hinrichs, T., 225, 1104  
 Hinrichs, T. R., 357, 1095  
 Hintikka, J., 358, 1098  
 Hinton, G. E., 35, 44, 136, 161, 726, 836–839, 872, 900, 905, 1029, 1072, 1086, 1092, 1098, 1099, 1102, 1103, 1107, 1111–1113  
 HIPAA, 1041  
 Hipp, J. D., 48, 1104, 1113  
 Hirschberg, J., 45, 1114

- Hirth, M., 1068, 1099  
 Hitachi, 377  
 HMM (hidden Markov model), 43, 479, 491, 491–496, 503, 515, 795, 881  
 Ho, J., 873, 985, 1095, 1112  
 Ho, M. K., 872, 1099  
 Ho, T. K., 736, 737, 1097, 1099  
 Ho, Y.-C., 836, 1089  
 Hoane, A. J., 222, 1089  
 Hobbes, T., 24  
 Hobbs, J. R., 359, 361, 904, 1097, 1099  
 Hochreiter, S., 838, 1099  
 Hodges, J. L., 735, 1095  
 Hoff, M. E., 39, 870, 1116  
 Hoffman, G., 986, 1114  
 Hoffman, M., 476, 667, 668, 798, 903, 1090, 1099, 1115  
 Hoffman, S. C., 1047, 1060, 1087  
 Hoffmann, J., 372, 374, 398–401, 1099  
 Hofleitner, A., 516, 1099  
 Hofmann-Wellenhof, R., 48, 1104  
 Hogan, N., 985, 1099  
 Hoiem, D., 1021, 1030, 1099, 1101  
 Holenstein, R., 517, 1085  
 Holland, J. H., 161, 1099, 1106  
 Hollerbach, J. M., 985, 1094  
 Holte, R. C., 106, 114, 126, 127, 637, 1085, 1087, 1090, 1094, 1099, 1103  
 Holzmann, G. J., 327, 1099  
 Homan, K. A., 1064, 1102  
 homeostatic, 34  
*Homo sapiens*, 19, 874  
 Hood, A., 29, 1099  
 Hooker, J., 191, 1099  
 Hoos, H. H., 160, 190, 549, 737, 1088, 1099, 1115  
 Hopcroft, J., 738, 984, 1088, 1112  
 Hopfield, J. J., 839, 1099  
 Hopfield network, 839  
 Hopkins Beast, 983  
 HORIZON (reinforcement learning platform), 873  
 horizon, 990  
 infinite, 586  
 horizon (in MDPs), 555  
 horizon effect, 204  
 Horn, A., 266, 1099  
 Horn, B. K. P., 1030, 1099  
 Horn, K. V., 428, 1099  
 Horn, W., 334, 1093  
 Horn clause, 247  
 Horn form, 265, 266  
 Horning, J. J., 1099  
 Horowitz, M., 268, 1108  
 Horrocks, J. C., 428, 1092  
 Horsfall, P., 667, 1087  
 Horswill, I., 987, 1099  
 Horvitz, E. J., 80, 473, 516, 549, 1070, 1099, 1108  
 Hoseini, S. S., 161, 1096  
 Hößfeld, T., 1068, 1099  
 Hotelling, H., 838, 1099  
 Houde, S., 1047, 1060, 1087  
 HOUDINI (chess program), 222  
 Houlsby, N., 737, 1117  
 Houston, M., 48, 1103  
 Hovel, D., 473, 1099  
 Howard, J., 734, 738, 930, 1099, 1105  
 Howard, R. A., 534, 548–550, 586, 1099, 1106  
 Howe, A., 362, 398, 1096  
 Howe, D., 331, 1099  
 Howe, P., 737, 1061, 1106  
 Howson, C., 665, 1099  
 HPP (Heuristic Programming Project), 41  
 Hruschka, E., 901, 1106  
 HSCP (planning system), 400  
 Hsiao, K., 588, 1099  
 Hsieh, H.-P., 736, 1117  
 HSP (Heuristic Search Planner), 398  
 Hsu, D., 588, 734, 1086, 1087  
 Hsu, F.-H., 222, 1089, 1099  
 Hsueh, C.-H., 222, 1116  
 HTML, 891  
 HTN (hierarchical task network), 375, 397, 858  
 Hu, H., 928, 1090  
 Hu, J., 587, 638, 1090, 1099  
 Hu, Y.-T., 1023, 1093  
 Hua, Y., 1060, 1086  
 Huang, A., 37, 45, 48, 222, 223, 1112  
 Huang, L., 896, 904, 905, 1096, 1118  
 Huang, S., 48, 1104  
 Huang, T., 516, 666, 598, 1095, 1099  
 Huang, X., 47, 1117  
 Huang, Y., 838, 1110  
 Huang, Z., 837, 1111  
 Hubble Space Telescope, 167, 181, 402  
 Hubel, D. H., 837, 1029, 1031, 1099  
 Huber, M., 983, 1091  
 Hubert, T., 45, 48, 201, 220, 223, 224, 871, 1112  
 Huddleston, R. D., 904, 1099  
 Huet, G., 330, 1087  
 Huffman, D. A., 38, 1099  
 Huffman, S., 983, 1091  
 Hughes, B. D., 156, 1099  
 Hughes, G. E., 358, 1099  
 Hughes, M., 901, 1090  
 HUGIN (Bayes net system), 474, 516  
 Huhns, M. N., 80, 1099  
 Hui, F., 45, 48, 1112  
 human-computer interaction, 32  
 human-robot interaction, 971, 986  
 human-level AI, 50  
 human actions, 986  
 human judgment, 528  
 human performance, 19  
 Hume, D., 24, 1099  
 Humphrys, M., 1035, 1099  
 Hungarian algorithm, 658  
 Hunkapiller, T., 516, 1086  
 Hunsberger, L., 627, 639, 1099  
 Hunt, J. J., 986, 1104  
 Hunt, W., 330, 1099  
 Hunter, L., 799, 1099  
 Huq, A., 1060, 1091  
 Hur, C.-K., 668, 1099  
 Hurst, M., 906, 1099  
 Hurst, S., 871, 1111  
 Hurwicz, L., 639, 1099  
 Hussein, A. I., 700, 702, 1101  
 Hut, P., 48  
 Hutchinson, B., 1046, 1106  
 Hutchinson, S., 79, 985, 987, 1090  
 Huth, M., 297, 1099  
 Huttenlocher, D., 1029, 1099  
 Hutter, F., 399, 737, 838, 1094, 1099, 1112, 1115  
 Huygens, C., 426, 637, 1100  
 Huyn, N., 125, 1100  
 Huynh, V. A., 588, 1100  
 Hwa, R., 904, 1100  
 Hwang, C. H., 357, 1100  
 Hyafil, L., 734, 1100  
 HYBRID-WUMPUS-AGENT, 260  
 hybrid A\*, 953  
 hyperbolic reward, 588  
 hyperparameter, 684, 781  
 hypertext, 32  
 hypertree width, 191  
 hypothesis, 669, 671
  - approximately correct, 691
  - null, 681
  - prior, 773
  - space, 671
 Hyun, S., 984, 1092
- 
- I**
- i.i.d. (independent and identically distributed), 683, 773  
 Iagnemma, K., 984, 1089  
 Iatauro, M., 47, 1086  
 IBAL (probabilistic programming language), 667  
 Ibarz, J., 820, 985, 986, 1088, 1097, 1103  
 IBM, 37, 1059, 1068  
 identifiability, 794  
 identification in the limit, 734  
 identity matrix ( $I$ ), 1077  
 identity uncertainty, 648  
 idiot Bayes, 420  
 IEEE P7001, 1048  
 Ieong, S., 638, 1100  
 ignorance, 404

- ignore-delete-lists heuristic, **372**  
 ignore-preconditions heuristic, **371**  
 Iida, H., 222, *1111*  
 III, H. D., 1046, *1096*  
 IJCAI (International Joint Conference on AI), 53, 127  
 ILOG (constraint logic programming system), 330  
 ILQR (iterative LQR), **962**, 982, 985  
 image, **989**  
   formation, 989–995  
   segmentation, 1001–1002  
 image captioning, 48  
 ImageNet, **723**  
 ImageNet (image data set), 43, 44, 46, 1003, 1071  
 image transformation, **1020**  
 IMDB, 1042  
 imitation learning, **864**, 973  
 imperfect information, 220, 608, **609**  
 implementation (of a high-level action), **375**  
 implementation level, **228**  
 implication, **235**  
 implicative normal form, 317  
 implicit model, **831**  
 importance sampling, **457**  
   adaptive, 475  
   sequential, **509**  
 imputation, **617**  
 inadmissible heuristic, 107, **108**  
 incentive, **591**  
 incentive compatible, 625  
 inclusion-exclusion principle, **411**  
 Inclusive Images Competition, 1046  
 income inequality, 1038, **1051**  
 incomplete information game, **968**  
 incompleteness, 314  
   theorem, 27, 323, 1034  
 incremental search, **159**  
 independence, 415–417, **416**, 419, 425  
   absolute, 416, 419  
   conditional, **419**, 424, 426, 433, 436–445, 472, 487  
   context-specific, **438**  
   marginal, 416  
 independent subproblems, **183**  
 indexical, **897**  
 indexing, **302**, 302–304  
 India, 34, 188, 356  
 indicator variable, **791**  
 indifference, principle of, **427**  
 individually rational offer, **633**  
 individual rationality, **617**  
 individuation, **339**  
 induced width, 190  
 induction, **24**, 670  
   mathematical, 27  
 inductive learning, 671–674, 733  
 inductive logic programming (ILP), 428, 750, **767**  
 attribute-based learning algorithm, 759, 760  
 CHOOSE-LITERAL, 763  
 constructive induction algorithms, **760**  
 description, 758  
 family tree, 760  
 inverse entailment, **765**  
 inverse resolution, **763**, 764–766  
*Journal of Molecular Biology*, **766**  
 knowledge-based induction problem, 758  
 linear resolution, **765**  
 LINUS system, 765  
 molecular biology experiments, 767  
 NEW-LITERALS, 762–763  
 positive and negative examples, 758, 759  
 PROGOL system, 765, 766  
 Prolog Horn-clause deduction, 765  
 resolvent C, **763**  
 three-dimensional configuration, 758  
 top-down inductive learning methods, 761–763  
 Induráin, E., 548, *1087*  
 Indyk, P., 735, *1085*, *1096*  
 inference, **227**  
   probabilistic, **413**, 413–415, 430  
   inference by enumeration, 445  
   inference procedure, 291  
   inference rule, **241**, 265  
   inferential frame problem, **257**, 268  
   infinite branching factor, 137  
   infinite horizon, 586  
   influence diagram, 472, 518, **534**, 534–537, 547  
 INFORMATION-GATHERING-AGENT, 541  
 information extraction, **901**  
 information gain, **680**, 681, 964  
 information gain ratio, 683  
 information gathering, **59**, 963  
 information retrieval (IR), **901**, 905  
 information set, **610**  
 information theory, 679–680, 733  
 information value, **537**, 550  
 informed search, 81, **102**, 102–123  
 Ingerman, P. Z., 903, *1100*  
 Ingham, M., 267, *1116*  
 inheritance, **335**, 347  
   multiple, **348**  
 initial state, **83**, 86, 123, 193, 363  
 initial state model, 482  
 input gate (in LSTM), **826**  
 input resolution, **326**  
 inside–outside algorithm, 891  
 instance (of a schema), **135**  
 instance-based learning, **704**, 704–706  
 instant runoff voting, **630**  
 insurance premium, **525**  
 integrated information theory, 1058  
 intelligence, **19**, 54  
 intelligence augmentation, **32**  
 intelligent backtracking, 179–181, 252  
 interior-point method, 161  
 interleaved execution, **592**  
 interleaving, 152, 178, 398  
 internal state, **69**  
 International Joint Conference on AI (IJCAI), 53, 127  
 interpolation (of data), **686**  
 interpolation smoothing, 902  
 interpretability, **729**, 737  
 interpretation, **276**, 295  
   extended, **279**, 295  
   intended, **276**  
 interval, 342–343  
 Intille, S., 516, *1100*  
 intractability, 39  
 intrinsic property, **340**  
 introspection, **20**, 31  
 invariance, temporal, 811  
 inverse (of a matrix), **1077**  
 inverse dynamics, **958**  
 inverse kinematics, **947**  
 inverse reinforcement learning, **864**, 1054, 1065  
 inverted pendulum, **867**  
 Ioffe, S., 837, *1100*  
 IPL (programming language), 36  
 IQ test, 38  
 IR (information retrieval), **901**, 905  
 Irpan, A., 835, 871, 985, 1088, *1100*  
 irrationality, 19, 521, **550**  
 irreversible action, **154**, 850  
 Irving, G., 330, *1085*, *1104*  
 IS-A links, 359  
 Isard, M., 516, *1100*  
 Isbell, C., 359, 1065, *1104*  
 ISBN, 648  
 Isele, R., 357, *1103*  
 ISIS (planning system), 401  
 Islam, R., 871, *1095*  
 Isola, P., 930, 1021, 1022, *1100*, *1118*  
 iterated best response, 600  
 iterated game, 604  
**ITERATIVE-DEEPENING-SEARCH**, **99**  
 iterative deepening search, **98**, 98–100, 123, 125, 201, 204, 376  
 iterative expansion, **126**  
 iterative LQR (ILQR), **962**, 982, 985  
 Ivanov, V., 1043, *1088*  
 Iwasawa, S., 1057, *1109*  
 iWeb (language corpus), 876  
 IXTeT (planning system), 401  
 Iyyer, M., 930, *1109*

**J**

Jaakkola, T., 476, 838, 1100, 1111  
 Jabbari, S., 1060, 1087  
 JACK (bridge program), 224  
 Jackel, L., 44, 837, 1029, 1103  
 Jackson, C., 666, 798, 1105  
 Jacobi, C. G., 667  
 Jacobs, D., 1014, 1101  
 Jacobson, D. H., 985, 1100  
 Jacquard, J., 33  
 Jacquard loom, 33  
 Jaderberg, M., 48, 225, 838, 1100  
 Jaffar, J., 330, 1100  
 Jaggi, M., 838, 1117  
 Jaguar, 401  
 Jain, A., 48, 905, 923, 1104, 1108, 1115  
 Jain, B., 873, 1111  
 Jain, D., 666, 1100  
 Jaityl, N., 900, 905, 1090, 1099  
 Jakob, M., 357, 1103  
 James, G., 738, 1100  
 James, W., 31  
 janitorial science, 57  
 Jankowiak, M., 667, 1087  
 Janz, D., 737, 1113  
 Janzing, D., 476, 1109  
 Japan, 41, 983  
 Jarrett, K., 837, 1100  
 Jasra, A., 517, 1093  
 Jastrzebski, S., 734, 1086  
 Jaumard, B., 266, 1098  
 Jauvin, C., 929, 1087  
 Jaynes, E. T., 412, 427–429, 1100  
 Jeffrey, R. C., 427, 548, 1100  
 Jeffreys, H., 902, 1100  
 Jelinek, F., 902, 906, 931, 1089, 1100  
 Jenkin, M., 987, 1094  
 Jenkins, G., 515, 838, 1088  
 Jenkins, N. W., 48, 1093  
 Jennings, H. S., 31, 1100  
 Jennings, N. R., 638, 1110  
 Jenniskens, P., 390, 1100  
 Jensen, F., 473, 474, 550, 1085, 1093, 1107  
 Jensen, F. V., 473, 474, 478, 1085, 1100  
 Jentzsch, A., 357, 1103  
 Jeopardy, 44, 48  
 Jevons, W. S., 266  
 Ji, Z., 1060, 1100  
 Jiang, H., 1045, 1100  
 Jiang, K., 986, 1085  
 Jiao, J., 838, 1118  
 Jie, K., 48  
 Jimenez, P., 162, 401, 1100  
 Joachims, T., 736, 903, 1100  
 job, 393, 1051  
 Job, J., 1069, 1107  
 job-shop scheduling problem, 393  
 Johansen, A. M., 517, 1093

Johanson, M., 48, 224, 637, 1088, 1107, 1118  
 Johnson, C. R., 80, 1089  
 Johnson, D. S., 1076, 1080, 1096  
 Johnson, I., 737, 1116  
 Johnson, J., 1069, 1115  
 Johnson, M., 901, 904, 905, 1057, 1090, 1094, 1100, 1103  
 Johnson, S. M., 571, 587, 1088  
 Johnson, W. W., 124, 1100  
 Johnston, M. D., 160, 190, 402, 1100, 1106  
 joint action, 593  
 joint agent, 971  
 joint probability distribution, 410 full, 411, 425, 430, 432–436  
 join tree, 452  
 Jonathan, P. J. Y., 1035, 1100  
 Jones, D. M., 587, 1096  
 Jones, G., 476, 1089  
 Jones, L., 901, 919, 931, 1090, 1115  
 Jones, M., 550, 1029, 1100, 1115  
 Jones, R., 906, 1100  
 Jones, R. M., 329, 1100, 1117  
 Jones, T., 79, 1100  
 Jonsson, A., 47, 401, 1100  
 Jordan, M., 872, 1112  
 Jordan, M. I., 476, 516, 517, 588, 779, 799, 837, 838, 863, 868, 872, 896, 903, 986, 1088, 1100, 1104, 1107, 1111–1113, 1116  
 Joseph, A. D., 1061, 1086  
 Joshi, M., 927, 930, 1104  
 Jouannaud, J.-P., 330, 1100  
 Joulin, A., 903, 1100  
 Jouppi, N. P., 1069, 1100  
 Joy, B., 1061, 1100  
 Jozefowicz, R., 838, 929, 966, 986, 1085, 1100  
 Juang, B.-H., 515, 1110  
 Judah, K., 551, 1094  
 Juels, A., 161, 1060, 1100, 1115  
 Julesz, B., 1028, 1100  
 Julian, K. D., 588, 1100  
 Juliani, A., 873, 1100  
 Jung, M. W., 873, 1103  
 Junker, U., 330, 1100  
 Jurafsky, D., 891, 900, 903, 906, 1100, 1113  
 Just, M. A., 271, 1106  
 justification (in a JTMS), 354

**K**

k-anonymity, 1042  
 k-consistency, 172  
 k-DL (decision list), 693  
 k-DT (decision tree), 693  
 k-d tree, 706  
 k-fold cross-validation, 684

Kaack, L. H., 48, 1110  
 Kadane, J. B., 551, 637, 1100  
 Kaden, Z., 873, 1096  
 Kadian, A., 873, 1111  
 Kaelbling, L. P., 267, 516, 586–588, 668, 984, 1090, 1092, 1099, 1100, 1104, 1106, 1112, 1114  
 Kager, R., 902, 1100  
 Kaggle, 716  
 Kahn, H., 475, 872, 1100  
 Kahneman, D., 436, 528, 550, 1100, 1115  
 Kaindl, H., 127, 1100  
 Kaiser, L., 901, 919, 928, 930, 931, 1101, 1115  
 Kalakrishnan, M., 985, 1088  
 Kalchbrenner, N., 830, 838, 900, 1115  
 Kale, A. U., 48, 1104  
 Kaliszyk, C., 330, 331, 1097, 1104  
 Kalman, R., 497, 515, 1101  
 Kalman filter, 479, 497, 497–503, 515, 942 extended (EKF), 501, 942 switching, 502  
 Kalman gain matrix, 501  
 Kalra, N., 160, 1093  
 Kalyanakrishnan, S., 45, 1114  
 Kalyanpur, A. A., 48, 1094  
 Kamar, E., 45, 1114  
 Kamber, M., 738, 1098  
 Kambhampati, S., 162, 399–402, 1089, 1091, 1093, 1101  
 Kameya, Y., 666, 1111  
 Kaminka, G., 639, 1114  
 Kan, A., 125, 395, 402, 1103  
 Kanada, K., 48, 1104  
 Kanade, T., 46, 1029, 1030, 1101, 1111, 1115  
 Kanal, E., 1041, 1101  
 Kanal, L. N., 126, 1102  
 Kanazawa, A., 1014, 1020, 1101  
 Kanazawa, K., 516, 586, 588, 799, 1087, 1088, 1095, 1101, 1111  
 Kanefsky, B., 190, 1090  
 Kang, S. M., 1064, 1101  
 Kannan, A., 900, 1090  
 Kannan, K., 1047, 1060, 1087  
 Kannan, R., 738, 1088  
 Kanodia, N., 587, 1097  
 Kanoui, H., 296, 329, 1091  
 Kanoulas, E., 901, 1085  
 Kant, E., 329, 1089  
 Kant, I., 26  
 Kanter, J. M., 737, 1101  
 Kantor, G., 79, 985, 987, 1090  
 Kantorovich, L. V., 161, 1101  
 Kanwal, M. S., 734, 1086  
 Kaplan, D., 359, 1101  
 Kaplan, H., 126, 1096  
 Kaplow, R., 588, 1112

- Karaboga, D., 160, *1101*  
 Karafiát, M., 929, 930, *1106*  
 Karaletsos, T., 667, *1087*  
 Karamchandani, A., 475, *1101*  
 Karlin, S., 571, 587, *1088*  
 Karlsson, R., 516, *1098*  
 Karmarkar, N., 161, *1101*  
 Karmiloff-Smith, A., 905, *1094*  
 Karp, R.M., 27, 125, 127, 1080, 1098, *1101*  
 Karpas, E., 399, *1098*  
 Karpathy, A., 837, 913, 930, *1101, 1111*  
 Karpatne, A., 738, *1114*  
 Karras, T., 831, *1101*  
 Karsch, K., 1021, *1101*  
 Kartam, N. A., 401, *1101*  
 Kasami, T., 886, 904, *1101*  
 Kasif, S., 474, *1118*  
 Kasparov, G., 48, 222  
 Kassirer, J. P., 428, *1097*  
 Kataoka, T., 837, *1106*  
 Katehakis, M. N., 587, *1101*  
 Katriel, I., 189, *1115*  
 Katz, B., 901, *1101*  
 Katz, S., 191, *1091*  
 Kaufmann, M., 331, *1101*  
 Kautz, D., 402, *1092*  
 Kautz, H., 160, 190, 266, 267, 399, 474, *1097, 1101, 1111, 1112*  
 Kautz, J., 930, *1104*  
 Kavraki, L., 79, 985, 987, 1090, *1101*  
 Kavukcuoglu, K., 830, 835, 837, 838, 841, 871, 873, 900, 1100, *1104, 1106, 1115*  
 Kawczyński, M. G., 48, *1093*  
 Kay, A.R., 29, *1108*  
 Kaynama, S., 872, *1085*  
 Kazemi, S. M., 668, *1101*  
 KB (knowledge base), *227, 264*  
 KB-AGENT, *227*  
 Keane, M. A., 161, *1102*  
 Keane, P. A., 48, *1104*  
 Kearns, M., 587, 588, 735, 736, 738, 871, *1060, 1087, 1101*  
 Kebeasy, R. M., 700, 702, *1101*  
 Kedzier, D., 871, *1111*  
 Keeney, R. L., 529, 534, 549, 550, *1101*  
 keepaway, *859*  
 Kegelmeyer, W. P., 725, 1046, *1090*  
 Keil, F. C., 21, 1058, *1116*  
 Kelcey, M., 985, *1088*  
 Kelley, H. J., 40, 836, *1101*  
 Kelly, J., 639, 798, 799, *1089, 1090*  
 Kelly, K., 1056  
 Kembhavi, A., 931, *1112*  
 Kemp, C., 550, *1097*  
 Kemp, M., 1027, *1101*  
 Kempe, A. B., 1034, *1101*  
 Kenley, C. R., 473, *1112*  
 Kephart, J. O., 79, *1101*  
 Kepler, J., 1027  
 Keras (machine learning software), 738, *1072*  
 Kern, C., 48, *1104*  
 kernel (in neural networks), *811*  
 kernel (in regression), 709  
 kernel function, *712, 787*  
 kernelization, *714*  
 kernel machine, 710–714, *735*  
 kernel trick, *710, 713, 735*  
 Kernighan, B. W., 125, *1104*  
 Kersting, K., 666, 668, *1101, 1106*  
 Keskar, N. S., 931, *1101*  
 keyframe, *975*  
 Keynes, J. M., 427, 1049, *1101*  
 key vector (in transformers), *920*  
 Khairy, K., 734, *1105*  
 Khanna, R., 737, *1101*  
 Khare, R., 357, *1101*  
 Khatib, O., 985, 987, *1101, 1110, 1112*  
 Khorsand, A., 127, *1100*  
 Khosla, A., 837, *1111*  
 Khot, T., 901, 927, 931, 1017, *1091, 1097*  
 Khudanpur, S., 929, 930, *1106*  
 Kichkaylo, T., 47, *1086*  
 killer move, *201*  
 Kim, B., 737, *1093, 1101*  
 Kim, H. J., 868, 872, 986, *1107*  
 Kim, J.-H., 1033, *1101*  
 Kim, J. H., 472, *1101*  
 Kim, T. W., 1061, *1101*  
 Kimmig, A., 668, *1101*  
 Kinect, 935  
 kinematic state, *958*  
 kinesthetic teaching, *975*  
 King, H., 871, 873, *1087, 1106*  
 Kingma, D. P., 838, *1101*  
 King Midas problem, *51, 1054*  
 Kingsbury, B., 900, 905, *1099*  
 Kinsey, E., 124  
 kinship domain, 284–286  
 Kirchlechner, B., 666, *1100*  
 Kirchner, C., 330, *1100*  
 Kirk, D. E., 79, *1101*  
 Kirk, J. T., 1058  
 Kirkpatrick, S., 160, 190, *1101*  
 Kirman, J., 587, *1092*  
 Kiros, J. R., 837, *1086*  
 Kirubarajan, T., 79, *1086*  
 Kishimoto, A., 223, *1111*  
 Kisiel, B., 901, *1106*  
 Kisynski, J., 668, *1101*  
 Kitaev, N., 904, 928, 930, *1101*  
 Kitani, K. M., 868, *1101*  
 Kitano, H., 983, *1101*  
 Kitchin, D. E., 399, *1115*  
 Kjaerulff, U., 516, *1101*  
 Klarman, H. E., 549, *1101*  
 Klein, D., 896, 903, 904, 928, *1101, 1104, 1108*  
 Kleinberg, J. M., 1045, 1060, *1101*  
 Klempner, P., 639, *1101*  
 Klempner, G., 473, *1107*  
 Kneser, R., 903, *1101*  
 Knight, B., 39, *1088*  
 Knoblock, C. A., 126, 362, 398, 400, *1090, 1096, 1101*  
 KNOWITALL (information extraction system), 906  
 knowledge  
   acquisition, 41, *290*  
   and action, 25, 344–346  
   autonomous learning agent, 748  
   background, *227, 320, 749*  
   base (KB), *227, 264*  
   commonsense, 37  
   cumulative learning process, 748  
   diagnostic, 418  
   engineering, *289, 289–295, 433*  
   entailment constraint, *747*  
   explanation-based learning, *749*  
   inductive logic programming, 750  
   inferential behavior, 748–749  
   knowledge-based inductive learning, *750*  
   level, *228, 265*  
   model-based, 418  
   prior, 58, 59, *670*  
   relevance-based learning, *749, 750*  
   knowledge-based agents, *226*  
   knowledge-based inductive learning (KBIL), *750, 767*  
   knowledge-based system, 40–42, 870  
   knowledge compilation, *768*  
   knowledge representation, *20, 35, 37, 41, 226, 269–274, 332–361*  
   for everything, 332  
   language, *227, 264, 269*  
   uncertain, 430–432  
 knowledge state, 405  
 Knuth, D. E., 86, 221, 267, 330, 985, *1097, 1101*  
 Ko, J., 48, *1094*  
 Kober, J., 986, *1101*  
 Kobilarov, G., 334, 357, *1088*  
 Koch, C., 1058, *1091, 1101*  
 Kochenderfer, M. J., 588, 873, *1100, 1102*  
 Kociemba, H., 124, *1110*  
 Kocsis, L., 222, 587, *1102*  
 Koditschek, D., 986, *1102*  
 Koehn, P., 931, *1102*  
 Koelsch, S., 1027, *1102*  
 Koenderink, J. J., 1030, *1102*  
 Koenig, S., 163, 399, 401, 586, 984, *1098, 1102, 1113*  
 Kohlberger, T., 48, *1104*  
 Kohli, P., 668, *1102*

- Kolesky, D. B., 1064, *1102, 1107*  
 Kollar, T., 986, *1114*  
 Koller, D., 222, 428, 473, 478, 516, 517,  
   551, 587, 611, 637, 666–668, 798,  
   799, 903, 984, *1086–1088, 1090,*  
   *1095, 1097, 1099, 1101, 1102,*  
   *1106, 1108, 1109, 1111, 1114*  
 Kolmogorov's axioms, **411**  
 Kolmogorov, A. N., 427, 428, 515, 734,  
   *1102*  
 Kolmogorov complexity, **734**  
 Kolobov, A., 588, 667, *1105, 1106*  
 Kolter, J. Z., 868, *1102*  
 Koltun, V., 873, *1111*  
 KOMODO (chess program), 222  
 Kondrak, G., 189–191, *1102*  
 Konečný, J., 1043, *1102*  
 Konolige, K., 190, 360, 636, 984, 985,  
   987, 1088, 1089, *1102*  
 Kononova, O., 923, *1115*  
 Kontokostas, D., 357, *1103*  
 Koopmans, T. C., 586, *1102*  
 Korb, K. B., 478, *1102*  
 Koren, S., 735, *1087*  
 Koren., Y., 985, *1088*  
 Korf, R. E., 116, 124–127, 163, 221, 399,  
   *1094, 1102, 1109*  
 Kortenkamp, D., 983, *1091*  
 Koss, F., 983, *1091*  
 Kotthoff, L., 737, *1099*  
 Koutsoupias, E., 160, 266, *1102*  
 Kovacs, D. L., 398, *1102*  
 Kowalski, R., 296, 312, 317, 329, 358,  
   *1102, 1111*  
 Kowalski form, 317  
 Koyama, M., 837, *1106*  
 Koyejo, O. O., 737, *1101*  
 Koyejo, S., 1023, *1093*  
 Koza, J. R., 161, *1102*  
 Krakovna, V., 873, 1054, *1102, 1103*  
 Kramer, S., 666, *1101*  
 Kraska, T., 329, *1102*  
 Kraus, S., 45, 639, 640, *1102, 1114*  
 Kraus, W. F., 161, *1104*  
 Krause, A., 551, *1102*  
 Krause, J., 837, *1111*  
 Krauss, P., 665, *1112*  
 Krawiec, K., 161, *1113*  
 Kreitmann, P., 1060, *1087*  
 Kretch, K. S., 852, *1085*  
 Kreuter, B., 1043, *1088*  
 Kriegspiel, **214**  
 Krikun, M., 47, 834, 901, 916, *1117*  
 Kripke, S. A., 358, *1102*  
 Krishna, V., 639, *1102*  
 Krishnamurthy, V., 588, *1102*  
 Krishnan, A., 516, *1104*  
 Krishnan, T., 799, *1106*  
 Krishnanand, K., 160, *1102*  
 Krizhevsky, A., 44, 837, 838, 986, 1029,  
   *1102, 1103, 1113*  
 Krogh, A., 516, 838, *1102*  
 Krueger, D., 734, *1086*  
 Kruppa, E., 1028, *1102*  
 Ktesibios of Alexandria, 33  
 Kübler, S., 904, *1102*  
 Kuffner, J. J., 985, *1102*  
 Kuhlmann, G., 873, *1114*  
 Kuhn, H. W., 658, 667, 637, *1102*  
 Kuipers, B. J., 360, 984, *1102*  
 Kulkarni, T., 668, *1102*  
 Kullback-Leibler divergence, **809**  
 Kumar, M. P., 734, *1087*  
 Kumar, P. R., 79, *1102*  
 Kumar, S., 33, 901, 1069, *1102, 1117*  
 Kumar, V., 126, 738, 1041, *1090, 1102,*  
   *1114*  
 Kumaran, D., 45, 48, 201, 220, 223, 871,  
   *873, 1106, 1112*  
 Kuniyoshi, Y., 983, *1101*  
 Kuo, W.-C., 1015, *1095*  
 Kuppuswamy, N., 1033, *1101*  
 Kuprel, B., 48, *1094*  
 Kurakin, A., 838, *1090*  
 Kurien, J., 162, *1102*  
 Kurth, T., 48, *1103*  
 Kurzweil, R., 30, 1055, 1056, 1061, *1103*  
 Küttler, H.,  
 Kwok, C., 905, *1103*
- 
- L**
- L-BFGS (optimization algorithm), 735  
 label (in machine learning), **671**  
 label (in plans), **143**  
 Laborie, P., 402, *1103*  
 Lacoste-Julien, S., 734, *1086*  
 Ladkin, P., 358, *1103*  
 Lafferty, J., 906, *1103*  
 Lagoudakis, M. G., 637, *1097*  
 Laguna, M., 160, *1096*  
 Lahiri, S., 838, *1109*  
 Lai, J. C., 929, *1089*  
 Lai, M., 45, 48, 201, 220, 223, 871, *1112*  
 Lai, T. L., 575, 587, *1103*  
 Laine, S., 831, *1101*  
 Laird, J., 310, 329, 400, *1100, 1103, 1117*  
 Laird, N., 515, 799, *1093*  
 Laird, P., 160, 190, *1106*  
 Lake, B., 667, *1103*  
 Lake, R., 223, *1111*  
 Lakemeyer, G., 984, *1089*  
 Lakoff, G., 357, 904, 1057, *1103*  
 Lally, A., 48, *1094*  
 Lam, J., 225, *1103*  
 Lamarck, J. B., 136, *1103*  
 Lample, G., 930, *1103*  
 Lamure, M., 549, *1087*  
 Lanctot, M., 45, 48, 201, 220, 223, 871, *1112*  
 Landauer, T. K., 903, 929, *1093*  
 Landhuis, E., 528, *1103*  
 landmark (recognizable feature), **940**  
 landmark point, **120**  
 land mine, 1038  
 Landolin, J. M., 735, *1087*  
 landscape (in state space), **128**  
 Lang, J., 639, *1088*  
 Langdon, W., 162, *1103, 1109*  
 Lange, D., 873, *1100*  
 Langton, C., 161, *1103*  
 language, 884  
   formal, 874  
   model, **875**, 911  
     in disambiguation, 899  
     masked, **925**  
   natural, 22, 270, 874  
   processing, 35, 874–931  
   source, **915**  
   target, **915**  
     understanding, 38, 41  
 language identification, **877**  
 Lao, N., 334, *1093*  
 LaPaugh, A. S., 125, *1103*  
 Laplace, P., 26, 426, 427, 878, 902, *1103*  
 Laplace smoothing, 878  
 Larkey, P. D., 637, *1100*  
 Larochelle, H., 690, 838, *1103, 1113*  
 Larsen, B., 474, *1103*  
 Larson, K., 638, *1111*  
 Larson, S. C., 734, *1103*  
 Laruelle, H., 401, *1096*  
 Laskey, K. B., 667, 550, *1103*  
 Lassez, J.-L., 330, *1100*  
 Lassila, O., 357, *1087*  
 late move reduction, **206**  
 latent Dirichlet allocation, 903  
 latent semantic indexing, 903  
 latent variable, **788**  
 Latham, D., 873, *1105*  
 Latombe, J.-C., 402, 984–986, *1093,*  
   *1101, 1103, 1118*  
 lattice theory, 331  
 Laugherty, K., 904, *1097*  
 Laurent, C., 900, *1118*  
 Lauritzen, S., 473, 474, 548, 549, 798,  
   *799, 1091, 1103, 1108, 1113*  
 LaValle, S., 402, 985, 987, *1102, 1103*  
 Lave, R. E., 588, *1111*  
 Lavie, A., 904, *1111*  
 Lawler, E. L., 125, 126, 395, 402, *1103*  
 laws of robotics, 1058  
 laws of thought, 21–22  
 layer (in neural networks), **801**  
   hidden, **805**  
   input, 807  
   mixture density, **810**  
   output, **805**

- Lazanas, A., 986, **1103**  
 laziness, **404**  
 La Mettrie, J. O., 1052, 1057, **1103**  
 La Mura, P., 549, **1103**  
 LCF (Logic for Computable Functions), 296  
 Le, Q. V., 47, 736, 834, 838, 900, 901, 916, 930, 931, 1072, **1092**, **1109**, **1110**, **1112**, **1114**, **1117**, **1118**  
 Le, T. A., 475, 668, **1103**  
 Leacock, C., 356, **1089**  
 leak node, **439**  
 Leaper, D. J., 428, **1092**  
**LEARN-DECISION-TREE**, **678**  
 learned index structures, 329  
 learning, **59**, 64, 78, 228, 235, 669, 874  
   apprenticeship, 1054  
   arch-learning program, 743  
   assessing performance of, 683–684  
   attributes, **739**  
   Bayesian, 719, **773**, 773–774, 797  
   Bayesian network, 785–786  
   in the blocks world, 38  
   boundary-set representation, **744**, 745–746  
   candidate elimination algorithm, **744**  
   in checkers, 37  
   computational theory of, 690  
   current-best-hypothesis search, 741–743  
   decision list, 692–694  
   decision tree, 675–679, 740  
   declarative bias, **757**  
   deep, **44**, 716, 801–839  
   dropping conditions, **743**  
   element, **74**  
   ensemble, **714**, 714–720  
   explanation-based learning  
    branching factor, **753**  
    definition, **750**  
    efficiency, 752–754  
    general rules, 751–752  
    logic programming implementation, 751  
    memoization, **751**  
    operationality, **754**  
    prune, **752**  
   extensions, **740**  
   false negative, **740**  
   false positive, **741**  
   functional dependencies/  
    determinations, 755  
   in game playing, 866–867  
   generalization, **741**, 742  
   generalization hierarchy, **747**  
   grammar, 905  
   G-set, **744**, 745, 746  
   heuristics, 122  
   in hidden Markov models, 795  
   hidden variables, 792–794  
   inductive, 671–674, 733, 740, 741  
   inductive logic programming  
    attribute-based learning algorithm, 759, 760  
    **CHOOSE-LITERAL**, 763  
    constructive induction algorithms, 760  
    description, 758  
    family tree, 760  
    inverse entailment, **765**  
    inverse resolution, **763**, 764–766  
    *Journal of Molecular Biology*, **766**  
    knowledge-based induction problem, 758  
    linear resolution, **765**  
    LINUS system, 765  
    molecular biology experiments, 767  
    NEW-LITERALS, 762–763  
    positive and negative examples, 758, 759  
    PROGOL system, 765, 766  
    Prolog Horn-clause deduction, 765  
    resolvent C, **763**  
    three-dimensional configuration, 758  
    top-down inductive learning  
       methods, 761–763  
   instance-based, **704**, 704–706  
   knowledge  
    autonomous learning agent, 748  
    background knowledge, 749  
    cumulative learning process, 748  
    entailment constraint, **747**  
    explanation-based learning, **749**  
    inductive logic programming, 750  
    inferential behavior, 748–749  
    knowledge-based inductive learning, 750  
    relevance-based learning, **749**, **750**  
   large-scale, **688**  
   least-commitment search, 744–747  
   MAP, 774–775  
   maximum likelihood, 776–780  
   Meta-DENDRAL system, **747**  
   metalevel, **122**  
   mixtures of Gaussians, 789–791  
   naïve Bayes, 778  
   neural network, 35  
   noise in, 681–682  
   nonparametric, 704  
   online, **721**, 855  
   PAC, 691, 735  
   parameter, **775**, 781–783  
   *Patrons* condition, 743  
   prior knowledge, **739**  
   Q, **841**, 853, 861, 932  
   rate of, **696**, 846  
   reinforcement, 28, 210, 585, **671**, **840**, 840–873, 986  
   relational, **871**  
   relevance information, 754–757  
   in the restaurant problem, 674  
   in search, 121–122, **121**  
   to search, 121  
   specialization, **741**, 742  
   statistical, 772–775  
   S-set, **744**, 745, 746  
   unsupervised, **671**, 789–791  
   version space, **744**  
**VERSION-SPACE-LEARNING**  
   algorithm, 746–747  
   weak, **718**  
   learning curve, **679**  
**Learning to Learn**, 737  
 least-constraining-value heuristic, **177**  
 leave-one-out cross-validation (LOOCV), 684  
 Lebedev, M.A., 29, **1103**  
 Lecoutre, C., 191, **1103**  
 LeCun, Y., 35, 44, 736, 837–839, 903, 1029, 1067, 1069, **1087**, **1090**, **1100**, **1103**, **1115**, **1118**  
 Lederberg, J., 40, 41, 356, **1094**, **1104**  
 Ledsam, J. R., 48, **1104**  
 Lee, A. X., 985, **1112**  
 Lee, B. K., 1027, **1111**  
 Lee, C.-H., 1033, **1101**  
 Lee, D., 476, 668, 798, 873, **1090**, **1103**  
 Lee, G. A., 29, **1118**  
 Lee, J. D., 837, **1093**  
 Lee, K., 48, 930, **1093**, **1104**, **1109**, **1110**  
 Lee, K.-F., 19, 53, **1103**  
 Lee, K.-H., 1033, **1101**  
 Lee, K. C., 986, **1093**  
 Lee, M. S., 798, **1107**  
 Lee, R. C.-T., 331, **1090**  
 Lee, S.-I., 737, **1104**  
 Lee, S.-W., 516, **1114**  
 Lee, T.-M., 29, **1108**  
 Lee, W. S., 588, **1086**  
 Leech, G., 903, 904, **1103**, **1110**  
 LEELAZERO (game-playing program), 218  
 Lefrancq, A., 873, 1054, **1087**, **1103**, **1114**  
 Legendre, A. M., 735, **1103**  
 Legg, S., 871, 873, 1054, **1087**, **1103**, **1106**  
 Lehmann, D., 639, **1102**  
 Lehmann, J., 334, 357, **1088**, **1103**  
 Lehrer, J., 550, **1103**  
 Lehtinen, J., 831, **1101**  
 Leibniz, G. W., 24, 138, 265, 427, 637  
 Leibo, J. Z., 873, **1087**  
 Leighton, M. J., 126, **1089**  
 Leike, J., 873, 1054, **1103**

- Leimer, H., 474, 1103  
 Leipzig, 30  
 Leiserson, C. E., 125, 1080, 1091  
 Lelis, L., 127, 1103  
 Lenat, D. B., 334, 357, 636, 1103  
 lens, 991  
 Lenstra, J.K., 125, 190, 395, 402, 1085, 1103  
 Lenzerini, M., 359, 1089  
 Leonard, H. S., 358, 1103  
 Leonard, J., 984, 1088, 1103  
 Leone, N., 191, 360, 1094, 1097  
 Lepage, G. P., 475, 1103  
 Lerman, K., 1060, 1106  
 Lerner, U., 473, 1103  
 Lesh, N., 401, 1094  
 Leśniewski, S., 358, 1103  
 Lesser, V. R., 636, 639, 1094, 1103  
 Letz, R., 330, 1103  
 Levasseur, L. P., 670, 1098  
 level of abstraction, 84  
 Lever, G., 48, 225, 1100  
 Levesque, H. J., 160, 267, 359, 361, 636, 1088, 1091, 1103, 1112  
 Leviathan, 24  
 Levin, D. A., 515, 1103  
 Levine, S., 737, 841, 872, 985, 986, 1085, 1088, 1095, 1103, 1111, 1112  
 Levinstein, B., 1053, 1085  
 Levitt, R. E., 401, 1101  
 Levitt, T. S., 984, 1102  
 Levskaya, A., 928, 930, 1101  
 Levy, D., 224, 1033, 1057, 1103  
 Levy, O., 927–931, 1097, 1103, 1104, 1116  
 Lew, A. K., 667, 668, 1091  
 Lewis, A., 160, 1106  
 Lewis, J. A., 1064, 1102, 1107  
 Lewis, M., 927, 930, 1104  
 lexical-functional grammar (LFG), 904  
 lexical category, 880  
 lexicalized PCFG, 892, 904  
 lexicon, 886, 904  
 Leyton-Brown, K., 638, 640, 737, 1103, 1112, 1115  
 Le Truong, H., 331, 1097  
 LFG (lexical-functional grammar), 904  
 Li, B., 900, 1090  
 Li, C. M., 266, 1103  
 Li, D., 901, 1085  
 Li, H., 222, 837, 1093, 1116  
 Li, K., 44, 838, 1093, 1103  
 Li, L., 517, 668, 1094, 1117  
 Li, L.-J., 44, 1093  
 Li, M., 735, 1104  
 Li, P., 476, 516, 668, 798, 1090, 1104  
 Li, S., 725, 1046, 1098  
 Li, T.-M., 1067, 1104  
 Li, W., 930, 985, 1104, 1110  
 Li, X., 160, 1104  
 Li, X.-R., 79, 1086  
 Li, Y., 837, 873, 930, 1085, 1114  
 Li, Z., 516, 1104  
 Liang, G., 473, 1090  
 Liang, J., 137, 161, 1106  
 Liang, P., 896, 931, 1104, 1110  
 Liang, Y., 873, 1096  
 LIBBi (probabilistic programming language), 668  
 Liberatore, P., 267, 1104  
 Libratus (poker program), 218, 224, 612  
 lidar, 1024  
 Lidar, D., 1069, 1107  
 Liebana, D. P., 222, 1089  
 Lifchits, A., 905, 1108  
 life, value of statistical, 523  
 life insurance, 529  
 Lifschitz, V., 360, 361, 1096, 1104, 1115  
 lifting, 300–304, 301, 668  
     in probabilistic inference, 648  
 lifting lemma, 321, 324  
 light, 993–995  
     ambient, 993  
 Lighthill, J., 39, 1104  
 Lighthill report, 39, 41  
 likelihood, 773  
 LIKELIHOOD-WEIGHTING, 458  
 likelihood weighting, 457, 472, 509  
 Likhachev, M., 163, 1102  
 Lillicrap, T., 45, 48, 224, 986, 1098, 1104, 1112  
 Lim, G., 334, 1113  
 LIME (explainable machine learning system), 730, 737  
 limited rationality, 22, 346  
 limit of means, 605  
 Lin, D., 905, 1108  
 Lin, J., 901, 905, 1086  
 Lin, S., 125, 639, 1104, 1116  
 Lin, T., 334, 1113  
 Lin, Y., 516, 1117  
 Linares López, C., 399, 1115  
 Lindley, D. V., 550, 1104  
 Lindsay, R. K., 356, 1104  
 Lindsten, F., 517, 1104  
 linear-Gaussian, 440, 473, 497, 656, 779  
 linear algebra, 1076–1077  
 linear constraint, 167  
 linear function, 694  
 linearization, 942  
 linear programming, 139, 159, 161, 167, 562, 603  
 linear quadratic regulator (LQR), 962, 982  
 linear regression, *see* regression (in machine learning)  
 linear resolution, 326  
 linear separability, 700  
 linear separator, 712  
 linear temporal logic, 346  
 line search, 139  
 linguistics, 34–35, 874  
 Linnaeus, 357  
 Lipton, Z. C., 1060, 1100  
 liquid (naive physics of), 360  
 Lisp, 37, 277, 362  
 lists (logical axioms for), 287  
 Lisý, V., 48, 224, 1107  
 literal (sentence), 235  
 literal, watched, 266  
 Littman, M. L., 161, 359, 401, 586, 587, 638, 872, 873, 1065, 1085, 1090, 1099, 1100, 1104, 1105  
 Lituiev, D., 48, 1093  
 Liu, B., 871, 1104  
 Liu, H., 838, 1027, 1104, 1105  
 Liu, J., 516, 517, 871, 873, 1104, 1111  
 Liu, L. T., 1060, 1104  
 Liu, M.-Y., 930, 1104  
 Liu, P. J., 930, 1110  
 Liu, W., 798, 1090  
 Liu, X., 48, 1104  
 Liu, Y., 48, 927, 930, 1104, 1113  
 Livescu, K., 516, 1104  
 Ljung, G. M., 515, 838, 1088  
 Lloyd, J., 737, 1113  
 Lloyd, S., 1069, 1071, 1098, 1104  
 Lloyd, W. F., 639, 1104  
 Llull, R., 24, 1104  
 Lo, H.-Y., 736, 1117  
 LoBue, V., 852, 1085  
 local-sensing vacuum world, 148  
 local beam search, 133  
 local consistency, 170  
 localist representation, 77  
 locality, 257  
 locality-sensitive hash (LSH), 707  
 localization, 151, 494, 939  
     Markov, 984  
 locally structured system, 435  
 locally weighted regression, 709  
 local search, 128–137, 160, 190, 253–254, 265, 266  
 location sensor, 935  
 Locke, J., 24  
 Lockhart, E., 224, 1112  
 Loebner Prize, 1057  
 Loftus, E., 271, 1104  
 log-normal distribution discrete, 650  
 Logemann, G., 251, 266, 1092  
 logic, 21, 232–235  
     atoms, 278  
     default, 352, 356, 359  
     equality in, 282  
     first-order, 269, 269–297  
         inference, 298–300  
         semantics, 274

syntax, 274  
 first-order (FOL), 269  
 fuzzy, 232, 273, 477  
 higher-order, 273  
 inductive, 428  
 interpretations in, 275–277  
 model preference, 352  
 models, 274–275  
 nonmonotonic, 243, 351, 351–353, 359  
 notation, 21  
 propositional, 226, 235–240, 264, 269  
     inference, 238–254  
     semantics, 236–238  
     syntax, 235–236  
 quantifier, 278–282  
 resolution, 243–247  
 sampling, 475  
 temporal, 273  
     terms, 277–278  
 logical agent, 404  
 logical connective, 35, 235, 264, 278  
 logical inference, 233, 298–331  
 logically exhaustive, 404  
 logical minimization, 337  
 logical omniscience, 346  
 logical piano, 266  
 logical positivism, 25  
 logical reasoning, 240–255  
 logicism, 21  
 logic programming, 248, 296, 312–316  
     constraint, 316, 330  
     probabilistic, 665  
     tabled, 315  
 Logic Theorist (LT), 36, 266  
 LOGISTELLO (Othello program), 206  
 logistic function, 442, 735  
 logistic regression, 703, 735, 881  
 logit model, inverse, 442  
 log likelihood, 776  
 Lohia, P., 1047, 1060, 1087  
 Lohn, J. D., 161, 1104  
 long-distance dependency, 897  
 long-term memory, 310  
 long short-term memory (LSTM), 826,  
     914  
 long tail, 730  
 Longuet-Higgins, H. C., 1030, 1104  
 LOOCV (leave-one-out cross-validation),  
     684  
     “Look, Ma, no hands,” 36  
 loopy belief propagation, *see* belief  
     propagation, loopy  
 Loos, S., 327, 330, 1086, 1104  
 loosely coupled system, 593  
 Lopez, A. M., 873, 1110  
 Lopez, P., 737, 1098  
 Lopez de Segura, R., 23, 1104  
 Lopyrev, K., 931, 1110  
 Lorentz, R., 222, 1104

Losey, D. P., 974, 1086  
 loss function, 687  
 Lotem, A., 399, 1116  
 lottery, 520  
     standard, 523  
 Lou, J.-K., 736, 1117  
 love, 1033  
 Love, B. C., 550, 1100  
 Love, N., 225, 1104  
 Lovejoy, W. S., 587, 588, 1104  
 Lovelace, A., 33, 1104  
 Loveland, D., 251, 266, 330, 1092, 1104  
 low-dimensional embedding, 944  
 Lowe, D., 1029, 1104  
 Löwenheim, L., 296, 1104  
 Lowerre, B. T., 160, 1104  
 lowest path cost, 83  
 low impact, 1053  
 Lowrance, J. D., 477, 1111  
 Lowry, M., 327, 330, 1098, 1104  
 Loyd, S., 124, 1104  
 Lozano-Perez, T., 588, 984–986, 1089,  
     1099, 1104, 1116  
 LQR (linear quadratic regulator), 962,  
     982  
 LRTA\*-AGENT, 158  
 LRTA\*-COST, 158  
 LRTA\*, 157, 163, 383, 570  
 LSH (locality-sensitive hash), 707  
 LSTM (long short-term memory), 826,  
     914  
 LSVRC, 46  
 LT (Logic Theorist), 36, 266  
 Lu, C., 872, 1117  
 Lu, F., 984, 1104  
 Lu, P., 223, 735, 1089, 1111  
 Lu, Y., 737, 1117  
 Luan, D., 930, 1110  
 Lubberts, A., 223, 1104  
 Luby, M., 132, 475, 1092, 1104  
 Lucas, J. R., 1034, 1104  
 Lucas, P., 428, 1104  
 Lucas, S. M., 222, 1089  
 Luce, D. R., 28, 637, 1104  
 Luddite, 1049  
 Luehr, N., 48, 1103  
 Lugosi, G., 736, 1090  
 Lukasiewicz, T., 665, 1104  
 Lum, K., 1060, 1094  
 LUNAR (question-answering system),  
     905  
 Lundberg, S. M., 737, 1104  
 Lunn, D., 666, 798, 1105  
 Luo, S., 1027, 1105  
 Luong, Q.-T., 1030, 1094  
 Lusk, E., 331, 1117  
 Luu, C., 1060, 1087  
 Lygeros, J., 79, 1090  
 Lyman, P., 737, 1105

LYNA (medical diagnosis system), 48  
 Lynch, K., 79, 985, 987, 1090, 1105

---

## M

---

MA\* search, 113, 113, 127  
 Ma, S., 734, 837, 1087, 1111  
 Maas, A. L., 868, 1118  
 MacDonald, R., 48, 1113  
 MacGlashan, J., 359, 872, 1065, 1099,  
     1104  
 Macherey, K., 47, 834, 901, 916, 1117  
 Macherey, W., 47, 834, 901, 916, 1090,  
     1117  
 Machina, M., 550, 1105  
 machine evolution, 39  
 machine learning, 20, 22, 53, 220, 669,  
     669–987  
 machine translation (MT), 900, 915, 931  
     unsupervised, 831  
 Machover, M., 297, 1087  
 Macià, N., 737, 1097  
 MacKay, D. J. C., 476, 738, 1105, 1106  
 MacKenzie, D., 331, 1105  
 Mackworth, A.K., 79, 171, 188–190,  
     1105, 1109  
 Macready, W. G., 733, 1117  
 macrop (macro operator), 400, 768  
 madaline, 836  
 Madams, T., 48, 1097  
 Maddison, C.J., 37, 45, 222, 223, 1112  
 Madhavan, R., 986, 1105  
 Madigan, C. F., 190, 266, 1107  
 Madry, A., 838, 1090  
 magic set, 310, 329  
 Magnini, B., 931, 1091  
 Magron, V., 331, 1097  
 Mahadevan, S., 871, 1104  
 Mahalanobis distance, 706  
 Maharaj, T., 734, 1086  
 Mahaviracarya, 426  
 Mahendirian, T., 48, 1104  
 Mahesh, A., 48, 1103  
 Maheswaran, R., 191, 1108  
 Mahlmann, T., 48, 225, 1094  
 Maier, D., 190, 329, 1086, 1087  
 Mailath, G., 637, 1105  
 Majercik, S. M., 401, 1105  
 makespan, 393  
 Makov, U. E., 799, 1115  
 Maksymets, O., 873, 1111  
 Malave, V. L., 271, 1106  
 Malhotra, P., 1041, 1105  
 Malik, D., 638, 1105  
 Malik, J., 516, 838, 873, 986, 1001, 1014,  
     1015, 1020, 1028, 1030, 1085,  
     1092, 1095, 1096, 1099, 1101,  
     1103, 1105, 1111, 1112  
 Malik, S., 190, 266, 1107  
 Malone, T. W., 1062, 1105

- MAML (Model-Agnostic Meta-Learning), 737  
 manager (of tasks), 623  
 Manchak, D., 358, 1115  
 Mandal, S., 734, 1087  
 Mané, D., 1061, 1085  
 Maneva, E., 267, 1105  
 Manhattan distance, 118  
 Manhattan heuristic, 116  
 manipulator, 933  
 Manna, Z., 296, 297, 1105  
 Manne, A. S., 586, 1105  
 Manning, C., 890, 901, 903–905, 923, 929, 931, 1088, 1090, 1101, 1105, 1108, 1109  
 Mannion, M., 296, 1105  
 Manolios, P., 331, 1101  
 Mansinghka, V. K., 667, 668, 1091, 1097, 1102, 1105, 1111  
 Mansour, Y., 587, 588, 872, 873, 1101, 1114  
 map, 726  
 MAP (maximum a posteriori), 774, 797, 822  
 mapping problem, 153  
 Marais, H., 905, 1113  
 Marbach, P., 872, 1105  
 Marcedone, A., 1043, 1088  
 March, J. G., 549, 1098  
 Marcinkiewicz, M. A., 880, 903, 1105  
 Marcot, B., 474, 1109  
 Marcus, G., 550, 1105  
 Marcus, M. P., 880, 903, 1105  
 Marcus, S. I., 587, 1090  
 margin, 711  
 marginal contribution, 618  
 marginal contribution net, 620  
 marginalization, 414  
 Marin-Urias, L. F., 986, 1113  
 Marinescu, R., 475, 1105  
 Mari Aparici, C., 48, 1093  
 Markov, A., 481, 515, 902, 1105  
 Markov assumption, 481, 515, 824  
 Markov blanket, 437, 460, 461  
 Markov chain, 460, 481, 877  
 Markov chain Monte Carlo (MCMC), 459, 459–466, 472, 475, 509, 668, 783  
 decayed, 517  
 Markov decision process (MDP), 28, 553, 585, 587, 840  
 factored, 587  
 partially observable (POMDP), 578, 578–585, 587  
 relational, 587  
 structural estimation of, 872  
 Markov game, 638  
 Markov network, 474  
 Markov process, 481  
 second-order, 481  
 time-homogeneous, 481–483, 514  
 Markov property, 489, 514, 553  
 Markov reward process, 571  
 Marler, R. T., 137, 1105  
 Maron, M. E., 428, 1105  
 Márquez, L., 903, 1105  
 Marr, D., 837, 1031, 1105  
 Marriott, K., 189, 1105  
 Marris, L., 48, 225, 1100  
 Mars Exploration Rover,  
 Marshall, A.W., 872, 1100  
 Marshall, P. J., 670, 1098  
 Marsland, S., 738, 1105  
 Martelli, A., 125, 162, 1105  
 Marthi, B., 400, 517, 653, 666, 667, 873, 1105, 1106, 1108  
 Martic, M., 873, 1054, 1103  
 Martin, D., 1028, 1105  
 Martin, F. G., 1051, 1105  
 Martin, J. H., 900, 903, 904, 906, 1100, 1105  
 Martin, N., 329, 1089  
 Martin, S., 88, 1087  
 Martino, J., 1047, 1060, 1087  
 Marx, G., 898  
 masked language model (MLM), 925  
 Maskell, S., 517, 1086  
 Maskin, E., 639, 1092  
 Mason, M., 162, 400, 985–987, 1094, 1104, 1105  
 Mason, R. A., 271, 1106  
 mass noun, 340  
 mass spectrometer, 40  
 Mataric, M. J., 987, 1105  
 matching pennies, 599  
 Mateescu, R., 191, 474, 1092  
 Mateis, C., 360, 1094  
 Matena, M., 930, 1110  
 materialism, 24  
 material science, 923  
 material value, 203  
 Materzynska, J., 873, 1110  
 Mates, B., 265, 1105  
 mathematical induction schema, 323  
 mathematics, history of, 26–27  
 Matheson, J. E., 534, 548, 550, 1099, 1106  
 Matheson, M., 48, 1103  
 Mathieu, M., 48, 225, 837, 1069, 1090, 1115  
 matrix, 1077  
 matrix form, 492  
 Matsubara, H., 222, 1095  
 Mattar, M., 873, 1100  
 Matuszek, C., 357, 1105  
 Mauchly, J., 32  
 Mausam., 402, 588, 1091, 1105  
 MAVEN (Scrabble program), 225  
 MAX-VALUE, 196, 200  
 maximin, 601  
 maximin equilibrium, 603  
 maximum  
     global, 129  
     local, 130  
 maximum a posteriori (MAP), 774, 797, 822  
 maximum expected utility (MEU), 405, 519, 565  
 maximum likelihood, 775, 776–780, 797  
 maximum margin separator, 710, 711  
 maximum mean discrepancy, 737  
 max norm, 564  
 Maxwell, J., 34, 426, 1027, 1105  
 Mayer, A., 127, 1098  
 Mayer, J., 734, 1105  
 Mayne, D. Q., 516, 985, 1098, 1100  
 Mayor, A., 1057, 1105  
 Maziarz, K., 736, 1072, 1112  
 MBP (planning system), 401  
 McAfee, A., 1062, 1089  
 McAleer, S., 124, 1085  
 McAllester, D. A., 43, 162, 221, 190, 360, 398, 399, 667, 872, 873, 889, 1085, 1094, 1101, 1102, 1105, 1114  
 McArthur, N., 1057, 1092  
 MCC (Microelectronics and Computer Technology Corporation), 41  
 McCallum, A., 668, 905, 906, 1091, 1095, 1100, 1103, 1105, 1109, 1114  
 McCann, B., 931, 1101  
 McCarthy, J., 35, 36, 50, 78, 265, 267, 296, 335, 359, 399, 1105  
 McCawley, J. D., 904, 1105  
 McClelland, J. L., 42, 836, 1111  
 McClure, M., 516, 1086  
 McCorduck, P., 1058, 1105  
 McCulloch, W. S., 34, 35, 38, 267, 801, 836, 1105  
 McCune, W., 326, 331, 1105  
 McDermott, D., 162, 329, 347, 358, 359, 362, 398, 401, 1090, 1096, 1105, 1106  
 McDermott, J., 41, 310, 329, 1106  
 McDonald, R., 272, 890, 904, 1102, 1108  
 McEliece, R. J., 476, 1106  
 McGregor, J. J., 189, 1106  
 McGrew, B., 966, 986, 1085  
 McGuinness, D., 350, 357, 359, 1086, 1088, 1113  
 McIlraith, S., 296, 1106  
 McKenzie, D., 476, 1106  
 McKinney, W., 738, 1106  
 McLachlan, G. J., 799, 1106  
 McLaughlin, S., 331, 1097  
 McMahan, B., 551, 1102

- McMahan, H. B., 1043, 1060, 1088, 1102, 1106  
 MCMC (Markov chain Monte Carlo), 459, 459–466, 472, 475, 509, 668, 783  
 McMillan, K. L., 399, 1106  
 McPhee, N., 162, 1109  
 MCTS (Monte Carlo tree search), 207  
 McWhorter, J. H., 296, 1106  
 MDL (minimum description length), 689, 734, 775  
 MDP (Markov decision process), 28, 553, 585, 587, 840  
 mean-field approximation, 476  
 measure, 337  
 measurement, 337–339  
 mechanism  
     strategy-proof, 625  
 mechanism design, 590, 622  
 Medea, 982  
 medical diagnosis, 41, 428, 436, 537  
 Medina-Lara, A., 549, 1087  
 Meehan, J., 329, 1090  
 Meehl, P., 356, 1097, 1106  
 Meek, C., 474, 1116  
*Meet* (interval relation), 342  
 meganode, 453  
 Megarian school, 265  
 megavariable, 491  
 Meggido, N., 611, 637, 1102  
 Mehrabi, N., 1060, 1106  
 Mehta, G. B., 548, 1087  
 Mehta, S., 1046, 1047, 1060, 1061, 1087, 1098  
 Mellish, C. S., 330, 1091  
 Melo, F. S., 225, 1112  
 memoization, 328  
 memorization (of data), 686  
 memory (in neural networks), 824  
 memory cell (in LSTMs), 826  
 memory requirements, 95, 98  
 MEMS (micro-electromechanical system), 1064  
 MENACE (learning algorithm), 223  
 Mendel, G., 136, 1106  
 Mendes, P. N., 357, 1103  
 Meng, X.-L., 476, 1089  
 Mengüç, Y., 1064, 1107  
 meningitis, 417  
 mental model, in disambiguation, 899  
 mental object, 344–346  
 Mercer's theorem, 713  
 Mercer, J., 713, 1106  
 Mercer, R. L., 902, 929, 931, 1089, 1100  
 Merel, J., 873, 1114  
 mereology, 358  
 Merkhofer, M. M., 548, 1106  
 Merleau-Ponty, M., 1057, 1106  
 Mertz, C., 970, 986, 1118  
 Meshulam, R., 127, 1094  
 Meta-DENDRAL system, 747  
 metalearning, 737, 967, 986  
 metalevel reasoning system (MRS), 316  
 metalevel state space, 121  
 metaphor, 899, 904  
 metareasoning, 219  
     decision-theoretic, 1070  
 metarule, 316  
 meteorite, 390, 403  
 metonymy, 898  
 Metropolis, N., 160, 222, 475, 1106  
 Metropolis-Hastings, 460  
 Metropolis algorithm, 160, 475  
 Metz, L., 828, 1110  
 Metzen, J. H., 838, 1094  
 Metzler, D., 901, 905, 1091  
     *et al.*, 48, 225, 1093, 1100  
 MEU (maximum expected utility), 405, 519, 565  
 Meyerson, E., 137, 161, 1106  
 Mézard, M., 160, 1106  
 MGONZ (chatbot), 1035  
 MGSS\*, 221  
 MGU (most general unifier), 302, 304, 324  
 MHT (multiple hypothesis tracker), 667  
 Mian, I. S., 516, 1102, 1107  
 Michael, J., 930, 931, 1116  
 Michalak, T. P., 638, 1110  
 Michaylov, S., 330, 1100  
 Michel, V., 738, 1109  
 Michie, D., 125, 126, 162, 221, 223, 867, 871, 872, 983, 1089, 1093, 1106, 1111  
 micro-electromechanical system (MEMS), 1064  
 microarray gene expression, 716  
 Microelectronics and Computer Technology Corporation (MCC), 41  
 micromort, 523, 549  
 Microsoft, 473, 1037, 1059, 1068  
 microworld, 38, 39  
 Middleton, B., 439, 473, 1109  
 Miesenböck, G., 29, 1118  
 Miikkulainen, R., 137, 161, 223, 1104, 1106  
 Mikolov, T., 903, 909, 929, 930, 1100, 1106  
 Milch, B., 653, 666–668, 1105, 1106, 1108, 1109  
 milestone, 953  
 Milgrom, P., 639, 1106  
 Milios, E., 984, 1104  
 Mill, J. S., 26, 1106  
 Miller, A. C., 548, 1106  
 Miller, D., 401, 1092  
 Miller, T., 737, 1061, 1106  
 million-queens problem, 181  
 MILLIONAIRE (mechanical calculator), 27  
 Millstein, T., 399, 1094  
 Milner, A. J., 296, 1097  
 MIN-CONFLICTS, 182  
 min-conflicts heuristic, 181, 190  
 MIN-VALUE, 196, 200  
 Minami, R., 1064, 1117  
 mind  
     dualistic view, 1057  
     as physical system, 24  
     theory of, 20  
 minibatch, 697  
 minimal model, 352  
 MINIMAX-SEARCH, 196  
 minimax algorithm, 195–196, 220, 221, 601  
 minimax decision, 195  
 minimax search, 194, 194–198  
 minimax value, 194  
 minimum  
     global, 129  
     local, 130  
 minimum-remaining-values, 177, 308  
 minimum description length (MDL), 689, 734, 775  
 minimum slack, 396  
 minimum spanning tree (MST), 127  
 Minka, T., 667, 798, 1098, 1106  
 Minker, J., 329, 361, 1095, 1106  
 Minkowski distance, 705  
 Minsky, M. L., 35, 37, 40, 41, 50, 359, 472, 636, 836, 1056, 1058, 1061, 1106  
 Minton, S., 160, 190, 400, 1090, 1106  
 Miranker, D. P., 190, 1086  
 Mirhoseini, A., 736, 1072, 1112  
 Mirjalili, S. M., 160, 1106  
 Mirza, M., 838, 1097  
 Misak, C., 296, 1106  
 Mishra, B. D., 927, 931, 1091  
 missing attribute values, 682  
 missing fluent, 390  
 missionaries and cannibals, 356  
 MIT, 32, 36, 37, 41, 1028, 983  
 Mitchell, D., 160, 267, 1091, 1112  
 Mitchell, M., 53, 161, 162, 1046, 1106  
 Mitchell, T. M., 80, 271, 901, 905, 906, 1088, 1089, 1091, 1106  
 MITSUKU (chatbot), 1058  
 Mittelstadt, B., 1038, 1106  
 Mitten, L. G., 550, 1106  
 mixed strategy, 596  
 mixing number, 134  
 mixing rate, 464  
 mixing time, 486  
 mixture distribution, 790  
 mixture of Gaussians, 790, 792, 793

- Miyake, S., 837, 1095  
 Miyato, T., 837, 1106  
 Mnih, V., 835, 841, 871, 873, 1106  
 MNIST, 43, 686, 1004  
 mobile robot, 934  
 Mobileye, 1025, 1031  
 modal logic, 345  
 modal operators, 345  
 model  
     causal, 418  
     class, 671  
     counting, 452  
         weighted, 452  
     selection, 684, 797  
     sensor, 492, 499, 514  
     transition, 70, 83, 123, 140, 193, 256,  
         479, 482, 510, 514, 553, 585, 939  
 model (abstract description of reality), 84  
     small-scale, 31  
 model (in logic), 232, 364, 272, 295, 345  
 model (in machine learning), 669, 671  
 model (in probability theory), 407  
 model-based  
     reflex agents, 78  
     reinforcement learning, 841, 966  
     vision, 988  
**MODEL-BASED-REFLEX-AGENT**, 71  
 model-free agent, 74  
 Model-free reinforcement learning, 841  
**MODEL-SELECTION**, 685  
 model checking, 233, 265  
 model predictive control (MPC), 963, 971  
 model theory (in logic), 296  
 Modus Ponens, 241, 265, 326, 328  
     Generalized, 300, 300–301  
 Moffat, A., 903, 1117  
 Mohamed, A. R., 900, 905, 1099  
 Mohamed, S., 838, 1110  
 Mohr, R., 189, 1029, 1106, 1112  
 Mohtashamian, A., 48, 1104  
 Moir, T. J., 1027, 1112  
 Mojsilovic, A., 1046, 1047, 1060, 1061,  
     1087, 1098  
 momentum (in optimization), 817  
 monitoring (in machine learning), 730  
 monitoring (state estimation), 150  
 monotone condition, 125  
 monotonic concession protocol, 634  
 monotonicity  
     of a logical system, 243, 351  
     of path costs, 107  
     of preferences, 520  
 Monro, S., 735, 1110  
 Montague, P. R., 873, 1112  
 Montague, R., 358, 359, 904, 1101, 1106  
 Montanari, U., 162, 188, 1088, 1105,  
     1106  
**MONTE-CARLO-LOCALIZATION**, 941  
**MONTE-CARLO-TREE-SEARCH**, 209  
 Monte Carlo algorithm, 453  
     sequential, 517  
 Monte Carlo localization, 941  
 Monte Carlo search, pure, 207  
 Monte Carlo tree search (MCTS), 207  
 Montemerlo, M., 984, 1106  
 Montezuma’s Revenge, 867  
 Mooney, R., 905, 1106, 1117  
 Moore’s law, 32  
 Moore, A. M., 798, 1107  
 Moore, A. W., 160, 586, 798, 871, 986,  
     1086, 1088, 1107  
 Moore, E. F., 125, 1107  
 Moore, J. D., 1060, 1107  
 Moore, J. S., 327, 330, 331, 1088, 1101  
 Moore, R. C., 359, 361, 735, 1099, 1107  
 Moore machine, 637  
 Moraes, G., 48, 1104  
 Moravčík, M., 48, 224, 1107  
 Moravec, H. P., 983, 984, 1061, 1107  
 Morscos, A. S., 48, 225, 1100  
 More, T., 36  
 Morgan, C. L., 161, 1107  
 Morgan, J., 636, 1091  
 Morgan, T. J. H., 136, 161, 1107  
 Morgenstern, J., 1046, 1096  
 Morgenstern, L., 359, 361, 1092  
 Morgenstern, O., 28, 521, 548, 638, 1116  
 Moricz, M., 905, 1113  
 Moritz, P., 872, 986, 1112  
 Morjaria, M. A., 473, 1107  
 Morrill, D., 48, 224, 1107  
 Morris, P., 47, 401, 1086, 1100  
 Morrison, E., 221, 1107  
 Morrison, P., 221, 1107  
 Morsey, M., 357, 1103  
 Morstatter, F., 1060, 1106  
 Moses, Y., 359, 1094  
 Moskewicz, M. W., 190, 266, 1107  
 Mossel, E., 267, 1105  
 most general unifier (MGU), 302, 304,  
     324  
 most likely explanation, 514  
 most likely state, 963  
 motion, 1011–1012  
 motion model, 939  
 motion parallax, 1027  
 motion planning, 938, 945, 949, 984  
 motion primitive, 967  
 Mott, A., 1069, 1107  
 Mott, J., 903, 1087  
 Motwani, R., 627, 735, 1085, 1096  
 Motzkin, T. S., 836, 1107  
 Mountney, P., 1069, 1093  
 Moutarlier, P., 984, 1107  
 move, 193  
 movies  
     2001: *A Space Odyssey*, 472, 1036  
     AI, 1033, 1051  
 Centennial Man, 1051  
*Her*, 1033  
*The Matrix*, 1052  
*Rogue One*, 1022  
*The Terminator*, 933, 1052  
*Wall-E*, 1033  
 MPC (model predictive control), 963, 971  
 MPE, *see* explanation, most probable  
 MPI (mutual preferential independence),  
     533  
 MRS (metalevel reasoning system), 316  
 MST (minimum spanning tree), 127  
 MT (machine translation), 900, 915, 931  
 MUC (Message Understanding  
     Conference), 905  
 Mudd, H. F., 1058  
 Mudigonda, M., 48, 1103  
 Mueller, E. T., 334, 358, 1107, 1113  
 Muggleton, S. H., 905, 1107  
 MUI (mutual utility independence), 534  
 Muldal, A., 873, 1114  
 Mullainathan, S., 1045, 1060, 1101  
 Müller, M., 223, 1107, 1111  
 Muller, U., 44, 837, 1029, 1107  
 multi-query planning, 954  
 multiagent environment, 73, 589–595  
 multiagent planning problem, 589  
 multiagent system, 79, 589  
 multiattribute utility theory, 530, 549  
 multibody planning, 589, 591–594  
 multiheaded attention, 920  
 multiple hypothesis tracker (MHT), 667  
 multiplexer, 646  
 multiply connected network, 451  
 multitask learning, 833  
 multivariable linear regression, 697  
 Mumford, D., 1028, 1107  
 Mundy, J., 1030, 1107  
 MUNIN (medical diagnosis system), 473  
 Munos, R., 872, 1107  
 Murdock, J. W., 48, 1094  
 Murphy, K., 334, 515–517, 667, 738, 800,  
     838, 984, 1087, 1093, 1095, 1107,  
     1114, 1115  
 Murphy, R., 987, 1107  
 Murray, I., 838, 1103  
 Murray, L. M., 668, 1107  
 Murray, R. M., 985, 1064, 1095, 1107  
 Murray-Rust, P., 358, 1107  
 Murthy, C., 331, 1107  
 Musat, C., 838, 1117  
 Muscettola, N., 47, 401, 402, 1100, 1107  
 Musk, E., 51  
 Muslea, I., 905, 1107  
 mutation, 39, 159  
 mutation rate, 134  
 Muth, J. T., 1064, 1107  
 mutual preferential independence (MPI),  
     533

mutual utility independence (MUI), 534  
 MUZERO (game-playing program), 224  
 MYCIN (expert system), 41, 477  
 Myers, R. H., 429, 1116  
 Myers, S.L., 429, 1116  
 Myerson, R., 638, 639, 1107  
 myopic best response, 600  
 myopic policy, 540

**N**

n-armed bandit, 571  
 n-gram model, 830, 877, 883, 902, 903, 907  
 Nachum, O., 1045, 1100  
 Naddaf, Y., 873, 1087  
 Naderan, M., 161, 1096  
 Nagar, S., 1047, 1060, 1087  
 Nair, P., 474, 1109  
 Nair, A. V., 986, 1085  
 Nair, R., 1046, 1061, 1098  
 Nair, V., 837, 1107  
 naive Bayes, 420, 426, 428, 778, 792, 793, 797, 875  
 naked triples, 175  
 Nalwa, V. S., 30, 1107  
 Nangia, N., 931, 1116  
 Narang, S., 930, 1100  
 Narayanan, A., 1042, 1060, 1107  
 Narayanan, V., 873, 1096  
 Narayanaswamy, A., 48, 1097  
 Nardi, D., 359, 1086, 1089  
 NAS (neural architecture search), 821, 838  
 NASA, 47, 360, 373, 402, 473  
 Nash's theorem, 599  
 Nash, J., 598, 637, 1107  
 Nash, P., 587, 1107  
 Nash equilibrium, 598, 635  
 Nash folk theorems, 607  
 NATACHATA (chatbot), 1035  
 Natarajan, S., 551, 1094  
 Natarajan, V., 48, 1104  
 naturalism, 24  
 natural kind, 338  
 natural language inference, 931  
 natural language processing (NLP), 20, 874, 874–931  
 natural numbers, 286  
 natural stupidity, 347  
 Nau, D. S., 222, 224, 399, 400, 402, 1094, 1096, 1108, 1113, 1116  
 Navruzyan, A., 137, 161, 1106  
 Nayak, P., 162, 360, 402, 1102, 1107  
 NBS (search algorithm), 126  
 nearest-neighbor filter, 658  
 nearest-neighbors, 705, 735, 787  
 nearest-neighbors regression, 709  
 neat vs. scruffy, 42  
 Nebel, B., 398, 1099

Neches, R., 1060, 1107  
 needle in a haystack, 234  
 negation, 235  
 negative example, 675  
 negative literal, 235  
 negative side effects, 1038  
 negotiation set, 631  
 Neil, M., 548, 1094  
 Neiswanger, W., 838, 1116  
 NELL (Never-Ending Language Learning), 901  
 Nelson, B., 1061, 1086  
 Nelson, G., 47, 1110  
 Nelson, P. Q., 48, 1104  
 Nemhauser, G. L., 551, 1107  
 Nemirovski, A., 161, 1087, 1107  
 Nesterov, Y., 161, 1107  
 Netflix Prize, 1042  
 network tomography, 473  
 Neumann, M., 930, 1109  
 neural architecture search (NAS), 821, 838  
 neural network, 35, 38, 42, 801, 750–839  
     convolutional, 44  
     expressiveness, 35  
     feedforward, 802  
     hardware, 35  
     learning, 35  
     multilayer, 40  
     recurrent, 802, 823–826, 911–915  
 NeurIPS, 45, 53, 478, 668, 738, 800, 839, 873, 1031, 1046  
 neurobiology, 1031  
 NEUROGAMMON (backgammon program), 866  
 neuron, 29, 35, 801  
 neuroscience, 29, 29  
 Newborn, M., 126, 1108  
 Newcomb, S., 1032  
 Newell, A., 20, 36, 37, 79, 124, 125, 265, 266, 310, 329, 398, 400, 1103, 1107, 1113  
 Newman, P., 984, 1088, 1093  
 Newton, I., 19, 67, 138, 160, 483, 735, 1107  
 Newton-Raphson method, 139  
 NEXTKB (knowledge base), 357  
 Ney, H., 515, 903, 931, 1101, 1108  
 Ng, A. Y., 586–588, 735, 736, 779, 863, 868, 872, 903, 986, 1055, 1085, 1088, 1091, 1101, 1102, 1107  
 Ng, M., 29, 1118  
 Nguyen, P., 900, 905, 1090, 1099  
 Nguyen, T. T., 331, 1097  
 Nicholson, A., 478, 516, 587, 1092, 1102, 1107  
 Nicolelis, M. A., 29, 1103  
 Nielsen, E., 731, 737, 1089  
 Nielsen, M. A., 738, 1107  
 Nielsen, P. E., 329, 1100  
 Nielsen, T., 550, 1107  
 Niemelä, I., 360, 1107  
 Nigam, K., 905, 906, 1091, 1100  
 Nikolaidis, S., 986, 1107  
 Niles, I., 357, 1107, 1109  
 Nilsson, D., 548, 1108  
 Nilsson, N. J., 50, 53, 79, 124, 125, 162, 296, 297, 321, 330, 398, 400, 401, 665, 666, 836, 983, 1095, 1096, 1098, 1108  
 Niranjan, M., 516, 871, 1092, 1111  
 Nisan, N., 638, 639, 1108  
 Niv, Y., 873, 1092, 1108  
 Nivre, J., 890, 904, 1102, 1108  
 Nixon, R., 352  
 Nixon diamond, 352  
 Niyogi, S., 296, 1114  
 NLP (natural language processing), 20, 874, 874–931  
 no-good, 180  
 no-regret learning, 722  
 NOAH (planning system), 398, 400  
 Nobel Prize, 28, 29, 40  
 Nocedal, J., 735, 1088, 1089  
 Noda, I., 983, 1101  
 node (in search trees), 89, 91  
 node consistency, 170  
 Nodelman, U., 516, 1108  
 Noe, A., 1057, 1108  
 noise, 688  
 noise (in images), 996  
 noise (in training data), 677, 681–682, 772  
 noisy-OR, 438  
 nominative case, 892  
 non-cooperative game, 591, 635  
 nondeterminism  
     angelic, 379  
     demonic, 379  
 nondeterministic environment, 63, 128  
 NONLIN (planning system), 398  
 NONLIN+ (planning system), 401  
 nonlinear constraints, 167  
 nonlinear dynamical system, 501  
 nonmonotonicity, 351  
 nonmonotonic logic, 243, 351, 351–353, 359  
 Nono, 305  
 nonstationary environment, 638, 730  
 nonstationary policy, 555  
 nonterminal symbol, 1081  
 Nordfors, D., 1062, 1108  
 Nori, A. V., 668, 1091, 1099  
 Normal-Wishart, 782  
 normal distribution, 1078  
     standard, 1079  
 normal equation, 698  
 normal form game, 595

normalization (of a probability distribution), 414, 418  
 normalization (of attribute ranges), **706**  
 normative theory, **528**  
 Norouzi, M., 47, 48, 834, 901, 916, 1027, *1094, 1104, 1117*  
 North, O.  
 North, T., 39, *1095*  
 Norvig, P., 44, 329, 339, 358, 515, 737, 902, 904, *1097, 1108, 1111*  
 notation  
   infix, **286**  
   logical, 21  
   prefix, **286**  
 noughts and crosses, 193, 223  
 Nourbakhsh, I., 162, *1096*  
 Novoa, R. A., 48, *1094*  
 Nowak, R., 473, *1090*  
 Nowatzyk, A., 222, *1099*  
 Nowick, S. M., 268, *1108*  
 Nowlan, S. J., 136, 161, *1098*  
 NP (hard problems), 1075–1076  
 NP-complete, 27, 124, 241, 266, 359, 452, 734, **1075**, 1076  
 NP-hard, **1076**  
 NQTHM (theorem prover), 330  
 number statement, **649**  
 number variable, **650**  
 NumPy, 738  
 NUPRL (theorem prover), 331  
 Nuro, 1031  
 Nvidia, 1031  
 Nyberg, E., 48, *1094*  
 Nyberg, L., 19, *1089*

---

**O**

---

*O()* notation, 1075  
 O’Malley, K., 639, *1116*  
 O’Malley, M. K., 974, *1086*  
 O’Neil, C., 736, 1060, *1108*  
 O'Reilly, T., 1066  
 O'Reilly, U.-M., 161, *1108*  
 O-PLAN (planning system), 377, 401, 402  
 Oaksford, M., 550, *1090*  
 Obermeyer, F., 667, *1087*  
 object, 272, 277  
   composite, **336**  
 object-level state space, **121**  
 object-oriented programming, 33, 348  
 objective case, 892  
 objective function, **128**  
 objectivism, **426**  
 object model, 988  
 observable environment, 62  
 observation model, 481  
 observation sentence, 25  
 occupancy grid, **984**  
 occur check, **302**, 313

Och, F. J., 515, 903, 931, *1088, 1108, 1118*  
 Ockham’s razor, 673, 733, 734, 775  
 Ockham, W., 673, 733  
 odometry, **936**  
 off-switch problem, 613  
 Office Assistant, 473  
 offline search, **152**  
 Ogasawara, G., 516, *1099*  
 Ogawa, S., 29, *1108*  
 Oglesby, F., 331, *1097*  
 Oh, M.-S., 475, *1108*  
 Oh, S., 667, *1108*  
 Ohashi, T., 225, *1115*  
 Oizumi, M., 1058, *1108*  
 Olah, C., 1061, *1085*  
 Olalainty, B., 402, *1095*  
 Olesen, K. G., 473, 474, *1085, 1108*  
 Oliver, N., 516, *1108*  
 Oliver, R. M., 548, *1108*  
 Olshen, R. A., 734, *1089*  
 Olson, N., 48, *1104*  
 Olteanu, A., 1046, 1061, *1098*  
 Olum, P., 1028, *1096*  
 omniscience, **58**  
 Omohundro, S., 51, 1061, *1108*  
 one-hot encoding, **725**, 808, 908  
 One Hundred Year Study on AI, 45  
 Ong, D., 667, *1106*  
 Ong, J., 47, *1086*  
 ONLINE-DFS-AGENT, **155**  
 online gradient descent, **697**  
 online learning, **721**, 855  
 online planning, 383  
 online replanning, **963**  
 online search, **152**, 152–159, 162–163  
 ontological commitment, **272**, 295, 404  
 ontological engineering, **332**, 332–334  
 ontology, **290**, 293  
   general, 335–346  
   upper, 355  
 open-loop, **82**, 958  
 open-world assumption, 385  
 OpenAI, 1059  
 OpenAI Gym (simulated environment), 873  
 open class, **886**  
 OPENCYC (knowledge base), 357  
 open list, *see* frontier  
 OPENMIND (knowledge base), 334  
 open universe probability model  
   (OUPM), **649**  
 operations research, **28**, 79, 125, 126  
 Oppacher, F., 161, *1108*  
 OPS-5 (logical reasoning system), 310, 329  
 optical flow, **1000**, 1028  
 optimal brain damage, 838  
 optimal control theory, 160  
 optimality (of a search algorithm), **93**  
 optimality theory (in linguistics), 902  
 optimally efficient algorithm, **108**  
 optimal solution, **83**  
 optimism under uncertainty, **157**  
 optimistic description (of an action), **380**  
 optimistic prior, 849  
 optimization, **684**  
   convex, **140**, 159  
 optimizer’s curse, **527**, 549  
 OPTIMUM-AIV (planning and scheduling system), 402  
 optogenetics, **19**  
 order-of-magnitude distribution, **650**  
 orderability, **520**  
 order statistic, **526**  
 ordinal utility, **522**  
 Organon (Aristotle), 265, 357  
 origin function, **649**  
 OR node, **141**  
 Orseau, L., 873, 1054, *1103*  
 Ortega, P. A., 873, 1054, *1103*  
 Osawa, E., 983, *1101*  
 Osborne, M. A., 1050, *1095*  
 Osborne, M. J., 638, *1108*  
 Osherson, D. N., 734, *1108*  
 Osindero, S., 837, 838, *1099, 1100*  
 Osman, I., 126, *1110*  
 Ostland, M., 666, 667, *1108*  
 Ostrom, E., 639, 1061, *1098, 1108*  
 Ostrovski, G., 871, 873, *1106*  
 Othello, 224  
 Ott, M., 927, 930, *1104*  
 Otten, L., 474, *1112*  
 OTTER (theorem prover), 331  
 out-of-bag error, **716**  
 out-of-vocabulary word, **878**  
 outcome, **405**, 596, 616  
 outlier, **725**  
 output gate (in LSTM), **826**  
 overhypotheses, **768**  
 over-sample, **725**  
 Overbeek, R., 331, *1117*  
 overfitting, **673**, 681–682, 698, 772, 774  
 overgeneration, **886**  
 Overmars, M., 985, *1101*  
 overriding, **349**  
 Owen, M. P., 588, *1100*  
 Owens, A. J., 161, *1095*  
 OWL (description logic), 357  
 Ozair, S., 838, *1097*

---

**P**

---

**P** (probability vector), 409, 410  
 $P(s'|s, a)$  (transition model), 553, 842  
 PAC (probably approximately correct), **691**, 693, 735  
 pace of change, **1050**  
 Pachocki, J., 966, 986, *1085*

- Padgham, L., 79, *1108*  
 Page, L., 905, *1089*  
 Paige, B., 516, 668, *1108*  
 Palacios, H., 400, *1108*  
 Palaniappan, M., 638, *1105*  
 Paleo, B. W., 331, *1087*  
 Palmer, J., 271, *1104*  
 Palmer, S., 1031, *1108*  
 Palmieri, G., 836, *1096*  
 Pan, X., 872, *1117*  
 Pandas, 738  
 Pandas (data analysis software), 727  
 Panini, 34, 903  
 Papadimitriou, C. H., 160, 162, 266, 586,  
     588, 638, 1080, 1093, 1102, *1108*  
 Papadopoulou, T., 1030, *1094*  
 Papavassiliou, V., 871, *1108*  
 paperclip game, **614**  
 Papernot, N., 838, *1090*  
 Papert, S., 40, 836, *1106*  
 paradox, 359, 529  
     Allais, 528, 550  
     Condorcet, 629  
     Ellsberg, 528, 550  
     Girard, 331  
     St. Petersburg, 547  
     Zeno, 83  
 parallel distributed processing (PDP), 836  
 parallel jaw gripper, **936**  
 parallel lines, 990  
 parameter, **430**, 776  
 parameter independence, **783**  
 parametric model, **704**  
 paramodulation, **325**, 330  
 parent node, **90**  
 Pareto optimality, **599**  
 Parikh, D., 873, 1017, 1097, *1111*  
 Parisi, D., 905, *1094*  
 Parisi, G., 160, 476, *1106*, *1108*  
 Parisi, M. M. G., 267, *1108*  
 Park, F. C., 987, *1105*  
 Park, J. D., 475, *1108*  
 Park, S., 327, *1098*  
 Park, T., 930, 1022, *1118*  
 Parker, A., 222, *1108*  
 Parker, D. B., 836, *1108*  
 Parker, L. E., 987, *1108*  
 Parmar, N., 901, 919, 931, *1090*, *1115*  
 Parr, R., 551, 586, 587, 637, 873, *1085*,  
     *1090*, *1097*, *1102*, *1108*  
 Parrod, Y., 402, *1085*  
 Parsey McParseface, 904  
 parsing, **886**, 886–891  
 part-of-speech tagging, **881**  
 Partee, B. H., 904, *1109*  
 partial-order planning, **370**  
 partial assignment, **165**  
 partial bracketing, **891**  
 partial evaluation method, **768**  
 partial observability, 128, 214, 578  
 partial program, **859**  
 partial solution, **165**  
 PARTICLE-FILTERING, **510**  
 particle filtering, **510**, 515  
     Rao-Blackwellized, **514**, 516, 984  
 particle MCMC, **517**  
 partition, **336**, 616  
 Partnership on AI, 53, 1059  
 part of speech (POS), **880**  
 Parzen, E., 799, *1108*  
 Parzen window, 799  
 Pasca, M., 905, 906, *1094*, *1108*  
 Pascal's wager, 427  
 Pascal, B., 24, 26, 426  
 PASCAL Challenge, 931  
 PASCAL VOC (image data set), 1029  
 Pascanu, R., 837, *1092*  
 Pasero, R., 296, 329, *1091*  
 Paskin, M., 666, *1108*  
 PASSIVE-ADP-LEARNER, **845**  
 PASSIVE-TD-LEARNER, **846**  
 passive learning agent, **842**  
 passive sensing, 988  
 Pastor, P., 985, 986, *1088*, *1103*  
 Pasula, H., 517, 653, 666, 667, *1105*,  
     *1108*  
 Patel, S., 1043, *1088*  
 Patel-Schneider, P., 357, 359, *1086*, *1108*  
 path, **83**, 123, 394, **945**, 949  
     loopy, **92**  
     redundant, **92**, 92–93  
 path consistency, **172**, 188  
 PATHFINDER (medical diagnosis system),  
     473  
 path integral, **956**  
 Patil, N., 1069, *1100*  
 Patil, R., 904, *1090*  
 Patrick, B. G., 126, *1108*  
 pattern database, **119**, 127, 374  
     disjoint, **119**  
 pattern matching, **308**  
 Patterson, D. A., 670, 1069, *1092*, *1100*  
 Paul, R. P., 985, *1108*  
 Paulin-Mohring, C., 330, *1087*  
 Paull, M., 266, *1095*  
 Pauls, A., 904, *1108*  
 Pavlovic, V., 474, *1118*  
 payoff function, 193, **595**  
 payoff matrix, **595**  
 payoff vector, **616**  
 Paz, A., 473, *1115*  
 Pazzani, M., 428, 798, *1093*  
 PBT (population-based training), **690**  
 PCA (principal components analysis),  
     828  
 PCFG, **884**, 884–886, 903  
     lexicalized, **892**, 904  
 P controller, **959**  
 PD controller, **960**  
 PDDL (Planing Domain Definition  
     Language), 362  
 PDP (parallel distributed processing), 836  
 Peano, G., 296, *1108*  
 Peano axioms, **286**, 296, 307  
 Pearce, J., 191, *1108*  
 Pearl, J., 35, 43, 82, 108, 125, 127, 160,  
     190, 221, 222, 431, 472–478, 798,  
     799, *1092*, *1096*, *1097*, *1100*,  
     *1101*, *1108*, *1115*  
 Pearlmutter, B. A., 1067, *1113*  
 Pearson, K., 798, 838, *1108*  
 PEAS description, **60**, 62  
 Pease, A., 357, *1107*, *1109*  
 Pecheur, C., 327, *1098*  
 pedigree analysis, **473**  
 Pednault, E. P. D., 398, 636, *1109*  
 Pedregosa, F., 738, *1109*  
 Peek, M. Y., 1064, *1095*  
 PEGASUS (reinforcement learning  
     algorithm), 868  
 Pei, J., 738, *1098*  
 Peirce, C. S., 188, 296, 347, 359, *1109*  
 Peleg, B., 638, *1109*  
 Pelikan, M., 161, *1109*  
 Pell, B., 402, *1107*  
 Pemberton, J. C., 163, *1109*  
 penalty, **75**  
 Penberthy, J. S., 398, *1109*  
 Peng, J., 871, *1109*  
 Peng, L., 48, *1097*, *1104*, *1113*  
 Penix, J., 327, *1098*  
 Pennachin, C., 51, *1096*  
 Pennington, J., 923, 929, *1109*  
 Pennsylvania, Univ. of, 32  
 Penn Treebank, **880**, 890, 928  
 Penrose, R., 1034, *1109*  
 Pentagon Papers, 550  
 people prediction, **938**  
 Peot, M., 401, 475, *1109*, *1112*  
 percept, **54**  
     possible, 148  
 perception, 54, 288, 988–1026  
 perceptron, 39, 836  
     convergence theorem, 39  
     learning rule, **701**  
     representational power, 40  
 percept schema, **384**  
 percept sequence, **54**, 58, 59  
 Pereira, F., 44, 312, 358, 736, 737, 889,  
     891, 906, 928, *1085*, *1093*, *1097*,  
     *1103*, *1109*, *1110*  
 Peres, Y., 515, 517, *1103*, *1105*  
 perfect information, **193**, **608**  
 perfect recall, 609  
 performance element, **74**  
 performance measure, **57**, 58, 60, 78, 403,  
     519

- Perkins, T., 334, *1113*  
 Perov, Y., 668, *1105*  
 perplexity, **930**  
 Perrault, C. R., 636, *1091*  
 persistence arc, **507**  
 persistent failure model, **506**  
 persistent variable, 1082  
 personal agent, **1066**  
 perspective, 1027  
 perspective projection, **990**  
 Persson, K. A., 923, *1115*  
 Pesch, E., 402, *1088*  
 Peshkin, M., 162, *1116*  
 pessimistic description (of an action), **380**  
 Peters, J., 476, 986, *1101*, *1109*  
 Peters, M. E., 930, *1109*, *1111*  
 Peters, S., 904, *1093*  
 Petersen, S., 871, 873, *1087*, *1106*  
 Peterson, C., 476, *1109*  
 Peterson, K., 970, 986, *1118*  
 Petosa, N., 224, *1109*  
 Petrie, K., 191, *1096*  
 Petrie, T., 515, 799, *1086*  
 Petrik, M., 871, *1104*  
 Petron, A., 966, 986, *1085*  
 Petrov, S., 890, 904, 1040, *1085*, *1108*  
 Pettersson, L., 873, *1089*  
 Pezeshki, M., 900, *1118*  
 Pfeffer, A., 222, 643, 666–668, 637,  
*1102*, *1109*  
 Pfeifer, G., 360, *1094*  
 Pfeifer, R., 1057, *1109*  
 Pham, H., 838, *1109*  
 phase transition, 267  
 Philips, A. B., 160, 190, *1106*  
 Phillippe, A. M., 735, *1087*  
 Phillips, E. H., 48, *1103*  
 Philo of Megara, 265  
 philosophy, 24–26, 79, 1032–1062  
   moral, 57  
 photogrammetry, 1028  
 photometry, 993–995  
 photosensitive spot, 1024  
 phrase structure, **884**  
 physical game, 225  
 physicalism, 24  
 physical symbol system, **37**  
 piano mover’s problem, **949**  
 Piantino, S., 1069, *1115*  
 Piccione, C., 637, *1118*  
 PICTURE (probabilistic programming language), 668  
 PID controller, **960**  
 Pieper, G., 331, *1117*  
 Pierson, E., 1060, *1091*  
 Pietra, V. J. D., 929, *1089*  
 pigeons, 31  
 Pineau, J., 588, 871, 986, *1095*, *1109*,  
*1112*  
 Pinedo, M., 402, *1109*  
 ping-pong, 225  
 pinhole camera, **990**, 989–991  
 Pinkas, G., 190, *1109*  
 Pinker, S., 270, 271, 296, 905, *1109*, *1111*  
 Pinto, D., 906, *1109*  
 Pinto, L., 986, *1109*  
 Pisa, tower of, 75  
 pit, bottomless, 228  
 Pitassi, T., 474, 1046, 1060, *1086*, *1094*,  
*1117*  
 Pitts, W., 34, 35, 38, 267, 801, 836, *1105*  
 pixel, **989**  
 PL-FC-ENTAILS?, **249**  
 PL-RESOLUTION, **246**  
 PLAN-ERS1 (planning and scheduling system), 402  
 PLAN-ROUTE, **260**  
 PLANEX (planning agent), 401  
 Plankalkül, 32  
 plan monitoring, 390  
 PLANNER (logic programming language),  
   41, 329, 636  
 planning, 220, 362–402  
   and acting, 383–385  
   as refinement, 370  
   as satisfiability, 369  
   blocks world, 38  
   case-based, **400**  
   classical, **362**  
   conformant, 383, 385–388, 397, **400**  
   contingency, 383, 388–389, 397  
   decentralized, **589**  
   hierarchical, 374–383, 397  
   hierarchical task network, 375  
   history of, 398  
   linear, **398**  
   Monte Carlo, **562**  
   multibody, **589**, 591–594  
   multieffector, **589**  
   online, 383  
   reactive, **401**, 975  
   regression, 368–369  
   route, 47  
   search space, 366–374  
   sensorless, 383, 385–388  
 planning graph, **370**  
 plan recognition, **595**  
 PlanSAT, 402  
   bounded, 402  
 Plappert, M., 966, 986, *1085*  
 plateau (in local search), **131**  
 Plato, 265, 358  
 Platt, J., 736, *1109*  
 player (in a game), **595**  
 playout, **207**  
 playout policy, **207**  
 Playter, R., 47, *1110*  
 Plotkin, G., 330, *1109*  
 Plummer, M., 798, *1109*  
 Plunkett, K., 905, *1094*  
 plurality voting, **630**  
 Pluribus (poker program), 224, 612  
 ply, **194**  
 Pnueli, A., 359, *1109*  
 poetry, 19  
 Poggio, T., 837, *1105*  
 Pohl, I., 125, 126, *1109*  
 pointwise product, **448**  
 Poisson distribution, **650**  
 poker, 224, 638  
**Poland**, 358  
 Poli, R., 162, *1103*, *1109*  
 policy, 206, **554**, 585  
   dominating, **577**  
   evaluation, **566**, 842  
   gradient, **862**  
   improvement, **566**  
   iteration, 562, **566**, 566–568, 586  
   asynchronous, **568**  
   modified, **568**  
   loss, **565**  
   optimal, **554**  
   proper, **556**  
   search, **841**, **861**, 861–863  
   stochastic, **862**  
   value, **862**  
**POLICY-ITERATION**, **567**  
 polite convention, **1036**  
 Pollack, M. E., 636, *1091*  
 Polosukhin, I., 901, 919, 931, *1115*  
 polysemy, 924  
 polytree, **451**, 452, 472, 488  
 Polyzotis, N., 329, *1102*  
**POMCP**, **584**  
 POMDP (partially observable MDP), **578**,  
   578–585  
**POMDP-VALUE-ITERATION**, **583**  
 Pomerleau, D. A., 1030, 984, *1109*  
 Ponce, J., 1031, *1095*  
 Ponte, J., 931, *1118*  
 Poole, B., 838, *1109*  
 Poole, D., 79, 474, 666, 668, 549, *1088*,  
*1101*, *1109*, *1118*  
 pooling (in neural networks), **813**  
 Popat, A. C., 903, 931, *1088*  
 Popescu, A.-M., 906, *1094*  
 Poppe, R., 1064, *1109*  
 Popper, K. R., 427, 734, *1109*  
 population-based training (PBT), **690**  
 Porphyry, 359  
 Port-Royal Logic, 547  
 Porter, B., 361, *1115*  
 portfolio, **397**  
 Portner, P., 904, *1109*  
 POS (part of speech), **880**  
 pose, **1012**, 946  
 Posegga, J., 330, *1087*

- position, **193**  
 positive example, **675**  
 positive literal, **235**  
 positivism, logical, **25**  
 possibility theory, **478**  
 possible percept, **148**  
 possible world, **232**, 264, 295, **345**, 406, **642**  
 Post, E. L., 266, **1109**  
 post-decision disappointment, **549**  
 POS tagging, **903**  
 posterior probability, *see* probability, conditional  
 potential (in MDPs), **560**  
 Potts, C., 896, 931, **1088**, **1104**  
 Pouget-Abadie, J., 838, **1097**  
 Poulton, C., 1064, **1109**  
 Poundstone, W., 637, **1109**  
 Pourret, O., 474, **1109**  
 Poverty of the Stimulus, **905**  
 Powell, G., 966, 986, **1085**  
 Powell, R., 48, 225, **1115**  
 Powers, R., 638, **1112**  
 Powley, E. J., 222, **1089**  
 PPCA (probabilistic principal components analysis), **827**  
 Prabhakar, A., 48, **1103**  
 Prabhavalkar, R., 900, **1090**  
 Prade, H., 478, **1093**  
 Pradhan, M., 439, 473, **1109**  
 Pradhan, N., 667, **1087**  
 Praet, J.-C., 549, **1087**  
 Prager, J., 48, **1094**  
 pragmatics, **897**  
 Pratt, L., 737, **1115**  
 Prawitz, D., 328, **1109**  
 precedence constraint, **166**  
 precomputation, **120**  
 precondition, **363**  
     missing, **390**  
 precondition axiom, **263**  
 predicate indexing, **302**  
 predicate symbol, **275**, **894**  
 prediction, 148, 485–486, **514**  
 predictive learning, **1068**  
 preference, **405**, **520**  
     monotonic, **524**  
     unknown, **543**–**546**  
 preference elicitation, **523**  
 preference independence, **533**  
 preference learning, **938**  
 preferred action, **373**  
 premise, **235**  
 president, **343**  
 Presley, E.,  
 Press, W. H., 160, **1109**  
 Presta, A., 904, **1085**  
 Preston, J., 1058, **1109**  
 pretraining, 908, **922**  
 Prettenhofer, P., 738, **1109**  
 Price, B., 587, **1088**  
 Price, E., 1060, **1098**  
 Price Waterhouse, 401  
 Prieditis, A. E., 118, 127, **1109**  
 Prince, A., 902, **1113**  
 principal components analysis (PCA), **828**  
*Principia Mathematica*, **36**  
 principle of insufficient reason, **427**  
 PRIOR-SAMPLE, **454**  
 prioritized sweeping, 586, **847**, 871  
 priority queue, **92**  
 prior knowledge, 58, 59, **670**  
 prior probability, **407**, 425  
 prismatic joint, **936**  
 prisoner’s dilemma, **596**  
 Pritzel, A., 986, **1104**  
 privacy, 729, 1060  
 private value, **624**  
 probabilistic agent, 404  
 probabilistically complete, **954**  
 probabilistic context-free grammar (PCFG), **902**  
 probabilistic principal components analysis (PPCA), **827**  
 probabilistic roadmap (PRM), **953**  
 probability, 21, **26**, 43, 403–478,  
     1078–1079  
     axioms of, **411**  
     conditional, **407**, 414, 417, 425, 434  
     conjunctive, **434**  
     density function, **409**, 1078  
     distribution, **409**, 445  
     history of, **429**  
     judgments, **436**  
     marginal, **413**  
     model, **407**, 1078  
         open-universe (OUPM), **649**  
     prior, **407**, 425  
         uninformative, **784**  
     theory, **274**, **404**, 547  
 probability logic, **665**  
 probability notation, 409  
 probably approximately correct (PAC), **691**, 693, 735  
 PROBCUT (game-tree search algorithm), **205**  
 probit, **442**, 473  
 problem, **83**, 123  
     assembly sequencing, 88, 125  
     bandit, **571**, 587, 849  
     conformant, **144**  
     constrained optimization, 139, **169**  
     8-puzzle, 115, 118  
     formulation, **82**, 84  
     generator, **75**  
     halting, 300  
     inherently hard, 1075–1076  
 million-queens, 181, 190  
*n*-queens, **254**  
 optimization, **128**  
 piano movers, 984  
 real-world, **84**  
 relaxed, **117**, 118, 371  
 robot navigation, 88  
 sensorless, **144**  
 solving, 40  
 standardized, **84**  
 touring, **88**  
 traveling salesperson, 125, 127  
 underconstrained, **254**  
 VLSI layout, 88, 133  
 Procaccia, A. D., 639, **1088**  
 procedural approach, **228**, 269  
 procedural attachment, **349**  
 PROCEED, **101**  
 PRODIGY (planning system), 400  
 production system, 68, **310**, 328, 329  
 product rule, **408**, 417  
 programming language, 269  
 progression (in planning), 367, 397  
 Prolog, 41, **312**, 329, 398, 765  
 Prolog Technology Theorem Prover (PTTP), 330  
 proof, **241**  
 proper policy, **556**  
 property (unary relation), **272**  
 proposal distribution, **465**  
     adaptive, **664**  
 proposition  
     probabilistic, 406–413  
     symbol, **235**  
 propositional attitude, **344**  
 propositionalization, **299**, 327  
 propositional logic, 226, 235–240, 264, **269**  
 proprioceptive sensor, **936**  
 Prosser, P., 189, **1109**  
 protein design, **88**  
 provably beneficial, 613  
 Provan, G. M., 439, 473, **1109**  
 PROVER9, 331  
 Provost, F., 1041, **1094**  
 Pruksachatkun, Y., 931, **1116**  
 pruning, **108**, **192**, 198, 681, 716  
     forward, **205**  
     futility, **222**  
     in contingency problems, 213  
 pseudocode, **1082**  
 pseudoexperience, **847**  
 pseudoinverse, **698**  
 pseudoreward, **858**  
 PSPACE, 402, 1076  
 psychological experiment, **20**  
 psychological reasoning, **361**  
 psychology, 30–32  
     experimental, 30

- psychophysics, 1031  
PTTP (Prolog Technology Theorem Prover), 330  
public key encryption, 327  
Puget, J.-F., 191, 1096  
Pullum, G. K., 296, 904, 905, 1099, 1109  
Puma (robot arm), 983  
Purdom, P., 191, 1089  
pure strategy, 596  
pure symbol, 251  
Puterman, M. L., 79, 586, 588, 1109  
Putnam, H., 251, 266, 321, 328, 428, 668, 1092, 1109  
Pyro (probabilistic programming language), 667  
Pyrros, A., 1023, 1093  
PySC2 (machine learning software), 873  
Python, 277  
PyTorch (machine learning software), 667, 738, 1072
- 
- Q**
- $Q(s, a)$  (value of action in state), 853  
Q-function, 535, 558, 841  
Q-learning, 841, 853, 861, 932  
Q-LEARNING-AGENT, 854  
Q-VALUE, 559  
QA3 (logical reasoning system), 296  
QALY (quality-adjusted life year), 524, 549  
Qian, H., 1060, 1087  
QUACKLE (Scrabble program), 225  
quadratic programming, 711  
Quake III, 46, 48  
qualia, 1036  
qualification problem, 259, 403, 1033  
qualitative physics, 339, 360  
qualitative probabilistic network, 478, 532  
quality-adjusted life year (QALY), 524, 549  
quantification, 896  
quantifier, 278, 296
  - existential, 280
  - in logic, 278–282
  - nested, 281
  - universal, 278–280, 298
quantum computing, 33  
quasi-logical form, 897  
query (logical), 284  
query (probabilistic), 413  
query variable, 445  
query vector (in transformers), 920  
question answering, 901, 931
  - visual (VQA), 46, 1017
queue, 92
  - FIFO, 92
  - LIFO, 92
  - priority, 92
Quevedo, T., 221
- quiescence, 204  
Quigley, A., 1064, 1117  
Quillen, D., 986, 1103  
Quillian, M. R., 359, 1109  
Quine, W. V., 297, 338, 357, 358, 1109  
Quinlan, J. R., 733, 734, 1109  
Quinlan, S., 985, 1110  
Quirk, R., 904, 1110  
QXTRACT (information extraction system), 906
- 
- R**
- R (statistical software), 738  
R1 (expert system), 41, 310, 329  
Rabani, Y., 161, 1110  
Rabe, M. N., 327, 1086  
Rabideau, G., 401, 1095  
Rabiner, L. R., 515, 1110  
Rabinovich, Y., 161, 1110  
Rabinowitz, N. C., 48, 225, 1100  
radar, 28, 935  
Radford, A., 828, 930, 1110  
Raedt, L. D., 666, 1101  
Raffel, C., 930, 1110  
Rafferty, A. N., 588, 1110  
Ragan-Kelley, J., 1067, 1104  
Raghavan, M., 1045, 1060, 1101  
Raghavan, P., 901, 905, 1105  
Raghu, M., 838, 1109  
Rahwan, T., 638, 1110  
Raibert, M., 47, 1110  
Raiffa, H., 28, 529, 549, 550, 637, 1101, 1104  
Rajamani, S. K., 668, 1091, 1099  
Rajan, K., 47, 401, 1100  
Raji, I. D., 1046, 1106  
Rajpurkar, P., 931, 1110  
Raju, B., 137, 161, 1106  
Ramage, D., 1043, 1088  
Ramamurthy, K. N., 1046, 1047, 1060, 1061, 1087, 1098  
Ramsey, F. P., 28, 427, 548, 1110  
Ramsundar, B., 517, 738, 1094, 1110  
RAND Corporation, 550, 586  
random-restart hill climbing, 131  
random forest, 715, 736  
randomForest (machine learning software), 716  
randomization, 55, 69  
randomized controlled trial, 471  
randomized weighted majority algorithm, 721  
random restart, 252  
random search, 690  
random variable, 408, 435
  - basic, 644
  - continuous, 409, 440, 473
  - indexed, 666
random walk, 156, 498
- range (of a random variable), 408  
range finder, 934
  - laser, 935
range sensor array, 940  
Ranzato, M., 837, 1100  
Rao, A., 80, 1117  
Rao, B., 516, 1099  
Rao, D. A. S., 1049, 1110  
Rao, K., 900, 1090  
Rao-Blackwellization, 514  
Raphael, B., 125, 329, 1097, 1098, 1110  
Raphson, J., 160, 735, 1110  
rapidly exploring random trees (RRT), 954  
rapid prototyping, 312  
Raschka, S., 738, 1110  
Raschke, U., 983, 1091  
Rashevsky, N., 35, 836, 1110  
Rasmussen, C. E., 799, 1110  
Rassenti, S., 639, 1110  
rating, 646  
Ratinov, L., 929, 1115  
Ratio Club, 34  
rational agent, 22, 21–22, 54, 57, 57–58, 73, 78, 79, 547  
rational decision, 405  
rationalism, 906  
rationality, 19, 57–58
  - Boltzmann, 865
  - individual, 617
  - limited, 22, 346
  - perfect, 22
rational thought, 21  
Ratliff, N., 958, 970, 973, 985, 986, 1110, 1118  
Ratnaparkhi, A., 903, 1110  
Ratner, D., 124, 1110  
rats, 31  
Rauber, J., 838, 1090  
Rauch, H. E., 515, 1110  
Rawal, A., 137, 161, 1106  
Ray, A., 966, 986, 1085  
Ray, B., 737, 1097  
Rayson, P., 903, 1103  
Rayward-Smith, V., 126, 1110  
Razavi, A., 838, 1100  
RBFS (recursive best-first search), 111–113, 123  
RDF (Resource Description Framework), 357  
Ré, C., 665, 1092  
reachable set, 379  
reached (states in search), 90  
reactive planning, 401, 975  
Real, E., 838, 1110  
real-world problem, 84  
realizability, 688  
reasoning, 21, 37, 226
  - default, 351–353, 477

- logical, 240–255  
 probabilistic, 413–415, 445–453  
     approximate, 453–467  
 receiver operating characteristic (ROC) curve, 728  
 receptive field, 813  
 Rechenberg, I., 161, 1110  
 Recht, B., 734, 1118  
 recombination, 133  
 recommendation, 643  
 recurrent neural network (RNN), *see* neural network, recurrent  
**RECURSIVE-BEST-FIRST-SEARCH**, 111  
 recursive best-first search (RBFS), 111–113, 123  
 recursive estimation, 484  
 Reddy, R., 35  
 reduction, 451, 1080  
 Rees, M., 51  
 Reeves, C., 126, 1110  
 Reeves, D., 639, 1116  
 reference class, 427  
 reference count, 110  
 referential transparency, 345  
 refinement (in hierarchical planning), 375  
 reflectance, 1012  
 reflection, 993  
**REFLEX-VACUUM-AGENT**, 67  
 reflex agent, 67, 67–69, 78, 554, 841  
 Reformer (natural language software), 928, 930  
 refutation, 241  
 refutation completeness, 321  
 Regan, N., 223, 1111  
 Regin, J., 189, 1110  
 region, 951  
 regional proposal network (RPN), 1007  
 region of interest, 1007  
 regression (in machine learning), 670  
     linear, 695, 735, 780  
         Bayesian, 783–785  
     tree, 683  
     regression (in planning), 368, 367–369, 397  
 regression search, 368  
 regret, 722  
 regularization, 689, 698, 720, 822, 1053  
 regularization function, 689  
 Reid, D. B., 667, 1110  
 Reid, M., 116, 125, 1102  
 Reif, J., 984, 986, 1089, 1110  
 reification, 335  
**REINFORCE** (reinforcement learning algorithm), 863  
 reinforcement, 840  
 reinforcement learning, 28, 210, 585, 671, 840, 789–873, 986  
     active, 842, 848–854  
     Bayesian, 851  
     deep, 835, 857  
     distributed, 636  
     generalization in, 854–861  
     hierarchical, 858, 1065  
     inverse, 52  
         cooperative, 638  
     model-based, 841  
     model-free, 841  
     multiagent, 636  
     off-policy, 853  
     on-policy, 853  
     passive, 842, 842–848  
     relational, 871  
 Reingold, E. M., 189, 1088  
 Reingold, O., 1060, 1094  
 Reinsel, G., 515, 838, 1088  
 Reiter, M. K., 1060, 1115  
 Reiter, R., 268, 359, 399, 587, 1088, 1110  
**REJECTION-SAMPLING**, 456  
 rejection sampling, 455  
     adaptive, 475  
 relation (in CSPs), 164  
 relation (in logic), 272  
 relational probability model (RPM), 643  
 relaxed problem, 117, 118, 371  
 relevance, 237  
 relevance-based learning (RBL), 749, 750, 767  
 relevant action, 368  
**ReLU**, 803  
 Remolina, E., 47, 1086  
 Remote Agent, 327, 373, 402  
 Remote Agent (planning agent), 47  
 Ren, S., 837, 1098  
 renaming, 305  
**RENDER-NOISY-IMAGE**, 660  
 rendering model, 988  
 Rényi, A., 428, 1110  
 repeated game, 598, 604  
 replanning, 383, 389–392, 401  
 representation, *see* knowledge representation  
     atomic, 77  
     contextual, 924  
     factored, 77  
     structured, 77  
 representation theorem, 533  
**REPRODUCE**, 137  
 resampling, 510  
 reserve bid, 624  
 residual (in neural networks), 815  
 residual network, 815  
 ResNet-50 model, 832  
 Resnick, C., 1027, 1094  
 Resnick, P., 47, 1110  
 resolution, 37, 39, 243–247, 244, 265, 316–328  
     closure, 246, 322  
     completeness proof for, 321  
     input, 326  
     linear, 326  
     strategies, 326–327  
 resolvent, 243, 318  
 resource constraint, 392  
 resources, 392–397  
 response, 31  
 restaurant hygiene inspector, 216  
 result, 363  
 Rete algorithm, 310  
 retiming, 958  
 retrograde, 207  
 reusable resource, 393  
 revelation principle, 625  
 revenue equivalence theorem, 626  
 Reversi, 224  
 revolute joint, 936  
 reward, 75, 553, 585, 840  
     additive, 556  
     discounted, 555  
     function, 1065  
     hyperbolic, 588  
     shaping, 858  
     sparse, 841  
 reward-to-go, 843  
 rewrite rule, 1081  
 Rezende, D. J., 838, 1110  
 RGB (red, green, blue), 995  
 Riazanov, A., 330, 331, 1110  
 Ribeiro, F., 225, 1115  
 Ribeiro, M. T., 737, 1110  
 Ribeiro-Neto, B., 901, 1086  
 Riccati equation, 962  
 Rice, T. R., 548, 1106  
 Richards, J. T., 1047, 1060, 1087  
 Richardson, K., 927, 931, 1091  
 Richardson, M., 666, 1110  
 Richardson, S., 475, 1096  
 Richtárik, P., 1043, 1102  
 Richter, S., 399, 1110  
 Riddell, A., 476, 668, 798, 1090  
 ridge (in local search), 130, 131  
 Ridley, M., 161, 1110  
 Riedmiller, M. A., 835, 841, 871, 873, 1106  
 Riesbeck, C., 41, 329, 1090, 1111  
 Riley, J., 639, 1110  
 Riley, P., 737, 1110  
 Riloff, E., 906, 1100, 1110  
 Ringgaard, M., 928, 1110  
 Rink, F. J., 473, 1107  
 Rintanen, J., 399–402, 1110  
 Ripley, B. D., 738, 1110  
 Rish, I., 474, 1092  
 risk aversion, 525  
 risk neutrality, 526  
 risk score, 1044  
 risk seeking, 525  
 Rissanen, J., 734, 1110

- Ristenpart, T., 1060, *1115*  
Ritov, Y., 666, 667, *1108*  
Rivest, R., 125, 734, 735, 1080, *1091*, *1100*, *1110*  
Rivest, Shamir, and Adelman (RSA), 327  
RMS (root mean square), 1080  
RNN (recurrent neural network), *see* neural network, recurrent  
Robbins, H., 576, 587, 735, *1103*, *1110*  
Robbins algebra, **331**  
ROBERTA (natural language system), 832, 926–928, 930  
Roberts, A., 930, 1027, *1094*, *1110*  
Roberts, G., 1057, *1094*  
Roberts, L. G., 1028, *1110*  
Robertson, N., 190, *1110*  
Robertson, S. E., 428, *1110*  
Robins, J., 476, *1110*  
Robinson, A., 297, 328, 331, *1110*  
Robinson, G., 330, *1117*  
Robinson, J. A., 37, 266, 296, 321, 329, *1059*  
Robinson, S., 125, *1110*  
RoboCup, 983  
robopocalypse, **1052**  
robot, **932**, 932–934, 982  
  hexapod, 975  
  mobile, 934  
  navigation, 88  
  rights, 1061–1062  
  soccer, 192, 640  
robotics, **20**, 505, 932–987  
robotic soccer, 225  
robust control theory, **852**  
Roche, E., 905, *1110*  
Rochester, N., 36, 37  
Rock, D., 1049, *1089*  
Rock, I., 1031, *1110*  
Rodríguez, H., 903, *1105*  
Rodríguez, K., 1064, *1117*  
Röger, G., 125, 399, *1098*, *1112*  
Rogue One, 1022  
Rohlfshagen, P., 222, *1089*  
ROI pooling, **1007**  
Rokicki, T., 124, *1110*  
Rolf, D., 266, *1110*  
Rolf, E., 1060, *1104*  
Rolland, N., 667, *1097*  
rollout, **207**  
Rolnick, D., 48, 838, *1110*  
Romania, 81, 165  
Romanovskii, I., 637, *1110*  
Romero, J., 48, *1103*  
Rong, Z., 923, *1115*  
Roossin, P., 931, *1089*  
root mean square (RMS), 1080  
Ros, G., 873, *1110*  
Rosen, C., 983  
Rosenblat, A., 1060, *1090*  
Rosenblatt, F., 39, 735, 836, 1028, *1088*, *1096*, *1110*  
Rosenblatt, M., 799, *1110*  
Rosenblitt, D., 398, *1105*  
Rosenbloom, P. S., 310, 329, 400, *1103*  
Rosenblueth, A., 34, *1110*  
Rosenbluth, A., 160, 475, *1106*  
Rosenbluth, M., 160, 475, *1106*  
Rosenschein, J. S., 639, 640, *1111*, *1114*  
Rosenschein, S. J., 79, 267, 268, *1100*, *1111*  
Rosenthal, G. D., 549, *1101*  
Ross, G., 27, *1111*  
Ross, S., 428, 986, 1080, *1111*  
Rosse, C., 334, *1113*  
Rossi, F., 188, 189, 191, *1086*, *1088*, *1111*  
Roth, A., 1060, *1087*, *1090*, *1094*  
Roth, D., 475, 668, *1092*, *1111*  
Roughgarden, T., 638, *1108*  
Roussel, P., 296, 329, *1091*, *1111*  
route finding, 87  
rover, **934**  
Roveri, M., 399–401, *1087*, *1091*  
Rowat, P. F., 984, *1111*  
Roweis, S. T., 473, 516, *1111*  
Rowley, H., 1029, *1111*  
row player, **596**  
Roy, B. V., 586, *1092*  
Roy, D., 667, *1085*, *1097*  
Roy, N., 588, 986, *1100*, *1111*, *1114*  
Rozonoer, L., 735, *1085*  
RPM (relational probability model), **643**  
RRT (rapidly exploring random trees), 954, 985  
RRT\*, **955**  
RSA (Rivest, Shamir, and Adelman), 327  
Ruan, P., 587, *1113*  
Rubik's Cube, 116, 118, 124  
Rubin, D., 476, 515, 516, 799, *1093*, *1096*, *1111*  
Rubin, J., 1060, *1115*  
Rubinstein, A., 588, 637, 638, 640, *1085*, *1108*, *1111*  
Rubinstein, B. I., 1060, *1107*  
Ruder, S., 926, 930, *1099*, *1111*  
Ruderman, A., 48, 225, *1100*  
rule, **235**  
  causal, 436  
  condition-action, **68**  
  default, **352**  
  diagnostic, 436  
  if-then, 50, 235  
  implication, 235  
  situation-action, 68  
rule-based system, 477  
Rumelhart, D. E., 42, 836, *1111*  
Rumelhart, D. E., 42, 836, *1111*  
Ruml, W., 126, 127, *1089*, *1098*, *1114*, *1117*  
Rummery, G. A., 871, *1111*  
Rush Hour (puzzle), 86  
Ruspini, E. H., 477, *1111*  
Russakovsky, O., 837, *1111*  
Russell, A., 125, *1093*  
Russell, B., 25, 35, 36, 328, *1116*  
Russell, J. G. B., 549, *1111*  
Russell, J. R., 331, *1107*  
Russell, S. J., 80, 126, 127, 162, 221, 222, 267, 316, 339, 400, 515–517, 551, 586–588, 637, 638, 653, 666–668, 799, 871–873, 984, 1070, 1071, 1085–1087, 1092–1095, 1097–1099, *1101*, *1105*–*1108*, *1111*, *1113*, *1114*, *1117*, *1118*  
Russo, C., 667, *1097*  
Rustagi, J. S., 476, *1111*  
Rusu, A. A., 871, 873, *1106*  
Ruzzo, W. L., 904, *1097*  
Ryan, M., 297, *1099*  
Rzepa, H. S., 358, *1107*
- 
- S**
- Saad, F., 667, 668, *1091*, *1111*  
Sabharwal, A., 266, 399, 901, *1091*, *1097*, *1099*  
Sabin, D., 189, *1111*  
Sabri, K. E., 331, *1111*  
Sacerdoti, E. D., 398, 400, *1111*  
Sackinger, E., 44, 837, 1029, *1103*  
Sadeghi, F., 986, *1111*  
Sadeh, N. M., 639, *1086*  
Sadigh, D., 971, 986, *1111*  
Sadiq, A., 871, 873, *1087*, *1106*  
Sadler, M., 223, *1111*  
Sadri, F., 358, *1111*  
Saeed, M., 737, *1097*  
safe exploration, 967  
safety engineering, **1052**  
Sagae, K., 904, *1111*  
Sagiv, Y., 329, *1086*  
Saha, D., 1047, 1060, *1087*  
Sahami, M., 903, *1102*, *1111*  
Sahin, N. T., 271, *1111*  
Sainath, T., 900, 905, *1090*, *1099*  
SAINT (mathematical problem solver), 38, 162  
St. Petersburg paradox, 547  
Sakuta, M., 222, *1111*  
Salakhutdinov, R., 667, 838, 930, *1103*, *1113*, *1117*  
Salesforce, 1068  
Salib, M., 731, 737, *1089*  
Salido, M. A., 189, 191, *1086*  
Salisbury, J., 985, *1105*  
Salmond, D. J., 516, *1097*  
Salomaa, A., 903, *1111*  
Salvatier, J., 46, *1097*  
Salzmann, M., 838, *1117*  
SAM (theorem prover), 331

- Samadi, M., 127, *1111*  
 Samet, H., 735, *1111*  
 Sammut, C., 871, *1111*  
 Samothrakis, S., 222, *1089*  
 sample complexity, **692**  
 sample size disparity, **1046**  
 sample space, **406**  
 sampling, 453–459, **1080**  
   adaptive, 475  
   correlated, **863**, 872  
   direct, 454  
   rejection, 455  
 Samsung, 1031  
 Samuel, A., 35–37, 51, 80, 223, 733, 870, 871, *1111*  
 Samuel, S., 668, *1099*  
 Samuelson, L., 637, *1105*  
 Samuelson, W., 639, *1110*  
 Sanchez-Lengeling, B., 1027, *1111*  
 Sanders, P., 126, *1093*  
 Sandholm, T., 48, 224, 637–639, 841, 1089, *1096*, *1111*  
 Sang, T., 474, *1111*  
 Sanna, R., 836, *1096*  
 Sanskrit, 356, 903  
 Santorini, B., 880, 903, *1105*  
 SAPA (planning system), 401  
 Sapir, E., 875, *1111*  
 Sapir-Whorf hypothesis, 270  
 Saraswat, V., 189, *1115*  
 Sarawagi, S., 906, *1111*  
 Sargent, T. J., 872, *1111*  
 SARSA (state-action-reward-state-action), **853**  
 Sartre, J.-P., 636, *1111*  
 Sastry, G., 872, *1111*  
 Sastry, S., 79, 667, 868, 872, 971, 986, 1098, *1107*, *1108*, *1111*  
 SAT, **241**  
 Satheesh, S., 837, *1111*  
 Satia, J. K., 588, *1111*  
 satisfaction (in logic), **232**  
 satisfiability, **240**, 266  
 satisfiability threshold conjecture, **254**  
 sacrificing, **28**, 108  
 SATMC (logical reasoning system), 268  
 Sato, T., 329, 666, *1111*, *1114*  
 SATPLAN, **262**  
 Sattigeri, P., 1047, 1060, *1087*  
 saturation, **322**  
 SATZ (logical reasoning system), 266  
 Saul, L. K., 476, 838, *1100*, *1111*  
 Saund, E., 903, *1111*  
 Saunders, W., 872, *1111*  
 Saurous, R. A., 667, *1115*  
 Savage, L. J., 412, 427, 548, *1111*  
 Savani, Y., 838, *1116*  
 Savva, M., 873, *1111*  
 Saxe, A. M., 837, *1097*  
 Saxe, R., 872, *1086*  
 Saxena, N., 1060, *1106*  
 Sayre, K., 1033, *1111*  
 scaled orthographic projection, **992**  
 scanning lidar, **935**  
 Scarcello, F., 191, 360, *1094*, *1097*  
 scene, **989**  
 Scha, R., 891, *1088*  
 Schaal, S., 986, *1086*, *1109*  
 Schabes, Y., 891, 905, *1109*, *1110*  
 Schaeffer, J., 106, 126, 127, 223, 637, 1085, *1087*, *1091*, *1094*, *1111*  
 Schäfer, A. M., 872, *1098*  
 Schank, R. C., 41, *1111*  
 Schapire, R. E., 718, 736, 903, *1095*, *1111*, *1112*  
 Scharir, M., 984, *1112*  
 Scharre, P., 1059, *1112*  
 Schaub, T., 359, *1093*  
 Schauenberg, T., 637, *1087*  
 scheduling, **392**, 392–396, **394**  
 Scheines, R., 798, *1113*  
 schema (in a genetic algorithm), **135**  
 schema acquisition, **768**  
 Schervish, M. J., 428, *1093*  
 Schickard, W., 24  
 Schlaefer, N., 48, *1094*  
 Schlenoff, C. I., 986, *1105*  
 Schmid, C., 1029, *1112*  
 Schmid, M., 48, 224, *1104*, *1107*  
 Schmidhuber, J., 838, 839, *1096*, *1099*, *1112*  
 Schmidt, G., 402, *1088*  
 Schmitt, S., 224, *1112*  
 Schneegäf, D., 872, *1098*  
 Schneider, J., 868, 873, 986, *1085*, *1086*, *1089*  
 Schnitzius, D., 402, *1092*  
 Schnizlein, D., 638, *1116*  
 Schoenberg, I. J., 836, *1107*  
 Schoenick, C., 901, *1091*  
 Schofield, M., 222, *1112*  
 Schölkopf, B., 476, 736, *1091*, *1092*, *1108*, *1112*  
 Schomer, D., 271, *1111*  
 Schön, T. B., 517, *1104*  
 Schöning, T., 266, *1112*  
 Schoppers, M. J., 401, *1112*  
 Schrag, R. C., 191, 266, *1086*  
 Schraudolph, N. N., 223, *1112*  
 Schrauwen, B., 47, *1115*  
 Schrittewieser, J., 45, 48, 201, 220, 223, 224, 871, 873, *1087*, *1112*  
 Schröder, E., 266, *1112*  
 Schrödl, S., 127, *1094*  
 Schubert, L. K., 357, *1100*  
 Schulman, J., 872, 873, 985, 986, 1061, 1085, *1089*, *1095*, *1112*  
 Schultes, D., 126, *1093*  
 Schultz, K., 668, *1105*  
 Schultz, W., 873, *1112*  
 Schulz, D., 667, 984, *1089*, *1112*  
 Schulz, S., 330, 331, *1112*, *1114*  
 Schumann, J., 330, *1093*, *1103*  
 Schuster, M., 47, 834, 901, 916, 929, 1090, *1100*, *1117*  
 Schutt, R., 736, *1108*  
 Schütze, H., 901, 903, 905, *1105*, *1112*  
 Schwartz, J. T., 984, *1112*  
 Schwartz, R., 928, *1097*  
 Schwartz, S. P., 357, *1112*  
 Schwartz, W. B., 428, *1097*  
 Schwinger, A., 1016, 1023, *1085*, *1093*  
 scientific discovery, 734  
 Scikit-Learn (machine learning software), 738, 1072  
 SciPY (scientific software), 1072  
 Sciuto, C., 838, *1117*  
 Scott, D., 665, *1112*  
 Scrabble, 225  
 scruffy vs. neat, 42  
 Sculley, D., 731, 737, *1089*  
 search, 40, 71  
   A\*, 103–108  
   algorithm, **89**  
   alpha-beta, 198–201, 220, 221  
   B\*, 221  
   backtracking, 98, 179–181, 183, 187  
   beam, **110**, 124, **133**, 205, 882, 887, 919  
   local, **133**  
   stochastic, **133**  
   best-first, **91**, 123  
   bidirectional, **100**, 114–115, 127  
   breadth-first, **94**, 94–95, 123, 376  
   conformant, 144–145  
   continuous space, 137–140, 160  
   cutting off, 204–206  
   depth-first, **96**, 96–98, 123, 376  
   depth-limited, **98**, 98  
   greedy best-first, **103**, 103  
   heuristic, 125  
   hill-climbing, 129–132, 156  
   in a CSP, 175–183  
   incremental belief-state, **147**  
   informed, 81, **102**, **102**–123  
   iterative deepening, **98**, 98–100, 123, 125, 201, 204, 376  
   iterative deepening A\*, 126  
   learning to, 121  
   local, 128–137, 160, 190, 253–254, 265, 266  
   greedy, **130**  
   local, for CSPs, 181–183  
   local beam, 133  
   memory-bounded, 110–113, 126  
   memory-bounded A\*, **113**, 113, 127  
   minimax, **194**, 194–198

- nondeterministic, 140–143  
 online, 152, 152–159, 162–163  
 partially observable, 144–152  
 policy, 841, 861, 861–863  
 problem, 949  
 quiescence, 204  
 real-time, 145, 202–206  
 recursive best-first (RBFS), 111–113, 126  
 simulated annealing, 132–133  
 stochastic beam, 133  
 systematic, 93  
 tabu, 160, 182  
 tree, 89, 193  
 uniform-cost, 95, 95–96, 123  
 uninformed, 63, 94–102, 123, 125  
 weighted A\*, 109  
 search tree, 89, 193  
 Searle, J.R., 29, 1036, 1058, 1112  
 secure aggregation, 1043  
 Sedgewick, R., 1080, 1112  
 Sedol, L., ix, 864  
 Sefidgar, Y. S., 974, 986, 1112  
 Segal, A., 1043, 1088  
 Segaran, T., 639, 738, 1112  
 segmentation (of an image), 1001  
 Seipp, J., 399, 1112  
 Sejnowski, T., 223, 839, 1098, 1112  
 selection (in evolutionary algorithms), 134  
 selection policy (in game tree search), 207  
 selection problem, 576  
 Self, M., 799, 1090  
 self-attention, 919  
 self-occlusion, 1003  
 Selfridge, O. G., 36  
 Sellart, L., 873, 1110  
 Selman, B., 160, 190, 266, 267, 359, 399, 1097, 1101, 1112  
 Selsam, D., 668, 1105  
 semantic interpretation, 894–895, 904  
 semantic network, 347, 347–349, 355, 359  
 semantics, 232, 884, 892  
 database, 283, 315, 363, 643  
 logical, 264  
 Semantic Web, 357  
 semi-Markov decision process, 860  
 semidecidable, 300, 328  
 semidynamic environment, 64  
 semisupervised learning, 723  
 semisupervised parsing, 891  
 Senges, M., 1062, 1108  
 Seni, G., 736, 1112  
 Senior, A., 830, 838, 900, 905, 1099, 1115  
 sensitivity analysis, 542  
 sensor, 54, 61, 989  
 active, 934  
 lidar, 61, 1024, 935  
 passive, 934  
 sonar, 934  
 tactile, 935  
 ultrasound, 61  
 sensor array, 940  
 sensor failure, 505  
 sensorless planning, 383, 385–388  
 sensor model, 70, 479, 482, 492, 499, 514, 578, 939  
 Sensory Graphplan (SGP), 401  
 sentence  
   atomic, 235, 278, 278, 282  
   complex, 235, 278  
   in a KB, 227, 264  
   natural language, 874, 885  
   as physical configuration, 234  
 sentiment analysis, 877, 914  
 Seo, H., 873, 1103  
 Seo, M., 931, 1112  
 separator (in Bayes net), 420  
 Seppi, K. D., 586, 1117  
 sequence-to-sequence model, 916  
 sequence form, 611  
 sequential decision problem, 552–562, 586  
 sequential environment, 63  
 sequential importance sampling with resampling (SIR), 516  
 serendipity, 392  
 Sergot, M., 358, 1102  
 serializable subgoals, 373  
 Servant of Philon, 983  
 Seshia, S. A., 971, 986, 1111  
 set (in first-order logic), 287  
 set-cover problem, 371  
 Seth, K., 1043, 1088  
 SETHEO (theorem prover), 330  
 set of support, 326  
 set partitioning problem, 621  
 Settle, L., 331, 1097  
 Settlers of Catan, 197  
 Seung, H. S., 986, 1114  
 Severini, S., 1069, 1093  
 Severyn, A., 904, 1085  
 SEXTANT (deep space navigation system), 47  
 Seymour, P. D., 190, 1110  
 SGD (stochastic gradient descent), 697, 816  
 SGP (Sensory Graphplan), 401  
 Sha, F., 928, 1090  
 Shachter, R. D., 473–475, 523, 548, 586, 1112, 1114  
 Shadow Dexterous Hand, 937  
 Shafer, G., 477, 1112  
 shaft decoder, 936  
 Shafro, P., 588, 1110  
 Shah, J., 1028, 986, 1107  
 Shahrzad, H., 137, 161, 1106  
 Shaked, T., 906, 1094  
 Shaker, G., 1064, 1117  
 Shakey, 20, 61, 162, 398, 401, 983, 986  
 Shalata, A., 474, 1112  
 Shalla, L., 330, 1117  
 Shamdas, M., 48, 1104  
 Shanahan, M., 341, 358, 1061, 1112  
 Shani, G., 588, 1112  
 Shankar, N., 331, 1112  
 Shannon, C. E., 36, 201, 221, 679, 733, 738, 902, 1112  
 SHAP (explainable machine learning system), 737  
 shape from shading, 1030  
 shaping theorem, 559  
 Shapley, L. S., 638, 1112  
 Shapley, S., 586, 638, 1112  
 Shapley value, 618, 619, 635  
 Sharan, R. V., 1027, 1112  
 shared model, 1068  
 Sharir, M., 985, 1097  
 Sharon, G., 126, 1099  
 Sharp, D. H., 836, 1091  
 Shatkay, H., 984, 1112  
 Shaw, J. C., 124, 266, 1107  
 Shawe-Taylor, J., 736, 1091  
 Shazeer, N., 736, 901, 919, 929–931, 1072, 1100, 1100, 1112, 1115  
 Shehory, O., 638, 1111  
 Shelley, M., 1052, 1112  
 Shelton, C., 516, 1108  
 Shen, Q., 799, 1092  
 Shepard, D., 549, 1117  
 Sheppard, B., 225, 1112  
 Shet, V., 820, 1097  
 Shi, E., 1060, 1107  
 Shi, J., 1001, 1028, 1112  
 Shieber, S., 1057, 1112  
 shift-reduce parsing, 889  
 Shimelevich, L. I., 516, 1117  
 Shimony, S. E., 475, 587, 1070, 1098, 1112  
 Shin, M. C., 586, 1109  
 Shinkareva, S. V., 271, 1106  
 Shmakov, A., 124, 1085  
 Shmatikov, V., 1042, 1060, 1086, 1107  
 Shmoys, D. B., 125, 395, 402, 1103  
 Shoar, S., 161, 1096  
 Shoham, Y., 79, 225, 330, 549, 638–640, 1085, 1091, 1100, 1103, 1112  
 short-term memory, 310  
 shortcut, 121  
 Shortliffe, E. H., 477, 1089, 1112  
 shoulder (in state space), 131  
 Shpitser, I., 653, 666, 1108  
 SHRDLU (natural language system), 38, 41, 364, 905  
 Shreve, S. E., 79, 1087

- Shroff, G., 1041, *1105*  
 Shtark, O., 474, *1112*  
 Siciliano, B., 987, *1112*  
 Siddiqui, N., 1023, *1093*  
 sideways move (in state space), **131**  
 Siebel, W. A., 1027, *1102*  
 Sievers, S., 399, *1112*  
 Sifre, L., 45, 48, 201, 220, 223, 224, 871, *1112*  
 SIGAI, 53  
 SIGART, 53  
 Sigaud, O., 588, *1112*  
 sigmoid, 703, **803**  
 Sigmund, K., 25, *1112*  
 signed distance field, **956**  
 significance test, **681**  
 Silberstein, M., 474, *1112*  
 Silva, R., 225, *1112*  
 Silver, D., 37, 45, 48, 201, 220, 222–225, 588, 835, 841, 871, 873, 986, *1104*, *1106*, *1112*, *1115*  
 Silverman, B. W., 838, *1112*  
 Silverstein, C., 905, *1113*  
 sim-to-real, 933, **966**, 985, 986  
 Sim'a'n, K., 891, *1088*  
 Simard, P., 44, 837, 1029, *1087*, *1103*  
 Simchowitz, M., 1060, *1104*  
 Simeon, T., 986, *1113*  
 Simmons, R., 516, 984, *1113*, *1115*  
 Simon's predictions, 39  
 Simon, D., 79, *1113*  
 Simon, H.A., 20, 28, 36, 37, 40, 79, 124, 125, 266, 327, 398, 551, *1100*, *1107*, *1113*  
 Simon, J. C., 267, *1113*  
 Simonis, H., 189, *1113*  
 Simons, P., 360, *1107*  
 Simonyan, K., 27, 30, 224, 830, 838, 900, 1027, *1094*, *1100*, *1104*, *1112*, *1115*  
**SIMPLE-REFLEX-AGENT**, **69**  
 simple majority vote, **630**  
 simplex algorithm, 161  
**SIMULATED-ANNEALING**, **133**  
 simulated annealing, 132–133, 159, 160, 395, 460  
 simulation, **207**  
 simultaneous localization and mapping (SLAM), **512**, **942**  
 Sin, B.-K., 516, *1114*  
 Sinclair, A., 132, 161, *1104*, *1110*  
 Singer, P. W., 1059, *1113*  
 Singer, Y., 516, 903, *1095*, *1111*  
 Singh, A., 930, 931, *1116*  
 Singh, M., 80, 1047, 1060, *1087*, *1099*  
 Singh, N., 48, *1104*  
 Singh, P., 50, 334, *1106*, *1113*  
 Singh, R., 667, *1087*  
 Singh, S., 163, 668, 587, 737, 871–873, 984, *1086*, *1089*, *1101*, *1105*, *1108*, *1114*  
 singly connected network, **451**  
 singular extension, **205**  
 singularitarianism, 1061  
 singularity, **30**, **1055**  
 singular matrix, **1077**  
 sins, seven deadly, 130  
 SIPE (planning system), 401  
 SIR (sequential importance sampling with resampling), 516  
 SIS, *see* importance sampling, sequential  
 Sisbot, E. A., 986, *1113*  
 Siskind, J. M., 1067, *1113*  
 Sistla, A. P., 399, *1113*  
 Sittler, R. W., 666, 667, *1113*  
 situated agent, 1033  
 situation calculus, **370**  
 Sjolander, K., 516, *1102*  
 Skinner, B. F., 34  
 skip-gram, **877**, 923  
 Skolem, T., 296, 328, *1113*  
 Skolem constant, **299**, 328  
 Skolem function, **318**, 328  
 skolemization, **317**  
 Skolnick, M. H., 474, *1089*  
 Skype, 47  
 slack (in scheduling), **394**  
 Slagle, J. R., 38, *1113*  
 SLAM (simultaneous localization and mapping), **512**, **942**  
 Slate, D. J., 125, *1113*  
 Slater, E., 221, *1113*  
 Slattery, S., 905, *1091*  
 sliding-tile puzzle, **86**, 120, 371  
 sliding window, **1006**  
 SLING (natural language system), 928  
 Slocum, J., 124, *1113*  
 Sloman, A., 50, *1106*  
 Slovic, P., 550, *1100*  
 small-scale learning, **688**  
 Smallwood, R. D., 587, *1113*  
 SMA\*, 123  
 Smith, A., 28, 782, *1087*  
 Smith, A. F. M., 516, 799, *1097*, *1115*  
 Smith, B., 47, 334, 358, 401, *1100*, *1113*  
 Smith, D. A., 904, *1113*  
 Smith, D. E., 162, 316, 329, 400, 401, *1089*, *1096*, *1102*, *1109*, *1113*, *1116*  
 Smith, E., 737, *1113*  
 Smith, G., 126, *1110*  
 Smith, J. E., 527, 549, *1113*  
 Smith, J. L., 48, *1104*  
 Smith, J. M., 161, 638, *1113*  
 Smith, J. Q., 548, *1108*, *1113*  
 Smith, K., 737, *1087*  
 Smith, M. K., 357, *1113*  
 Smith, N. A., 928, *1097*  
 Smith, R. C., 984, *1113*  
 Smith, R. G., 80, 639, *1089*, *1113*  
 Smith, S. J. J., 224, *1113*  
 Smith, T., 47, *1086*  
 Smith, V., 639, *1110*  
 Smith, W. D., 221, 473, *1086*, *1107*  
 Smith, W. E., 551, *1113*  
 SMODELS (logic programming system), 360  
 Smola, A. J., 736, *1112*  
 Smolensky, P., 902, 1068, *1113*  
 smoothing, 486–489, 514, 795, **878**, 998  
 complexity, 488  
 linear interpolation, **878**  
 online, 493  
 SMOTE (data generation system), 725, 1046  
 Smallyan, R. M., 297, *1113*  
 Smyth, P., 516, *1113*  
 SNARC (neural network hardware), 35  
 Snell, J., 428, *1097*  
 Snoek, J., 690, *1113*  
 Snyder, W., 330, *1086*  
 SOAR (cognitive architecture), 310, 329, 400  
 soccer, 225  
 robotic, 192, 640  
 Socher, R., 44, 923, 929, 931, 1072, *1093*, *1098*, *1101*, *1109*  
 social choice function, **629**  
 social choice theory, **628**  
 social law, **595**  
 social outcome, **629**  
 social preference order, 629  
 social welfare, **599**  
 social welfare function, **629**  
 societal bias, **1043**  
 Society for Artificial Intelligence and Simulation of Behaviour (AISB), 53  
 society of mind, 636  
 Socrates, 21  
 Soderland, S., 357, 901, 906, *1086*, *1094*  
 softbot, **61**  
 soft margin, **714**  
 softmax, **809**, 862  
 softplus, **803**  
 soft threshold, 441  
 software agent, **61**  
 software engineering, 1053  
 Sohl-Dickstein, J., 838, *1109*  
 Sohn, J. H., 48, *1093*  
 Soika, M., 984, *1088*  
 sokoban puzzle, **85**  
 solipsism, 1036  
 Solla, S., 838, *1103*  
 Solomonoff, R. J., 36, 734, *1113*  
 solution, **82**, 83, 123, **165**

- optimal, 83  
 solution concept, 596  
 soma, 30  
 sonar, 934  
 Sondik, E. J., 587, 1113  
 Sonenberg, L., 737, 1061, 1106  
 Song, Z., 837, 1085  
 Sonneveld, D., 124, 1113  
 Sontag, D., 667, 1106  
 Sophia (robot), 1051  
 Sorensen, T., 637, 1096  
 Sosic, R., 190, 1113  
 soundness (of inference), 234, 240, 248, 264, 306  
 Sowa, J., 361, 1113  
 Spaan, M. T. J., 588, 1113  
 space complexity, 93, 123  
 spacecraft assembly, 402  
 SPACY (parser), 904  
 spam detection, 877  
 Sparck Jones, K., 428, 1110  
 Sparrow, R., 1051, 1113  
 sparse model (in machine learning), 699  
 sparse reward (in reinforcement learning), 841  
 sparse system, 435  
 sparse transition model, 560  
 SPASS (theorem prover), 330  
 spatial reasoning, 360  
 species, 42, 334, 335, 357, 789, 1061  
 speech act, 897  
 speech recognition, 20, 43, 900, 905  
 speech synthesis, 47  
 speedy search, 110  
 sphe wasp, 59, 392  
 SPI (Symbolic Probabilistic Inference), 474  
 Spiegelhalter, D. J., 474–476, 666, 549, 798, 1091, 1096, 1103, 1105, 1113  
 SPIKE (planning system), 402  
 SPIN (logical reasoning system), 327  
 Spiropulu, M., 1069, 1107  
 Spirites, P., 798, 1113  
 Spirtkovsky, V. I., 891, 1113  
 Spitzer, E., 1046, 1106  
 split point, 682  
 Spronck, P., 222, 1090  
 Sproull, R. F., 80, 548, 1094  
 SQuAD (Stanford Question Answering Dataset), 43, 46, 931  
 square root, 67  
 Srebro, N., 1060, 1098  
 SRI, 38, 296, 398, 548, 983  
 Srinivasa, S., 958, 970, 985, 986, 1093, 1110, 1118  
 Srinivasan, M. V., 1064, 1095  
 Srivas, M., 327, 1113  
 Srivastava, N., 838, 1113  
 Srivastava, S., 587, 1113  
 SSS\*algorithm, 221  
 Staab, S., 357, 1113  
 stability  
     of a controller, 960  
     strict, 960  
 stack (data structure), 92  
 stacked generalization, 717  
 Stader, J., 402, 1085  
 STAGE (local search algorithm), 160  
 stage game, 604  
 Stallman, R. M., 189, 1113  
 STAN (probabilistic programming language), 476, 668  
 standard deviation, 219, 1079  
 standardized problem, 84  
 standardizing apart, 302, 368  
 standard model, 22  
 Stanfill, C., 735, 1113  
 Stanford Parser, 904  
 Stanford University, 37, 40, 41, 45, 296, 639  
 Stanhope Demonstrator, 266  
 Staniland, J. R., 428, 1092  
 Stanislawska, K., 161, 1113  
 Stanley (autonomous vehicle), 979, 980, 984  
 StarCraftII, 46, 48, 193, 218, 225, 873  
 START (question-answering system), 901  
 Star Trek, 1058  
 start symbol, 1081  
 state, 362  
     repeated, 92  
     world, 84  
 state-action-reward-state-action (SARSA), 853  
 state abstraction, 373, 872  
 state description, 86  
 state estimation, 150, 215, 260, 265, 483, 938  
     recursive, 150, 484  
 States, D. J., 799, 1099  
 state space, 83  
     graph, 123, 193  
     joint, 860  
     metalevel, 121  
 static environment, 63  
 stationarity, 683  
 stationarity (for preferences), 556  
 stationary distribution, 462, 486  
 statistics, 26, 1079  
 Statnikov, A. R., 737, 1097  
 steepest ascent, 129  
 Stefik, M., 361, 477, 1113  
 Stein, C., 125, 1080, 1091  
 Stein, J., 473, 1107  
 Steinbach, M., 738, 1114  
 Steinberg, R., 639, 1091  
 Steiner, D. F., 48, 1113  
 Steiner, W., 984, 1089  
 Steinhardt, J., 1061, 1085  
 Steinruecken, C., 737, 1113  
 Stensrud, B., 329, 1114  
 Stentz, A., 160, 1093  
 Stepleton, T., 872, 1107  
 step size, 139  
 stereogram, random dot, 1028  
 stereo vision, 934  
 Stergiou, K., 189, 1113  
 Stern, H. S., 799, 1096  
 Stickel, M. E., 266, 330, 1113, 1118  
 stiction, 959  
 stiff neck, 417  
 Stiller, L., 222, 1113  
 stimulus, 31, 988  
 Stob, M., 734, 1108  
 stochastic beam search, 133  
 stochastic dominance, 531, 547  
 stochastic environment, 163, 552  
 stochastic gradient descent, 697, 816  
 stochastic gradient descent (SGD), 697, 816  
 stochastic local search, 159  
 STOCKFISH (chess program), 206, 222, 223  
 Stockman, G., 221, 1113  
 Stoffel, K., 357, 1113  
 Stoic school, 265  
 Stokes, I., 402, 1085  
 Stolcke, A., 47, 1117  
 Stone, C. J., 734, 1089  
 Stone, M., 734, 1113  
 Stone, P., 45, 225, 639, 640, 873, 984, 1086, 1097, 1113, 1114  
 stopping time, 573  
 Stork, D. G., 738, 800, 1094  
 Storvik, G., 517, 1114  
 Story, W. E., 124, 1100  
 Stoyanov, V., 927, 930, 1104  
 Strachey, C., 32, 35, 221, 223, 1114, 1115  
 straight-line distance, 103  
 Strat, T. M., 477, 1111  
 strategic decision making, 590  
 strategy, 215, 596  
 strategy-proof, 625  
 strategy profile, 596  
 Stratонovich, R. L., 515, 550, 1114  
 Straub, J., 873, 1111  
 Straw, A. D., 1064, 1095  
 strawberries, enjoy, 1033  
 strictly stable, 960  
 stride, 812, 1007  
 Striebel, C. T., 515, 1110  
 string (in logic), 359  
 STRIPS (planning system), 398, 400, 401  
 Strohm, G., 401, 1095  
 Strohman, T., 901, 905, 1091  
 Strohmann, T., 334, 1093

strong AI, 1032, 1056, 1057  
 strong domination, 597  
 Stross, C., 30, 1093  
 structural equation, 468  
 structural estimation of MDPs, 872  
 structured light, 935  
 structured representation, 77, 81, 408  
 Stuckey, P. J., 189, 330, 1100, 1105  
 STUDENT (natural language system), 38  
 stuff, 339  
 Stuhlmüller, A., 668, 872, 1111, 1117  
 Stumpe, M. C., 48, 1097, 1104, 1113  
 stupid backoff, 879  
 stupid pet tricks, 59  
 Sturtevant, N.R., 106, 114, 126, 163,  
     1085, 1090, 1094, 1099, 1114  
 Stützle, J., 799, 1090  
 Stutzle, T., 160, 190, 1093, 1099  
 style transfer, 1022  
 Su, H., 837, 1111  
 Su, Y., 125, 1093  
 subcategory, 335, 892  
 subgame, 609  
 subgame perfect Nash equilibrium, 609  
 subgoal independence, 374  
 subjective case, 892  
 subjectivism, 427  
 subproblem, 118  
 Subrahmanian, V. S., 222, 1108  
 Subramanian, D., 360, 1071, 1111, 1114  
 substance, 339–340  
 substitutability (of lotteries), 520  
 substitution (in first-order logic), 284, 298  
 subsumption  
     in description logic, 349  
     in resolution, 326  
 subsumption architecture, 976  
 subsumption lattice, 303  
 successor-state axiom, 258, 268  
 successor node, 90  
 Suciu, D., 665, 668, 1092, 1097  
 Suderth, E., 653, 667, 1085  
 Sudholter, P., 638, 1109  
 Sudoku, 173  
 Suk, H.-I., 516, 1114  
 Sulawesi, 183  
 SUMMATION, 1074  
 Summers-Stay, D., 1017, 1097  
 summing out, 414, 449  
 sum of squared differences (SSD), 1000  
 Sun, A., 47, 1118  
 Sun, J., 837, 1098  
 Sun, R., 1023, 1093  
 Sun, S., 334, 1093  
 Sun, Y., 930, 1114  
 SUNFISH (chess program), 222  
 Sunstein, C., 550, 1114  
 Sunter, A., 666, 1094  
 superadditivity, 617

SUPERGLUE (natural language benchmark), 931  
 Superman, 270  
 supervised learning, 671, 840, 855  
 support vector machine (SVM), 710,  
     710–714  
 surely expanded nodes, 108  
 Suresh, A. T., 1043, 1102  
 sure thing, 525  
 survey propagation, 267  
 survival of the fittest, 516  
 Sussman, G. J., 189, 398, 1113, 1114  
 Sussman anomaly, 398  
 Sutcliffe, G., 330, 331, 1114  
 Sutherland, G. L., 40, 1089  
 Sutherland, I., 188, 1114  
 Sutphen, S., 223, 1111  
 Sutskever, I., 44, 837, 838, 900, 909, 930,  
     931, 1029, 1100, 1102, 1106,  
     1110, 1113, 1114  
 Suttner, C., 331, 1114  
 Sutton, C., 906, 1114  
 Sutton, R. S., 80, 586, 588, 871–873,  
     1086, 1114  
 Svartvik, J., 904, 1110  
 Svestka, P., 985, 1101  
 Svetnik, V. B., 516, 1117  
 SVM (support vector machine), 710,  
     710–714  
 Swade, D., 33, 1114  
 Swartout, W. R., 1060, 1107  
 Swartz, R., 356, 1089  
 Swayamdipta, S., 928, 930, 1097, 1111  
 Sweeney, L., 1042, 1060, 1114  
 Swerling, P., 515, 1114  
 Swersky, K., 1046, 1117  
 Swetter, S. M., 48, 1094  
 Swift, T., 329, 1114  
 switching Kalman filter, 502  
 sybil attack, 648  
 syllogism, 21, 24, 265  
 SYMBA\* (planning system), 399  
 Symbolic Probabilistic Inference  
     (SPI), 474  
 symmetric players, 619  
 symmetry breaking (in CSPs), 187  
 symmetry reduction, 372  
 synapse, 30  
 synchronization (in multiagent systems),  
     592  
 syntactic ambiguity, 898, 904  
 syntactic category, 884  
 syntactic sugar, 286  
 syntactic theory (of knowledge), 359  
 syntax, 41, 232, 235  
     of logic, 264  
     of natural language, 884  
     of probability, 410  
 synthesis (by theorem provers), 327

of algorithms, 327  
 SYNTHIA (simulated environment), 873  
 Syrjänen, T., 360, 1107  
 systematic search, 93  
 SYSTRAN (machine translation software),  
     931  
 Syverson, C., 1049, 1089  
 Szafron, D., 637, 638, 1087, 1088, 1116  
 Szathmáry, E., 161, 1113  
 Szegedy, C., 327, 330, 837, 838, 1085,  
     1086, 1100, 1104, 1114  
 Szeliski, R., 1031, 1114  
 Szepesvari, C., 222, 587, 873, 1102, 1114  
 Szerlip, P., 667, 1087  
 Szita, I., 222, 1090

---

**T**

*T* (fluent holds), 341  
 t-distributed stochastic neighbor  
 embedding (t-SNE), 727, 737  
 T-SCHED (planning system), 402  
 t-SNE (t-distributed stochastic neighbour  
 embedding), 727, 737  
 T4 (planning system), 401  
 T5 (natural language system), 928, 931  
 TABLE-DRIVEN-AGENT, 66  
 Tableau (data analysis software), 727  
 table lookup, 705  
 tabu search, 160, 182  
 tactile sensor, 935  
 Tadepalli, P., 551, 871, 1094, 1114  
 Tafjord, O., 901, 1091  
 tag (part of speech), 880  
 tagging system (for images), 1016  
 Tait, P.G., 124, 1114  
 Takusagawa, K. T., 668, 1109  
 Talos, 982, 1057  
 Talukdar, P., 901, 1106  
 Tamaki, H., 329, 1114  
 Tambe, M., 191, 1108  
 Tammelin, O., 48, 224, 1088  
 Tan, P., 738, 1114  
 Tang, E., 1069, 1114  
 Tang, J., 873, 1089  
 tanh, 804  
 Tank, D.W., 29, 1108  
 Tardos, E., 638, 1108  
 Tarricone, R., 549, 1087  
 Tarski, A., 296, 297, 904, 1114  
 Tash, J. K., 587, 1114  
 task-oriented domain (in multiagent  
     systems), 633  
 task announcement, 623  
 Taskar, B., 668, 1096  
 task environment, 60, 78  
 task planning (in robotics), 938  
 Tasmania, 183  
 Tassa, Y., 873, 986, 1098, 1104, 1114

- Tate, A., 377, 398, 400–402, 1085, 1087, 1114  
 Tatman, J. A., 586, 1114  
 Tattersall, C., 206, 1114  
 Tavener, S., 222, 1089  
 taxi, 60, 61, 403  
     automated, 75, 228, 403, 1070  
 taxonomic hierarchy, 41, 335, 357  
 Taylor, A. D., 638, 1114  
 Taylor, C., 1030, 1092  
 Taylor, G., 329, 1114  
 Taylor, M., 357, 1113  
 Taylor, P., 900, 1114  
 Taylor, R., 986, 1104  
 Taylor, W., 190, 1090  
 Taylor expansion, 942  
 TD-GAMMON (backgammon program), 37, 224, 867  
 technological unemployment, 983, 1049  
 teddy bear, 1033  
 Tedrake, R., 986, 1114  
 Tegmark, M., 838, 1110  
 Teh, Y. W., 516, 837, 1099, 1108  
 telephone, 914  
 telepresence robots, 978  
 Teller, A., 160, 475, 1106  
 Teller, E., 160, 475, 1106  
 Teller, S., 984, 986, 1088, 1114  
 Tellex, S., 986, 1114  
 Templeton Foundation, 1037  
 temporal-difference learning, 844–848, 869  
 temporal inference, 483–491  
 temporal invariance, 811  
 temporal logic, 273  
 temporal projection, 267  
 temporal reasoning, 134–142, 255–264, 340–343, 479–517  
 Tenenbaum, J. B., 225, 296, 667, 668, 550, 872, 1086, 1095, 1097, 1102, 1103, 1114  
 Tennenholz, M., 871, 1088  
 tennis, 593  
 tense, 897  
 tensor, 814  
 TensorFlow, 738  
 TensorFlow (machine learning software), 1060, 1072  
 tensor processing unit (TPU), 33, 45, 814  
 Teplyashin, D., 873, 1087  
 term (in logic), 277, 277–278  
 ter Meulen, A., 297, 1115  
 terminal state, 193  
 terminal symbol, 1081  
 terminal test, 193  
 termination condition, 964  
 term rewriting, 330  
 Tesauro, G., 37, 222, 224, 855, 866, 867, 871, 1114  
 Tesla, 1031, 980  
 test set, 672, 684  
 Tetlock, P. E., 46, 356, 1114  
 TETRAD (machine learning software), 798  
 Tetris, 562, 571  
 Teukolsky, S.A., 160, 1109  
 text-to-speech, 900  
 text classification, 421, 913  
 TEXTRUNNER (information extraction system), 334, 901, 906  
 texture, 999  
 texture gradient, 1028  
 Teyssier, M., 798, 1114  
 Thaler, R., 549, 550, 1114  
 Thayer, J. T., 127, 1114  
 thee and thou, 886  
 Theocharous, G., 516, 1114  
 theorem, 285  
     incompleteness, 27, 323, 1034  
 theorem proving, 39, 240, 316–327, 398  
 thermostat, 34  
 Theseus, 733  
 Thiele, T., 515, 1114  
 Thielscher, M., 222, 268, 358, 1112, 1114  
 thingification, 335  
 thinking humanly, 20  
 thinking rationally, 21  
 thinkism, 1056  
 Thirion, B., 738, 1109  
 Thitimajshima, P., 476, 1087  
 Thng, F., 48, 1113  
 Thomas, A., 476, 666, 798, 1096, 1105  
 Thomas, J., 738, 1091  
 Thomas, P. S., 872, 1114  
 Thomaz, A., 986, 1085, 1114  
 Thompson, B., 588, 1117  
 Thompson, E., 474, 1089, 1112  
 Thompson, K., 222, 1091, 1114  
 Thompson, W. R., 576, 587, 1114  
 Thompson sampling, 576  
 Thorndike, E., 870, 1114  
 Thornton, C., 737, 1115  
 Thorpe, C., 46, 1101  
 thought, 21, 37, 226  
     laws of, 21–22  
 thrashing, 113  
 3-SAT, 266, 267, 308, 451  
 threshold function, 701  
 Throop, T.A., 224, 1113  
 Thrun, S., 46, 48, 79, 516, 588, 737, 984–987, 1089, 1090, 1094, 1095, 1106, 1109, 1111, 1115  
 Thurstone, L. L., 667, 1115  
 Tian, H., 930, 1114  
 Tian, J., 473, 1115  
 Tiao, G. C., 798, 1088  
 Tibshirani, R., 735–738, 800, 1087, 1095, 1098, 1100  
 tic-tac-toe, 193, 221  
 Tikhonov, A. N., 734, 1115  
 time, 479–517  
 time (discrete), 480  
 time (in grammar), 897  
 time-of-flight camera, 935  
 time complexity, 93, 123  
 time interval, 358  
 time slice (in DBNs), 480  
 time well spent, 1066  
 Timofeev, A., 48, 1104  
 Tinsley, M., 223  
 Tipping, M. E., 838, 1115  
 Tirole, J., 638, 1095  
 Tishby, N., 516, 1095  
 Tit-for-Tat, 604  
 Titterington, D. M., 799, 1115  
 TMS (truth maintenance system), 206, 353, 353–355, 360  
 Tobarra, L., 268, 1085  
 Tobin, J., 986, 1085  
 Todorov, E., 985, 1104  
 Tohmé, F., 638, 1111  
 tokenization, 876  
 Tolpin, D., 587, 1070, 1098  
 Toma, P., 931, 1115  
 Tomasi, C., 1030, 1115  
 Tomlin, C. J., 872, 1085  
 Tononi, G., 1058, 1108  
 toothache, 404  
 Topcu, U., 359, 1065, 1104  
 top-down inductive learning methods, 761–763  
 Topol, E., 48, 1104, 1115  
 topological ordering, 434  
 topological sort, 183, 184  
 Torbica, A., 549, 1087  
 torque sensor, 936  
 Torralba, A., 399, 708, 1115  
 Torras, C., 162, 401, 1100  
 Toshev, A., 48, 930, 1115  
 total cost, 122  
 touring problem, 88  
 Toutanova, K., 930, 1093  
 Tovey, C. A., 639, 1086  
 TPTP (Thousands of Problems for Theorem Provers), 331  
 TPU (tensor processing unit), 33, 45, 814  
 tractability of inference, 27, 350  
 tragedy of the commons, 627, 1053  
 training, 684  
 training curve, 701  
 training set, 671, 684  
     weighted, 717  
 trajectory, 946  
 trajectory tracking control, 946, 985  
 Tramér, F., 1060, 1115  
 Tran, D., 667, 1115  
 Tran-Gia, P., 1068, 1099

transfer learning, 670, **832**, 922, 967, 1067  
 transformer decoder, **922**  
 transformer encoder, **922**  
 transformer model, 884, **919**  
 transhumanism, **1056**, 1061  
 transient failure, **505**  
 transient failure model, **506**  
 transition kernel, **462**  
 transition model, **70**, **83**, 123, **140**, **193**, 256, 479, 482, 510, 514, 553, 585, 939  
 sparse, 560  
 transitivity (of preferences), **520**  
 translation, *see* machine translation transparency transpose, **1077**  
 transposition (in a game), **201**  
 transposition table, **201**  
 Trappenberg, T., 839, **1115**  
 traveling salesperson problem (TSP), 88, 125, 127  
 Traverso, P., 399, 401, 402, **1087**, **1091**, 1096  
 treasure hunt, **541**  
 TREC (Text REtrieval Conference), 901  
 tree, 183  
 TREE-CSP-SOLVER, 184  
 tree-like search, **92**  
 treebank, 890, 902  
 Penn, **880**, 890, 928  
 tree decomposition, **185**, 188  
 tree width, **186**, 188, 190, 452  
 Treichler, S., 48, **1103**  
 trial (in reinforcement learning), **843**  
 triangle inequality, **106**  
 trichromacy, **995**  
 Trick, M. A., 639, **1086**  
 Triggs, B., 1029, **1092**  
 Trivedi, H., 48, **1093**  
 Troyanskii, P., 931  
 Truby, R.L., 1064, **1102**, **1107**  
 true concurrency, **592**  
 true majority rule voting, **630**  
 trust, 728, **1047**, 1060  
 trust region policy optimization, 872  
 Truszkowski, P., 48, **1113**  
 truth, **232**, 278  
 truth-preserving inference, **234**  
 truth-revealing, **625**  
 truth maintenance system (TMS), 190, **353**, 353–355, 360  
 assumption-based, **354**  
 justification-based, **354**  
 truth table, **237**, 266  
 truth value, **236**  
 Tsang, E., 190, 191, **1099**, **1115**  
 Tse, D., 838, **1118**  
 Tshitoyan, V., 923, **1115**

Tsipras, D., 838, **1090**  
 Tsitsiklis, J. N., 428, 586, 588, 856, 871–873, 1080, **1087**, **1105**, **1108**, 1115  
 TSP (traveling salesperson problem), 88, 125, 127  
 Tsuruoka, Y., 1072, **1098**  
 TT-CHECK-ALL, **239**  
 TT-ENTAILS?, **239**  
 Tuberg, D., 474, **1112**  
 Tukey, J. W., 726, 737, **1115**  
 Tulsiani, M., 1020, **1101**  
 Tumer, K., 639, **1115**  
 Tung, F., 515, **1110**  
 tuple, **274**  
 turbo decoding, **476**  
 Turian, J., 929, **1115**  
 Turing, A., 20, 27, 32, 35, 37, 51, 74, 221, 300, 328, 472, 733, 836, 870, 1032–1035, 1037, 1057, 1061, 1071, 1073, **1115**  
 Turing Award, 359, 670, 839, 1080  
 Turing machine, 27, 734  
 Turing test, **20**, 20, 22, 874, 1035, 1057  
 total, **20**  
 turtle, 983  
 Tversky, A., 436, 528, 550, **1095**, **1100**, 1115  
 two-finger Morra, 595  
 2001: A Space Odyssey, 472  
 Tygar, J. D., 1061, **1086**, **1115**  
 Type A strategy, **201**  
 Type B strategy, **201**  
 type signature, **644**  
 typical instance, 338  
 Tzemach, A., 474, **1112**

---

## U

---

*U* (utility), 519  
 $u_T$  (best prize), 523  
 $u_{\perp}$  (worst catastrophe), 523  
 UAV (unmanned aerial vehicle), 934  
 Uber, 1031  
 UCB1 (upper confidence bound), **209**  
 UCPOP (planning system), 398  
 UCT (game-tree search algorithm), **209**  
 Udluft, S., 872, **1098**  
 UI (Universal Instantiation), **298**  
 Ulam, S., 222, **1106**  
 Ullman, J.D., 329, **1086**, **1115**  
 Ullman, S., 1029, 1030, **1099**, **1115**  
 ULMFiT (natural language software), 930  
 ultimatum game, **632**  
 ultraintelligent machine, **1055**  
 unbalanced classes, **725**  
 unbiased (estimator), **526**  
 unbounded-cost search, **110**  
 uncertain reasoning, 404

uncertainty, 41, 333, 403–429  
 existence, **648**  
 identity, **648**  
 relational, **646**  
 rule-based approach to, 477  
 summarizing, 404  
 and time, 479–483  
 unconditional probability, *see* probability, prior underfitting, **673**  
 undergeneration, **886**  
 undersampling, **725**  
 unification, **301**, 301–302, 327, 328  
 and equality, 324  
 equational, **325**  
 unifier, **301**  
 most general (MGU), 302, 304, 324  
 UNIFORM-COST-SEARCH, **95**  
 uniform-cost search, **95**, 95–96, 123  
 uniform convergence theory, 735  
 uniform distribution, 409, **1080**  
 uniform prior, 775  
 UNIFY, **303**  
 UNIFY-VAR, **303**  
 UNIMATE (robot arm), 983  
 uninformative prior, **784**  
 uninformed search, 81, 94–102, 123, 125  
 unintended side effect, **1053**  
 unique names assumption, 282, 643  
 unit (in neural networks), **802**  
 unit clause, **244**, 251, 326  
 United Nations, 49  
 United States, 31, 32, 523, 1050  
 unit preference (in resolution), **326**  
 unit propagation, **251**  
 unit resolution, **243**, 326  
 units function, **337**  
 universal approximation, 803  
 Universal Dependencies (natural language data set), 890  
 universal grammar, **905**  
 Universal Instantiation (UI), **298**  
 universal plan, 401  
 universal quantifier, **279**  
 unknown word, 878  
 unmanned aerial vehicle (UAV), 934  
 unobservability, 69  
 UNPOP (planning system), 398  
 unrolling, 508, **647**  
 unsatisfiability, 264  
 unsupervised learning, **671**, 789–791  
 unsupervised parsing, **891**  
 UOSAT-II (satellite), 402  
 update (in temporal reasoning), 148, 150, 151, 484–486  
 upper confidence bound (UCB), **575**  
 upper ontology, 355  
 Urban, J., 330, **1085**  
 Urban Challenge, 984

- Urmson, C., 984, 1115  
 urn-and-ball model, 773  
 Usher, J. M., 357, 1095  
 Usher, N., 1069, 1093  
 Uszkoreit, J., 901, 919, 931, 1115  
 utilitarianism, 26  
 utilitarian social welfare, 599  
 utility, 25, 72, 405
  - adaptive, 551
  - axioms of, 521
  - estimation, 843
  - expected, 73, 80, 405, 518, 519, 524
  - function, 72, 73, 193, 519, 522–529
  - independence, 534
  - of money, 524–526
  - multiattribute, 530–534, 547
  - multiplicative, 534
  - node, 535
  - normalized, 523
  - ordinal, 522
  - theory, 10, 405, 519–522, 547
 utility-based agent, 73  
 utopia, 1073  
 UWL (planning system), 401
- 
- V**
- vacuum tube, 35  
 vacuum world, 55, 57
  - erratic, 140
  - kindergarten, 150
  - slippery, 143
 VAE, *see* autoencoder, variational  
 vagueness, 477, 875  
 Valdés, V., 873, 1087  
 Valiant, L., 735, 1115  
 validation set, 684  
 validity, 240, 264  
 Vallati, M., 399, 1115  
 value (of a variable), 77
  - in a CSP, 164**VALUE-ITERATION**, 563  
 value alignment problem, 23, 1054  
 value function, 522
  - additive, 533
 value iteration, 562, 562–566, 585
  - point-based, 588
 value node, *see* utility node  
 value of computation, 1070  
 value of information, 537–543, 547, 579, 1070  
 value of perfect information (VPI), 538  
 value symmetry, 187  
 value vector (in transformers), 920  
**VAMPIRE** (theorem prover), 330, 331  
 vanBeek, P., 188–191, 358, 399, 1086, 1102, 1111, 1115  
 van Benthem, J., 297, 1115  
 Vandenberghe, L., 161, 1088  
 van Doorn, A. J., 1030, 1102  
 van Harmelen, F., 361, 1115  
 van Heijenoort, J., 331, 1115  
 van Hoeve, W.-J., 189, 1115  
 Vanhoucke, V., 900, 905, 985, 1088, 1099  
 vanishing gradient, 914  
 vanishing point, 991  
 van Kleef, P., 357, 1103  
 van Lambalgen, M., 358, 1115  
 van Maaren, H., 267, 1087  
 van Nunen, J. A. E. E., 586, 1115  
 van Run, P., 191, 1086  
 Vanschoren, J., 737, 1099  
 Van den Broeck, G., 668, 1097, 1101  
 van den Driessche, G., 45, 48, 1112  
 van den Oord, A., 47, 830, 838, 900, 1115  
 van der Gaag, L., 428, 1104  
 van der Maaten, L., 1060, 1097  
 Van Gysel, C., 901, 1085  
 Van Hentenryck, P., 189, 1115  
 Van Ooyen, B., 903, 1089  
 Van Roy, B., 856, 871, 1115  
 Van Roy, P. L., 312, 329, 1115  
 Vapnik, V. N., 44, 735, 736, 738, 837, 1029, 1088, 1091, 1103, 1115  
 Varaiya, P., 79, 873, 1095, 1102  
 Vardi, M. Y., 359, 1094, 1115  
 variable, 77
  - atemporal, 256
  - elimination, 448, 448–451, 472, 474, 508
  - in a CSP, 164
  - in continuous state space, 138
  - indicator, 791
  - irrelevant, 451
  - in logic, 279
  - ordering, 177, 450
  - random, 408, 435
    - Bernoulli, 408
    - Boolean, 408
    - continuous, 409, 440, 473
    - unmodeled, 469
 Varian, H. R., 47, 639, 737, 1105, 1110, 1115  
 variance, 673, 688, 1079  
 variational approximation, 476  
 variational autoencoder, 838  
 variational lower bound, 829  
 variational parameter, 476  
 variational posterior, 829  
 Varoquaux, G., 738, 1109  
 Varshney, K. R., 1046, 1047, 1060, 1061, 1085, 1098  
 Varshney, L., 931, 1101  
 Varzi, A., 358, 1090  
 Vasilache, N., 1069, 1115  
 Vasserman, L., 1046, 1106  
 Vaswani, A., 901, 919, 931, 1115  
 Vaucanson, J., 983  
 Vaughan, J. W., 1046, 1096  
 Vazirani, U., 160, 738, 1085, 1101  
 Vazirani, V., 638, 1108  
 Vazquez, D., 873, 1110  
 VC dimension, 735  
 Vckay, E., 873, 1100  
 Veach, E., 475, 1115  
 Vecchi, M. P., 160, 190, 1101  
 Vecchione, B., 1046, 1096  
 vector, 1076  
 vector field histogram, 985  
 Veeramachaneni, K., 737, 1101  
 Veinott, A. F., 587, 1101  
 Veit, A., 1060, 1086  
 Veloso, M., 225, 986, 1085, 1112  
 Veness, J., 588, 871, 873, 1087, 1106, 1112  
 Venkataraman, S., 587, 1097  
 Venkatesh, S., 429, 1115  
 Venn, J., 427  
 Venugopal, A., 931, 1118  
 Venugopalan, S., 48, 1097, 1104  
 Vere, S. A., 401, 1115  
 verification, 327
  - hardware, 294
 verification and validation, 1047  
 Verma, S., 1060, 1115  
 Verma, T., 473, 474, 798, 1096, 1108  
 Verma, V., 516, 1115  
**VERSION-SPACE-LEARNING**
  - algorithm, 746–747
 version space method, 768  
 Verweij, G., 1049, 1110  
 Vetterling, W. T., 160, 1109  
 Vickrey, W., 626  
 Vickrey-Clarke-Groves mechanism (VCG), 628  
 video game, 225
  - Atari, 46, 224, 835, 867
 Viegas, E., 737, 1097  
 Viégas, F., 737, 1116  
 Vienna Circle, 25  
 Vig, L., 1041, 1105  
 Vihma, T., 161, 1113  
 Vincent, P., 929, 1087  
 Vinge, V., 30, 1115  
 Vinyals, O., 48, 225, 734, 830, 837, 838, 873, 900, 929–931, 1097, 1100, 1104, 1114, 1115, 1116, 1118  
 Viola, P., 1029, 1115  
 Virasoro, M., 160, 1106  
 Virochsiri, K., 873, 1096  
 virtual count, 782  
 visibility graph, 950, 985  
 vision, 20, 30, 38, 188, 989–1026  
 Visser, U., 225, 1115  
 Visser, W., 327, 1098  
 visual programming, 975  
 Vitali set, 411  
 Vitányi, P., 735, 1104

Viterbi, A. J., 515, 1115  
 Viterbi algorithm, 491, 881  
 Vlassis, N., 588, 640, 1113, 1116  
 Vlimant, J.-R., 1069, 1107  
 VLSI layout, 88, 125, 133  
 Vogt, D. M., 1064, 1107  
 Volk, K., 799, 1096  
 von Mises, R., 427, 1116  
 von Neumann, J., 28, 34, 35, 521, 548, 637, 638, 1116  
 von Stengel, B., 611, 637, 1102  
 von Winterfeldt, D., 548, 1116  
 von Linne, C., 357  
 Vorhees, E., 901, 1085  
 Voronkov, A., 297, 330, 331, 1110  
 Voronoi diagram, 951  
 Voronoi graph, 951  
 Vossen, T., 399, 1116  
 voted perceptron, 736  
 voting, 630  
 VPI (value of perfect information), 538  
 VQA (question answering visual), 46, 1017

---

W

---

Wadsworth, C. P., 296, 1097  
 wafer scale engine (WSE), 33  
 Wagner, D., 126, 1093  
 Wagner, S. K., 48, 1104  
 Wahba, G., 734, 1097  
 Wainwright, M., 267, 476, 873, 1087, 1105, 1116  
 Waldinger, R., 296, 297, 1105  
 Walker, E., 47, 1091  
 Walker, G., 838, 1116  
 Walker, H., 799, 1096  
 Walker, R. J., 189, 1116  
 WALKSAT, 253  
 Wall, R., 904, 1093  
 Wallace, A. R., 136, 1116  
 Wallach, H. M., 1046, 1096  
 Walpole, R.E., 429, 1116  
 Walras, L., 28  
 Walsh, M. J., 161, 1095  
 Walsh, T., 189, 191, 267, 1049, 1087, 1111, 1113, 1116  
 Walsh, W., 639, 1116  
 Walter, G., 983  
 Walter, M. R., 986, 1114  
 Waltz, D., 38, 188, 735, 1113, 1116  
 WAM (Warren Abstract Machine), 329  
 Wang, A., 930, 931, 1116  
 Wang, D. Z., 901, 906, 1089  
 Wang, E., 360, 1114  
 Wang, H., 930, 1114  
 Wang, J., 222, 1116  
 Wang, L., 837, 1093  
 Wang, S., 930, 1114  
 Wang, T., 33, 1069, 1117

Wang, Z., 872, 1117  
 Wanner, E., 271, 1116  
 Ward, T., 873, 1087  
 Warde-Farley, D., 838, 1097  
 Warmuth, M., 124, 735, 1088, 1110  
 Warner, C., 903, 1087  
 WARPLAN (planning system), 398  
 Warren, D. H. D., 312, 329, 398, 1109, 1116  
 Warren, D. S., 329, 1114  
 Warren Abstract Machine (WAM), 329  
 Washington, G., 343  
 wasp, sphex, 59, 392  
 Wasserman, L., 738, 1116  
 watched literal, 266  
 Watkins, C. J., 586, 871, 1116  
 Watson (question-answering system), 44  
 Watson, J., 31  
 Watson, J.-P., 190, 1087  
 Watson, J. D., 136, 1116  
 Watt, J., 34  
 Wattenberg, M., 161, 737, 1100, 1116  
 Watts, M., 1064, 1109  
 Waugh, K., 48, 224, 638, 1107, 1116  
 WaveNet, 830, 838  
 WaveNet (speech generation software), 900  
 Way, D. H., 48, 1104  
 Waymo, 46, 1031, 979, 980  
 Wayne, G., 986, 1098  
 Wayne, K., 1080, 1112  
 WBRIDGE5 (bridge program), 224  
 weak AI, 1032, 1056, 1057  
 weak domination, 597  
 weakly supervised learning, 723, 1068  
 weak method, 40  
 weapon  
   lethal autonomous, 1059  
 weapon, lethal autonomous, 49, 1038  
 Weaver, W., 679, 733, 738, 902, 931, 1112  
 Weber, J., 516, 1099  
 Webster, D. R., 48, 1104  
 Wefald, E.H., 127, 221, 1070, 1111  
 Wegbreit, B., 984, 1106  
 Weglarz, J., 402, 1088  
 Wei, J. N., 1027, 1111  
 Wei, X., 906, 1109  
 Weibull, J., 638, 1116  
 Weidenbach, C., 330, 1116  
 weight, 694  
 weight decay, 822  
 WEIGHTED-SAMPLE, 458  
 weighted A\* search, 109  
 weighted linear function, 203  
 weight space, 695  
 Weinstein, S., 734, 1108  
 Weiss, D., 904, 1085  
 Weiss, G., 80, 640, 1116  
 Weiss, R., 738, 900, 1090, 1109  
 Weiss, S., 903, 1085  
 Weiss, Y., 476, 517, 708, 1107, 1115–1117  
 Weissbrod, O., 474, 1112  
 Weissman, V., 296, 1097  
 Weizenbaum, J., 1052, 1116  
 Weld, D. S., 162, 357, 360, 398–402, 905, 906, 1059, 1091, 1094, 1103, 1109, 1113, 1116, 1117  
 Welinder, P., 986, 1085  
 well calibrated, 1044  
 Welling, M., 838, 1101  
 Wellman, M. P., 473, 478, 516, 666, 549, 586, 638, 639, 985, 1092, 1099, 1116  
 Welty, C., 48, 357, 1094, 1113  
 Wen, M., 359, 1065, 1104  
 Werbos, P., 586, 587, 836, 837, 871, 1116  
 Wermuth, N., 473, 1103  
 Werneck, R. F., 126, 1096  
 Wertheimer, M., 1028  
 Wesley, M. A., 985, 1116  
 West, Col., 305  
 West, D. M., 1062, 1116  
 West, M., 517, 1104  
 West, S. M., 1046, 1116  
 Westinghouse, 401  
 Weston, L., 923, 1115  
 Wexler, Y., 474, 1116  
 Weymouth, T., 983, 1091  
 Wheatstone, C., 1028, 1116  
 White, C., 838, 1116  
 White, J., 738, 1091  
 White, J. L., 327, 1098  
 Whitehead, A. N., 35, 328, 1116  
 Whitehouse, D., 222, 1089  
 Whiter, A. M., 401, 1114  
 Whitney, W. F., 225, 1095  
 Whittaker, M., 1046, 1116  
 Whittaker, W., 46, 984, 1101, 1115  
 Whittle, P., 571, 1116  
 Whorf, B., 270, 296, 1116  
 Widner, K., 48, 1097  
 Widrow, B., 39, 836, 870, 1116  
 Widrow–Hoff rule, 855  
 Wiedijk, F., 331, 1116  
 Wiegley, J., 162, 1116  
 Wiener, N., 34, 36, 51, 57, 221, 515, 836, 931, 1059, 1062, 1110, 1116  
 Wierstra, D., 835, 838, 841, 871, 873, 986, 1104, 1106, 1110  
 Wiesel, T. N., 837, 1029, 1099  
 wiggly belief state, 261  
 Wijmans, E., 873, 1111  
 Wikipedia, 334, 874, 876  
 Wilcox, S., 327, 1086  
 Wild, P. P., 475, 1096  
 Wildes, R. P., 1064, 1101

- Wilensky, R., 41, **1116**  
 Wilfong, G. T., 983, **1091**  
 Wilkins, D. E., 401, **1116**  
 Wilks, Y., 1061, 1062, **1116**  
 Williams, A., 931, **1116**  
 Williams, B., 267, 360, 402, **1107**, **1116**  
 Williams, C. K. I., 799, **1110**  
 Williams, J., 588, **1117**  
 Williams, R. J., 586, 836, 837, 863, 871, 872, **1109**, **1111**, **1116**  
 Williamson, J., 358, **1107**  
 Williamson, M., 401, **1094**  
 Willighagen, E. L., 358, **1107**  
 Wilmer, E. L., 515, **1103**  
 Wilson, A., 903, **1103**  
 Wilson, D. H., 1052, **1116**  
 Wilson, R., 188, **1116**  
 Wilson, R. A., 21, 1058, **1116**  
 Wilt, C. M., 126, **1117**  
 Wiltschko, A. B., 1027, **1111**  
 Windows (operating system), 473  
 Winfield, A., 160, 1048, **1089**, **1093**  
 Wingate, D., 668, 586, **1117**  
 Winikoff, M., 79, **1108**  
 Winker, S., 331, **1117**  
 Winkler, R. L., 527, 549, **1113**  
 winner's curse, **549**  
 Winograd, S., 38, **1117**  
 Winograd, T., 38, 41, 905, **1117**  
 Winograd Schema Challenge, 928  
 Winston, P. H., 38, 50, **1087**, **1117**  
 Wintermute, S., 329, **1117**  
 Winternitz, L., 47, **1117**  
 Witbrock, M., 357, **1105**  
 Witten, D., 738, **1100**  
 Witten, I. H., 738, 871, 902, 903, **1117**  
 Witten–Bell smoothing, 902  
 Wittgenstein, L., 25, 234, 266, 267, 338, 357, **1117**  
 Wohlhart, P., 985, **1088**  
 Wojciechowski, W. S., 327, **1117**  
 Wojcik, A. S., 327, **1117**  
 Wolf, A., 904, **1097**  
 Wolf, T., 930, **1111**  
 Wolfe, J., 162, 222, 400, **1105**, **1111**, **1117**  
 Wolpert, D., 639, 733, **1115**, **1117**  
 Wolsey, L. A., 551, **1107**  
 Wolski, F., 986, **1085**  
 Wong, C., 737, **1117**  
 Wong, K. W., 1035, **1100**  
 Wong, W.-K., 798, **1107**  
 Wood, D. E., 126, **1103**  
 Wood, F., 475, 516, 668, **1103**, **1108**  
 Wood, R. J., 1064, **1107**  
 Woodruff, A., 1060, **1087**  
 Woods, W. A., 359, 905, **1117**  
 Wooldridge, M., 79, 80, 638, 640, **1090**, **1091**, **1110**, **1117**
- Woolsey, K., 867  
 word, 421, 875  
     out-of-vocabulary, 878  
 WORD2VEC (word embedding software), 908, 926, 929  
 word embedding, 834, 879, **908**  
     positional, **921**  
 WordNet, **879**  
 WordNet (lexical database), 357  
 work, future of, 1049–1051  
 workspace, **946**  
 world model, in disambiguation, 899  
 world state, 84  
 World War II, 28, 472, 515  
 worst possible catastrophe, 523  
 Worswick, S., 1057  
 Wos, L., 330, 331, **1117**  
 Wossnig, L., 1069, **1093**  
 Wray, R. E., 329, **1117**  
 Wright, S., 161, 472, **1117**  
 WSE (wafer scale engine), 33  
 Wu, D., 48, **1097**  
 Wu, E., 901, 906, **1089**  
 Wu, F., 357, **1117**  
 Wu, H., 930, **1114**  
 Wu, I.-C., 222, **1116**  
 Wu, J., 930, **1110**  
 Wu, L., 47, **1117**  
 Wu, S., 1046, **1106**  
 Wu, Y., 47, 517, 668, 834, 837, 900, 901, 916, 929, 1046, **1090**, **1094**, **1100**, **1117**  
 wumpus world, **228**, 228–231, 238, 268, 288–289, 333, 422–426, 884  
 Wundt, W., 30, 1027  
 Wurman, P., 639, **1116**
- 
- X**
- XAI (explainable AI), 737, **1048**  
 XCON (expert system), 310  
 Xiong, C., 931, 1072, **1098**, **1101**  
 Xiong, W., 47, **1117**  
 XLM (multilingual language model), 930  
 XLNET (natural language system), 930  
 xor, 237  
 Xu, B., 838, **1097**  
 Xu, J., 329, **1117**  
 Xu, P., 903, 931, **1088**
- 
- Y**
- Yahtzee, 214  
 Yakimovsky, Y., 548, **1094**  
 Yale, 41  
 Yampolskiy, R. V., 1061, **1117**  
 Yan, D., 401, **1095**  
 Yang, B., 901, **1106**  
 Yang, G., 516, **1117**  
 Yang, X.-S., 160, **1117**  
 Yang, Y., 838, 930, **1104**, **1117**
- Yang, Z., 930, **1117**  
 Yannakakis, M., 162, 190, **1087**, **1108**  
 Yao, L., 47, **1118**  
 Yao, X., 160, **1104**  
 Yap, R. H. C., 330, **1100**  
 Yarowsky, D., 44, 906, **1117**  
 Yates, A., 906, **1094**  
 Ye, K. E., 429, **1116**  
 Ye, N., 588, **1086**  
 Ye, X., 873, **1096**  
 Ye, Y., 586, **1117**  
 Yedidia, J., 476, **1117**  
 Yeo, H.-S., 1064, **1117**  
 Ying, C., 33, 1069, **1117**  
 Yip, K. M.-K., 360, **1117**  
 Yngve, V., 904, **1117**  
 Yob, G., 268, **1117**  
 Yoon, C. J. M., 1060, **1095**  
 York, S., 873, **1087**  
 Yoshida, Y., 837, **1106**  
 Yoshikawa, T., 985, **1117**  
 You, Y., 872, **1117**  
 Young, C., 670, 1069, **1092**, **1100**  
 Young, H. P., 640, **1117**  
 Young, S., 588, **1117**  
 Young, T., 995  
 Younger, D. H., 886, 904, **1117**  
 Yu, B., 473, **1090**  
 Yu, D., 839, 900, 905, **1093**, **1099**, **1117**  
 Yu, F. X., 1043, **1102**  
 Yu, H.-F., 736, **1117**  
 Yu, K., 838, **1117**  
 Yudkowsky, E., 51, 1061, **1117**  
 Yule, G. U., 838, **1117**  
 Yule–Walker equations, **830**  
 Yvanovich, M., 402, **1092**
- 
- Z**
- Z-3 (early computer), 32  
 Zadeh, L. A., 477, 478, **1117**  
 Zadeh, R. B., 738, **1110**  
 Zahavi, U., 106, **1094**  
 Zaldivar, A., 1046, **1106**  
 Zapp, A., 46, 1030, 984, **1093**  
 Zaremba, W., 838, 873, 986, **1085**, **1089**, **1100**, **1114**  
 Zaritskii, V. S., 516, **1117**  
 Zaykov, Y., 667, **1106**  
 Zecchina, R., 267, **1108**  
 Zeckhauser, R., 549, **1117**  
 Zeeberg, A., 833, **1117**  
 Zeilinger, M. N., 872, **1085**  
 Zelle, J., 905, **1117**  
 Zemel, R., 1046, 1060, **1094**, **1117**  
 Zemelman, B. V., 29, **1118**  
 Zen, H., 830, 838, 900, **1115**  
 Zeng, H., 296, **1106**  
 Zermelo, E., 637, **1118**  
 zero-sum game, **193**, 590, **600**

- Zettlemoyer, L., 668, 896, 905, 927, 930, 1104, 1106, 1109, 1118  
Zeuthen strategy, 634  
Zhai, X., 837, 1093  
Zhang, B., 46, 1097  
Zhang, C., 734, 1118  
Zhang, F., 1060, 1115  
Zhang, H., 266, 1118  
Zhang, J., 931, 1110  
Zhang, L., 190, 266, 474, 1107, 1118  
Zhang, M., 904, 1118  
Zhang, N. L., 474, 1118  
Zhang, S., 29, 900, 1118  
Zhang, T. W., 986, 1114  
Zhang, W., 127, 334, 1093, 1102  
Zhang, X., 837, 903, 1098, 1118  
Zhang, Y., 900, 901, 904, 906, 1047, 1060, 1087, 1089, 1118  
Zhang, Z., 106, 1094  
Zhao, J., 903, 1118  
Zhao, K., 896, 904, 905, 1096, 1118  
Zhao, Y., 190, 266, 873, 1107, 1111  
Zhao, Z., 1023, 1093  
Zhou, K., 550, 1118  
Zhou, R., 127, 1118  
Zhou, T., 1021, 1100  
Zhou, Y., 930, 1110  
Zhu, B., 838, 1118  
Zhu, C., 735, 1089  
Zhu, D. J., 984, 1118  
Zhu, J., 904, 1118  
Zhu, J.-Y., 930, 1021, 1022, 1100, 1118  
Zhu, M., 904, 1118  
Zhu, T., 222, 1116  
Zhu, W. L., 334, 1113  
Ziebart, B. D., 868, 970, 986, 1118  
Zilberstein, S., 162, 389, 401, 1098  
Zilles, S., 114, 126, 127, 1085, 1090, 1094, 1103  
Zimdars, A., 637, 1111  
Zimmermann, H.-J., 477, 1118  
Zinkevich, M., 637, 973, 1110, 1118  
Zipf's Law, 929  
Zipf, G., 929, 1118  
Zipser, D., 837, 1116  
Zisserman, A., 734, 1030, 1031, 1087, 1098, 1107  
Zlot, R., 160, 1093  
Zlotkin, G., 639, 640, 1111  
Zobrist, A. L., 223, 1118  
Zollmann, A., 931, 1118  
Zoph, B., 838, 1109, 1118  
Zucker, M., 958, 985, 1110  
Zuckerman, D., 132, 1104  
Zufferey, J. C., 1064, 1095  
Zuse, K., 32, 221, 1118  
Zweben, M., 402, 1092  
Zweig, G., 516, 1118  
Zwicker, W. S., 638, 1114  
zyzzyva, 876